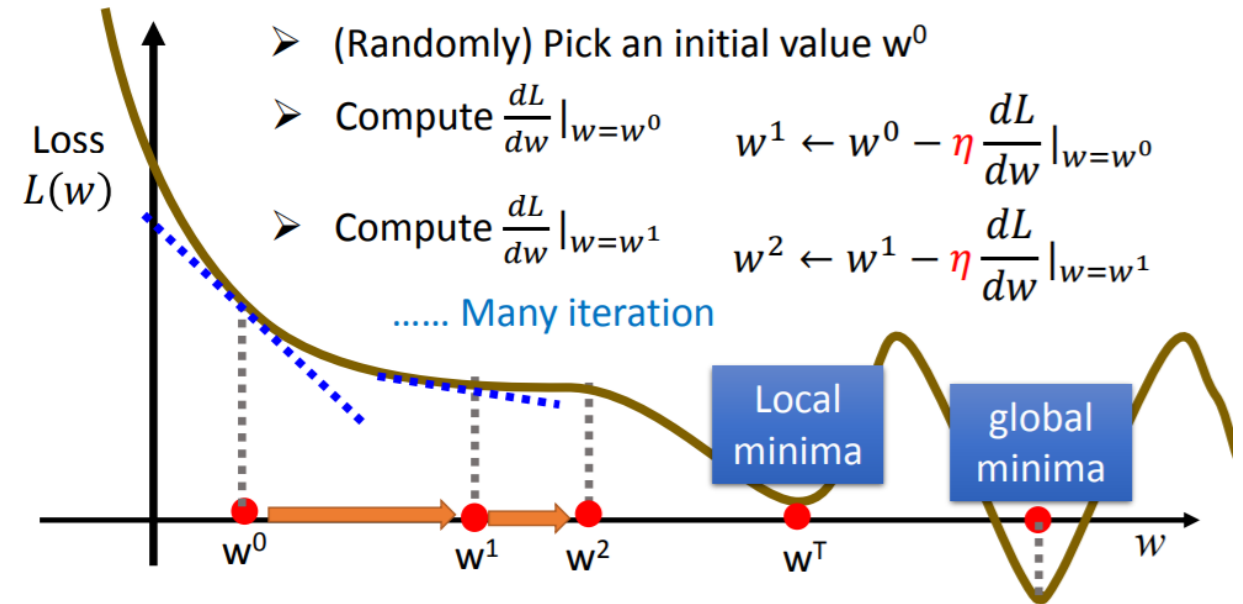
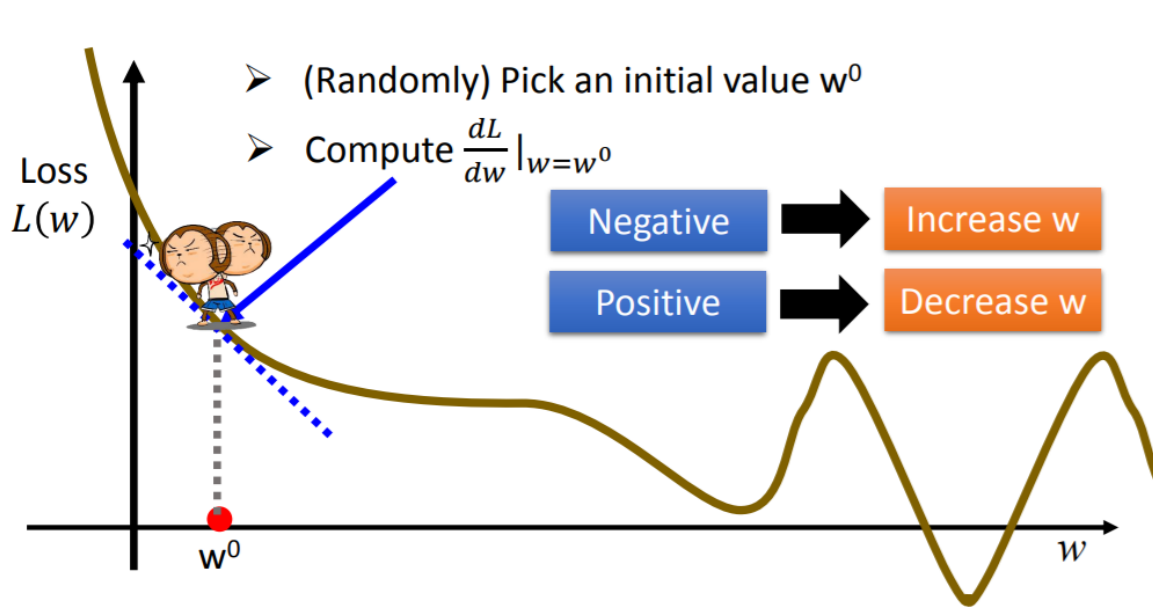


Review – How machine learns from data?

- Define a function to be learned: $y = f(x)$
- Define a loss function \mathcal{L} to describe the error between $\hat{y} = f(x)$ and y
- Find the optimal parameters of f that minimize \mathcal{L}

Review – Gradient descent to find optimal parameters that minimize the loss

2



Reference: 李弘毅 ML Lecture 1 <https://youtu.be/CXgbekl66jc>

Gradient descent to find optimal parameters that minimize the loss


Suppose that θ has two variables $\{\theta_1, \theta_2\}$

Randomly start at $\theta^0 = \begin{bmatrix} \theta_1^0 \\ \theta_2^0 \end{bmatrix}$


$$\nabla L(\theta) = \begin{bmatrix} \partial L(\theta_1)/\partial \theta_1 \\ \partial L(\theta_2)/\partial \theta_2 \end{bmatrix}$$

$$\begin{bmatrix} \theta_1^1 \\ \theta_2^1 \end{bmatrix} = \begin{bmatrix} \theta_1^0 \\ \theta_2^0 \end{bmatrix} - \eta \begin{bmatrix} \partial L(\theta_1^0)/\partial \theta_1 \\ \partial L(\theta_2^0)/\partial \theta_2 \end{bmatrix} \Rightarrow \theta^1 = \theta^0 - \eta \nabla L(\theta^0)$$

$$\begin{bmatrix} \theta_1^2 \\ \theta_2^2 \end{bmatrix} = \begin{bmatrix} \theta_1^1 \\ \theta_2^1 \end{bmatrix} - \eta \begin{bmatrix} \partial L(\theta_1^1)/\partial \theta_1 \\ \partial L(\theta_2^1)/\partial \theta_2 \end{bmatrix} \Rightarrow \theta^2 = \theta^1 - \eta \nabla L(\theta^1)$$



Review – Practice gradient calculation



$$\nabla L = \left[\frac{\partial L}{\partial w_{13}}, \frac{\partial L}{\partial w_{14}}, \frac{\partial L}{\partial b_3}, \dots, \frac{\partial L}{\partial w_{6y}}, \dots \right]$$

Reference: 許洺嘉(1071414)

```
for name, param in MyNet.named_parameters():
    if param.requires_grad:
        print(name, '\n', param.data, '\n', param.grad)
```

```
0.weight
tensor([[ 0.4343,  0.6134],
        [-0.2977,  0.1472],
        [-0.3772, -0.2806]])
None
0.bias
tensor([ 0.2081, -0.5165,  0.3692])
None
1.weight
tensor([[ 0.5137,  0.5441,  0.4459],
        [-0.0712,  0.0366, -0.2097]])
None
1.bias
tensor([0.2730, 0.5439])
None
2.weight
tensor([[0.5989, 0.3610]])
None
2.bias
tensor([-0.3608])
None
```

```
0.weight
tensor([[ 0.4343,  0.6134],
        [-0.2977,  0.1472],
        [-0.3772, -0.2806]])
tensor([[ -80.2851, -46.4245],
        [-96.5485, -55.8287],
        [-54.4879, -31.5074]])
0.bias
tensor([ 0.2081, -0.5165,  0.3692])
tensor([-10.9224, -13.1350, -7.4128])
1.weight
tensor([[ 0.5137,  0.5441,  0.4459],
        [-0.0712,  0.0366, -0.2097]])
tensor([[ -139.3767,  48.2291,  83.4421],
        [ -84.0035,  29.0681,  50.2912]])
1.bias
tensor([0.2730, 0.5439])
tensor([-23.2005, -13.9832])
2.weight
tensor([[0.5989, 0.3610]])
tensor([[ -24.1840, -30.7693]])
2.bias
tensor([-0.3608])
tensor([-38.7382])
```

Review – Practice gradient calculation

Loss function (MSE): $\mathcal{L} = (y - \hat{y})^2$

$$\hat{y} = w_{6y}n_6 + w_{7y}n_7$$

$$\therefore \mathcal{L} = (y - w_{6y}n_6 + w_{7y}n_7)^2$$

$$\therefore \nabla \mathcal{L}(w_6) = \frac{\partial \mathcal{L}}{\partial w_{6y}} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w_{6y}} \text{ (by chain rule)}$$

$$\frac{\partial \mathcal{L}}{\partial \hat{y}} = 2(y - \hat{y}) \cdot (-1)$$

$$\frac{\partial \hat{y}}{\partial w_{6y}} = n_6$$

$$\therefore \nabla \mathcal{L}(w_6) = 2(y - \hat{y}) \cdot (-1) \cdot n_6$$

<table><tr><th>n6</th><th>n7</th></tr><tr><td>0.69657858</td><td>0.51114952</td></tr><tr><td>0.859582</td><td>0.56898564</td></tr><tr><td>0.543156</td><td>0.90834304</td></tr></table>	n6	n7	0.69657858	0.51114952	0.859582	0.56898564	0.543156	0.90834304		<table><tr><th>W6Y</th><td>0.5989</td></tr><tr><th>W7Y</th><td>0.361</td></tr></table>	W6Y	0.5989	W7Y	0.361							
n6	n7																				
0.69657858	0.51114952																				
0.859582	0.56898564																				
0.543156	0.90834304																				
W6Y	0.5989																				
W7Y	0.361																				
<table><tr><th>X*W</th></tr><tr><td>0.601705888</td></tr><tr><td>0.720207476</td></tr><tr><td>0.653207966</td></tr></table>	X*W	0.601705888	0.720207476	0.653207966	+	<table><tr><th>BY</th></tr><tr><td>-0.3608</td></tr><tr><td>-0.3608</td></tr><tr><td>-0.3608</td></tr></table>	BY	-0.3608	-0.3608	-0.3608	=	<table><tr><th>Y_hat</th></tr><tr><td>0.24090589</td></tr><tr><td>0.35940748</td></tr><tr><td>0.29240797</td></tr></table>	Y_hat	0.24090589	0.35940748	0.29240797	<table><tr><th>Y</th></tr><tr><td>7</td></tr><tr><td>12</td></tr><tr><td>40</td></tr></table>	Y	7	12	40
X*W																					
0.601705888																					
0.720207476																					
0.653207966																					
BY																					
-0.3608																					
-0.3608																					
-0.3608																					
Y_hat																					
0.24090589																					
0.35940748																					
0.29240797																					
Y																					
7																					
12																					
40																					
<table><tr><th>GradientsPerData</th></tr><tr><td>-9.416480357</td></tr><tr><td>-20.01208761</td></tr><tr><td>-43.13483372</td></tr></table>	GradientsPerData	-9.416480357	-20.01208761	-43.13483372		<table><tr><th>Gradient n6</th></tr><tr><td>-24.1878006</td></tr></table>	Gradient n6	-24.1878006													
GradientsPerData																					
-9.416480357																					
-20.01208761																					
-43.13483372																					
Gradient n6																					
-24.1878006																					

More details about gradient descent

```
# initialize NN weights
for name, param in MyNet.named_parameters():
    if(param.requires_grad):
        torch.nn.init.normal_(param, mean=0.0, std=0.02)
loss_func = torch.nn.MSELoss()
optimizer = torch.optim.Adam(MyNet.parameters(), lr=0.0003)

# train NN
print("epoch", end=": ")
epoch_lossLst=[]
for epoch in range(1, 700):
    if(epoch%100 == 0):
        print(epoch, end=",")
    for (batchX, batchY) in loader:
        batchY_hat = MyNet(batchX)
        loss = loss_func(batchY_hat, batchY)
        epoch_lossLst.append(float(loss))
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```

Give different parameters
different lr?


Why mini-batch?

Fixed or adaptive
lr?

Learning rates

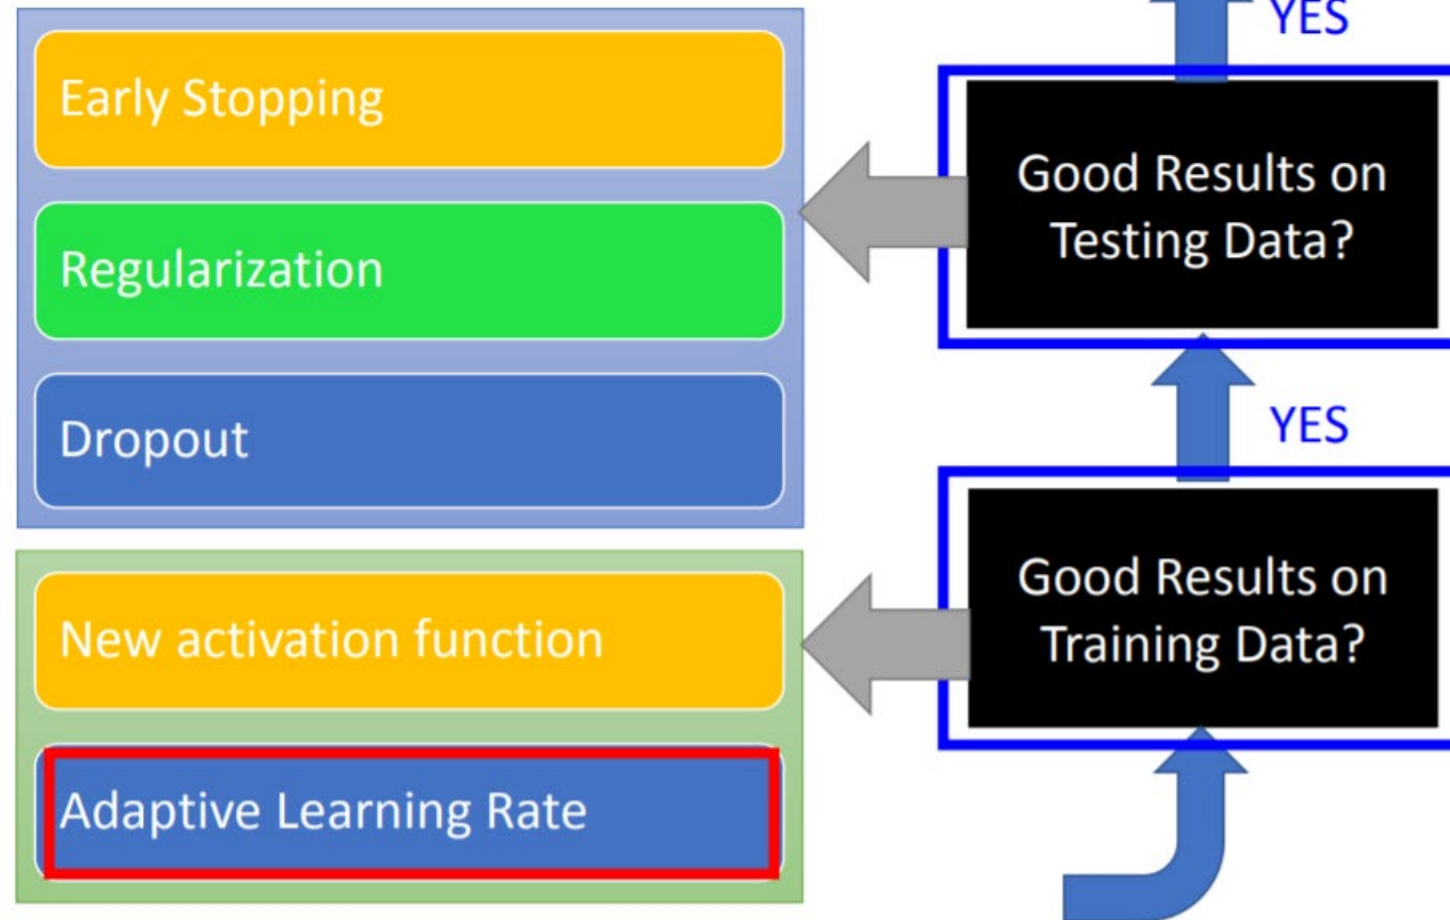
3. Gradient descent.ipynb

- 0.5
- 0.05
- 0.0003



Adaptive learning rates

Recipe of Deep Learning



Learning rate schedule

- Reduce the learning rate by some factor every few epochs.

At the beginning, we are far from the destination, so we use larger learning rate

After several epochs, we are close to the destination, so we reduce the learning rate

3. Gradient descent.ipynb

Adagrad

- Give different parameters different learning rates.
- Divide the learning rate of each parameter by the root mean square of its previous derivatives

$$w^{t+1} \leftarrow w^t - \frac{\eta^t}{\sigma^t} g^t \quad g^t = \frac{\partial L(\theta^t)}{\partial w}$$

η^t : adaptive lr

$$\begin{aligned} w^1 &\leftarrow w^0 - \frac{\eta^0}{\sigma^0} g^0 & \sigma^0 &= \sqrt{(g^0)^2} \\ w^2 &\leftarrow w^1 - \frac{\eta^1}{\sigma^1} g^1 & \sigma^1 &= \sqrt{\frac{1}{2} [(g^0)^2 + (g^1)^2]} \\ w^3 &\leftarrow w^2 - \frac{\eta^2}{\sigma^2} g^2 & \sigma^2 &= \sqrt{\frac{1}{3} [(g^0)^2 + (g^1)^2 + (g^2)^2]} \\ &\vdots & & \\ w^{t+1} &\leftarrow w^t - \frac{\eta^t}{\sigma^t} g^t & \sigma^t &= \sqrt{\frac{1}{t+1} \sum_{i=0}^t (g^i)^2} \end{aligned}$$


Why divided by root mean square of previous derivatives?

$$w^{t+1} \leftarrow w^t - \frac{\eta}{\sqrt{\sum_{i=0}^t (g^i)^2}} g^t$$

Use first derivative to estimate second derivative

3. Gradient descent.ipynb

- Adagrad




Stochastic Gradient Descent (SGD)

- Make the training faster
- Loss for only one sample

Gradient Descent

Update after seeing all examples

$$\theta^t = \theta^{t-1} - \eta \nabla L(\theta^{t-1})$$




$$\nabla L(\theta) = \begin{bmatrix} \partial L / \partial w_1 \\ \vdots \\ \partial L / \partial w_K \end{bmatrix}$$

Stochastic Gradient Descent

Update for each example

If there are 20 examples, 20 times faster.

$$\theta^t = \theta^{t-1} - \eta \nabla L^n(\theta^{t-1})$$




$$\nabla L(\theta) = \begin{bmatrix} \partial L^n / \partial w_1 \\ \vdots \\ \partial L^n / \partial w_K \end{bmatrix}$$

3. Gradient descent.ipynb

- SGD

RMS Prop

Error Surface can be very complex when training NN.



RMS Prop

$$w^1 \leftarrow w^0 - \frac{\eta}{\sigma^0} g^0 \quad \sigma^0 = g^0$$

$$w^2 \leftarrow w^1 - \frac{\eta}{\sigma^1} g^1 \quad \sigma^1 = \sqrt{\alpha(\sigma^0)^2 + (1 - \alpha)(g^1)^2}$$

$$w^3 \leftarrow w^2 - \frac{\eta}{\sigma^2} g^2 \quad \sigma^2 = \sqrt{\alpha(\sigma^1)^2 + (1 - \alpha)(g^2)^2}$$

\vdots

$$w^{t+1} \leftarrow w^t - \frac{\eta}{\sigma^t} g^t \quad \sigma^t = \sqrt{\alpha(\sigma^{t-1})^2 + (1 - \alpha)(g^t)^2}$$

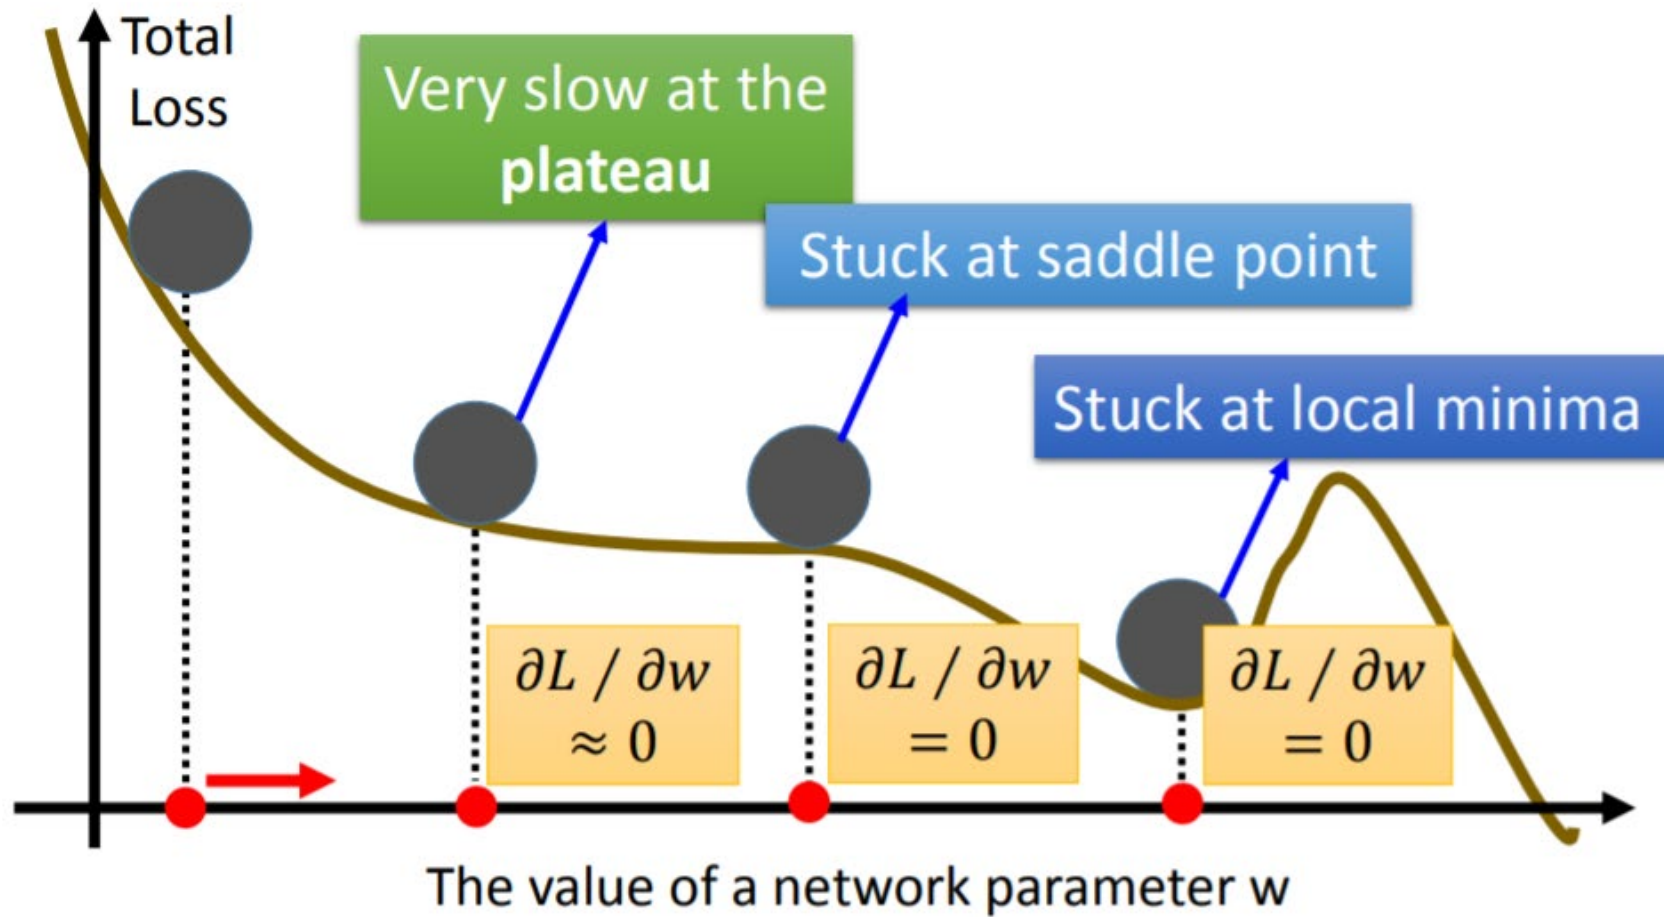
Root Mean Square of the gradients
with previous gradients being decayed

RMS Prop

3. Gradient descent.ipynb


- RMS Prop

Hard to find optimal parameters




Momentum

- Momentum



Vanilla gradient decent



Start at position θ^0

Compute gradient at θ^0

Move to $\theta^1 = \theta^0 - \eta \nabla L(\theta^0)$

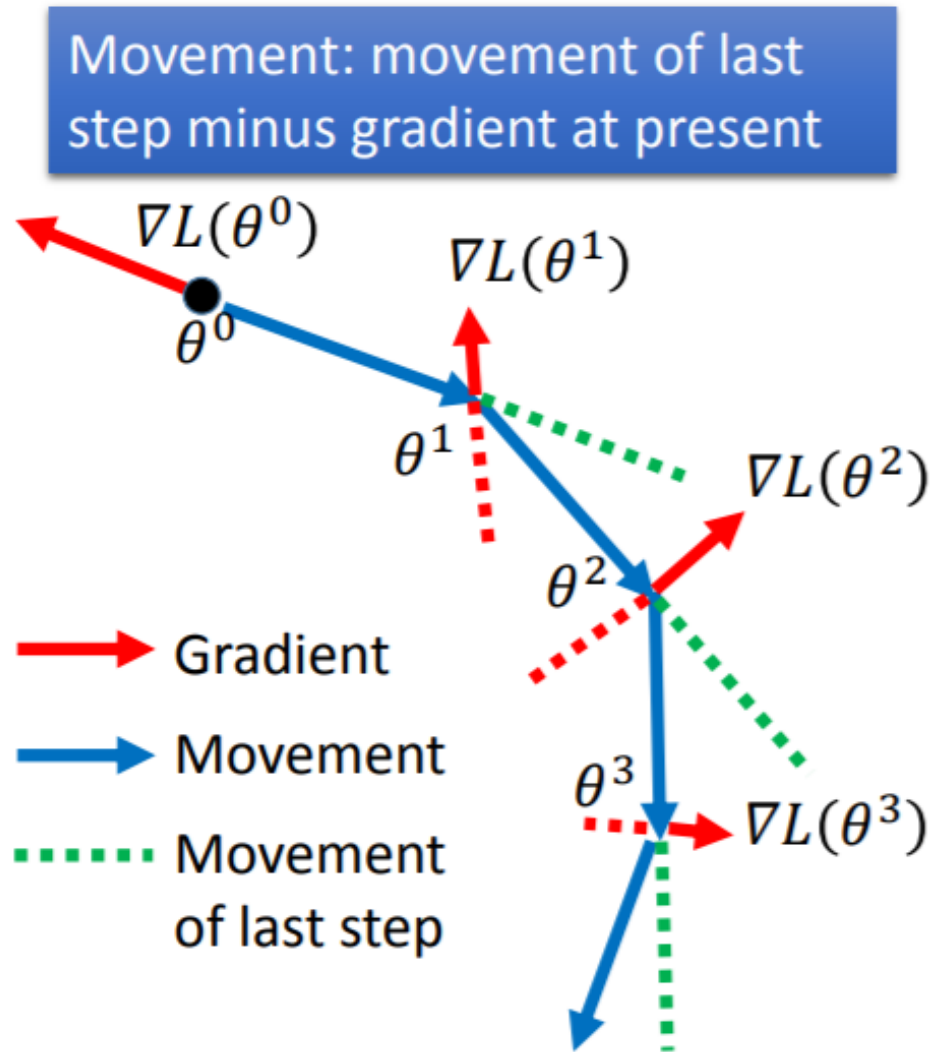
Compute gradient at θ^1

Move to $\theta^2 = \theta^1 - \eta \nabla L(\theta^1)$

...

Stop until $\nabla L(\theta^t) \approx 0$

Momentum



Start at point θ^0

Movement $v^0=0$

Compute gradient at θ^0

Movement $v^1 = \lambda v^0 - \eta \nabla L(\theta^0)$

Move to $\theta^1 = \theta^0 + v^1$

Compute gradient at θ^1

Movement $v^2 = \lambda v^1 - \eta \nabla L(\theta^1)$

Move to $\theta^2 = \theta^1 + v^2$

Movement not just based on gradient, but previous movement.

Momentum

Movement: movement of last step minus gradient at present

v^i is actually the weighted sum of all the previous gradient:

$$\nabla L(\theta^0), \nabla L(\theta^1), \dots \nabla L(\theta^{i-1})$$

$$v^0 = 0$$

$$v^1 = -\eta \nabla L(\theta^0)$$

$$v^2 = -\lambda \eta \nabla L(\theta^0) - \eta \nabla L(\theta^1)$$

\vdots

Start at point θ^0

Movement $v^0=0$

Compute gradient at θ^0

Movement $v^1 = \lambda v^0 - \eta \nabla L(\theta^0)$

Move to $\theta^1 = \theta^0 + v^1$

Compute gradient at θ^1


Movement $v^2 = \lambda v^1 - \eta \nabla L(\theta^1)$

Move to $\theta^2 = \theta^1 + v^2$

Movement not just based on gradient, but previous movement

Momentum

Still not guarantee reaching global minima, but give some hope



Algorithm 1: *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation. g_t^2 indicates the elementwise square $g_t \odot g_t$. Good default settings for the tested machine learning problems are $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$. All operations on vectors are element-wise. With β_1^t and β_2^t we denote β_1 and β_2 to the power t .

Require: α : Stepsize

Require: $\beta_1, \beta_2 \in [0, 1)$: Exponential decay rates for the moment estimates

Require: $f(\theta)$: Stochastic objective function with parameters θ

Require: θ_0 : Initial parameter vector

$m_0 \leftarrow 0$ (Initialize 1st moment vector) \rightarrow for momentum

$v_0 \leftarrow 0$ (Initialize 2nd moment vector) \rightarrow for RMSprop

$t \leftarrow 0$ (Initialize timestep)

while θ_t not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)

end while

return θ_t (Resulting parameters)

Adam

3. Gradient descent.ipynb

- Adam

Why gradient decent can find local minimum?



Why gradient decent can find local minimum?

Based on Taylor Series:

If the red circle is small enough, in the red circle


$$L(\theta) \approx L(a, b) + \frac{\partial L(a, b)}{\partial \theta_1}(\theta_1 - a) + \frac{\partial L(a, b)}{\partial \theta_2}(\theta_2 - b)$$

$$s = L(a, b)$$

$$u = \frac{\partial L(a, b)}{\partial \theta_1}, v = \frac{\partial L(a, b)}{\partial \theta_2}$$

$$L(\theta)$$

$$\approx s + u(\theta_1 - a) + v(\theta_2 - b)$$



Why gradient decent can find local minimum?

Red Circle: (If the radius is small)

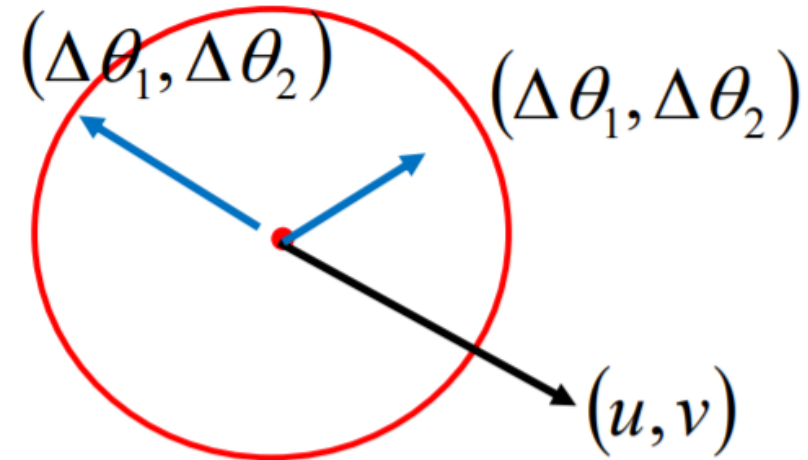
$$L(\theta) \approx \cancel{s} + u \frac{(\theta_1 - a)}{\Delta \theta_1} + v \frac{(\theta_2 - b)}{\Delta \theta_2}$$

Find θ_1 and θ_2 in the red circle
minimizing $L(\theta)$


$$\frac{(\theta_1 - a)^2}{\Delta \theta_1^2} + \frac{(\theta_2 - b)^2}{\Delta \theta_2^2} \leq d^2$$

To minimize $L(\theta)$

$$\begin{bmatrix} \Delta \theta_1 \\ \Delta \theta_2 \end{bmatrix} = -\eta \begin{bmatrix} u \\ v \end{bmatrix} \Rightarrow \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} = \begin{bmatrix} a \\ b \end{bmatrix} - \eta \begin{bmatrix} u \\ v \end{bmatrix}$$




Limitation of gradient decent



Mini-batches

We do not really minimize total loss!

Mini-batch




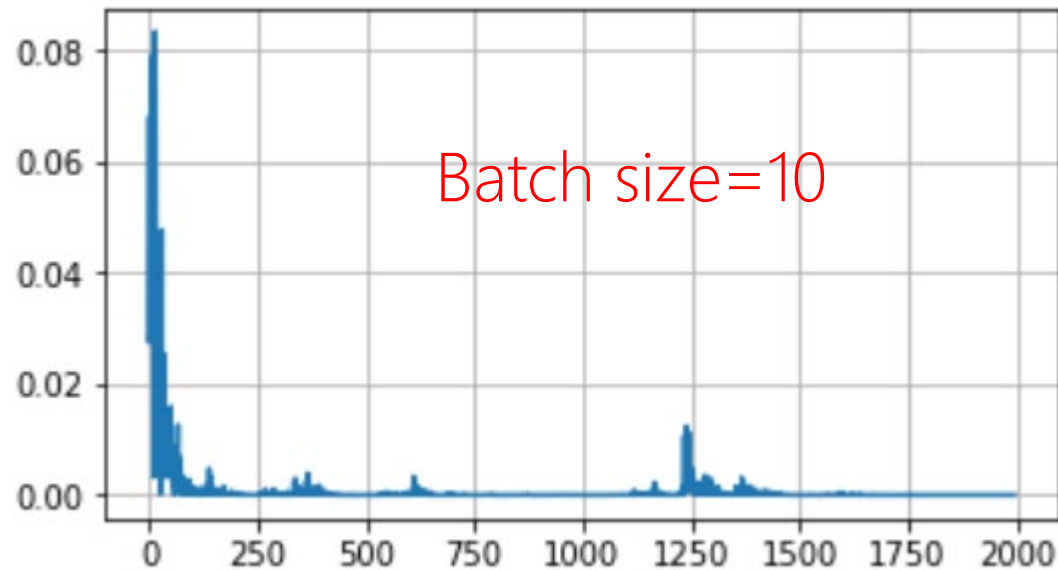
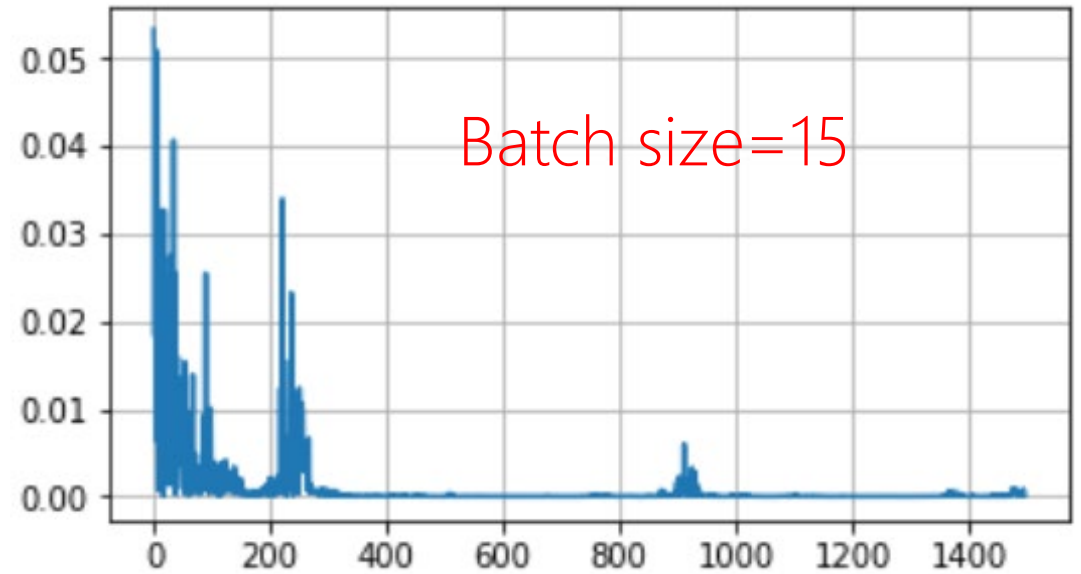

- Randomly initialize network parameters

- Pick the 1st batch
 $L' = l^1 + l^{31} + \dots$
Update parameters once
- Pick the 2nd batch
 $L'' = l^2 + l^{16} + \dots$
Update parameters once
⋮
- Until all mini-batches have been picked

one epoch

Repeat the above process

Run "3.3. Mini_Batch_Training"



Batch size influences both speed and performance. You have to tune it.

Speed

Very large batch size can yield worse performance

- Smaller batch size means more updates in one epoch
 - E.g. 50000 examples
 - batch size = 1, 50000 updates in one epoch 166s 1 epoch
 - batch size = 10, 5000 updates in one epoch 17s 10 epoch

