

Classification

4. Classification.ipynb

Cross entropy

Measures the differences between the true probability p_i and the predicted probability q_i

$$H(p, q) = - \sum_i p_i \ln(q_i)$$


Use Excel to verify
Excel formula: =LN(x)

動物	實際機率分佈	預測機率分佈	Entropy
Cat	0%	2%	$0\% * -\log(2\%) = 0$
Dog	0%	30%	$0\% * -\log(30\%) = 0$
Fox	0%	45%	$0\% * -\log(45\%) = 0$
Cow	0%	0%	$0\% * -\log(0\%) = 0$
Red Panda	100%	25%	$100\% * -\log(25\%) = 1.386$
Bear	0%	5%	$0\% * -\log(5\%) = 0$
Dolphin	0%	0%	$0\% * -\log(0\%) = 0$
總計: cross-entropy = 1.386			

Entropy

More information → more uncertain → larger entropy

$$Entropy = - \sum_i p_i \log_2(p_i)$$



$$\begin{aligned} &75\% \times 0.41 \\ &+ 25\% \times 2 \\ &= 0.81 \text{ bits} \end{aligned}$$



Softmax

```
In [15]: print(tensorY.shape, "\n", tensorY)
```

```
torch.Size([5, 2])  
tensor([[ -0.0180,  0.0855],  
        [ -0.0244,  0.0741],  
        [ -0.0187,  0.0850],  
        [ -0.0258,  0.0687],  
        [ -0.0267,  0.0617]], device='cuda:0', grad_fn=<
```

```
In [16]: # apply softmax  
tensorY = torch.softmax(tensorY, 1)  
print(tensorY.shape, "\n", tensorY)
```

```
torch.Size([5, 2])  
tensor([[0.4742, 0.5258],  
        [0.4754, 0.5246],  
        [0.4741, 0.5259],  
        [0.4764, 0.5236],  
        [0.4779, 0.5221]], device='cuda:0', grad_fn=<So
```

```
In [17]: MaxOfEachRow = torch.max(tensorY, 1)  
print(MaxOfEachRow)
```

```
torch.return_types.max(  
values=tensor([0.5258, 0.5246, 0.5259, 0.5236, 0.5221],  
              grad_fn=<MaxBackward0>),  
indices=tensor([1, 1, 1, 1, 1], device='cuda:0'))
```

`torch.softmax(tensor, 1)`

0: ↓

1: →

`torch.max(tensor, 1)`

$$\frac{e^{y_1}}{e^{y_1} + e^{y_2}}$$

$$\frac{e^{y_2}}{e^{y_1} + e^{y_2}}$$

Torch.max

```
tensor([[0.4742, 0.5258],
        [0.4754, 0.5246],
        [0.4741, 0.5259],
        [0.4764, 0.5236],
        [0.4779, 0.5221]],
```

`torch.max(tensor, 1)[1]`

[1]: The 2nd item of
torch.max results

```
In [17]: MaxOfEachRow = torch.max(tensorY, 1)
         print(MaxOfEachRow)

torch.return_types.max(
  values=tensor([0.5258, 0.5246, 0.5259, 0.5236, 0.5221], device='c
    grad_fn=<MaxBackward0>),
  indices=tensor([1, 1, 1, 1, 1], device='cuda:0'))
```

```
In [18]: MaxIdxOfEachRow = torch.max(tensorY, 1)[1]
         print(MaxIdxOfEachRow)
```

```
tensor([1, 1, 1, 1, 1], device='cuda:0')
```

```
In [19]: correct = 0
         MaxIdxOfEachRow = torch.max(tensorY, 1)[1]
         for i in range(batchY_hat.shape[0]):
             print(int(MaxIdxOfEachRow[i]), int(batchY_hat[i]), end=="==>")
             if (int(MaxIdxOfEachRow[i]) == int(batchY_hat[i])):
                 print("correct")
                 correct += 1
             else:
                 print("wrong")
         print(correct)
         accuracy = correct/batchY_hat.shape[0]
         print("%.2f" % accuracy)
```

```
1 0==>wrong
1 0==>wrong
1 0==>wrong
1 1==>correct
1 1==>correct
2
0.40
```

Bayesian's rule to understand $y_1 = p(C_1|x)$

$$y_1 = P(C_1|x) = \frac{P(x|C_1)P(C_1)}{P(x|C_1)P(C_1) + P(x|C_2)P(C_2)}$$

Generative Model $P(x) = P(x|C_1)P(C_1) + P(x|C_2)P(C_2)$

Probabilistic generative model

$$f_{\mu^1, \Sigma^1}(x) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma^1|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu^1)^T (\Sigma^1)^{-1} (x - \mu^1) \right\}$$

$$\mu^1 = \begin{bmatrix} 75.0 \\ 71.3 \end{bmatrix} \quad \Sigma^1 = \begin{bmatrix} 874 & 327 \\ 327 & 929 \end{bmatrix}$$

$$P(C_1) = 79 / (79 + 61) = 0.56$$

$$P(C_1|x) = \frac{P(x|C_1)P(C_1)}{P(x|C_1)P(C_1) + P(x|C_2)P(C_2)}$$

$$f_{\mu^2, \Sigma^2}(x) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma^2|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu^2)^T (\Sigma^2)^{-1} (x - \mu^2) \right\}$$

$$\mu^2 = \begin{bmatrix} 55.6 \\ 59.8 \end{bmatrix} \quad \Sigma^2 = \begin{bmatrix} 847 & 422 \\ 422 & 685 \end{bmatrix}$$

$$P(C_2) = 61 / (79 + 61) = 0.44$$

$$\text{If } P(C_1|x) > 0.5$$

Assuming x^n are sampled from a Gaussian distribution, then we can use maximum likelihood to find the best Gaussian distribution behind them.


Represent $y_1 = p(C_1|x)$ as a sigmoid function of z

$$y_1 = P(C_1|x) = \frac{P(x|C_1)P(C_1)}{P(x|C_1)P(C_1) + P(x|C_2)P(C_2)}$$

$$= \frac{1}{1 + \frac{P(x|C_2)P(C_2)}{P(x|C_1)P(C_1)}} = \frac{1}{1 + \exp(-z)} = \sigma(z)$$

Sigmoid function

$$z = \ln \frac{P(x|C_1)P(C_1)}{P(x|C_2)P(C_2)}$$



Represent $y_1 = p(C_1|x)$ as a sigmoid function of z

$$P(C_1|x) = \sigma(z)$$

Assuming the covariance matrices of the two classes are the same

$$z = \ln \frac{|\Sigma^2|^{1/2}}{|\Sigma^1|^{1/2}} - \frac{1}{2} x^T (\Sigma^1)^{-1} x + (\mu^1)^T (\Sigma^1)^{-1} x - \frac{1}{2} (\mu^1)^T (\Sigma^1)^{-1} \mu^1 + \frac{1}{2} x^T (\Sigma^2)^{-1} x - (\mu^2)^T (\Sigma^2)^{-1} x + \frac{1}{2} (\mu^2)^T (\Sigma^2)^{-1} \mu^2 + \ln \frac{N_1}{N_2}$$

$$\Sigma_1 = \Sigma_2 = \Sigma$$

$$z = \underbrace{(\mu^1 - \mu^2)^T \Sigma^{-1} x}_{\mathbf{w}^T} - \underbrace{\frac{1}{2} (\mu^1)^T \Sigma^{-1} \mu^1 + \frac{1}{2} (\mu^2)^T \Sigma^{-1} \mu^2 + \ln \frac{N_1}{N_2}}_b$$

$$y_1 = P(C_1|x) = \sigma(\mathbf{w} \cdot \mathbf{x} + b)$$

How about directly find \mathbf{w} and b ?

In generative model, we estimate $N_1, N_2, \mu^1, \mu^2, \Sigma$

Then we have \mathbf{w} and b

Logistic regression

If we use gradient decent to find optimal w and b for the posterior probability $y_1 = p(C_1|x) = \sigma(w \cdot x + b)$, then the problem becomes logistic regression.


We want to find $P_{w,b}(C_1|x)$

If $P_{w,b}(C_1|x) \geq 0.5$, output C_1

Otherwise, output C_2

$$P_{w,b}(C_1|x) = \sigma(z)$$

$$z = w \cdot x + b$$



Logistic regression vs regression

Logistic Regression

$$f_{w,b}(x) = \sigma \left(\sum_i w_i x_i + b \right)$$

Output: between 0 and 1

Linear Regression

$$f_{w,b}(x) = \sum_i w_i x_i + b$$

Output: any value

Maximum likelihood

Assuming the training data is generated from $y_1 = P_{w,b}(C_1 | x) = \sigma(w \cdot x +$

Training Data	x^1	x^2	x^3	$\dots \dots$	x^N
	C_1	C_1	C_2		C_1

$$\max L(w, b) = f_{w,b}(x^1) f_{w,b}(x^2) (1 - f_{w,b}(x^3)) \cdots f_{w,b}(x^N)$$

$$\min -\ln L(w, b) = -\ln f_{w,b}(x^1) - \ln f_{w,b}(x^2) - \ln (1 - f_{w,b}(x^3)) \cdots$$

\hat{y}^n : 1 for class 1, 0 for class 2

$$= \sum_n - \left[\hat{y}^n \ln f_{w,b}(x^n) + (1 - \hat{y}^n) \ln (1 - f_{w,b}(x^n)) \right]$$

Cross entropy between two Bernoulli distribution

Loss function

Training data: (x^n, \hat{y}^n)

\hat{y}^n : 1 for class 1, 0 for class 2

$$L(f) = \sum_n C(f(x^n), \hat{y}^n)$$

Training data: (x^n, \hat{y}^n)

\hat{y}^n : a real number

$$L(f) = \frac{1}{2} \sum_n (f(x^n) - \hat{y}^n)^2$$

Cross entropy:

$$C(f(x^n), \hat{y}^n) = -[\hat{y}^n \ln f(x^n) + (1 - \hat{y}^n) \ln(1 - f(x^n))]$$

Multi-class classification

