



Auto-encoder

- CNN Image Classifier – Convolution section + MLP classifier
- CNN Autoencoder – Convolution section + Deconvolution section to recover the input image

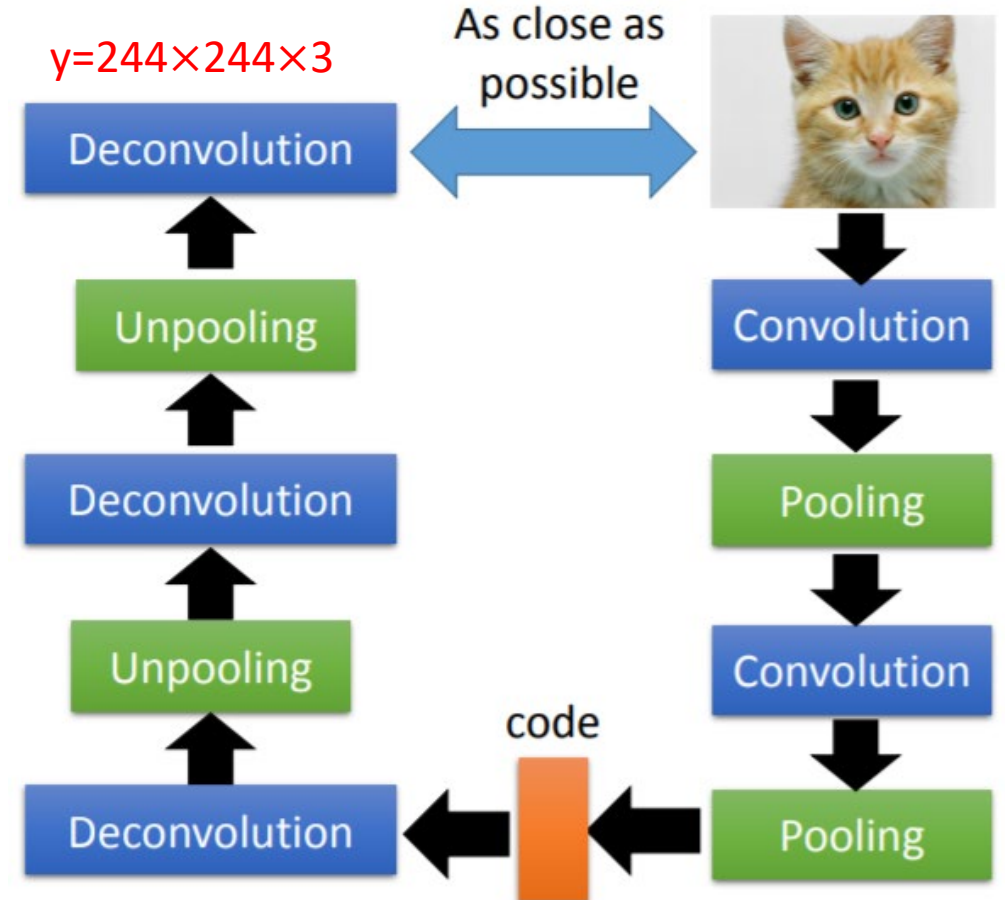
HW5 $y = f(x)$




$x=244 \times 244 \times 3$



HW6 $x = f(x)$

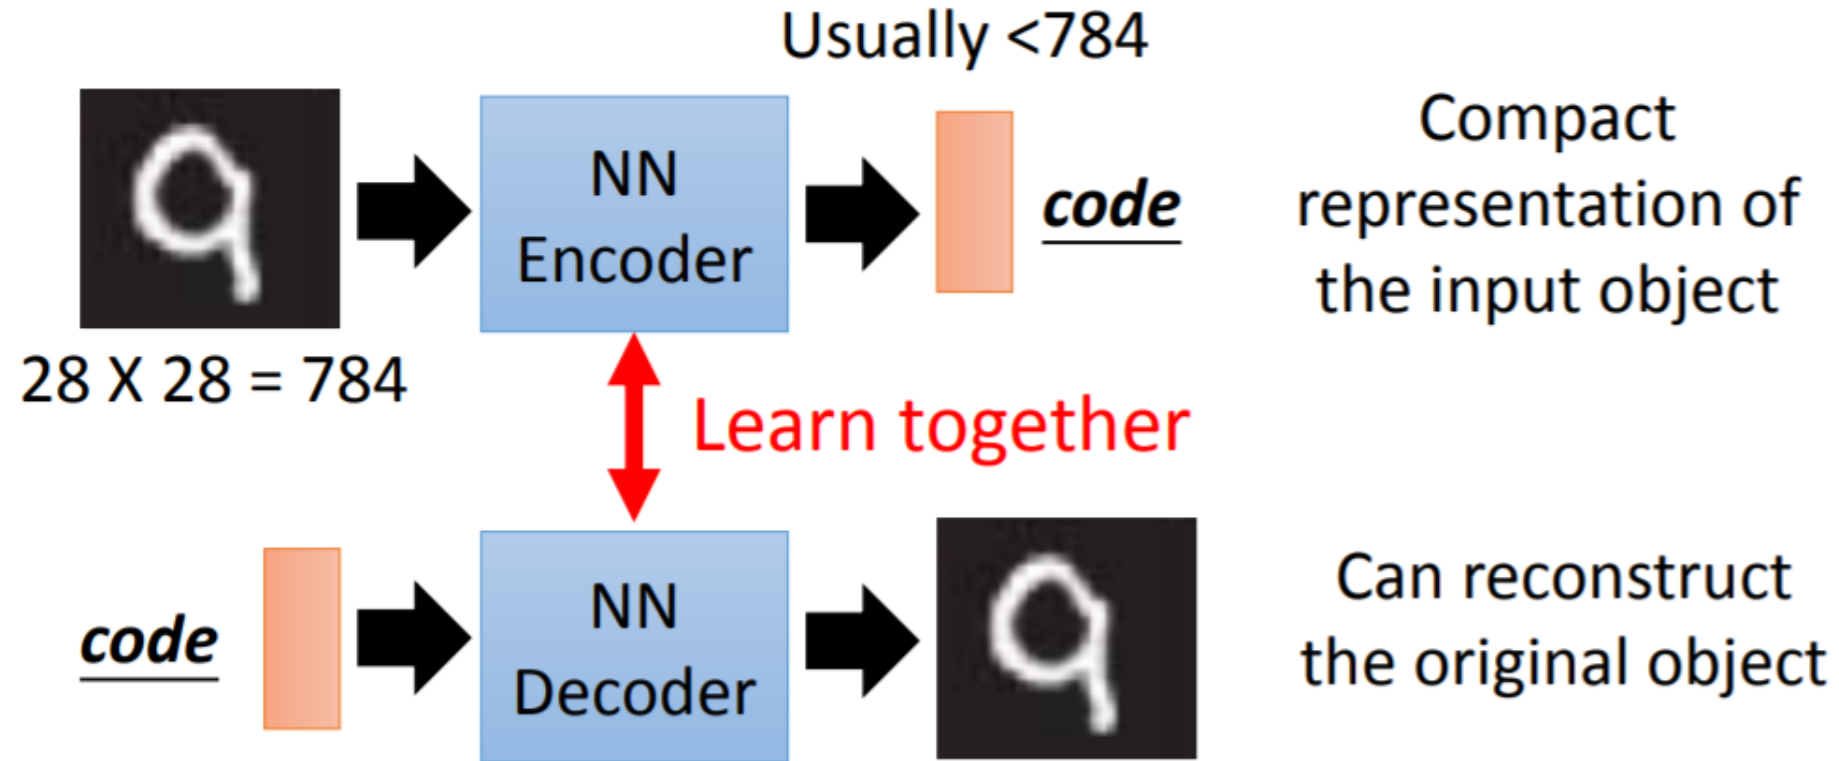


MLP based autoencoder



Reference: Hinton, Geoffrey E., and Ruslan R. Salakhutdinov. "Reducing the dimensionality of data with neural networks." *Science* 313.5786 (2006): 504-507

Autoencoder learns a compact representation of the input image




Practice

- Run "7.1.Conv_AE.ipynb"

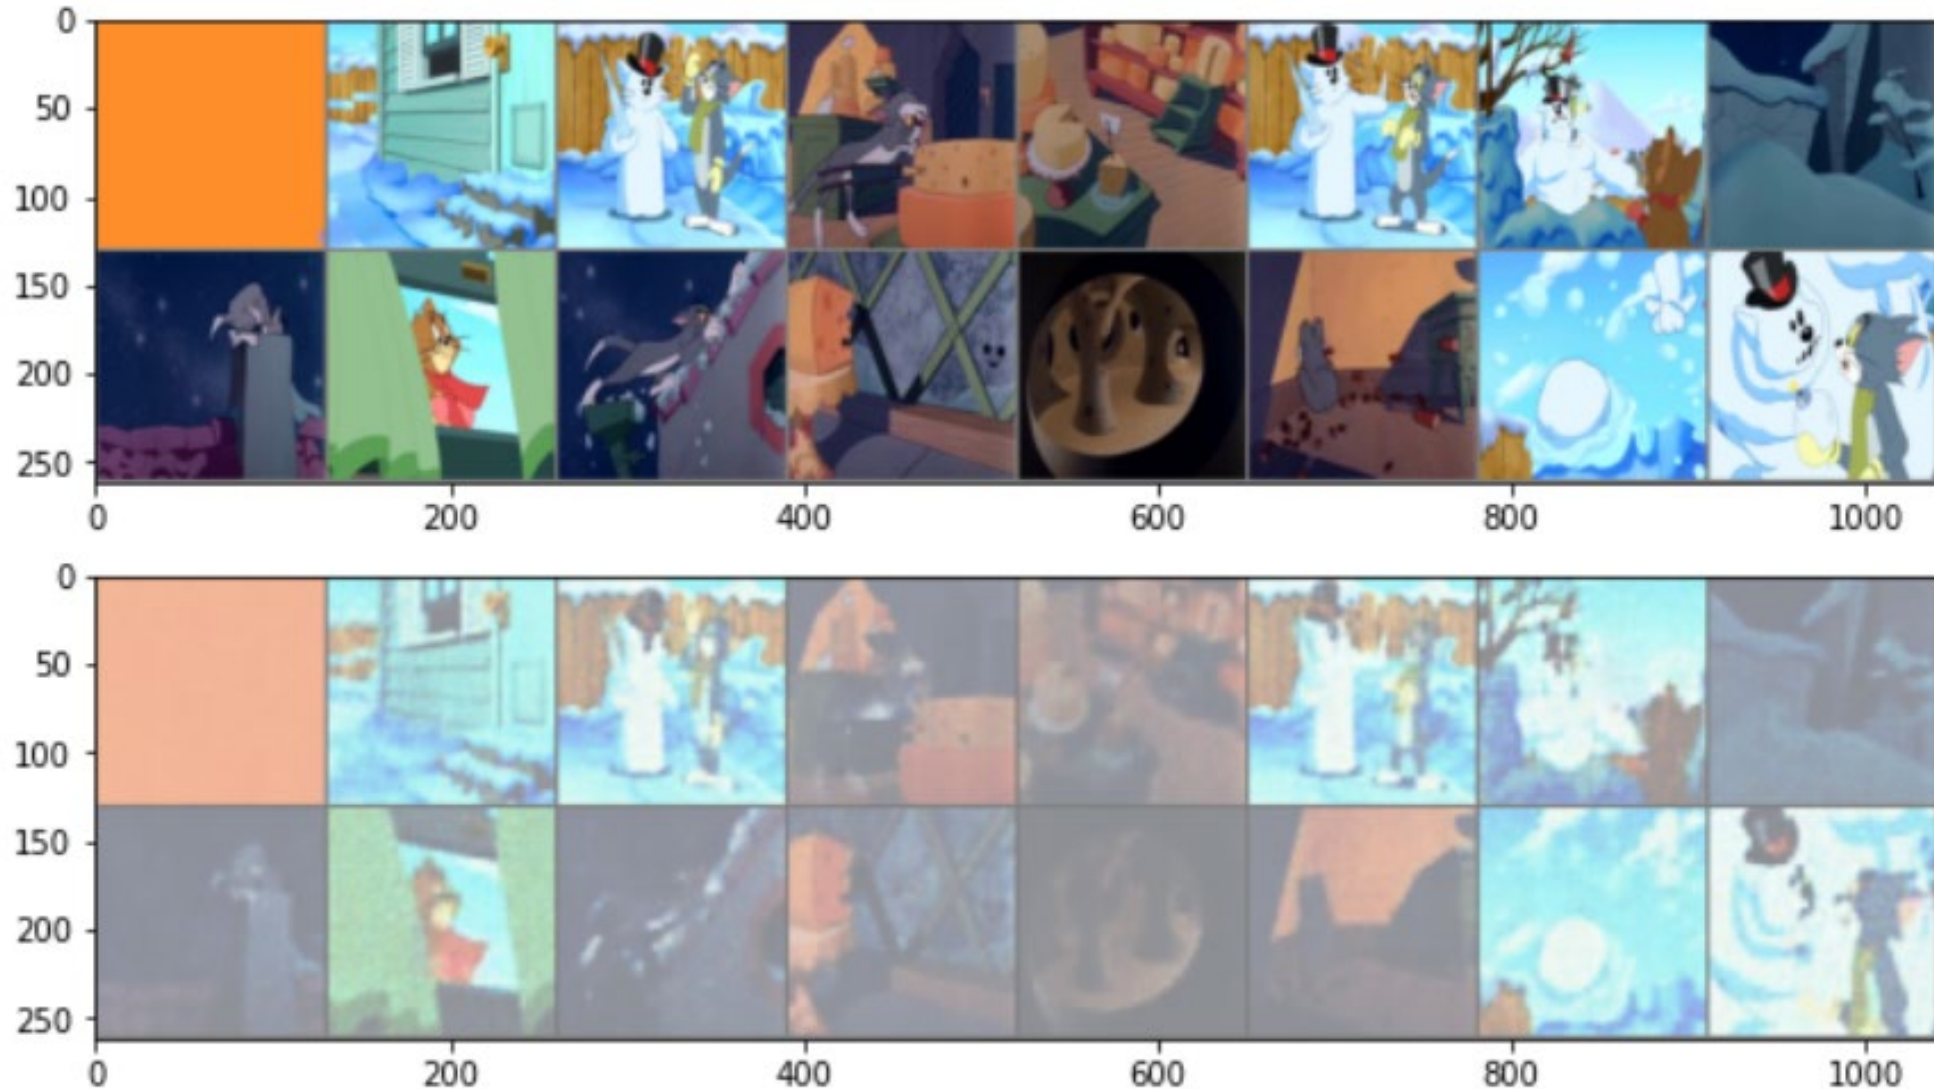


Train 200 epochs



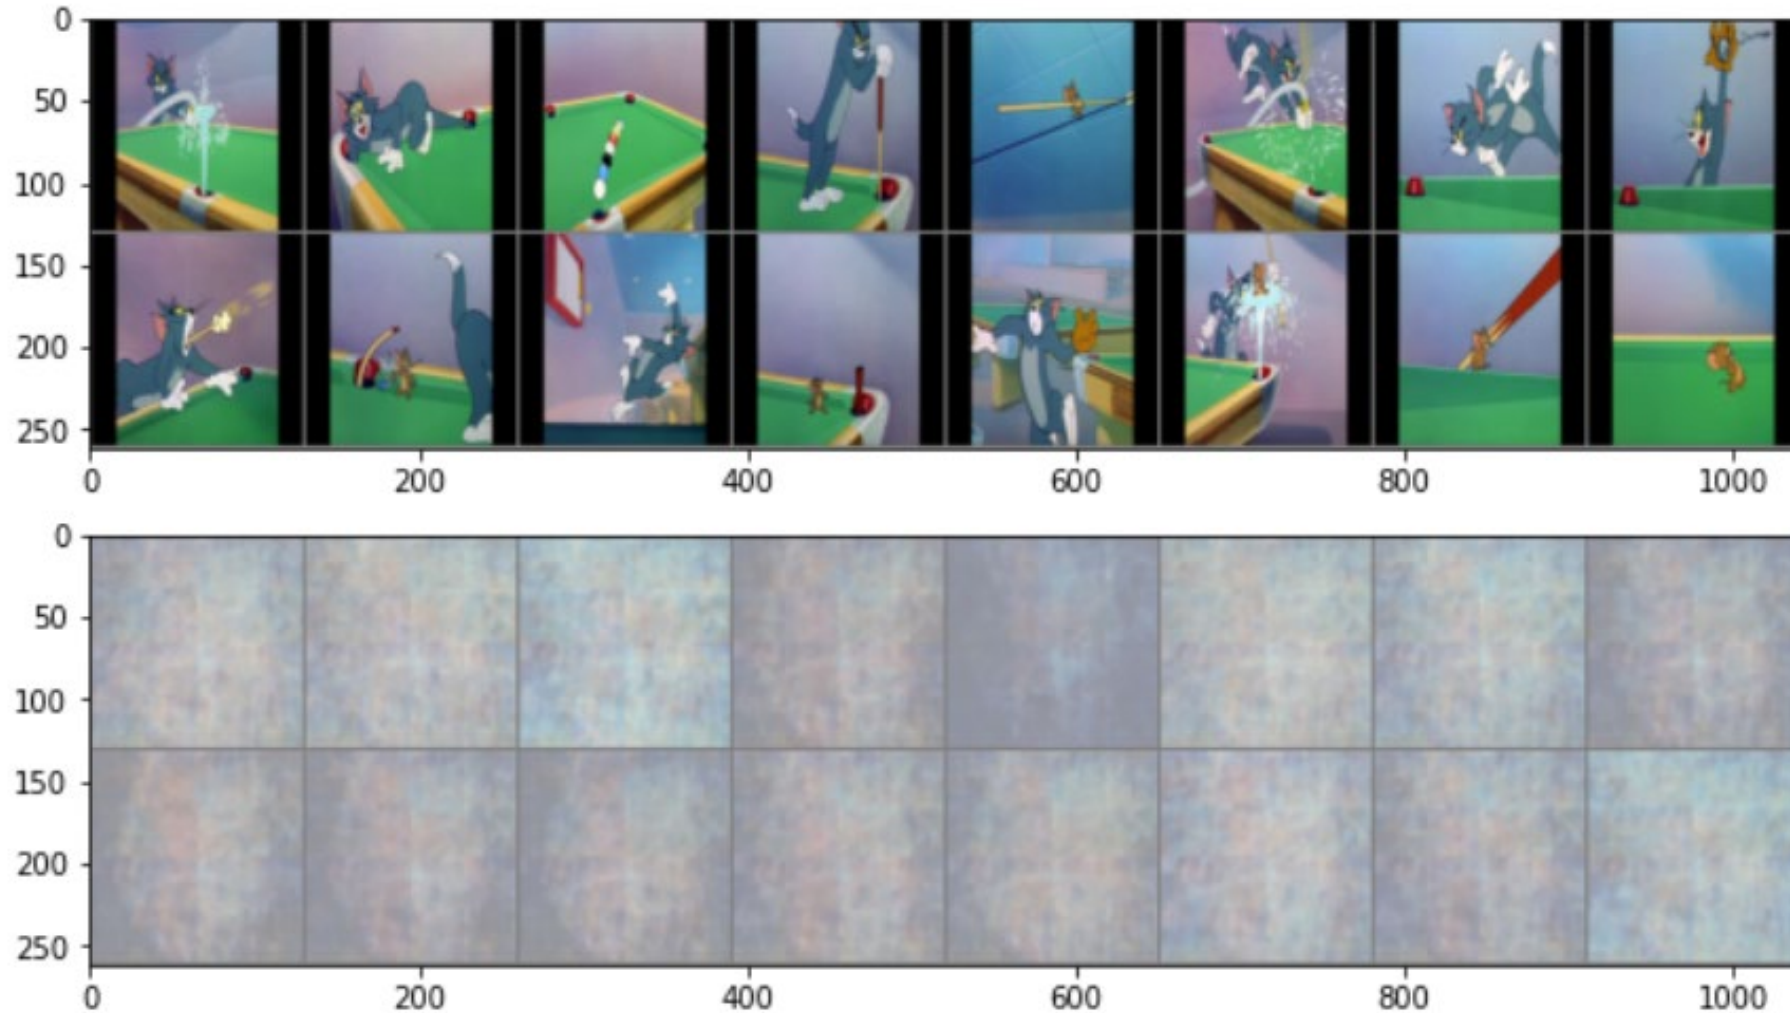
Train 200 epochs

Test on training images – the NN is able to recover more from the input images




Train 200 epochs

Test on un-seen images – fails to reconstruct the input images

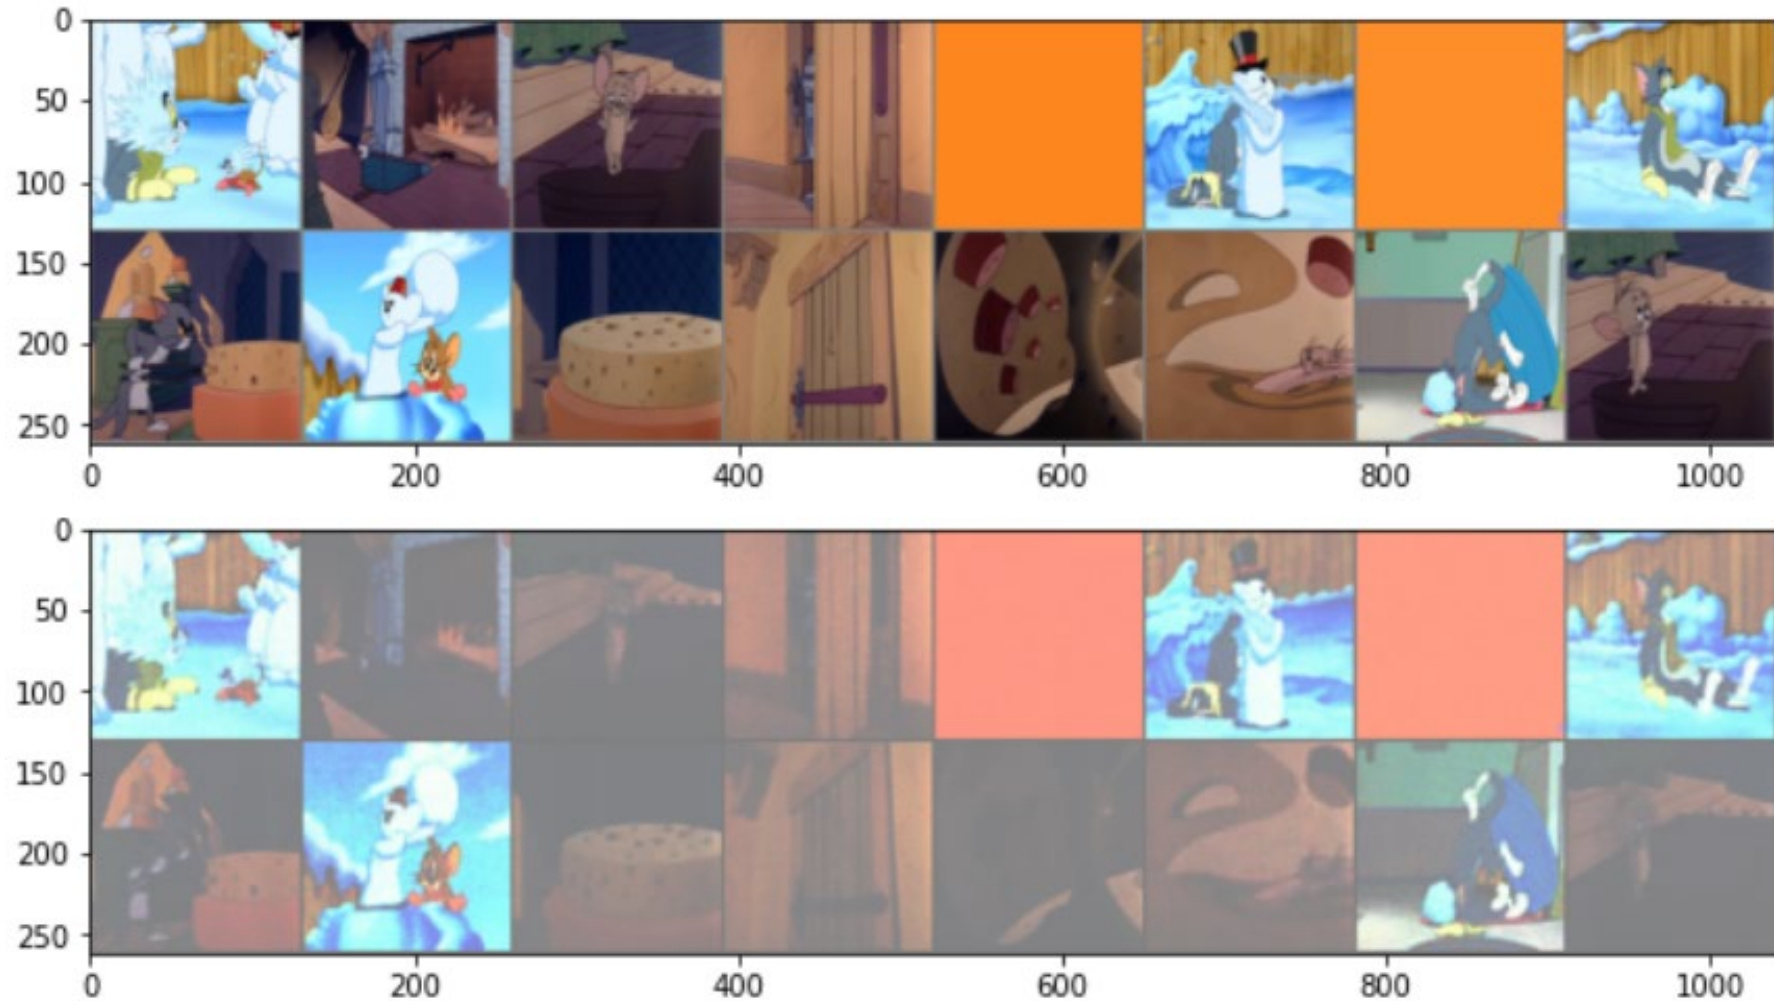


Train 800 epochs



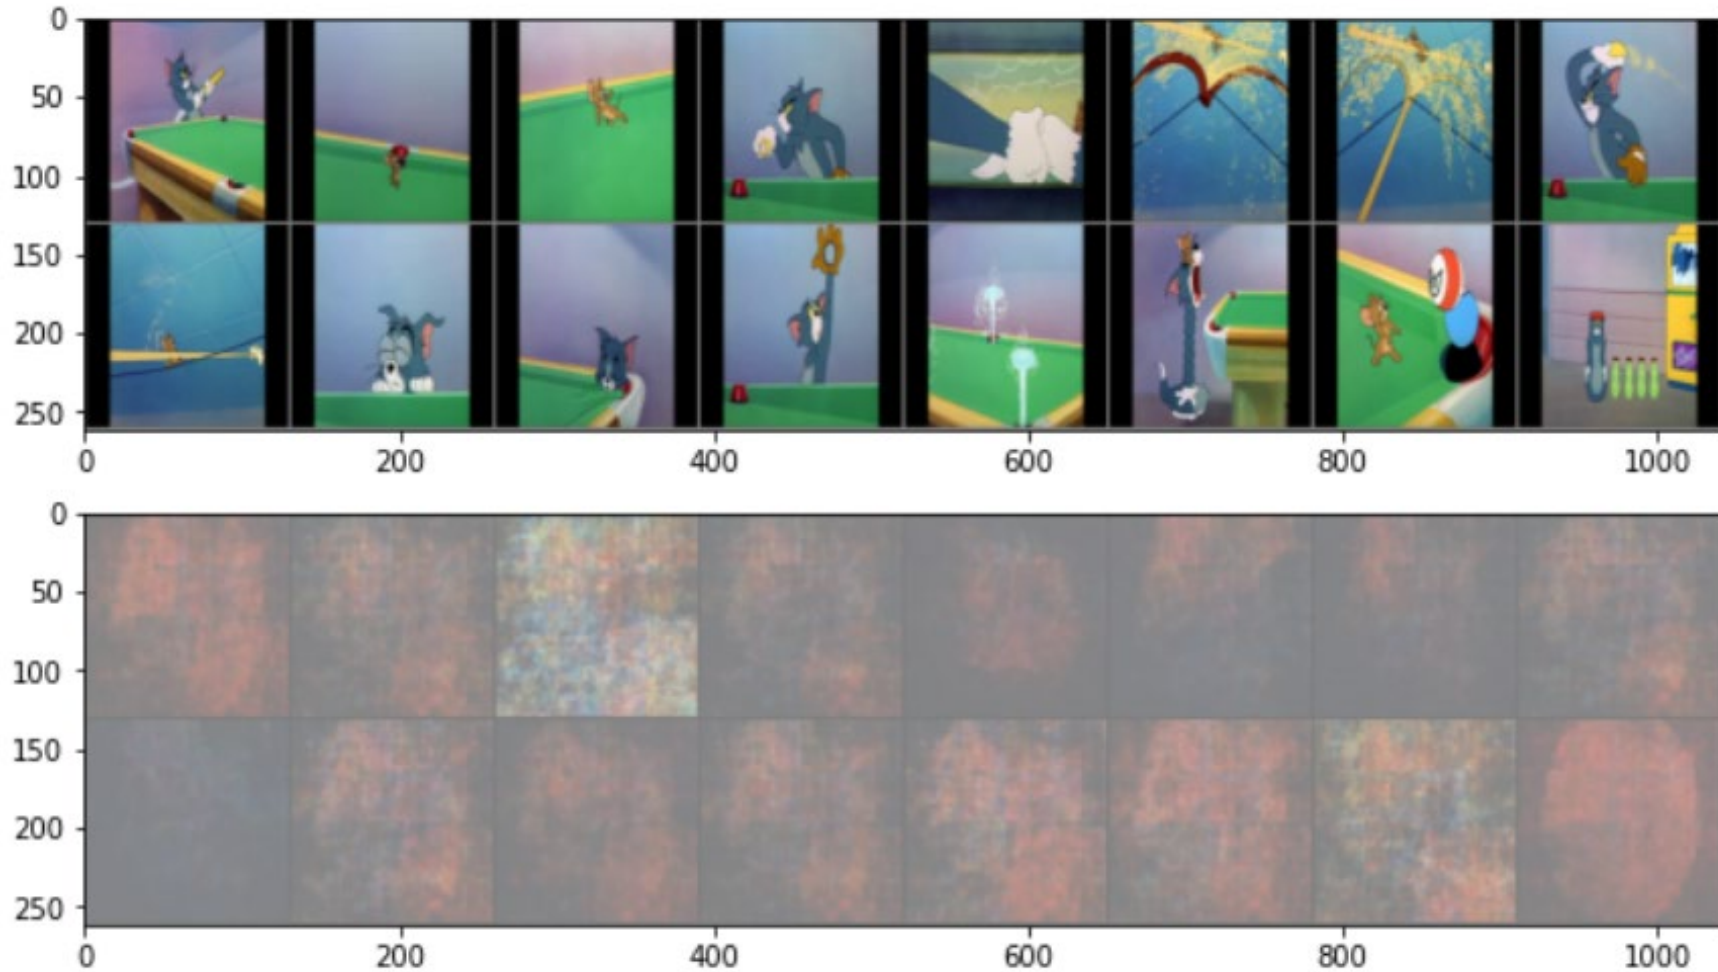
Train 800 epochs

Test on training images

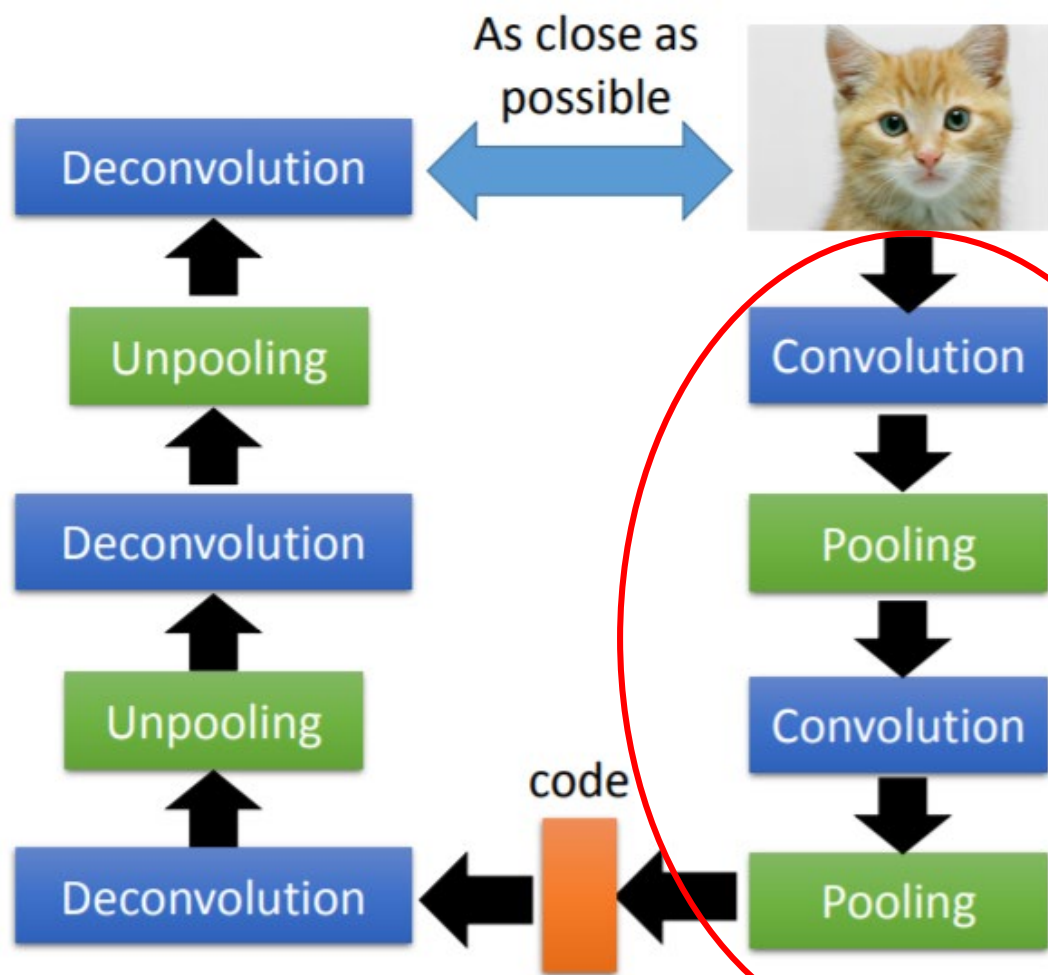


Train 800 epochs

Test on un-seen images – fails to reconstruct the input images



Encoder



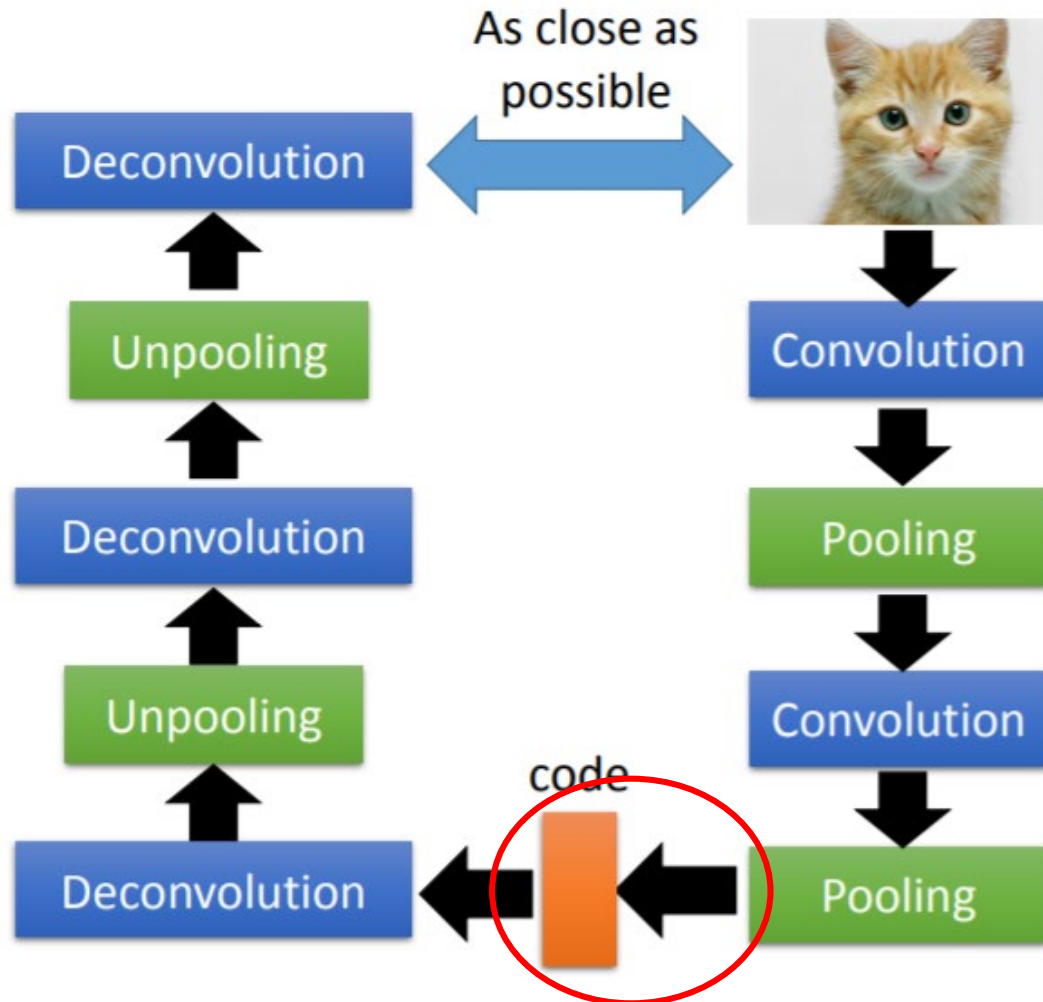
```
self.encoder = nn.Sequential(  
    nn.Conv2d(3, 32, kernel_size=2, stride=2),  
    nn.BatchNorm2d(32, eps=1e-05, momentum=0.1, af  
    nn.ReLU(),  
    nn.Conv2d(32, 64, kernel_size=2, stride=2),  
    nn.BatchNorm2d(64, eps=1e-05, momentum=0.1, af  
    nn.ReLU(),  
    nn.Conv2d(64, 128, kernel_size=2, stride=2),  
    nn.BatchNorm2d(128, eps=1e-05, momentum=0.1, a  
    nn.ReLU(),  
    nn.Conv2d(128, 256, kernel_size=2, stride=2),  
    nn.BatchNorm2d(256, eps=1e-05, momentum=0.1, a  
    nn.ReLU(),  
    nn.Conv2d(256, 512, kernel_size=2, stride=2),  
    nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, a  
    nn.ReLU(),  
    nn.Conv2d(512, 1024, kernel_size=2, stride=2),  
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1,  
    nn.ReLU(),  
    nn.Conv2d(1024, 1024, kernel_size=2, stride=2)  
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1,  
    nn.ReLU(),  
    Flatten(),  
    nn.Linear(in_features=i, out_features=o),  
)
```

Practice: Draw the feature maps of encoder

- Let input image = $224 \times 224 \times 3$
- Draw the feature maps (H, W, depth) after each convolution and max pooling
- What is the number of nodes after flatten?



Latent vector



```
class autoencoder(nn.Module):
    def __init__(self, i=1024, o=64):
        super(autoencoder, self).__init__()
        self.encoder = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=2, stride=
            nn.BatchNorm2d(32, eps=1e-05, momentum=
            nn.ReLU(),
```

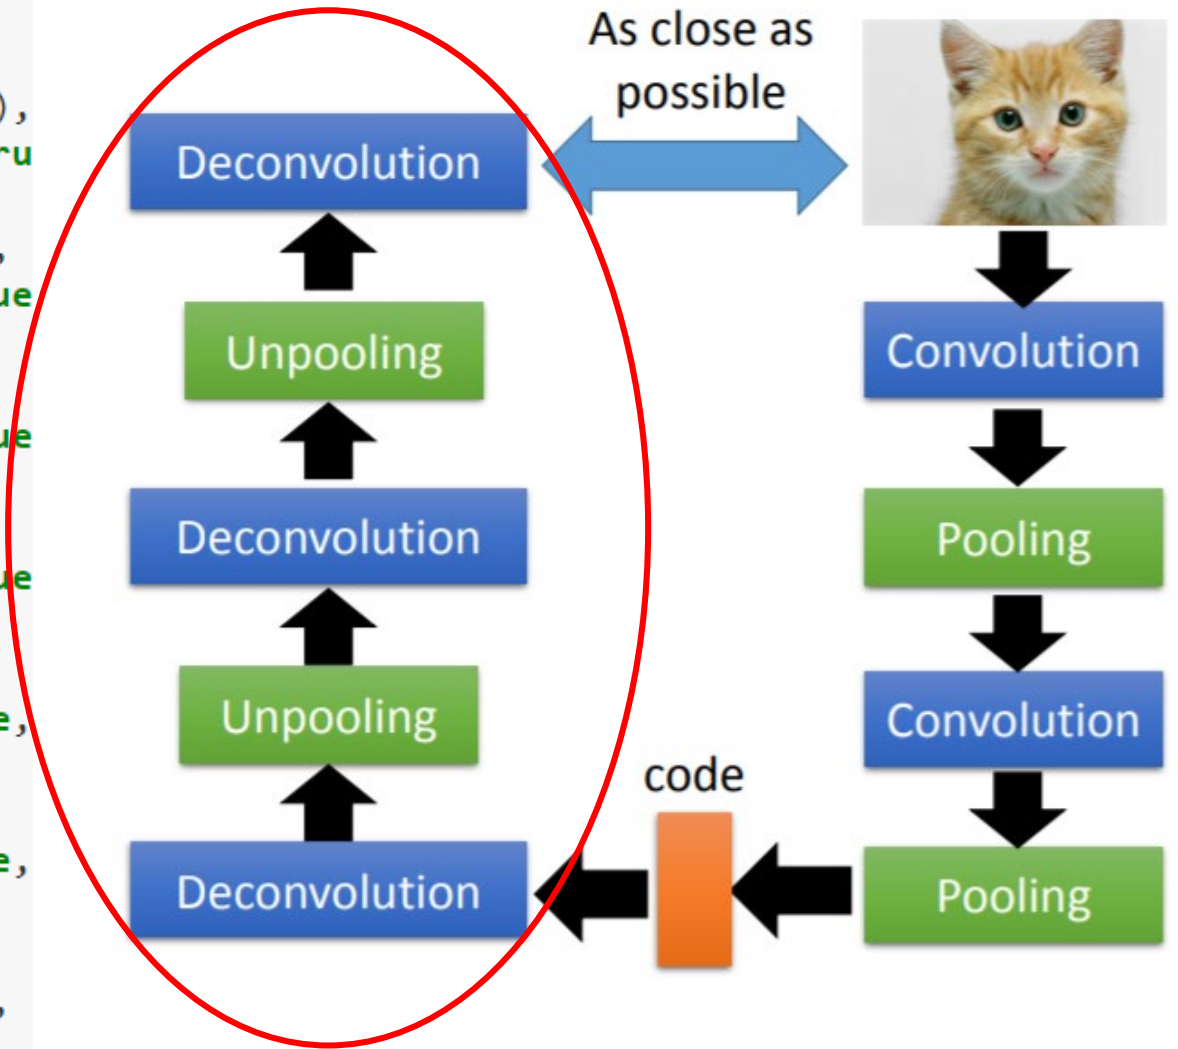
```
nn.BatchNorm2d(1024, eps=1e-05, momentum=0.
nn.ReLU(),
nn.Conv2d(1024, 1024, kernel_size=2, stride
nn.BatchNorm2d(1024, eps=1e-05, momentum=0.
nn.ReLU(),
Flatten(),
nn.Linear(in_features=i, out_features=o),
)
```

Flatten-22
Linear-23
Linear-24
UnFlatten-25

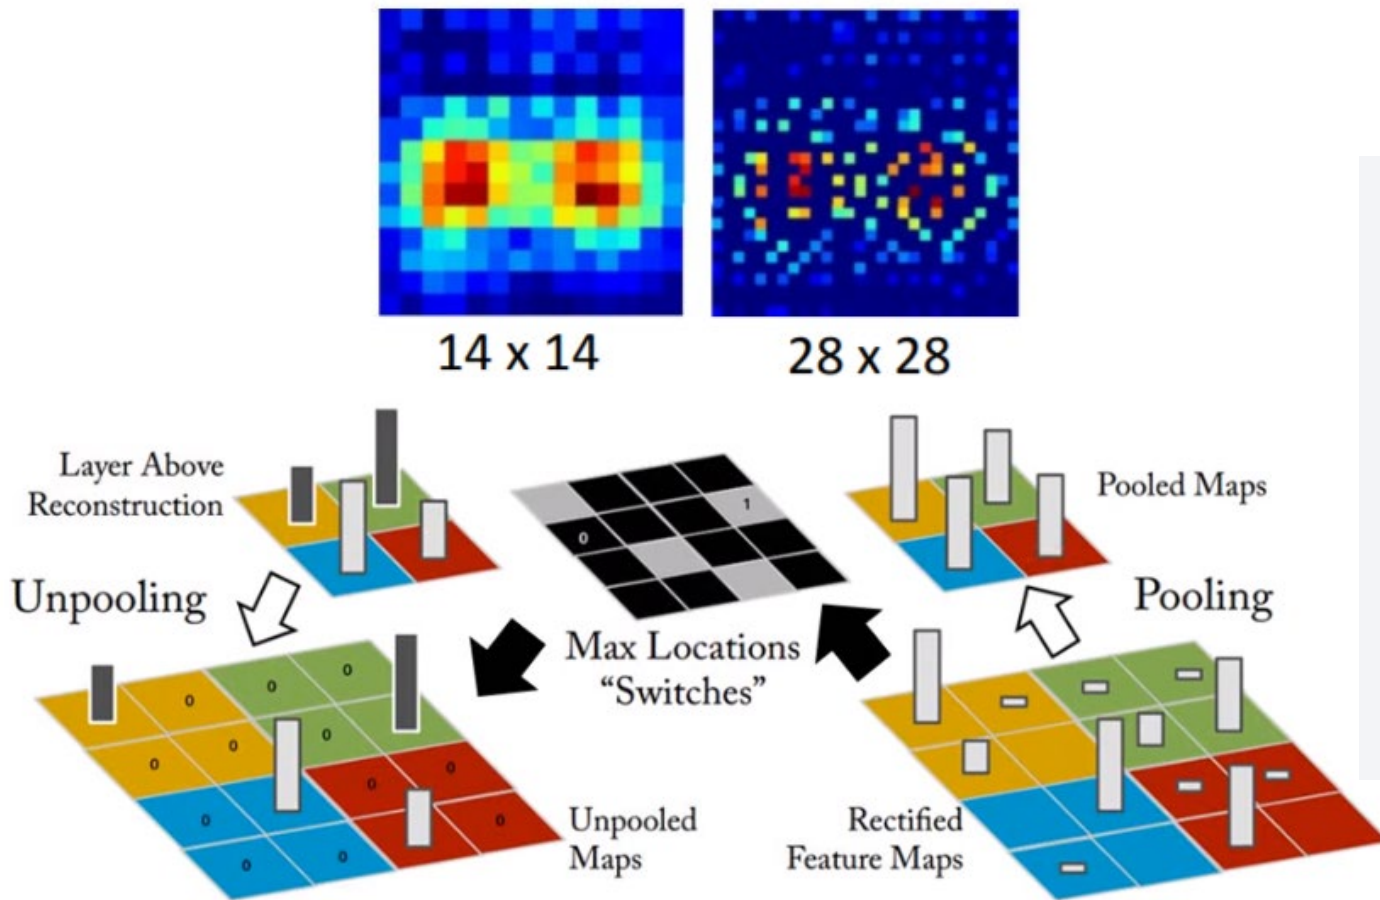
```
[-1, 1024]
[-1, 64]
[-1, 1024]
[-1, 1024, 1, 1]
```

Decoder

```
self.decoder = nn.Sequential(  
    nn.Linear(in_features=o, out_features=i),  
    UnFlatten(),  
    nn.ConvTranspose2d(1024, 1024, kernel_size=2, stride=2),  
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(1024, 512, kernel_size=2, stride=2),  
    nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(512, 256, kernel_size=2, stride=2),  
    nn.BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(256, 128, kernel_size=2, stride=2),  
    nn.BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(128, 64, kernel_size=2, stride=2),  
    nn.BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(64, 32, kernel_size=2, stride=2),  
    nn.BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(32, 3, kernel_size=2, stride=2),  
    nn.BatchNorm2d(3, eps=1e-05, momentum=0.1, affine=True),  
    nn.Sigmoid(),  
)
```



Unpooling




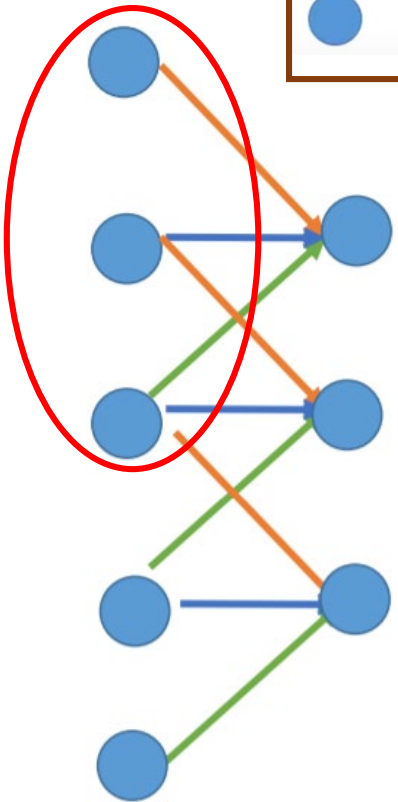


```
>>> pool = nn.MaxPool2d(2, stride=2, return_indices=True)
>>> unpool = nn.MaxUnpool2d(2, stride=2)
>>> input = torch.tensor([[[[ 1.,  2,  3,  4],
                             [ 5,  6,  7,  8],
                             [ 9, 10, 11, 12],
                             [13, 14, 15, 16]]]])

>>> output, indices = pool(input)
>>> unpool(output, indices)
tensor([[[[ 0.,  0.,  0.,  0.],
           [ 0.,  6.,  0.,  8.],
           [ 0.,  0.,  0.,  0.],
           [ 0., 14.,  0., 16.]]]]])
```



Reference: 李弘毅 ML Lecture 16 <https://youtu.be/Tk5B4seA-AU>

Deconvolution

1D convolution, $k=3$ 1D deconvolution, $k=3$



1D convolution, $k=3$



In this ConvAE example, we only use deconvolution for up sampling, no un-pooling is used.

```
self.decoder = nn.Sequential(  
    nn.Linear(in_features=o, out_features=i),  
    UnFlatten(),  
    nn.ConvTranspose2d(1024, 1024, kernel_size=2, stride=2),  
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(1024, 512, kernel_size=2, stride=2),  
    nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(512, 256, kernel_size=2, stride=2),  
    nn.BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(256, 128, kernel_size=2, stride=2),  
    nn.BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(128, 64, kernel_size=2, stride=2),  
    nn.BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(64, 32, kernel_size=2, stride=2),  
    nn.BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True),  
    nn.ReLU(),  
    nn.ConvTranspose2d(32, 3, kernel_size=2, stride=2),  
    nn.BatchNorm2d(3, eps=1e-05, momentum=0.1, affine=True),  
    nn.Sigmoid(),  
)
```


Practice: Draw the feature maps of decoder

- Input – the number of nodes after un-flatten
- Draw feature maps (H, W, depth) after each de-convolution and un-max pooling

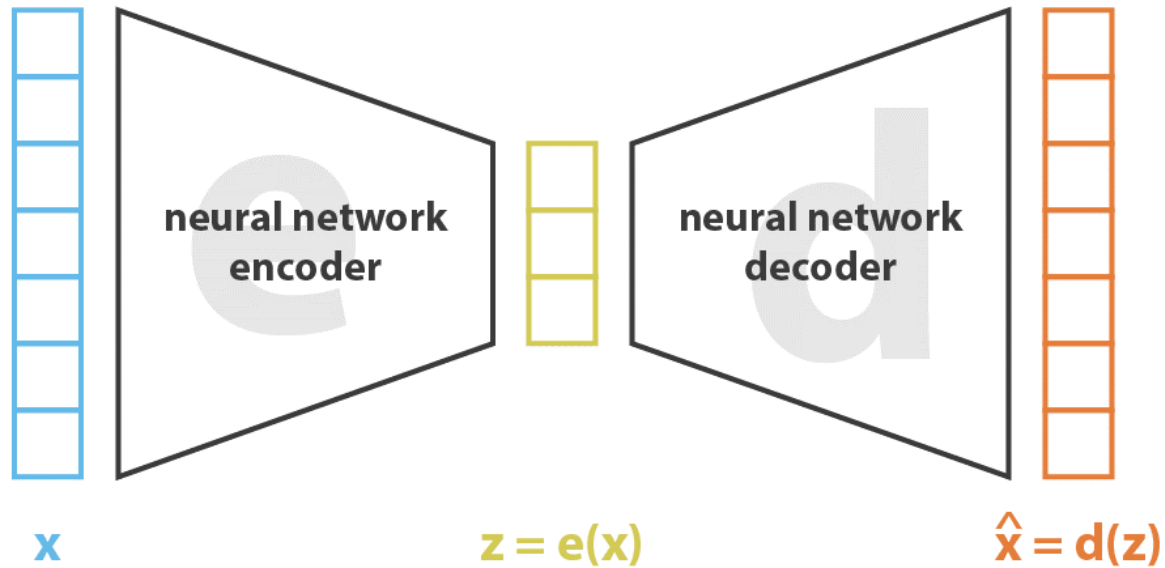


Deconvolution

```
(2): ConvTranspose2d(1024, 1024, kernel_size=(2, 2), stride=(2, 2))
(3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_r
(4): ReLU()
(5): ConvTranspose2d(1024, 512, kernel_size=(2, 2), stride=(2, 2))
(6): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_ru
(7): ReLU()
(8): ConvTranspose2d(512, 256, kernel_size=(2, 2), stride=(2, 2))
(9): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_ru
(10): ReLU()
```

ConvTranspose2d-26	[-1, 1024, 2, 2]
BatchNorm2d-27	[-1, 1024, 2, 2]
ReLU-28	[-1, 1024, 2, 2]
ConvTranspose2d-29	[-1, 512, 4, 4]
BatchNorm2d-30	[-1, 512, 4, 4]
ReLU-31	[-1, 512, 4, 4]
ConvTranspose2d-32	[-1, 256, 8, 8]
BatchNorm2d-33	[-1, 256, 8, 8]
ReLU-34	[-1, 256, 8, 8]
ConvTranspose2d-35	[-1, 128, 16, 16]
BatchNorm2d-36	[-1, 128, 16, 16]
ReLU-37	[-1, 128, 16, 16]
ConvTranspose2d-38	[-1, 64, 32, 32]
BatchNorm2d-39	[-1, 64, 32, 32]
ReLU-40	[-1, 64, 32, 32]

Loss function



$$\text{loss} = \|x - \hat{x}\|^2 = \|x - d(z)\|^2 = \|x - d(e(x))\|^2$$

Source: <https://towardsdatascience.com/understanding-variational-autoencoders-vaes-f70510919f73>

```
[13]: for batchX, _ in loader:
      break;
      print(batchX.shape)

      torch.Size([16, 3, 128, 128])

[14]: tensorY=model(batchX.to(device))
      print(tensorY.shape)

      torch.Size([16, 3, 128, 128])

[15]: loss = loss_func(tensorY, batchX.to(device))
      print(loss)

      tensor(0.6961, device='cuda:0', grad_fn=<Msel
```

Save and load PyTorch model

```
[27]: torch.save(model.state_dict(), "AE800.pt")
```

```
[28]: model=autoencoder() #build NN architecture  
      model.load_state_dict(torch.load("AE800.pt")) #Load model weights  
      model.to(device)  
      model.eval()
```



Save and load PyTorch model



The screenshot displays a Jupyter Notebook titled "3_AlexNet_(1).ipynb". The interface includes a top navigation bar with options like "檔案", "編輯", "檢視畫面", "插入", "執行階段", "工具", "說明", and "無法儲存變更". On the left, a file explorer sidebar shows a directory structure with folders "gdrive" and "sample_data", and files "AE800.pt" and "tSNE.csv". The "AE800.pt" file and its parent folder "sample_data" are circled in red. The main notebook area shows two code cells. Cell [27] contains the command `torch.save(model.state_dict(), "AE800.pt")`, where the filename is circled in red. Cell [28] contains code to load the model, build an autoencoder architecture, and evaluate it. The code is as follows:

```
[27] torch.save(model.state_dict(), "AE800.pt")

[28] model=autoencoder() #build NN architecture
model.load_state_dict(torch.load("AE800.pt")) #load
model.to(device)
model.eval()

autoencoder(
    (encoder): Sequential(
      (0): Conv2d(3, 32, kernel_size=(2, 2), stride=
      (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, a
      (2): ReLU()
```


Get latent vectors of all training images

```
[37]: for step, (batchX, batchY) in enumerate(loader):  
    tensorY = model.encoder(batchX.to(device))  
    if(step==0):  
        arrayX = np.array(tensorY.cpu().detach().numpy())  
        arrayY = batchY.cpu().detach().numpy()  
    else:  
        arrayX = np.concatenate((arrayX, tensorY.cpu().detach().numpy()))  
        arrayY = np.concatenate((arrayY, batchY.cpu().detach().numpy()))  
    print(arrayX.shape, arrayY.shape)
```

(298, 64) (298,)



Use t -SNE to reduce dimensions from 64 to 2

```
[38]: from sklearn.manifold import TSNE  
      tsne = TSNE(perplexity=5, n_components=2, init='pca', n_iter=5000)  
      # try perplexity = 5, 10, 30, 50
```


```
[39]: x=tsne.fit transform(arrayX)  
      print(x.shape)
```

```
(298, 2)
```

```
[40]: plt.figure(figsize=(18,9))  
      plt.scatter(x[:, 0], x[:, 1], c= arrayY)  
      plt.show()
```



Use t -SNE to reduce dimensions from 64 to 2



Save data to csv file

```
[42] print(x.shape, arrayY.shape)
```

```
(298, 2) (298,)
```

```
[43] arrayY1 = arrayY.reshape(arrayY.shape[0], 1)
      print(arrayY1.shape)
```

```
(298, 1)
```

```
[44] XYArray = np.hstack((x, arrayY1))
      print(XYArray.shape)
```

```
(298, 3)
```

```
[45] # Save data to excel for further Tableau visualization
      import pandas as pd
      pd.DataFrame(XYArray).to_csv("tSNE.csv")
```

Save data to csv file

CO 3_AlexNet_(1).ipynb

檔案 編輯 檢視畫面 插入 執行階段 工具 說明 無法儲存變更

檔案

gdrive
sample_data
AF800.pt
tSNE.csv

+ 程式碼 + 文字 複製到雲端硬碟

```
(2, 3) (2, 1) (2, 4)
```

```
[42] print(x.shape, arrayY.shape)
```

```
(298, 2) (298,)
```

```
[43] arrayY1 = arrayY.reshape(arrayY.shape[0], 1)
```

```
print(arrayY1.shape)
```

```
(298, 1)
```

```
[44] XYArray = np.hstack((x, arrayY1))
```

```
print(XYArray.shape)
```

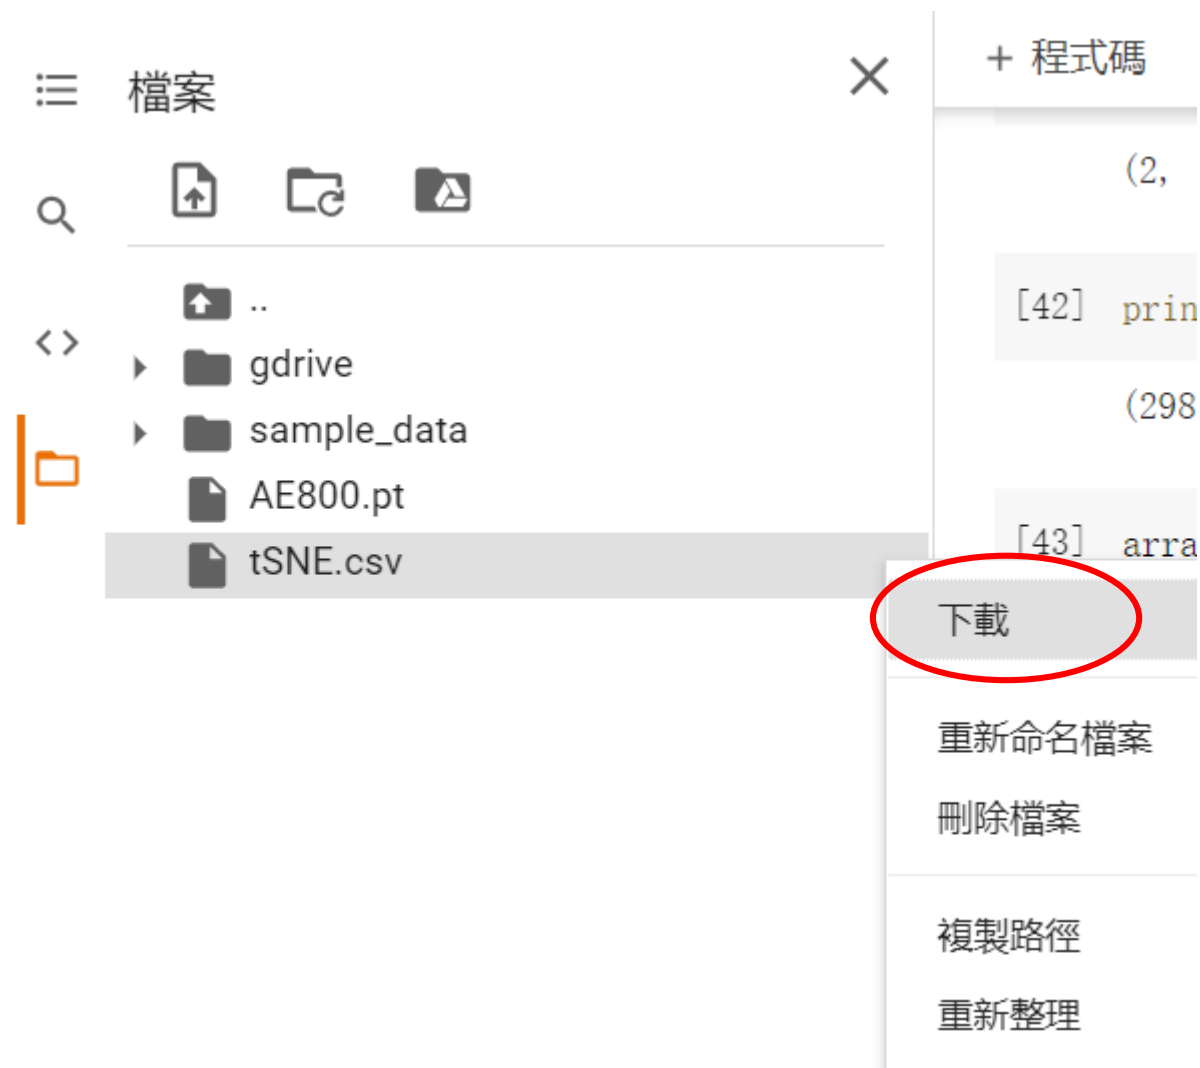
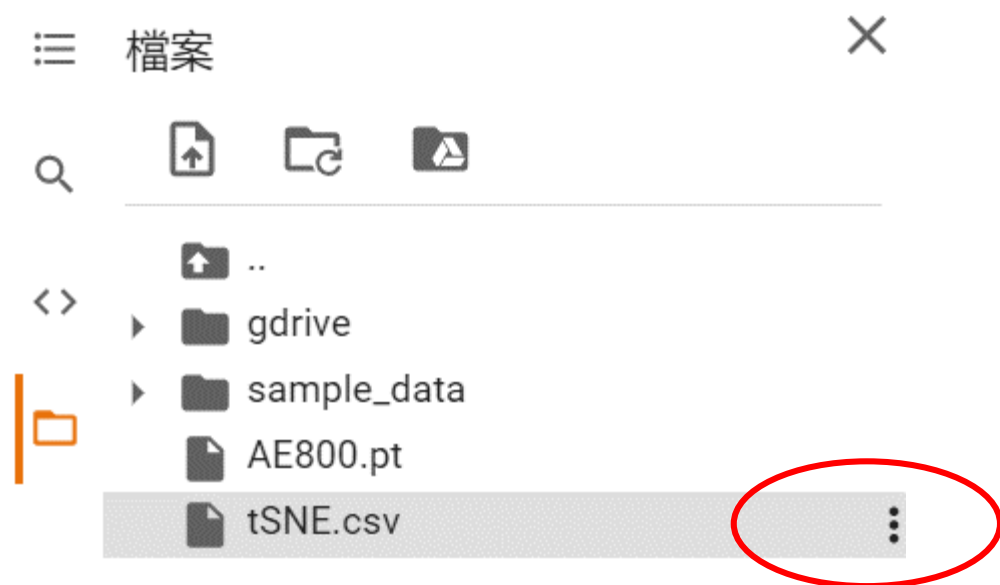
```
(298, 3)
```

```
[45] # Save data to excel for further Tableau visual
```

```
import pandas as pd
```

```
pd.DataFrame(XYArray).to_csv("tSNE.csv")
```


Download csv file



Visualize in Tableau public

ConnectionsAdd

tSNE (2)
Text file

Files

☐ Use Data Interpreter

Data Interpreter might be able to clean your Text file workbook.

HW1 lecture.txt

tSNE (1).csv

tSNE (2).csv

—

New Union

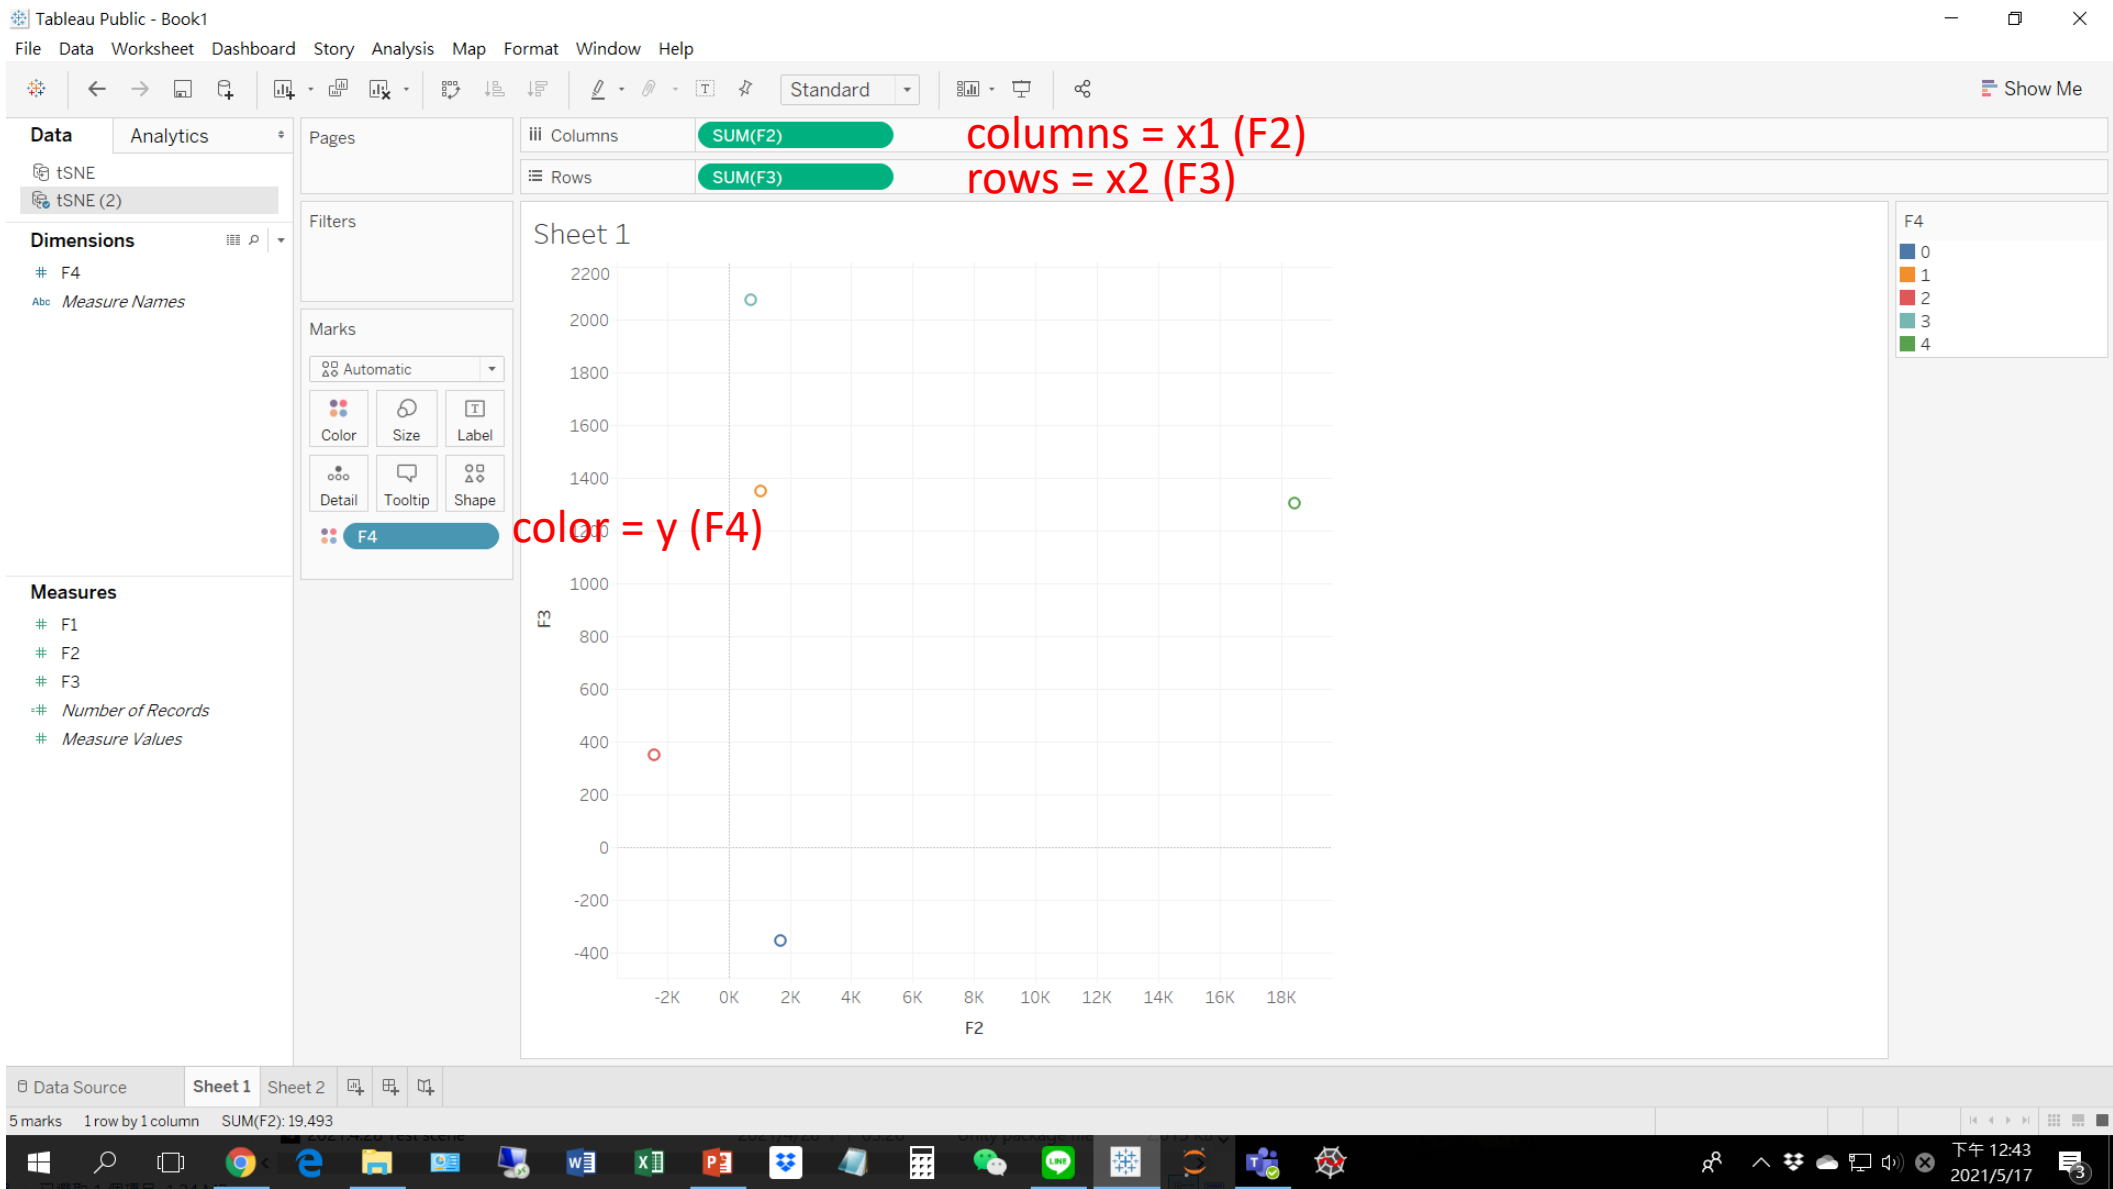
tSNE (2).csv

Sort fields

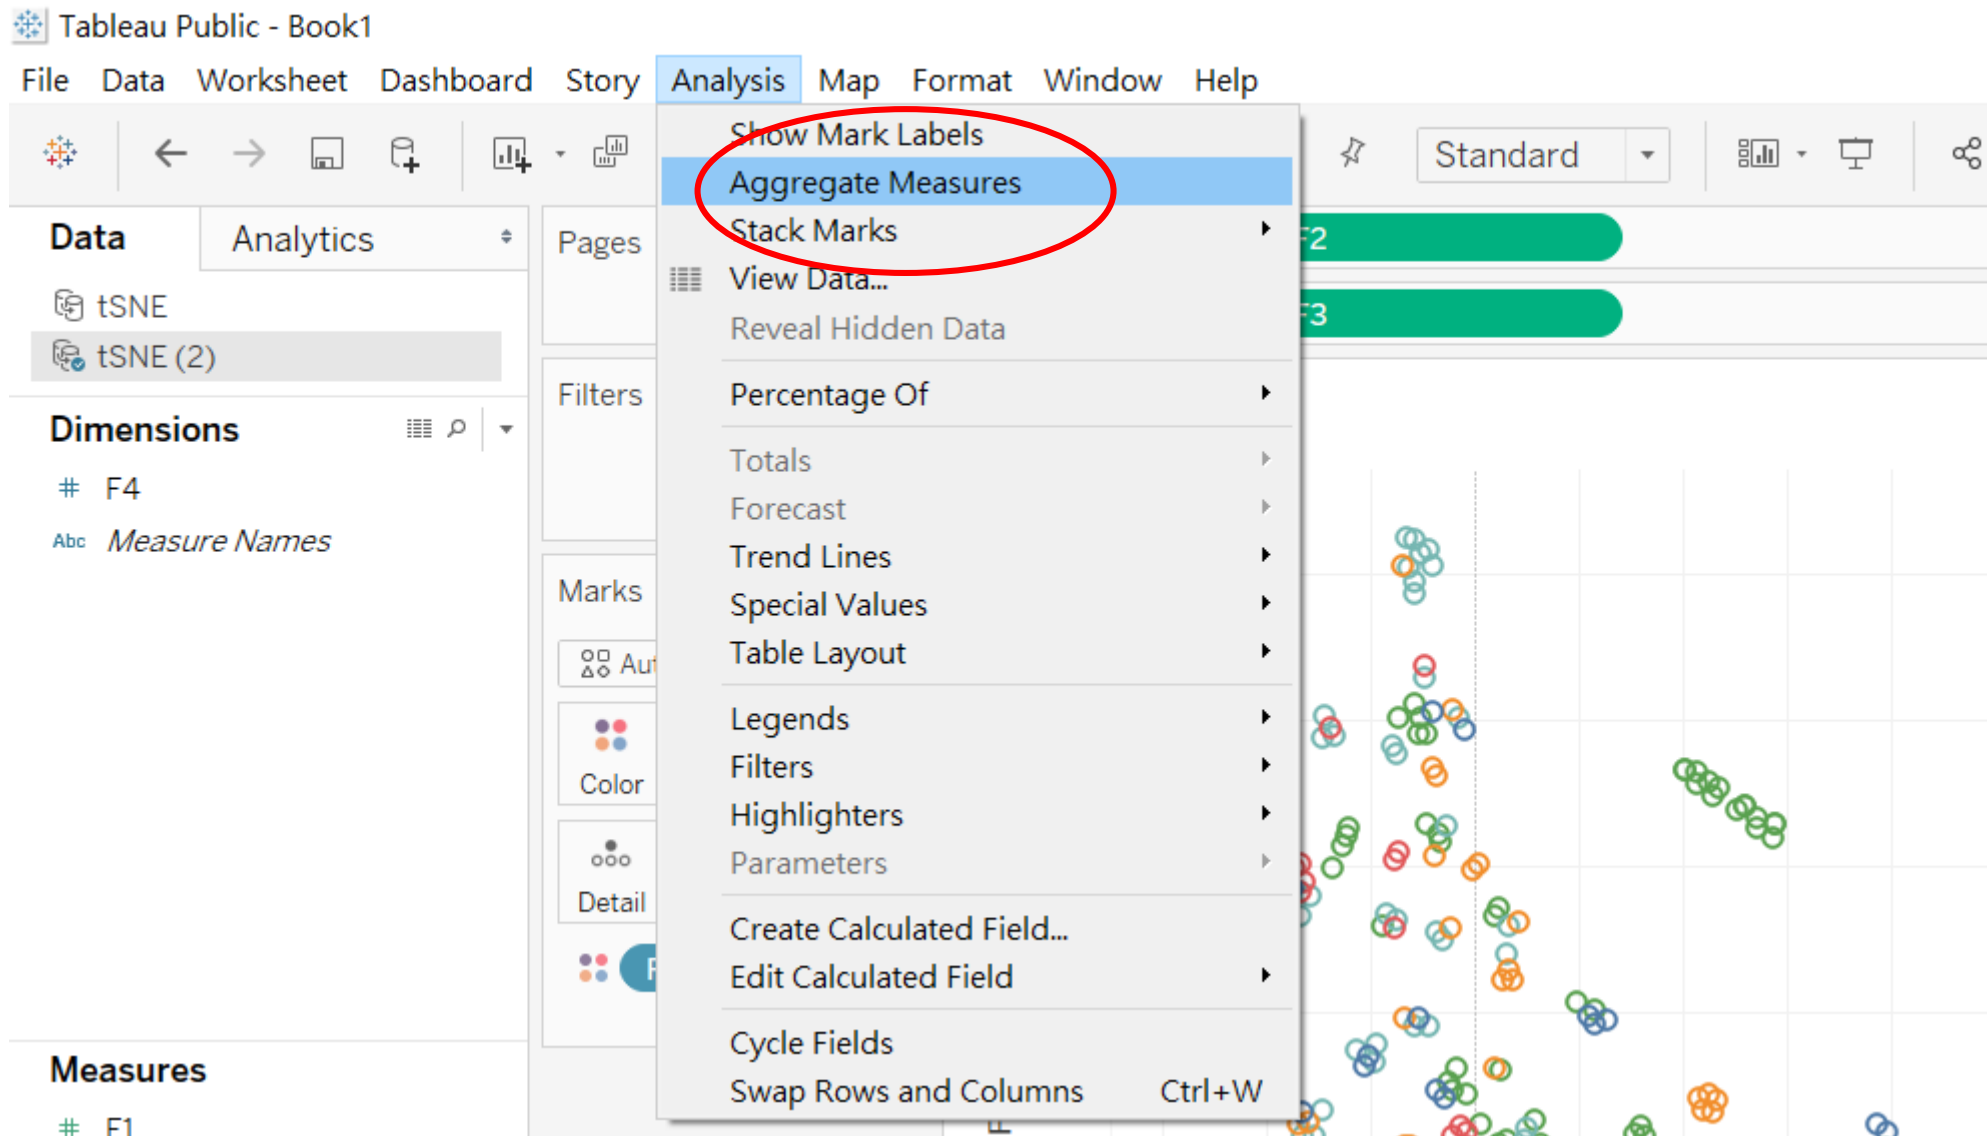
Data source order

#	# X1	# X2	# y
tSNE (...)	tSNE (2...)	tSNE (2)....	tSNE (2)...
F1	F2	F3	F4
2	-6.65	18.604	2.00000
3	-7.06	9.053	2.00000
4	-98.34	-372.219	0.00000
5	-165.39	31.733	0.00000
6	-156.86	179.328	3.00000

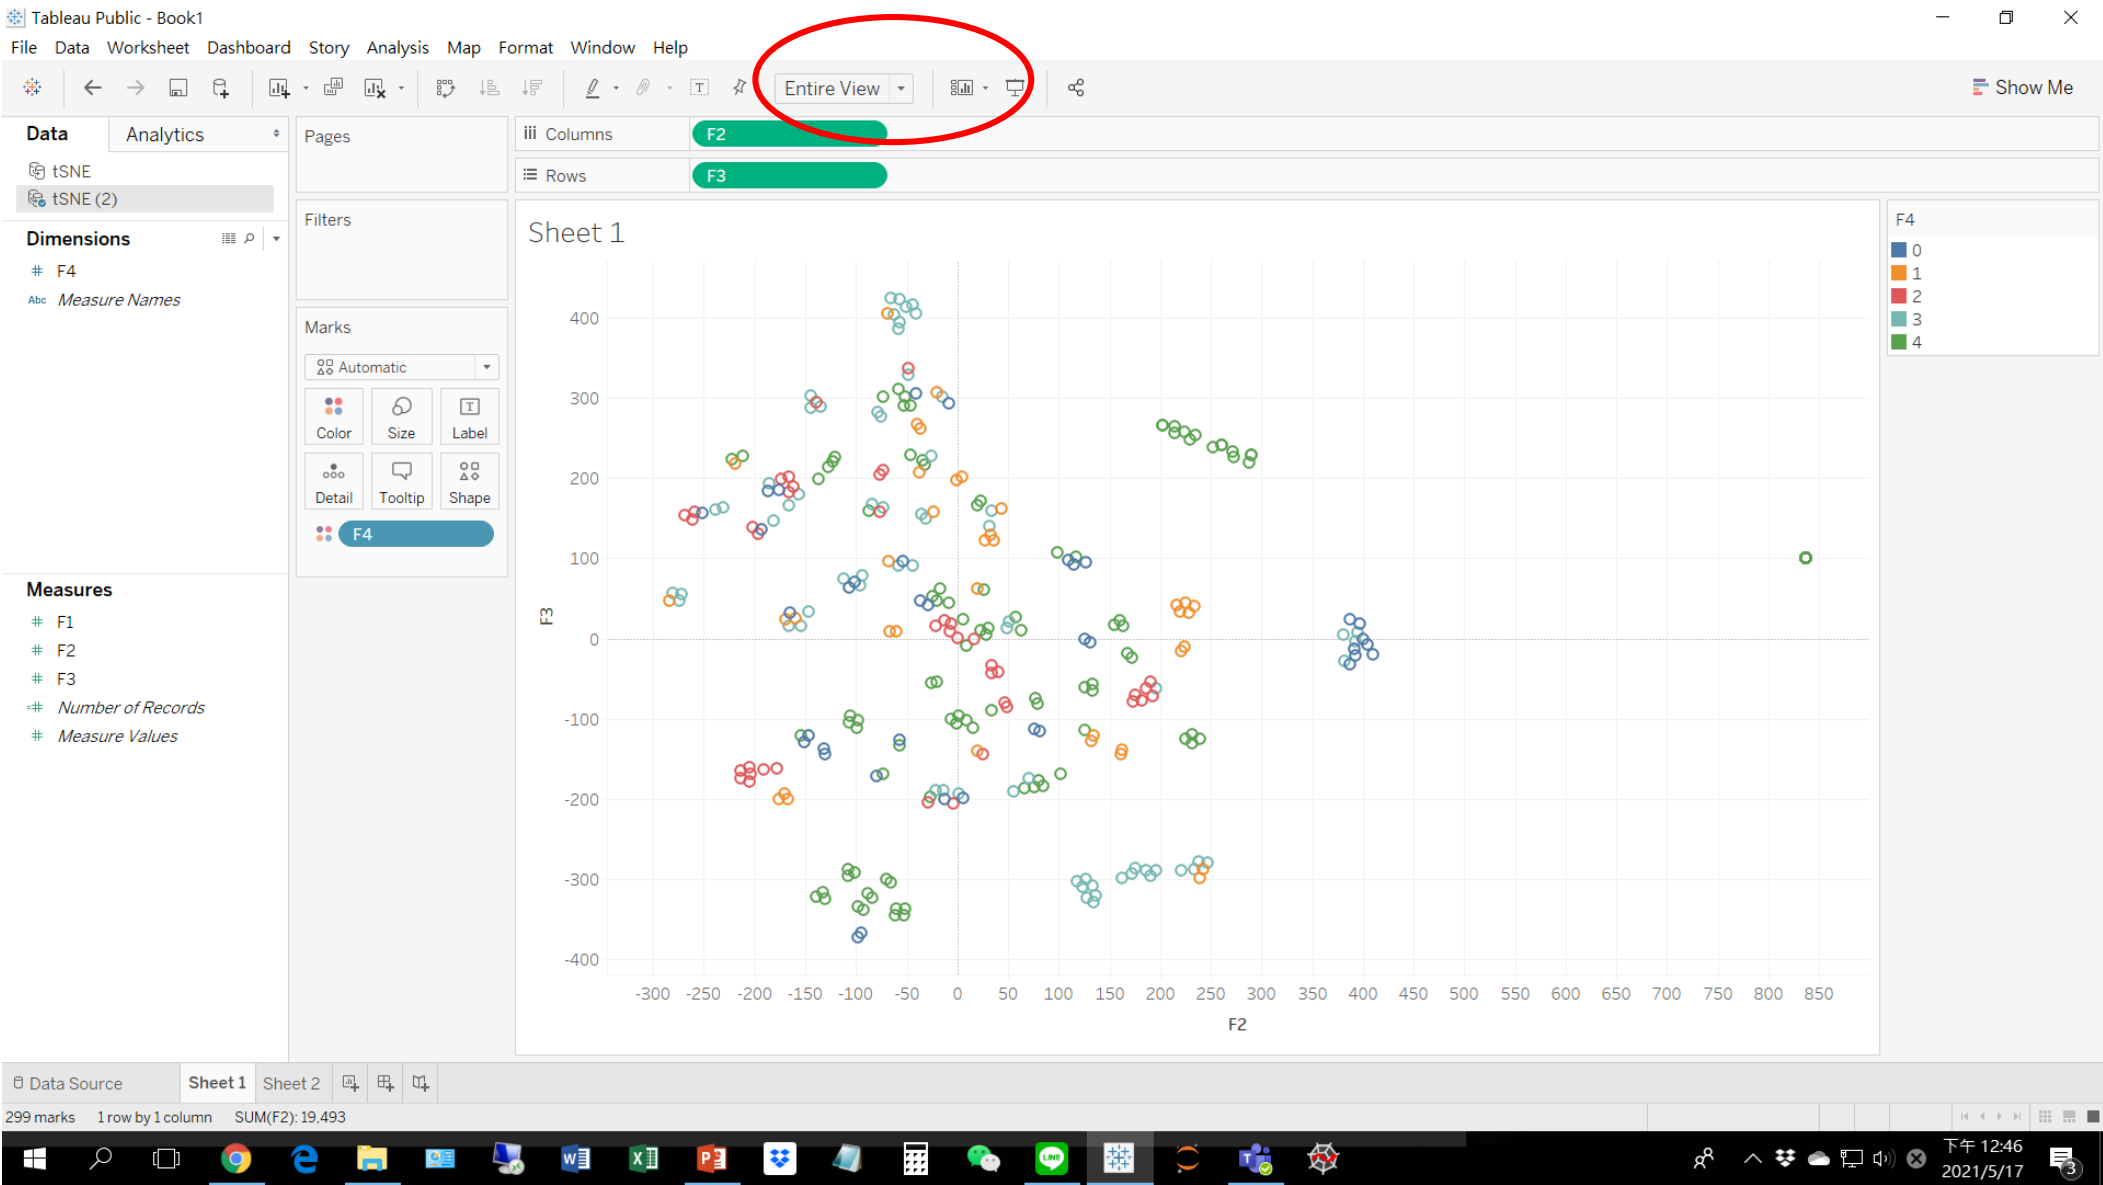
Visualize in Tableau public



Visualize in Tableau public



Visualize in Tableau public




HW6 (1)


- Train an AE to learn a compact representation (try latent vector of size 20, 30, 50) of your facial expression. Test with 10 happy and 10 angry faces.
- Show the recovered image.
- Send the latent vectors to t -SNE to see whether they form clusters.



Happy = ?, Angry = ?, Latent vector size = ?, 1160 epochs

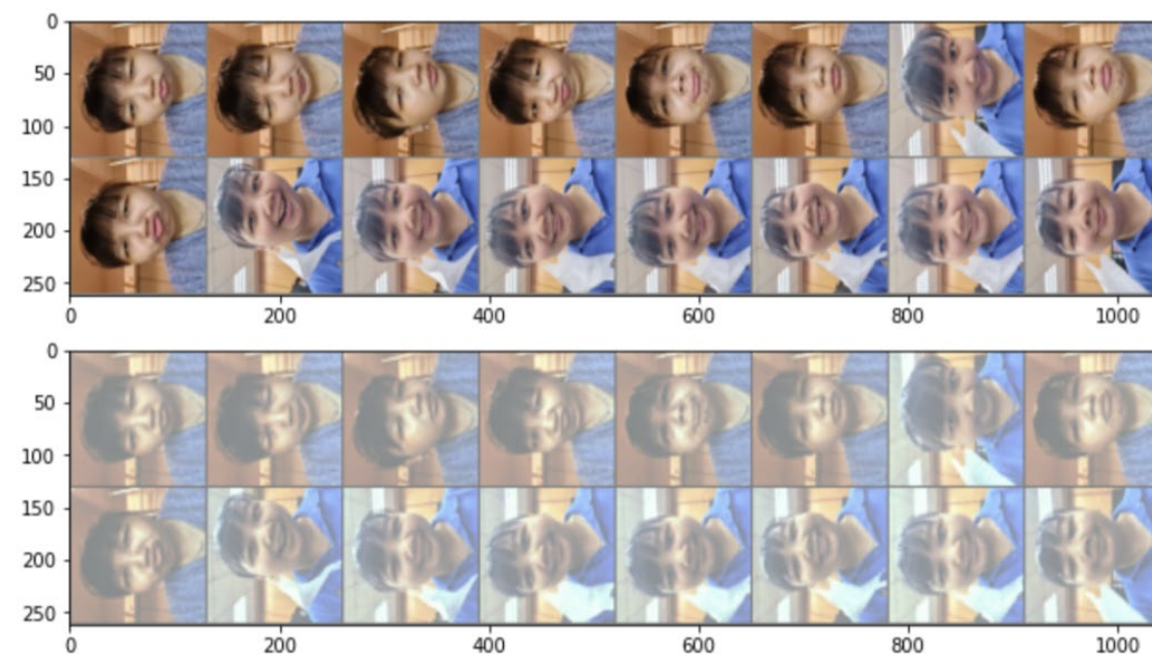


t-SNE (perplexity=?) results of training images




1061259 Keren


Recovered training images




Recovered un-seen test images




Happy = 12, Angry = 8, Latent vector size = 20, 1200 epochs



t-SNE (perplexity=?) results of training images




Recovered training images



Recovered un-seen test images



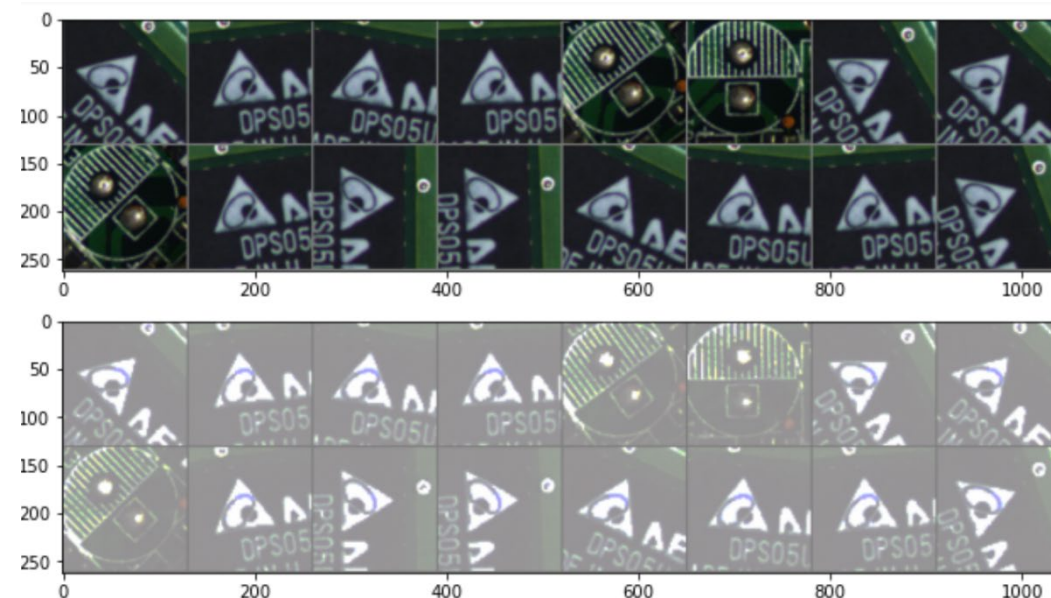
Class 1 = ?, Class 2 = ?, Latent vector size = ?, ? epochs



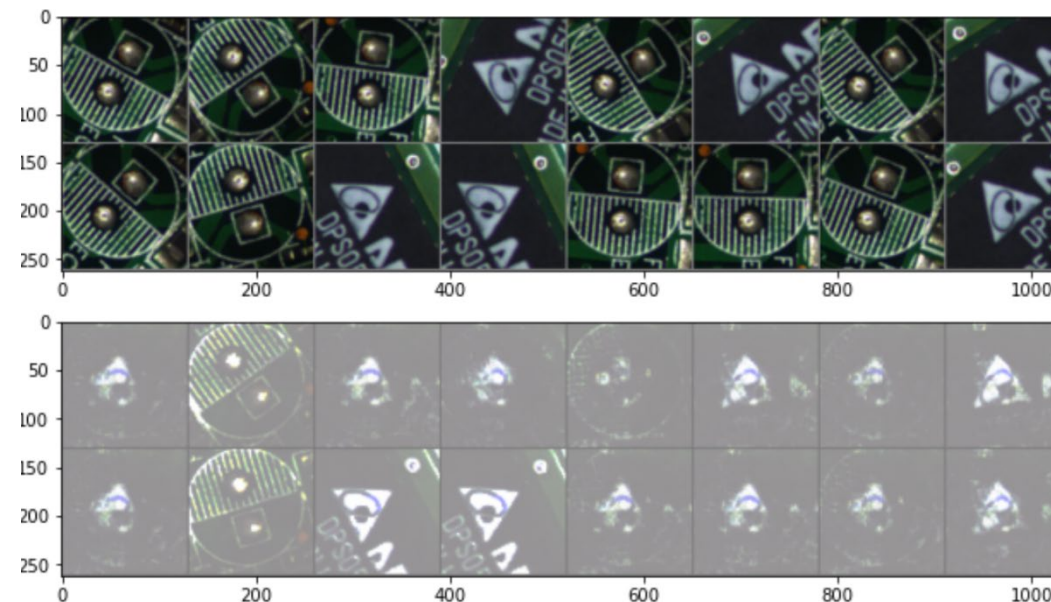
t-SNE (perplexity=?) results of training images

1085442 Carlos

Recovered training images



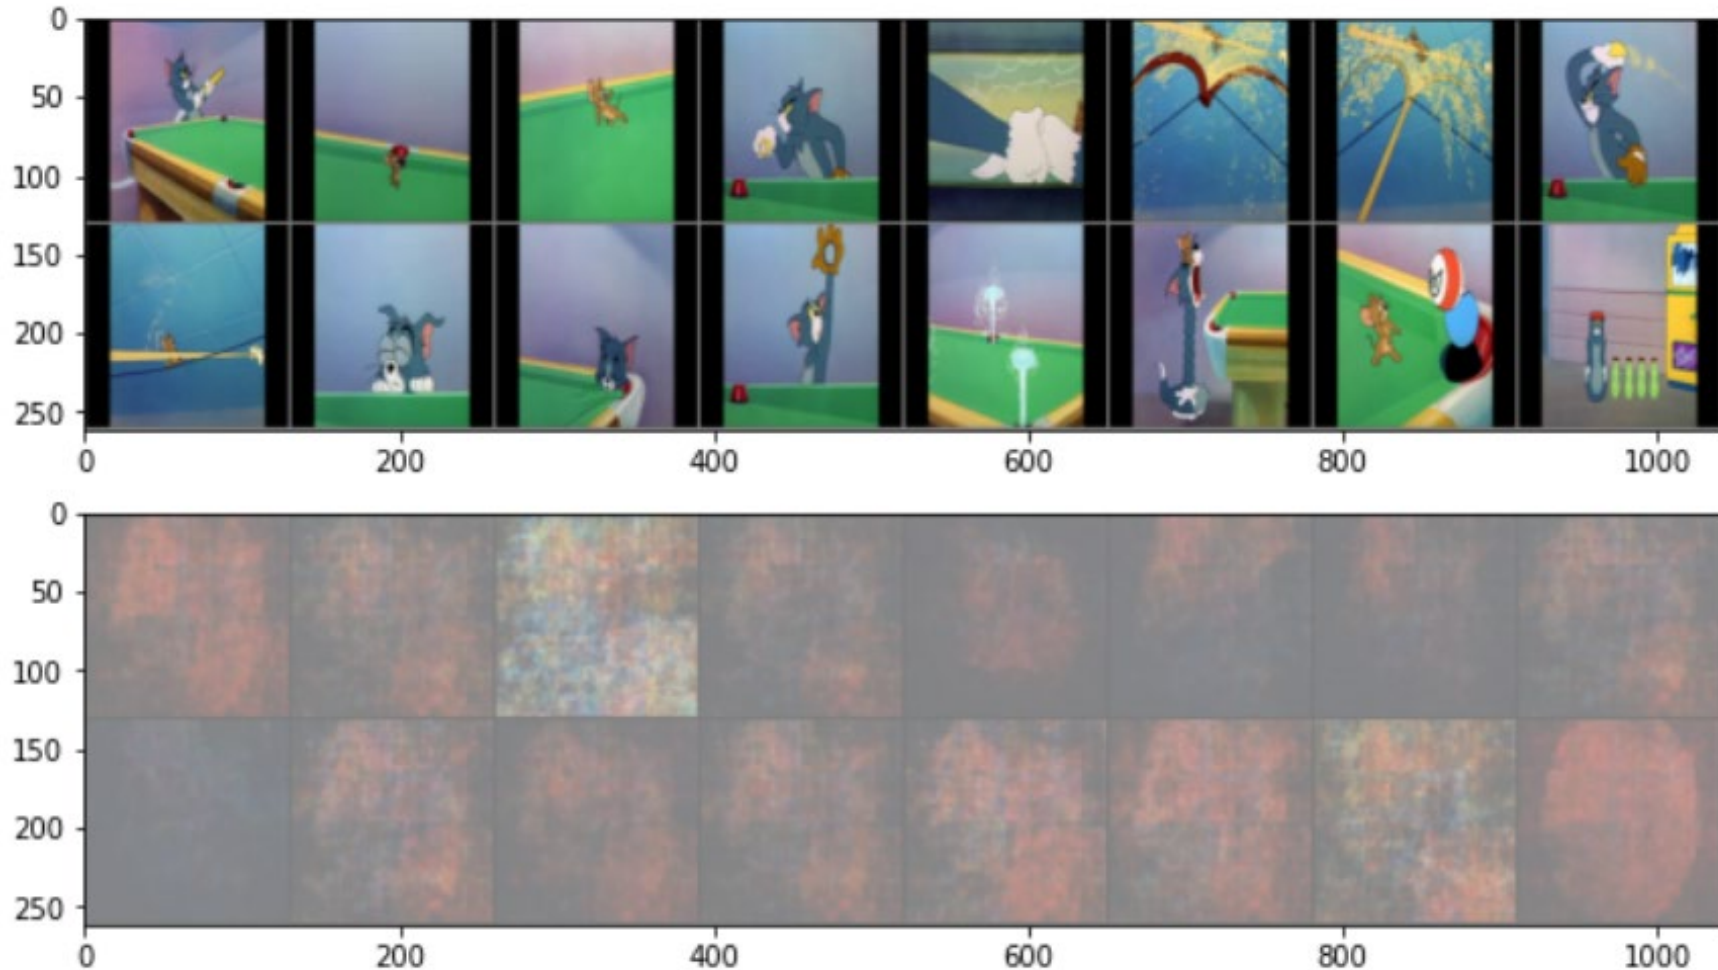
Recovered test images



How about Tom & Jerry?

Results are still not good after 1200 epochs

Test on un-seen images – fails to reconstruct the input images

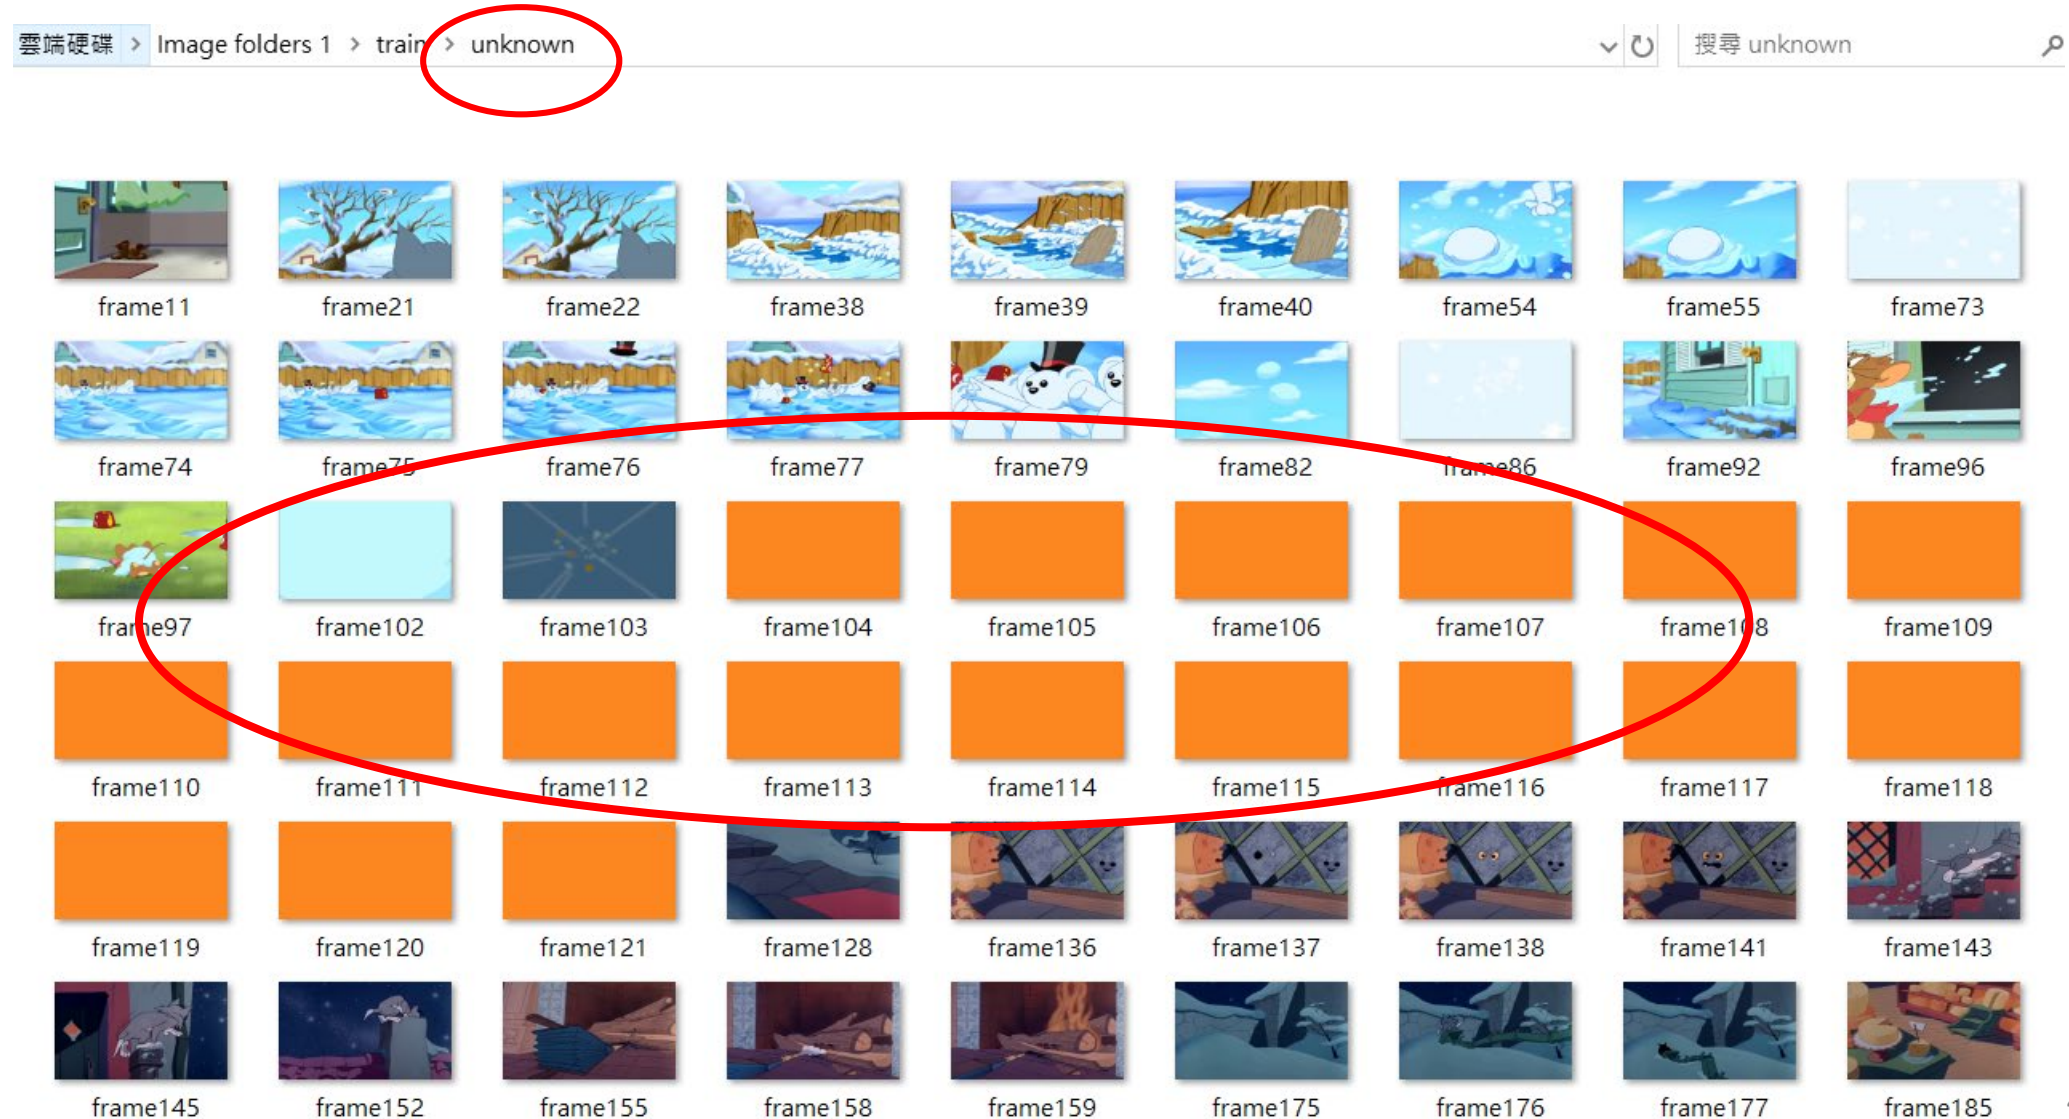


What to do if training is not successful?


- Check if the training data is too diversified? Your AI model is as good as your data, and as bad as your data too.
- Examine loss plot to understand the gradient decent process. Train with more epochs if the trend of the optimization process is good.
- Tune other hyper-parameters, e.g., batch size?



Your AI model is as good as your data, and as bad as your data too




Create a new folder "Image folder 1", remove the "Unknown" sub-folder

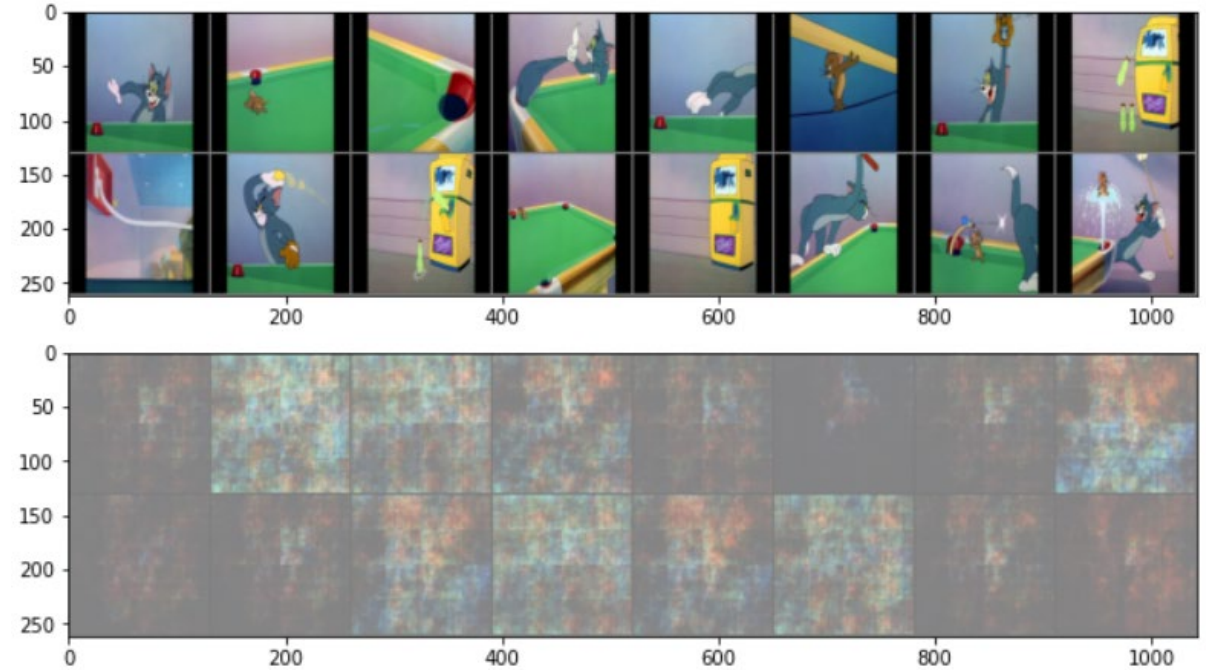
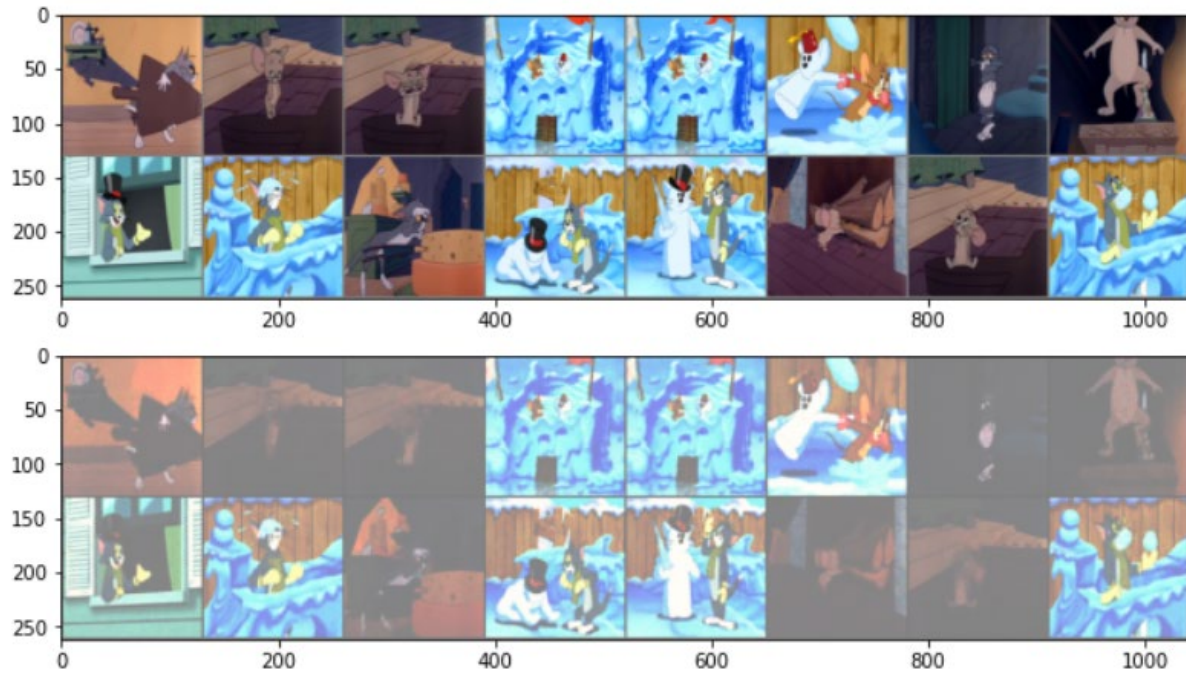


Test folder remains the same


Train 1200 epochs after removing the "unknown" folder




Train 1200 epochs after removing the "unknown" folder



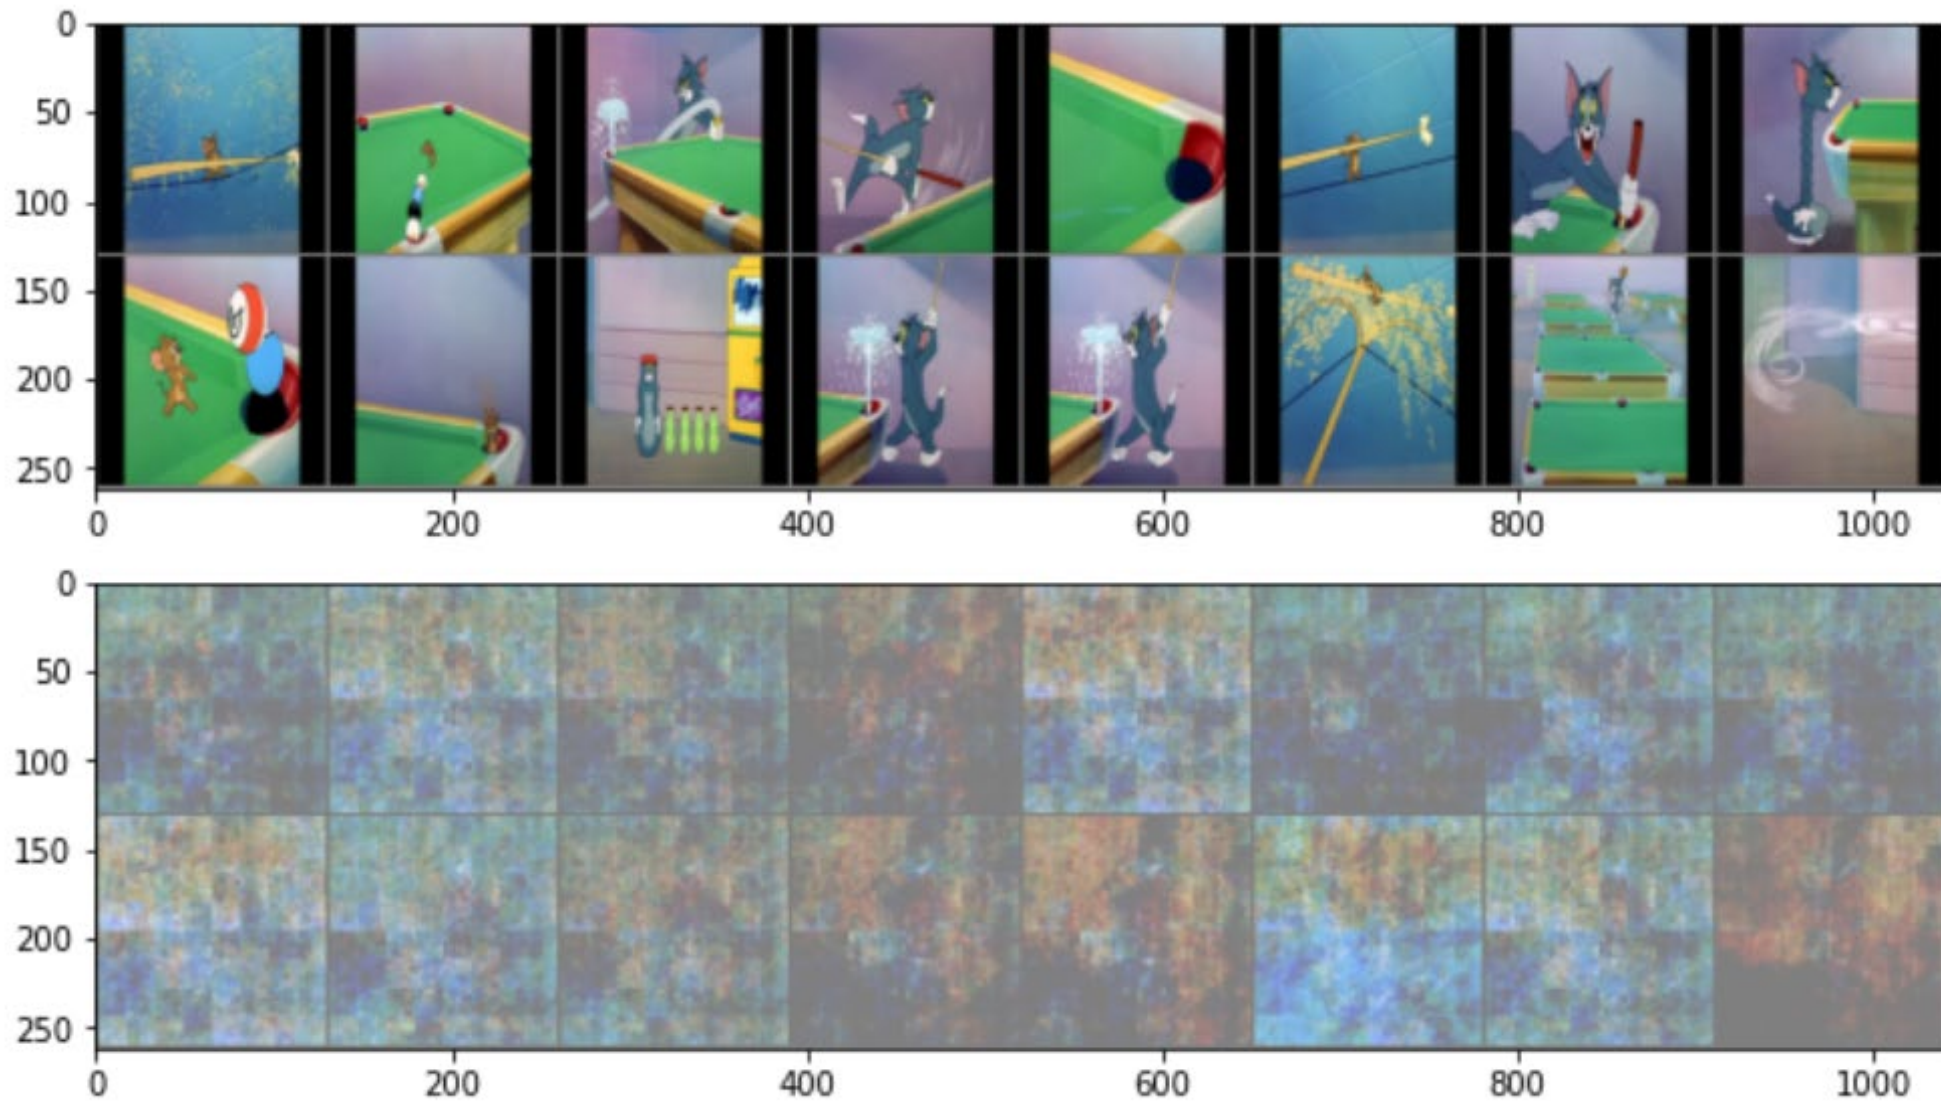
Examine the loss plot to understand the gradient decent process.



Results after training 20,000 epochs





Results after training 20,000 epochs



How about batch size? Increase or decrees ?


```
self.encoder = nn.Sequential(
    nn.Conv2d(3, 32, kernel_size=2, stride=2),
    nn.BatchNorm2d(32, eps=1e-05, momentum=0.1, af
    nn.ReLU(),
    nn.Conv2d(32, 64, kernel_size=2, stride=2),
    nn.BatchNorm2d(64, eps=1e-05, momentum=0.1, af
    nn.ReLU(),
    nn.Conv2d(64, 128, kernel_size=2, stride=2),
    nn.BatchNorm2d(128, eps=1e-05, momentum=0.1, a
    nn.ReLU(),
    nn.Conv2d(128, 256, kernel_size=2, stride=2),
    nn.BatchNorm2d(256, eps=1e-05, momentum=0.1, a
    nn.ReLU(),
    nn.Conv2d(256, 512, kernel_size=2, stride=2),
    nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, a
    nn.ReLU(),
    nn.Conv2d(512, 1024, kernel_size=2, stride=2),
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1,
    nn.ReLU(),
    nn.Conv2d(1024, 1024, kernel_size=2, stride=2)
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1,
    nn.ReLU(),
    Flatten(),
    nn.Linear(in_features=i, out_features=o),
)
```

```
[12]: import torch.utils.data as Data
      loader = Data.DataLoader(
          dataset=train_dataset,
          batch_size=16,
          shuffle=True)
```




epoch=1200, batch size=8

Train:

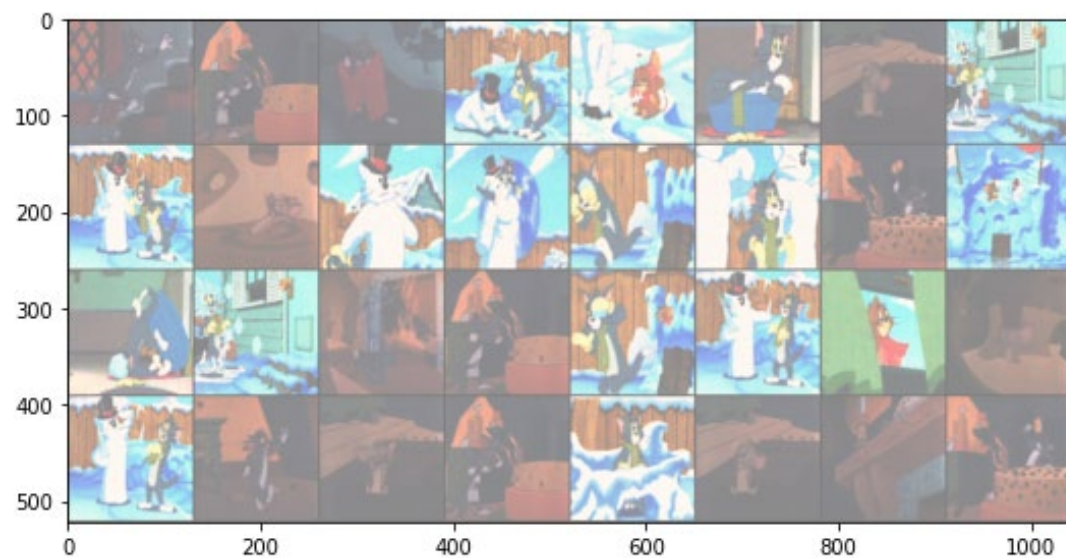
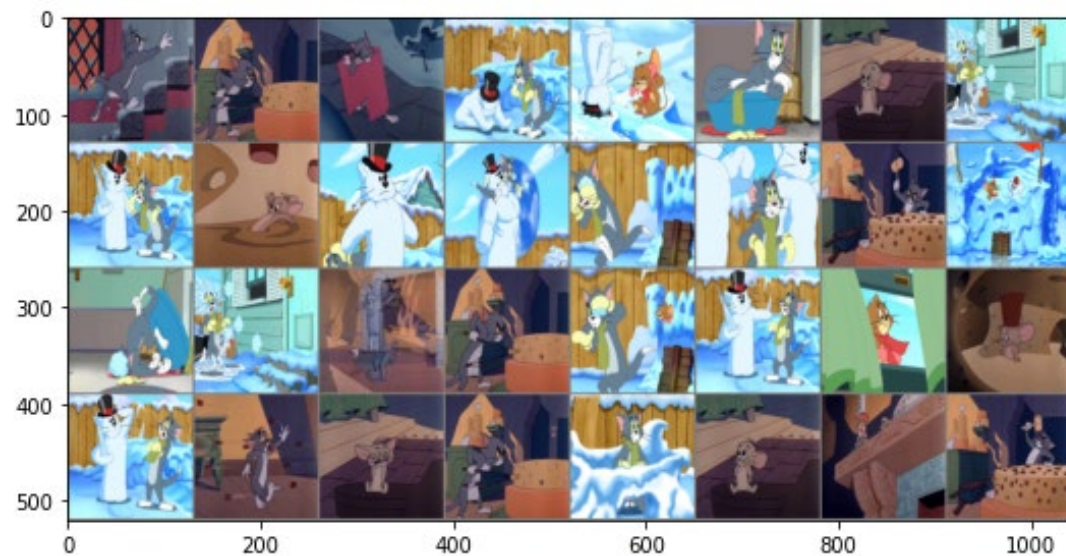


Test:





epoch=1200, batch size=32

Train:



Test:



2000 epochs using new training images