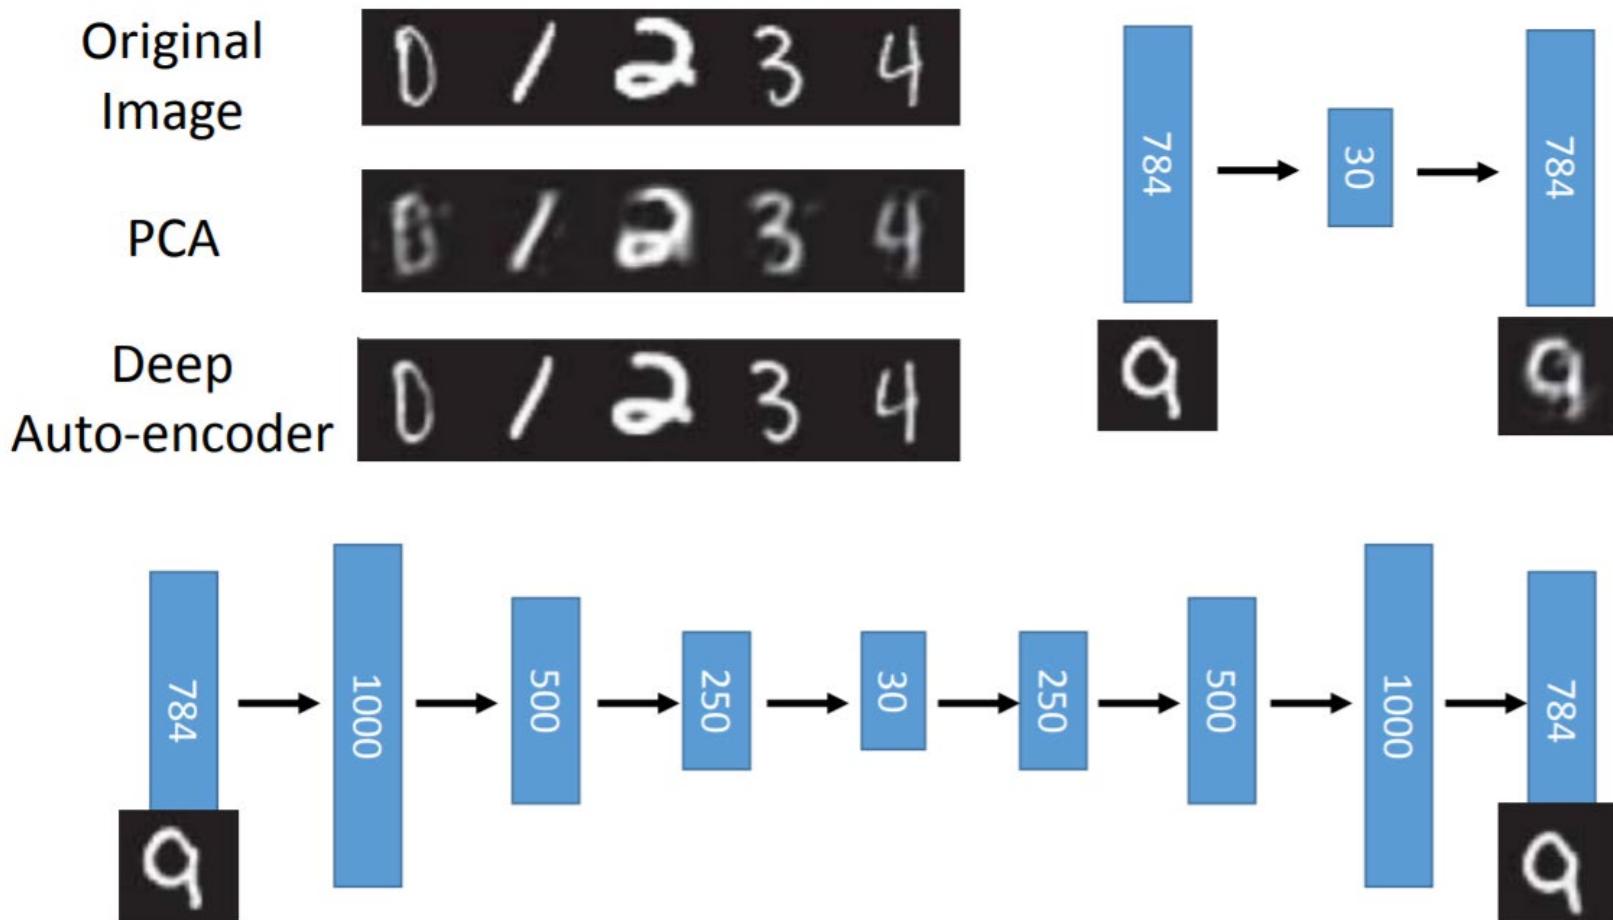


Auto-encoder

MLP based autoencoder



Reference: Hinton, Geoffrey E., and Ruslan R. Salakhutdinov. "Reducing the dimensionality of data with neural networks." *Science* 313.5786 (2006): 504-507

- CNN Image Classifier – Convolution section + MLP classifier
- CNN Autoencoder – Convolution section + Deconvolution section to recover the input image

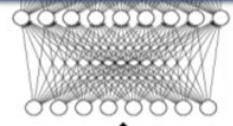
HW5 $y = f(x)$

$y=1, \dots, 1000$

cat dog

HW4

Fully Connected
Feedforward network



$x=244 \times 244 \times 3$



Convolution

Max Pooling

Convolution

Max Pooling

Can repeat
many times

HW6 $x = f(x)$

$y=244 \times 244 \times 3$

Deconvolution

Unpooling

Deconvolution

Unpooling

Deconvolution

As close as
possible

$x=244 \times 244 \times 3$



Convolution

Pooling

Convolution

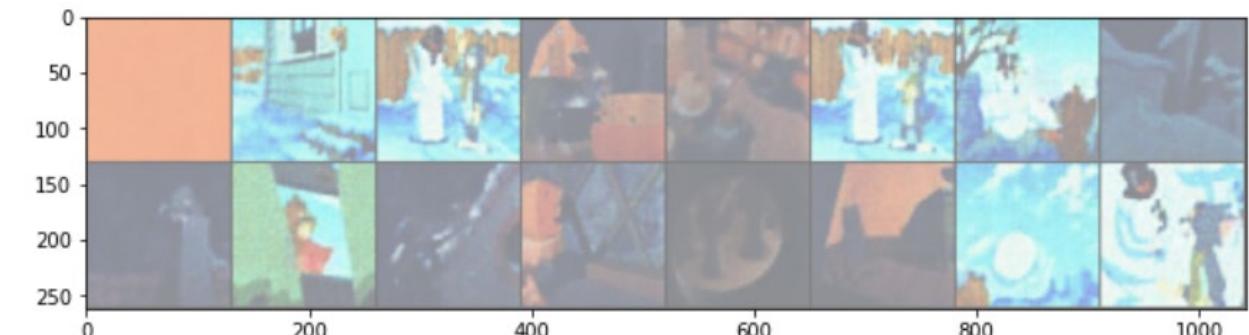
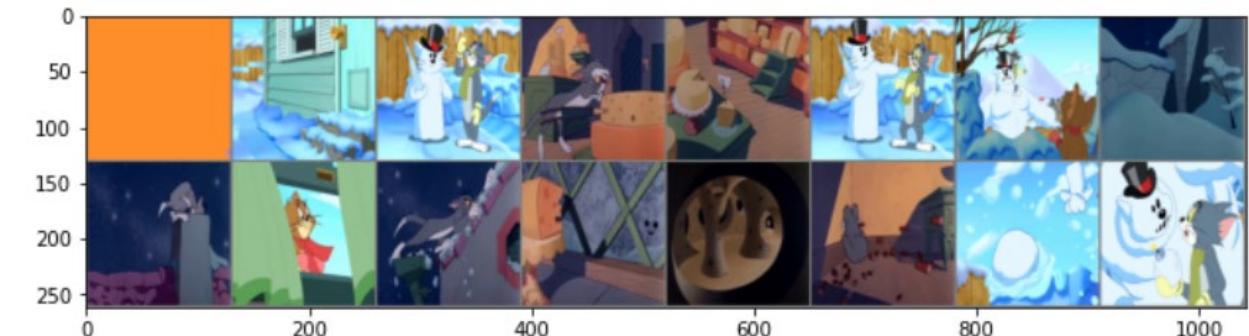
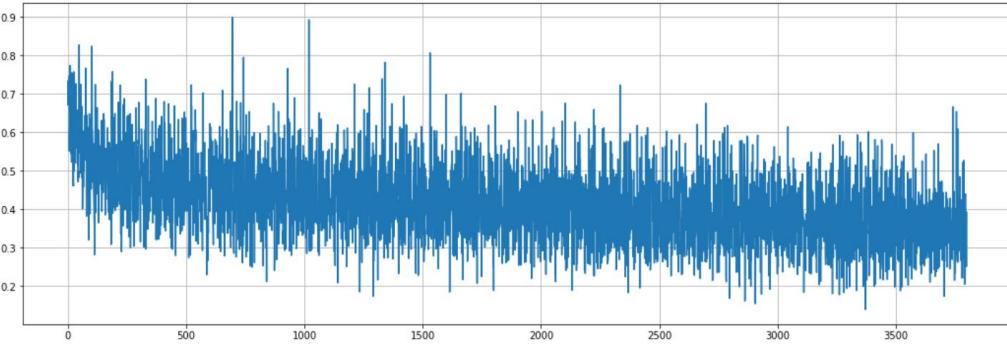
Pooling

Practice

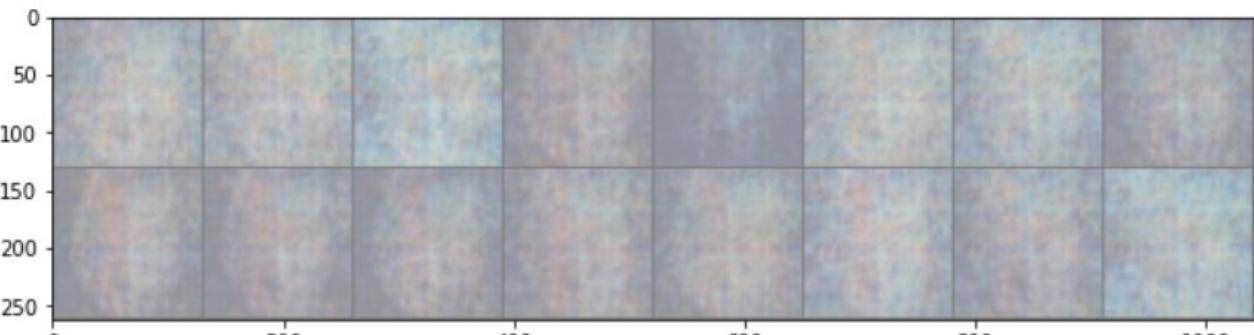
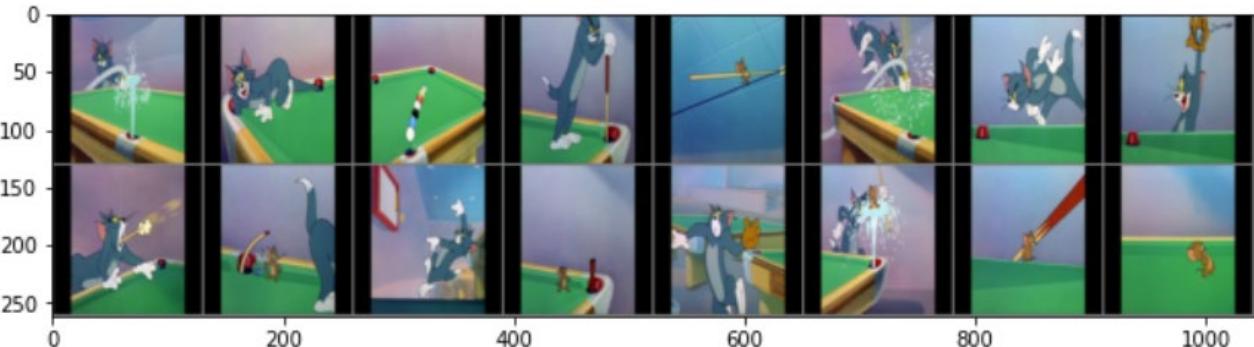
- Run "7.1.Conv_AE.ipynb"



Train 200 epochs

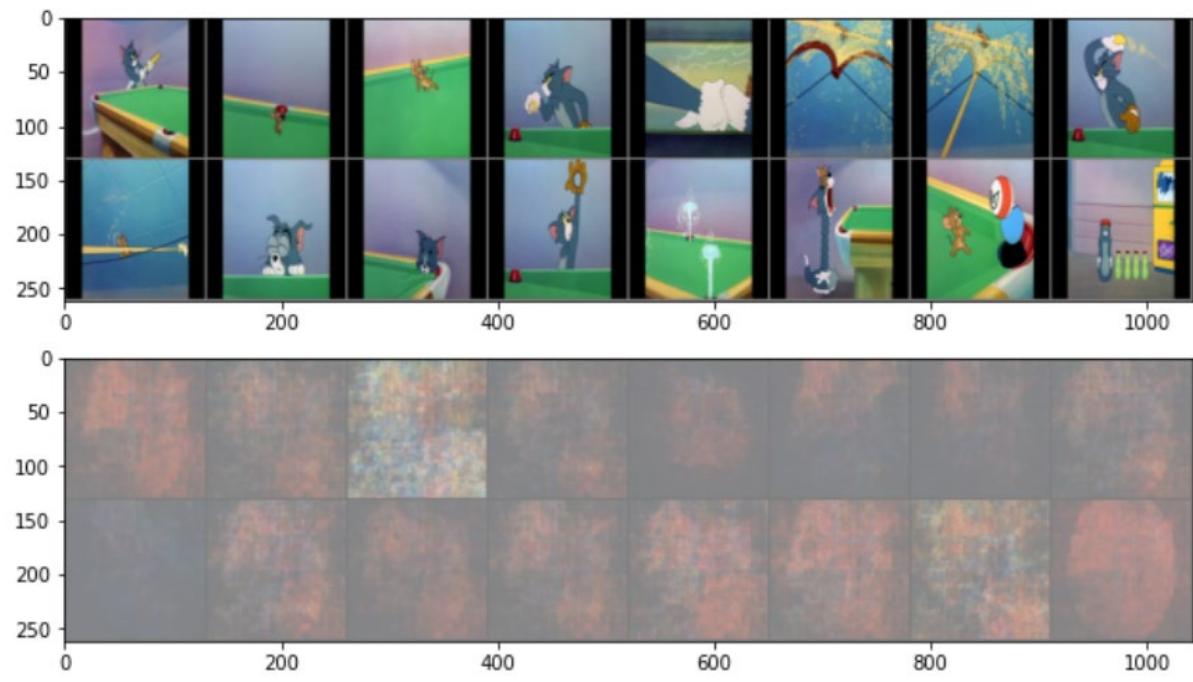
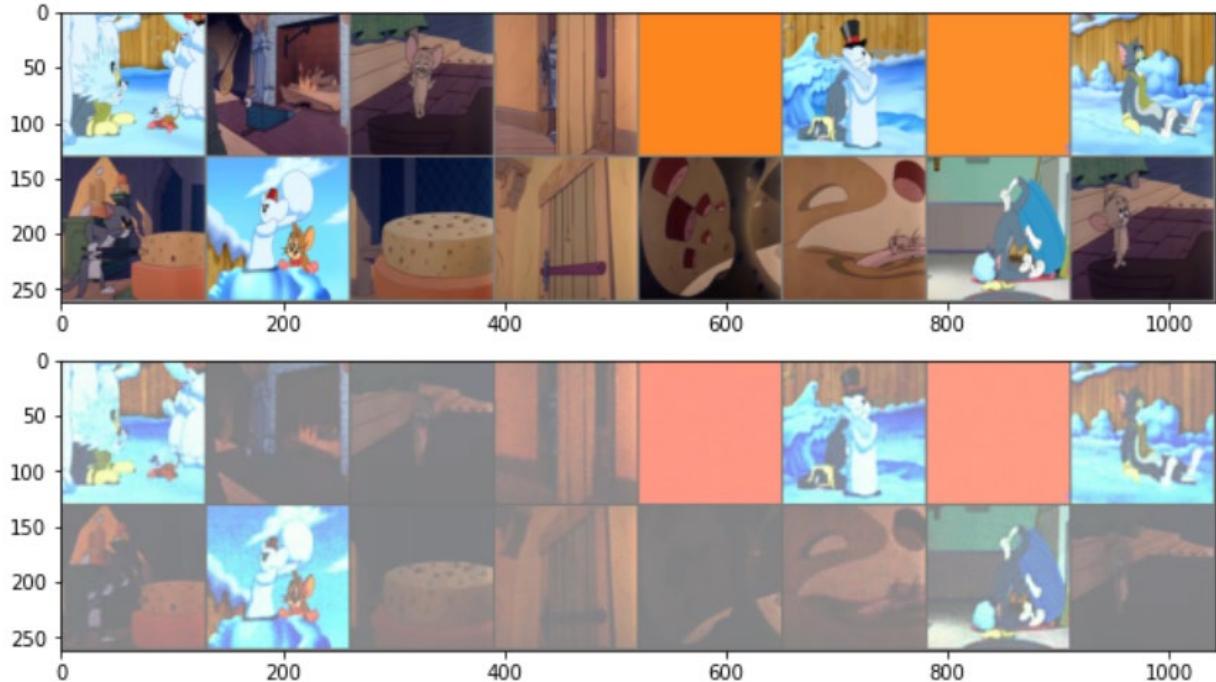
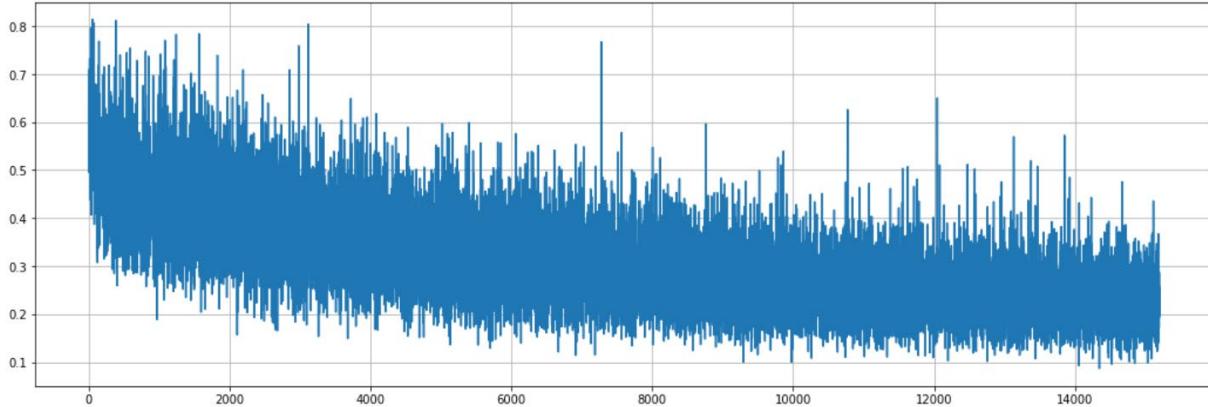


Able to recover the training images



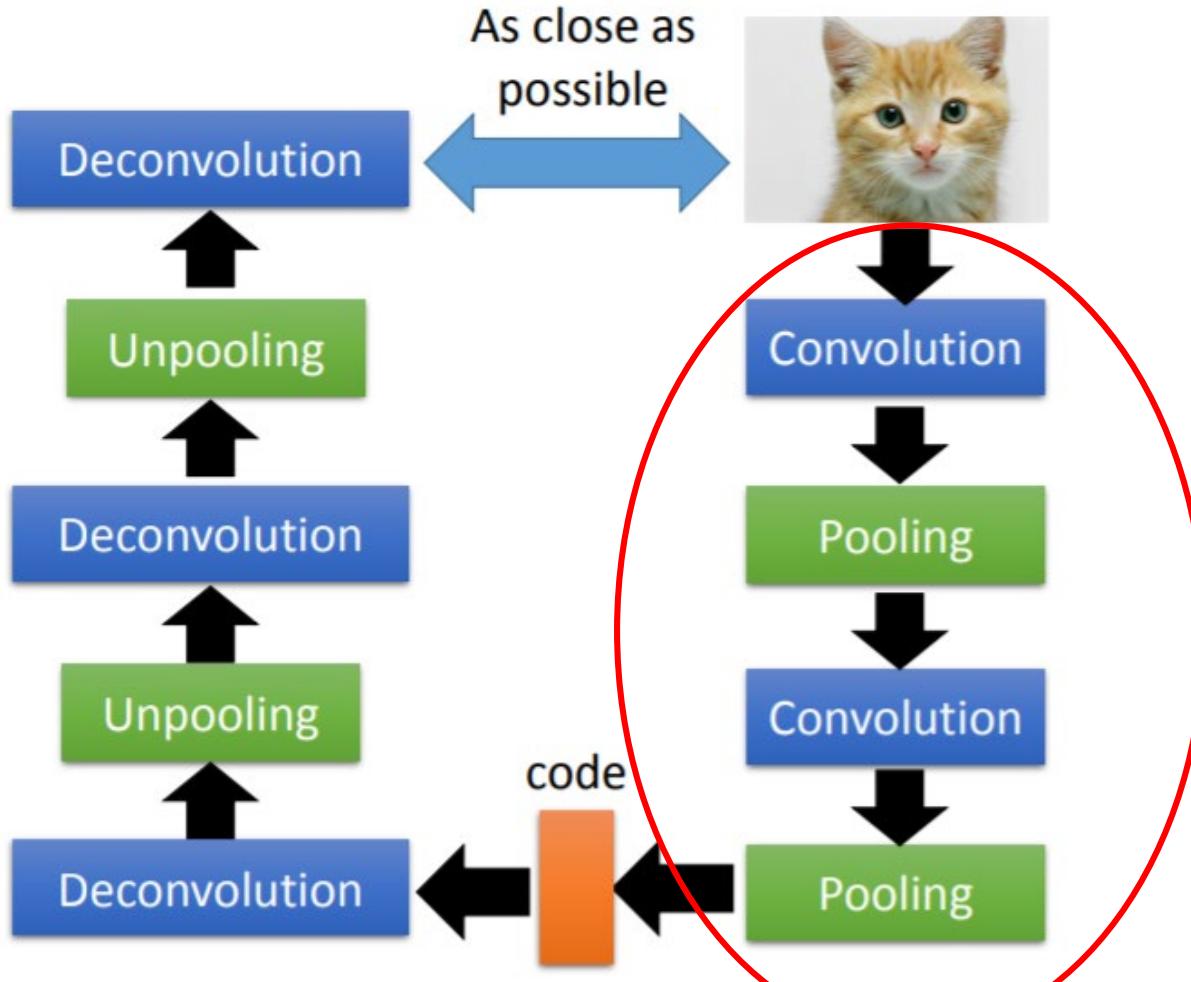
Fails to reconstruct the test images

Train 800 epochs



Still fails to reconstruct the test images

Encoder

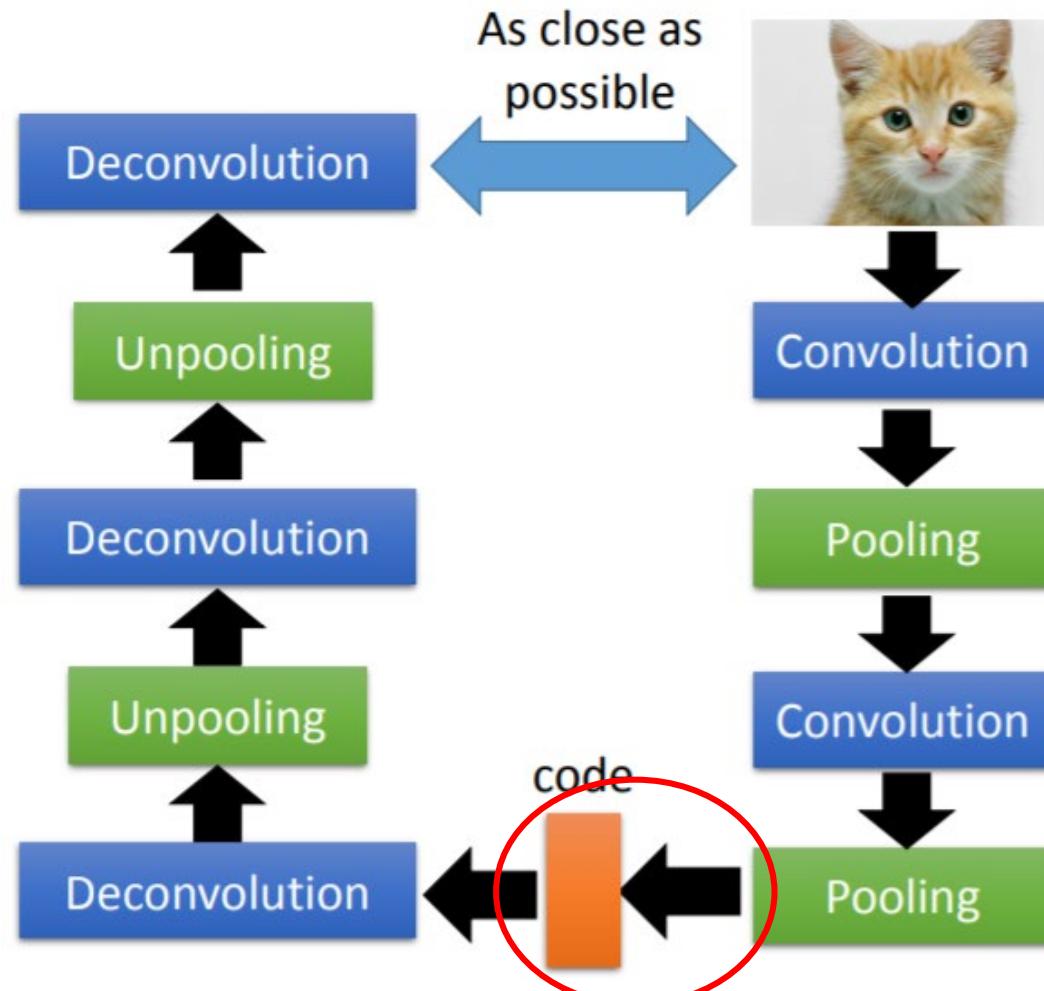


```
self.encoder = nn.Sequential(  
    nn.Conv2d(3, 32, kernel_size=2, stride=2),  
    nn.BatchNorm2d(32, eps=1e-05, momentum=0.1, af  
    nn.ReLU(),  
    nn.Conv2d(32, 64, kernel_size=2, stride=2),  
    nn.BatchNorm2d(64, eps=1e-05, momentum=0.1, af  
    nn.ReLU(),  
    nn.Conv2d(64, 128, kernel_size=2, stride=2),  
    nn.BatchNorm2d(128, eps=1e-05, momentum=0.1, a  
    nn.ReLU(),  
    nn.Conv2d(128, 256, kernel_size=2, stride=2),  
    nn.BatchNorm2d(256, eps=1e-05, momentum=0.1, a  
    nn.ReLU(),  
    nn.Conv2d(256, 512, kernel_size=2, stride=2),  
    nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, a  
    nn.ReLU(),  
    nn.Conv2d(512, 1024, kernel_size=2, stride=2),  
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1,  
    nn.ReLU(),  
    nn.Conv2d(1024, 1024, kernel_size=2, stride=2)  
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1,  
    nn.ReLU(),  
    Flatten(),  
    nn.Linear(in_features=i, out_features=o),  
)
```

Practice to write down the feature map size
and the results after flatten

- Let input image = $224 \times 224 \times 3$
- Draw the feature maps (H , W , depth) after each convolution and max pooling
- What is the number of nodes after flatten?

Latent vector



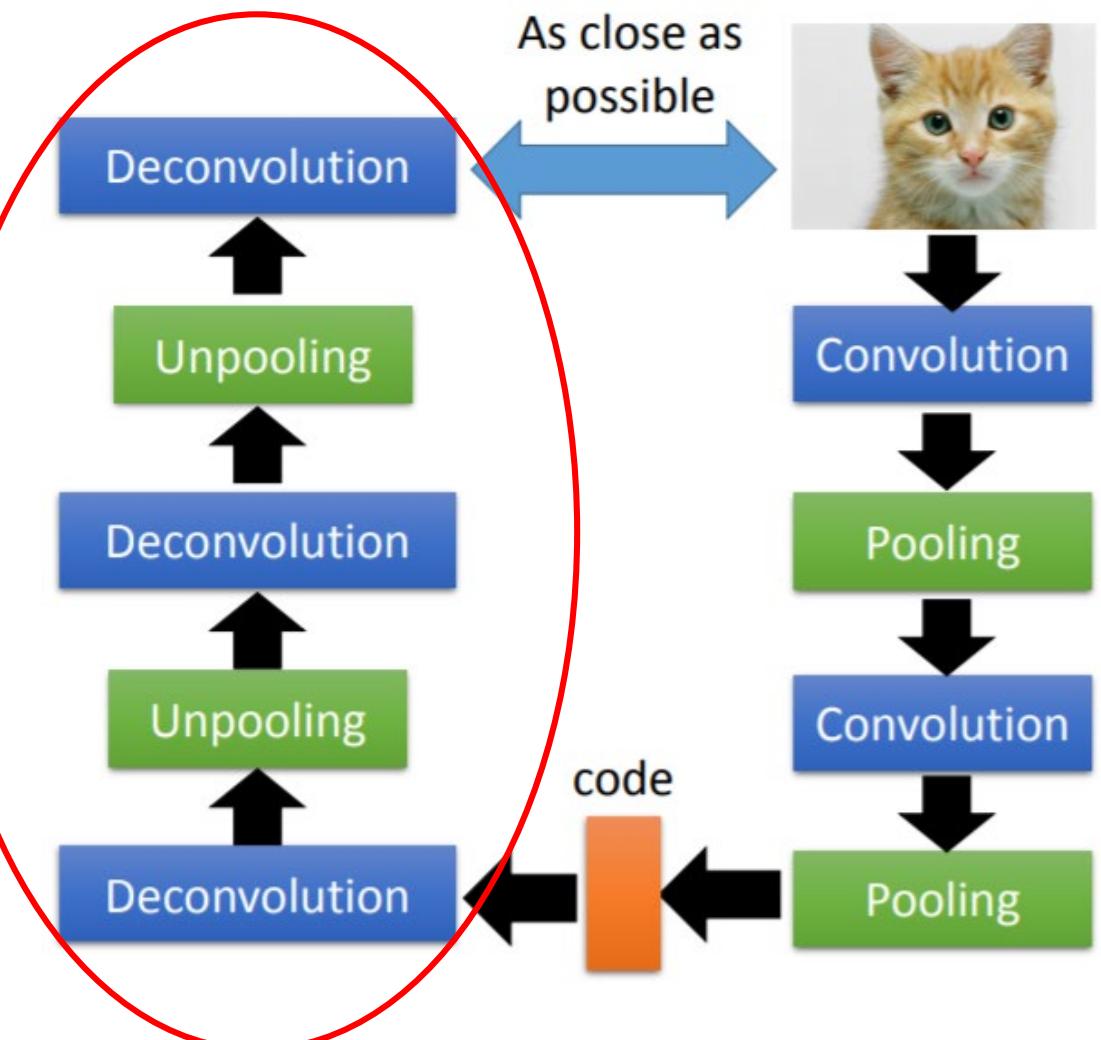
Reference: 李弘毅 ML Lecture 16 <https://youtu.be/Tk5B4seA-AU>

```
class autoencoder(nn.Module):
    def __init__(self,i=1024,o=64):
        super(autoencoder, self).__init__()
        self.encoder = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=2, stride=2),
            nn.BatchNorm2d(32, eps=1e-05, momentum=0.9),
            nn.ReLU(),
            nn.Conv2d(32, 64, kernel_size=2, stride=2),
            nn.BatchNorm2d(64, eps=1e-05, momentum=0.9),
            nn.ReLU(),
            nn.Conv2d(64, 128, kernel_size=2, stride=2),
            nn.BatchNorm2d(128, eps=1e-05, momentum=0.9),
            nn.ReLU(),
            nn.Conv2d(128, 256, kernel_size=2, stride=2),
            nn.BatchNorm2d(256, eps=1e-05, momentum=0.9),
            nn.ReLU(),
            nn.Conv2d(256, 512, kernel_size=2, stride=2),
            nn.BatchNorm2d(512, eps=1e-05, momentum=0.9),
            nn.ReLU(),
            nn.Conv2d(512, 1024, kernel_size=2, stride=2),
            nn.BatchNorm2d(1024, eps=1e-05, momentum=0.9),
            nn.ReLU(),
            nn.Conv2d(1024, 1024, kernel_size=2, stride=2),
            nn.BatchNorm2d(1024, eps=1e-05, momentum=0.9),
            nn.ReLU(),
            nn.Conv2d(1024, 1024, kernel_size=2, stride=2),
            nn.BatchNorm2d(1024, eps=1e-05, momentum=0.9),
            nn.ReLU(),
            nn.Conv2d(1024, 1024, kernel_size=2, stride=2),
            nn.BatchNorm2d(1024, eps=1e-05, momentum=0.9),
            nn.ReLU(),
            nn.Flatten(),
            nn.Linear(in_features=i, out_features=o),
        )
```

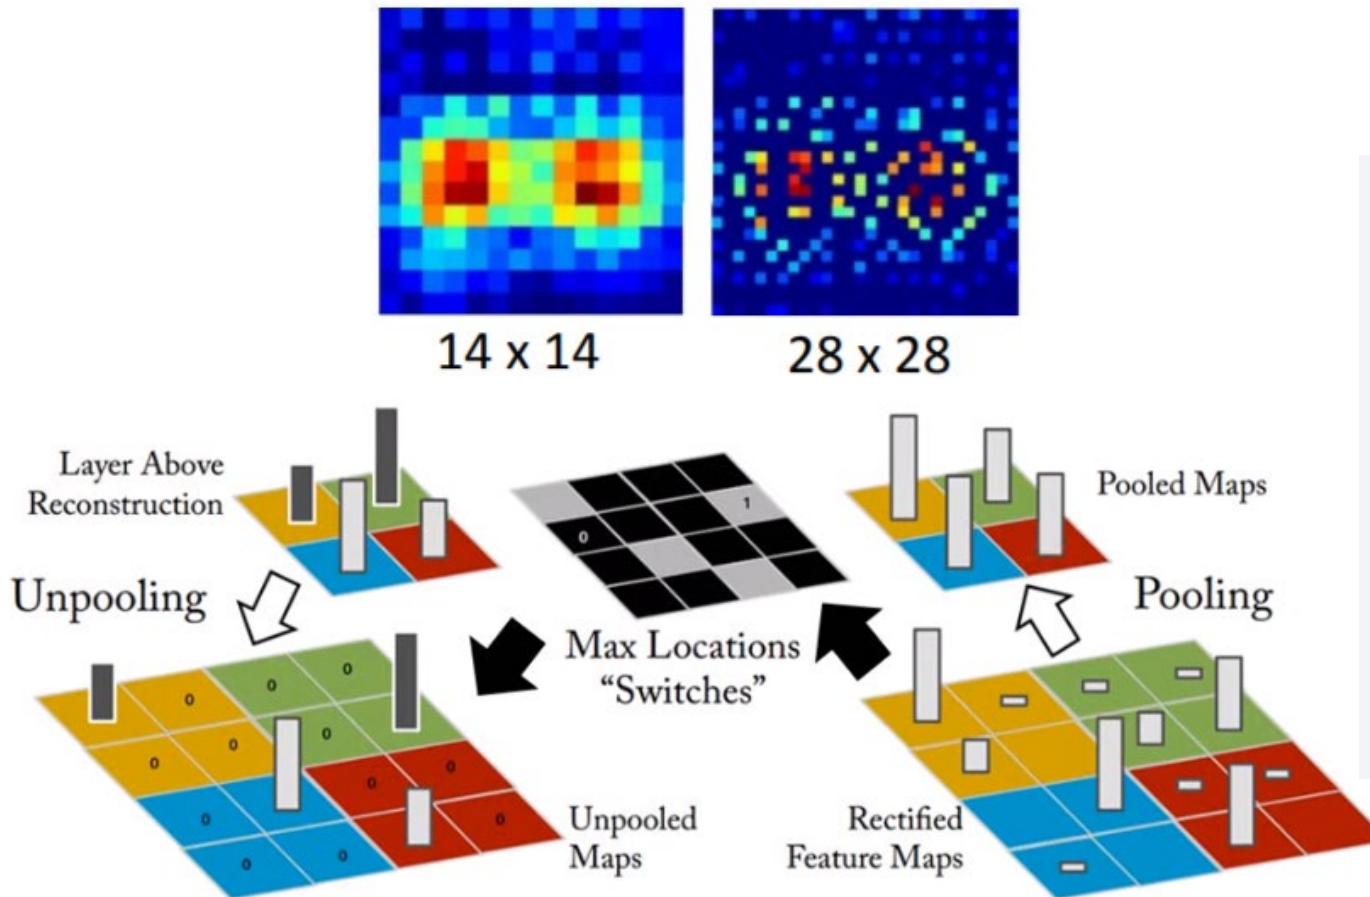
Flatten-22 [-1, 1024]
Linear-23 [-1, 64]
Linear-24 [-1, 1024]
UnFlatten-25 [-1, 1024, 1, 1]

Decoder

```
self.decoder = nn.Sequential(  
    nn.Linear(in_features=o, out_features=i),  
    UnFlatten(),  
    nn.ConvTranspose2d(1024, 1024, kernel_size=2, stride=2),  
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,  
    nn.ReLU(),  
    nn.ConvTranspose2d(1024, 512, kernel_size=2, stride=2),  
    nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,  
    nn.ReLU(),  
    nn.ConvTranspose2d(512, 256, kernel_size=2, stride=2),  
    nn.BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
    nn.ReLU(),  
    nn.ConvTranspose2d(256, 128, kernel_size=2, stride=2),  
    nn.BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,  
    nn.ReLU(),  
    nn.ConvTranspose2d(128, 64, kernel_size=2, stride=2),  
    nn.BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,  
    nn.ReLU(),  
    nn.ConvTranspose2d(64, 32, kernel_size=2, stride=2),  
    nn.BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,  
    nn.ReLU(),  
    nn.ConvTranspose2d(32, 3, kernel_size=2, stride=2),  
    nn.BatchNorm2d(3, eps=1e-05, momentum=0.1, affine=True,  
    nn.Sigmoid(),  
)
```



Unpooling

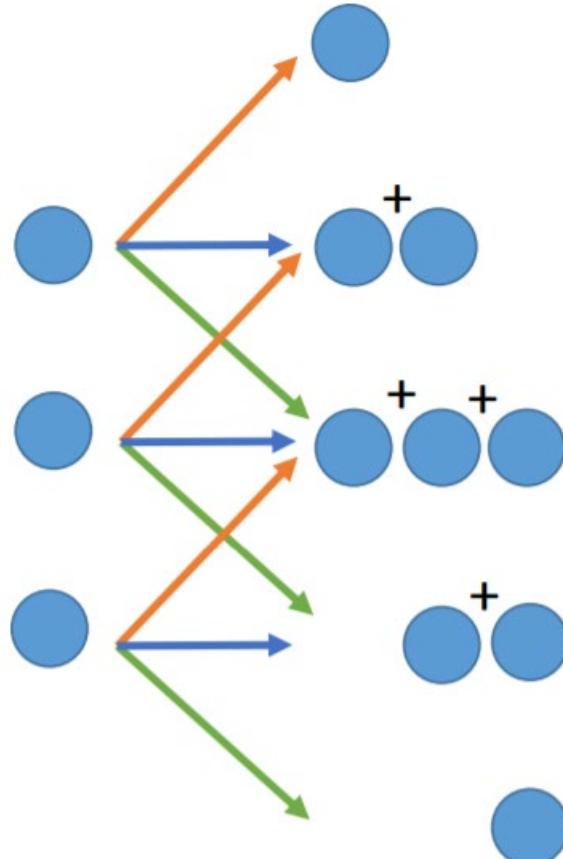
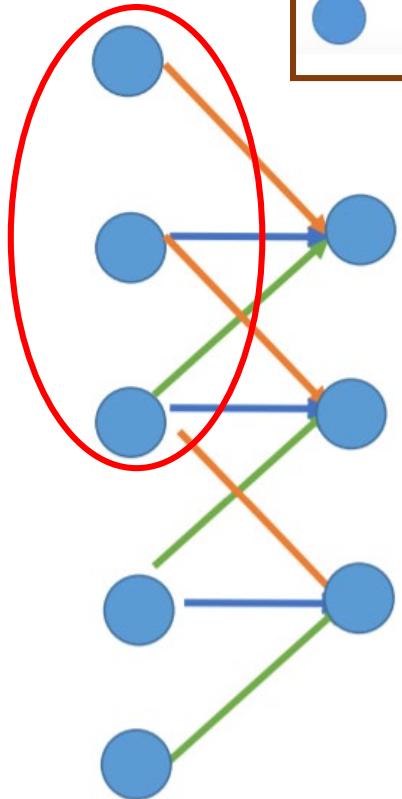
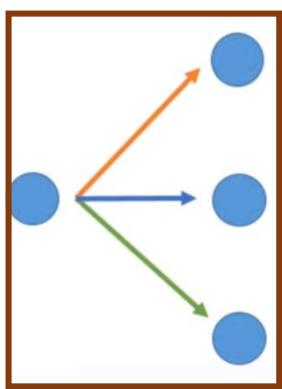
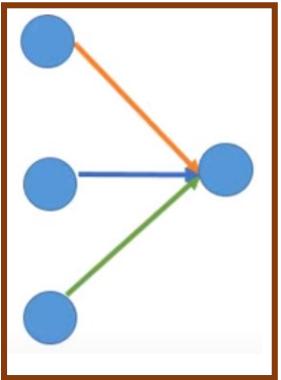


Reference: 李弘毅 ML Lecture 16 <https://youtu.be/Tk5B4seA-AU>

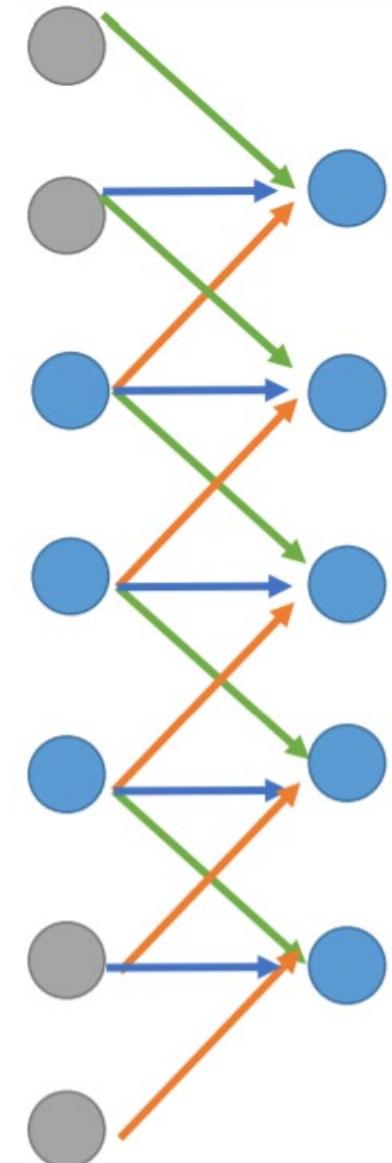
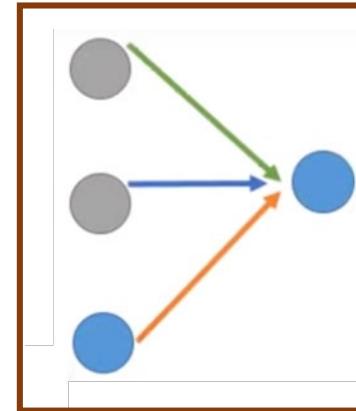
```
>>> pool = nn.MaxPool2d(2, stride=2, return_indices=True)
>>> unpool = nn.MaxUnpool2d(2, stride=2)
>>> input = torch.tensor([[[[ 1.,  2,  3,  4],
   [ 5,  6,  7,  8],
   [ 9, 10, 11, 12],
   [13, 14, 15, 16]]]]])
>>> output, indices = pool(input)
>>> unpool(output, indices)
tensor([[[[ 0.,  0.,  0.,  0.],
   [ 0.,  6.,  0.,  8.],
   [ 0.,  0.,  0.,  0.],
   [ 0., 14.,  0., 16.]]]])
```

Deconvolution

1D convolution, k=3 1D deconvolution, k=3



1D convolution, k=3



We only use deconvolution for up sampling, un-pooling is not used

```
self.decoder = nn.Sequential(  
    nn.Linear(in_features=o, out_features=i),  
    UnFlatten(),  
    nn.ConvTranspose2d(1024, 1024, kernel_size=2, stride=2),  
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,  
    nn.ReLU(),  
    nn.ConvTranspose2d(1024, 512, kernel_size=2, stride=2),  
    nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True  
    nn.ReLU(),  
    nn.ConvTranspose2d(512, 256, kernel_size=2, stride=2),  
    nn.BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True  
    nn.ReLU(),  
    nn.ConvTranspose2d(256, 128, kernel_size=2, stride=2),  
    nn.BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
    nn.ReLU(),  
    nn.ConvTranspose2d(128, 64, kernel_size=2, stride=2),  
    nn.BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,  
    nn.ReLU(),  
    nn.ConvTranspose2d(64, 32, kernel_size=2, stride=2),  
    nn.BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,  
    nn.ReLU(),  
    nn.ConvTranspose2d(32, 3, kernel_size=2, stride=2),  
    nn.BatchNorm2d(3, eps=1e-05, momentum=0.1, affine=True,  
    nn.Sigmoid(),  
)
```

Practice to write down the feature map size after deconvolution

- Input – the number of nodes after un-flatten
- Draw feature maps (H, W, depth) after each de-convolution and un-max pooling

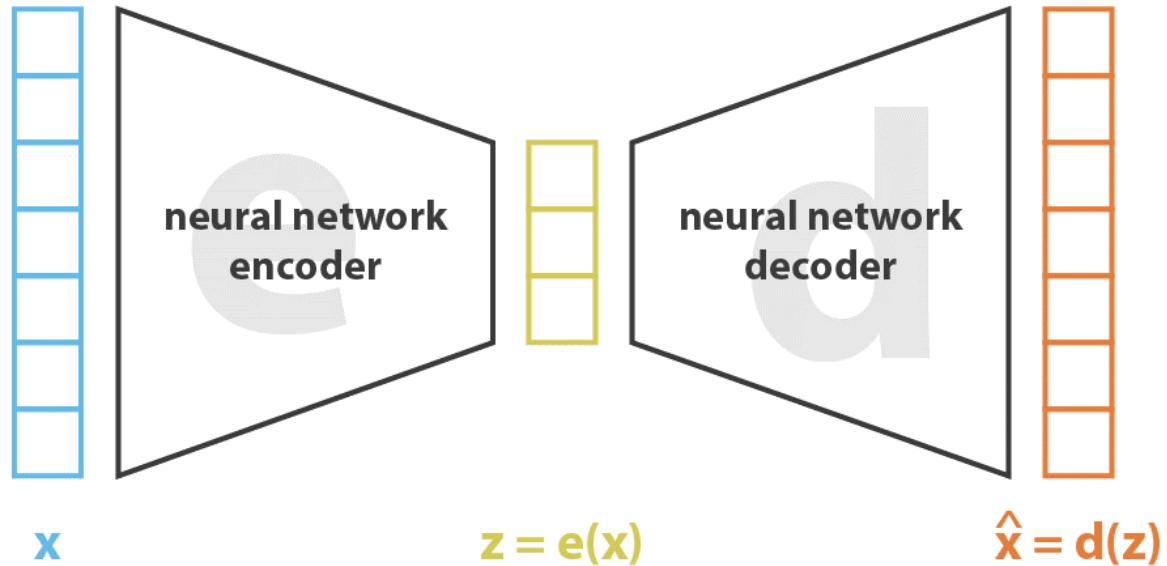


Feature map size after deconvolution

```
(2): ConvTranspose2d(1024, 1024, kernel_size=(2, 2), stride=(2, 2))
(3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(4): ReLU()
(5): ConvTranspose2d(1024, 512, kernel_size=(2, 2), stride=(2, 2))
(6): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(7): ReLU()
(8): ConvTranspose2d(512, 256, kernel_size=(2, 2), stride=(2, 2))
(9): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(10): ReLU()
```

ConvTranspose2d-26	[-1, 1024, 2, 2]
BatchNorm2d-27	[-1, 1024, 2, 2]
ReLU-28	[-1, 1024, 2, 2]
ConvTranspose2d-29	[-1, 512, 4, 4]
BatchNorm2d-30	[-1, 512, 4, 4]
ReLU-31	[-1, 512, 4, 4]
ConvTranspose2d-32	[-1, 256, 8, 8]
BatchNorm2d-33	[-1, 256, 8, 8]
ReLU-34	[-1, 256, 8, 8]
ConvTranspose2d-35	[-1, 128, 16, 16]
BatchNorm2d-36	[-1, 128, 16, 16]
ReLU-37	[-1, 128, 16, 16]
ConvTranspose2d-38	[-1, 64, 32, 32]
BatchNorm2d-39	[-1, 64, 32, 32]
ReLU-40	[-1, 64, 32, 32]

Loss function



$$\text{loss} = \|x - \hat{x}\|^2 = \|x - d(z)\|^2 = \|x - d(e(x))\|^2$$

Source: <https://towardsdatascience.com/understanding-variational-autoencoders-vaes-f70510919f73>

```
[13]: for batchX, _ in loader:  
    break;  
print(batchX.shape)  
torch.Size([16, 3, 128, 128])  
  
[14]: tensorY=model(batchX.to(device))  
print(tensorY.shape)  
torch.Size([16, 3, 128, 128])  
  
[15]: loss = loss_func(tensorY, batchX.to(device))  
print(loss)  
tensor(0.6961, device='cuda:0', grad_fn=<Msel
```

Save and load PyTorch model



The screenshot shows a Jupyter Notebook interface with the following details:

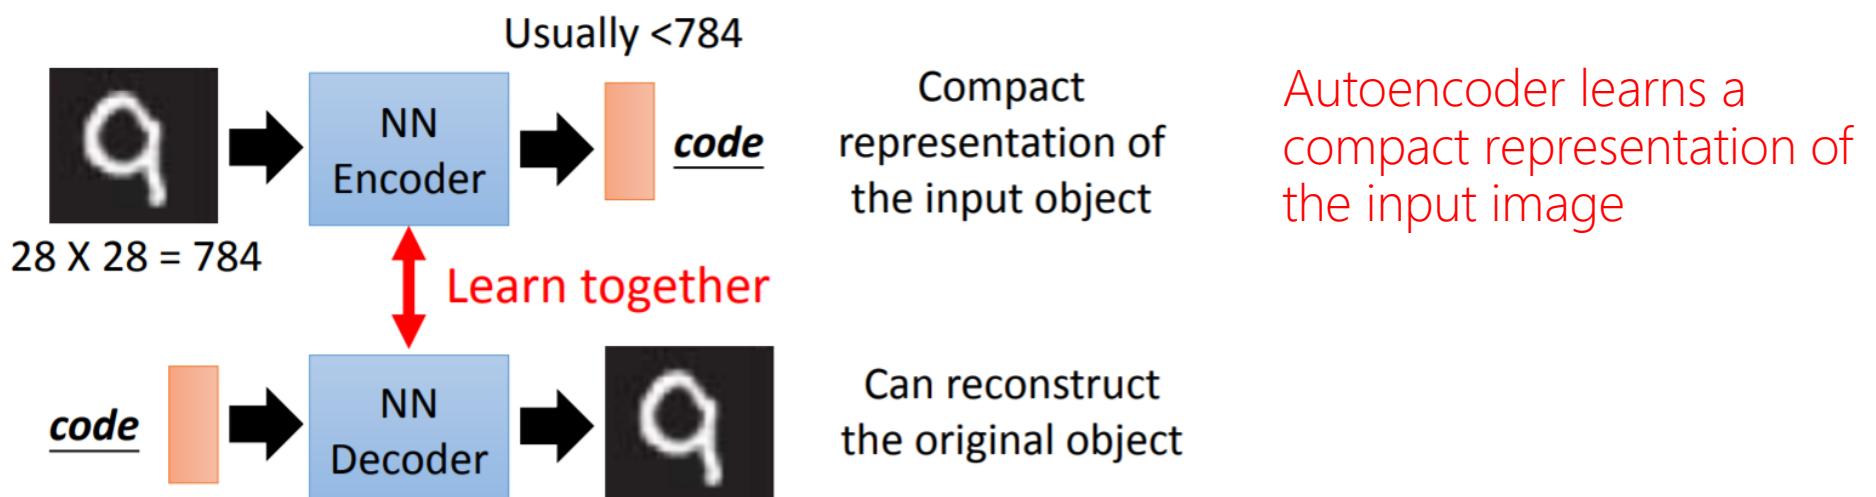
- Title Bar:** CO 3_AlexNet_(1).ipynb
- Menu Bar:** 檔案 編輯 檢視畫面 插入 執行階段 工具 說明 無法儲存變更
- Left Sidebar:** 檔案 (with a folder icon circled in red), 檔案, .., gdrive, sample data, AE800.pt (circled in red), tSNE.csv
- Right Main Area:**
 - Section Header:** Save and load a Pytorch model (II)
 - Code Cells:**
 - [27] `torch.save(model.state_dict(), "AE800.pt")` (The string "AE800.pt" is circled in red)
 - [28] `model=autoencoder() #build NN architecture`
`model.load_state_dict(torch.load("AE800.pt")) #load`
`model.to(device)`
`model.eval()`
 - Autoencoder Definition:**

```
autoencoder(
    encoder): Sequential(
        (0): Conv2d(3, 32, kernel_size=(2, 2), stride=1)
        (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True)
        (2): ReLU()
```

Pass images to AE to get their compact representation (latent vectors)

```
[37]: for step, (batchX, batchY) in enumerate(loader):
    tensorY = model.encoder(batchX.to(device))
    if(step==0):
        arrayX = np.array(tensorY.cpu().detach().numpy())
        arrayY = batchY.cpu().detach().numpy()
    else:
        arrayX = np.concatenate((arrayX, tensorY.cpu().detach().numpy()))
        arrayY = np.concatenate((arrayY, batchY.cpu().detach().numpy()))
    print(arrayX.shape, arrayY.shape)
```

(298, 64) (298,)



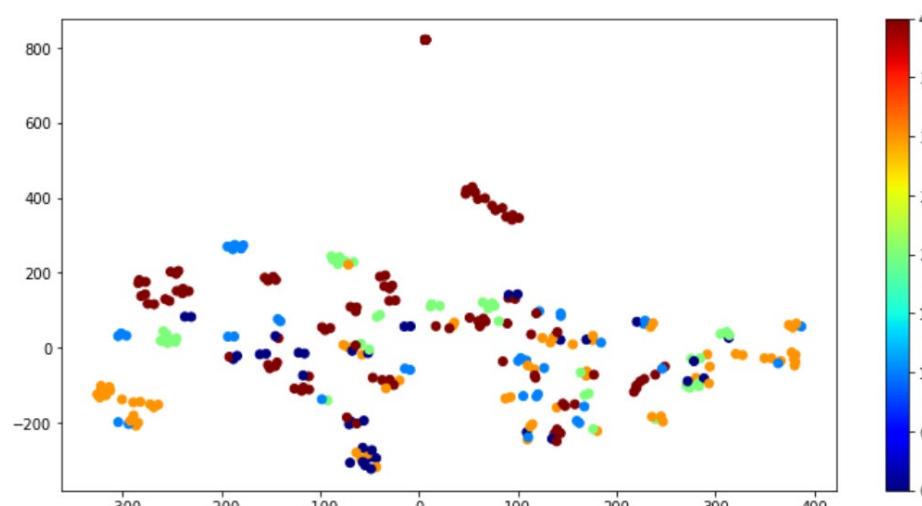
Use t -SNE to reduce the latent vector dimensions from 64 to 2

```
[38]: from sklearn.manifold import TSNE
tsne = TSNE(perplexity=5, n_components=2, init='pca', n_iter=5000)
# try perplexity = 5, 10, 30, 50
```

```
[39]: x=tsne.fit_transform(arrayX)
print(x.shape)
```

(298, 2)

```
[40]: plt.figure(figsize=(18,9))
plt.scatter(x[:, 0], x[:, 1], c= arrayY)
plt.show()
```



Save the results to csv file

CO 3_AlexNet_(1).ipynb

檔案 編輯 檢視畫面 插入 執行階段 工具 說明 無法儲存變更

+ 檔案 + 程式碼 + 文字 複製到雲端硬碟

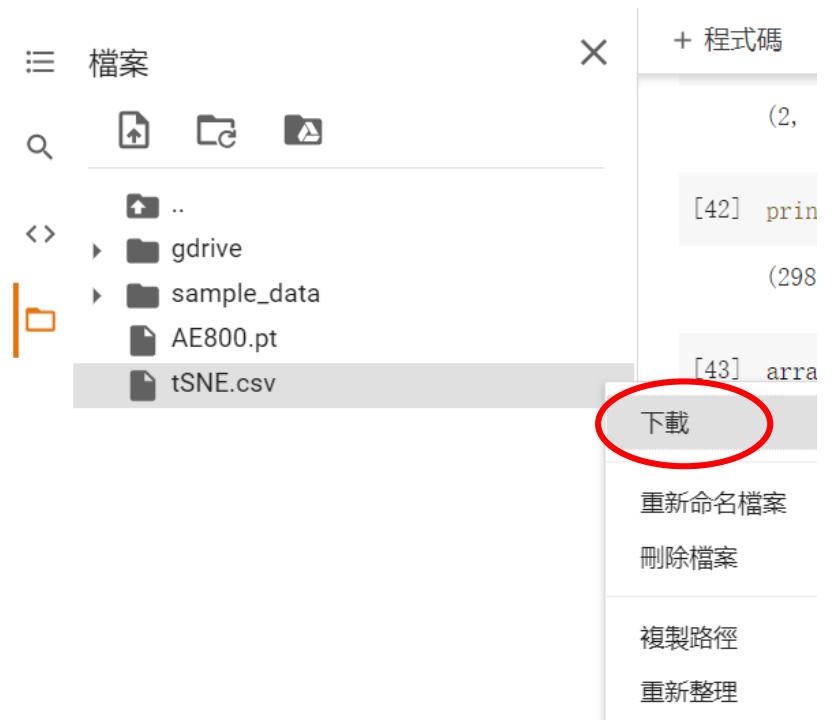
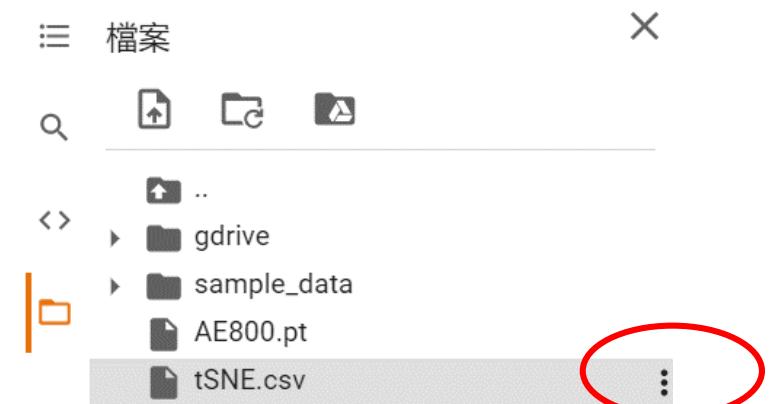
```
(2, 3) (2, 1) (2, 4)

[42] print(x.shape, arrayY.shape)
(298, 2) (298, )

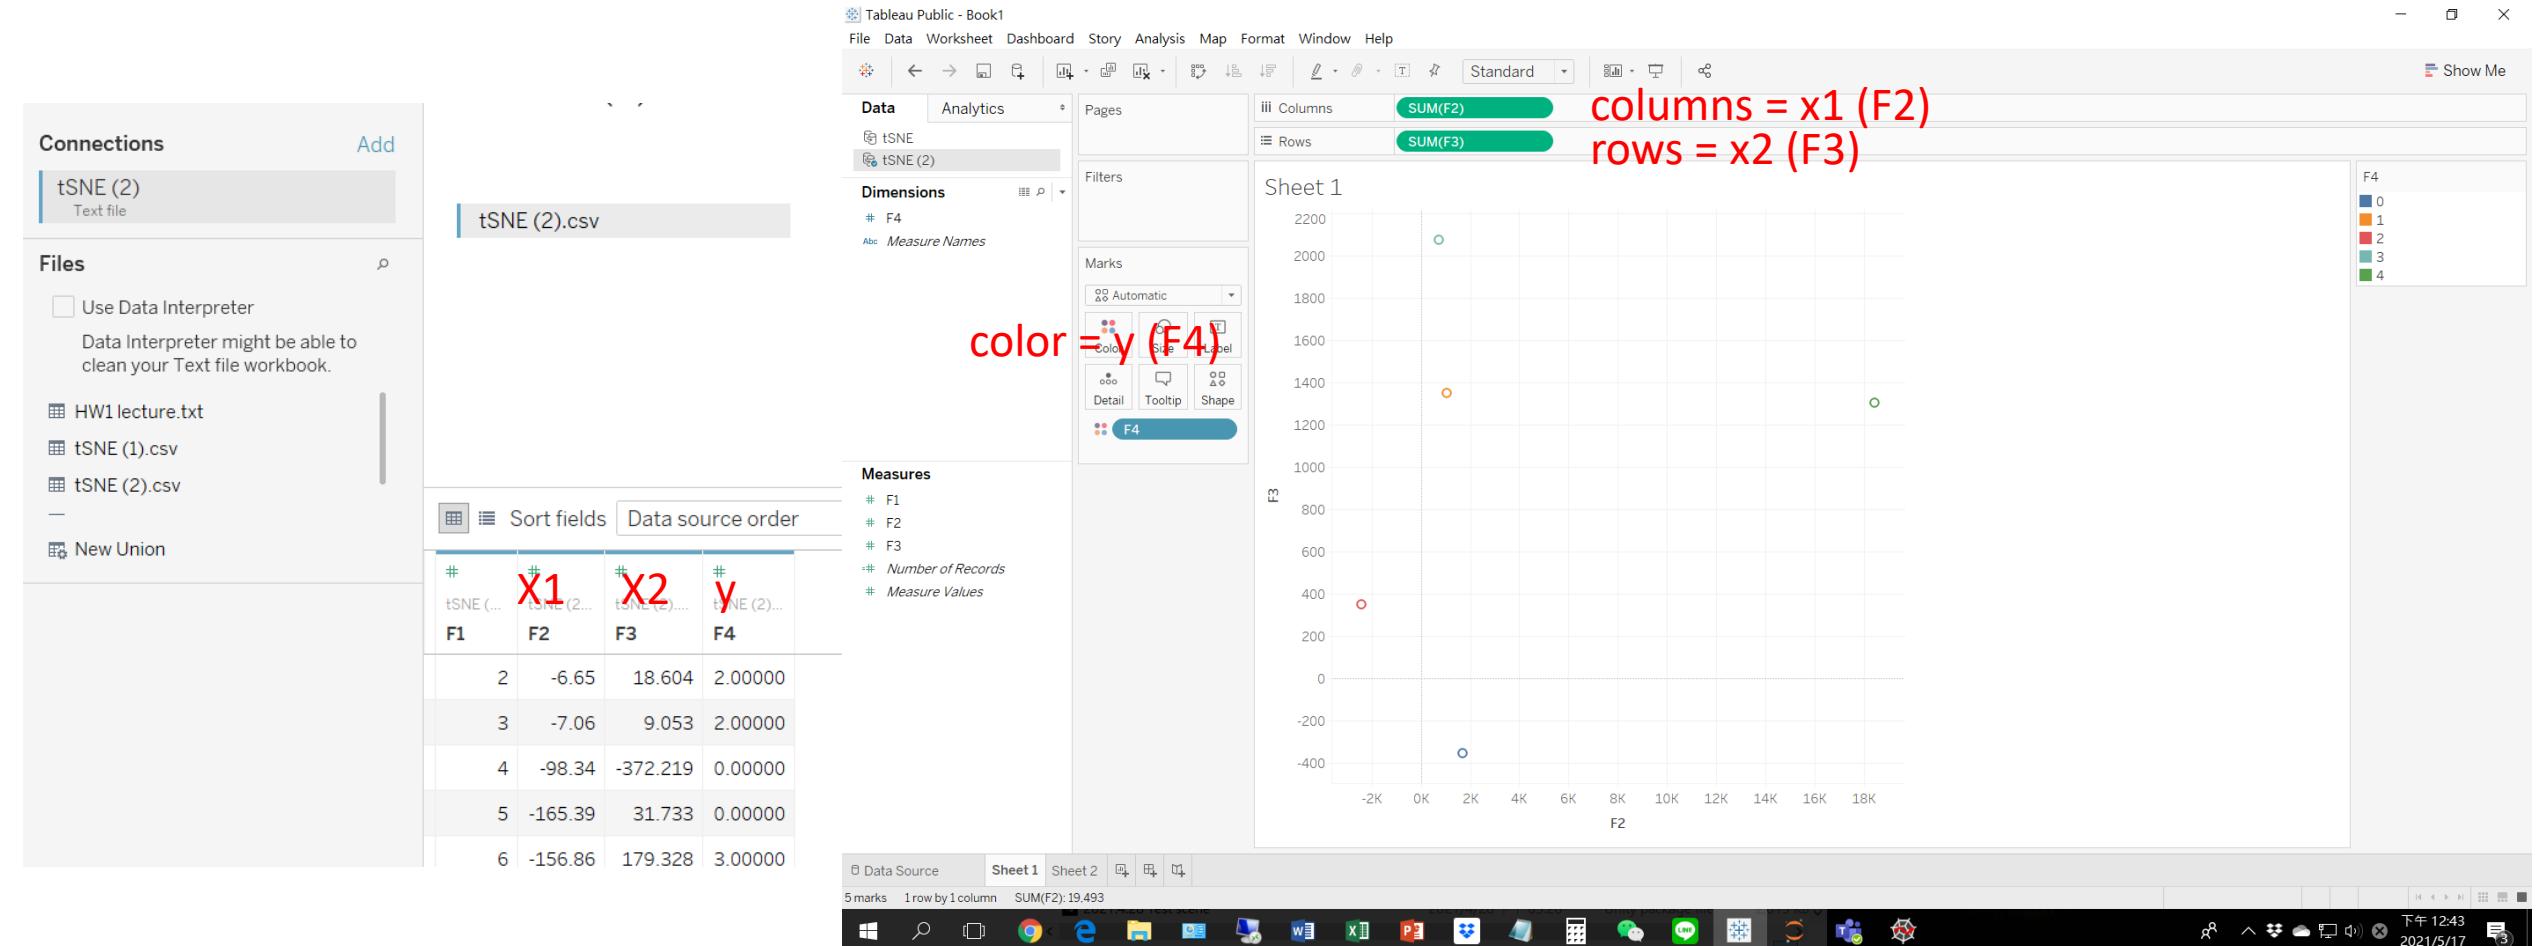
[43] arrayY1 = arrayY.reshape(arrayY.shape[0], 1)
print(arrayY1.shape)
(298, 1)

[44] XYArray = np.hstack((x, arrayY1))
print(XYArray.shape)
(298, 3)

[45] # Save data to excel for further Tableau visual
import pandas as pd
pd.DataFrame(XYArray).to_csv("tSNE.csv")
```

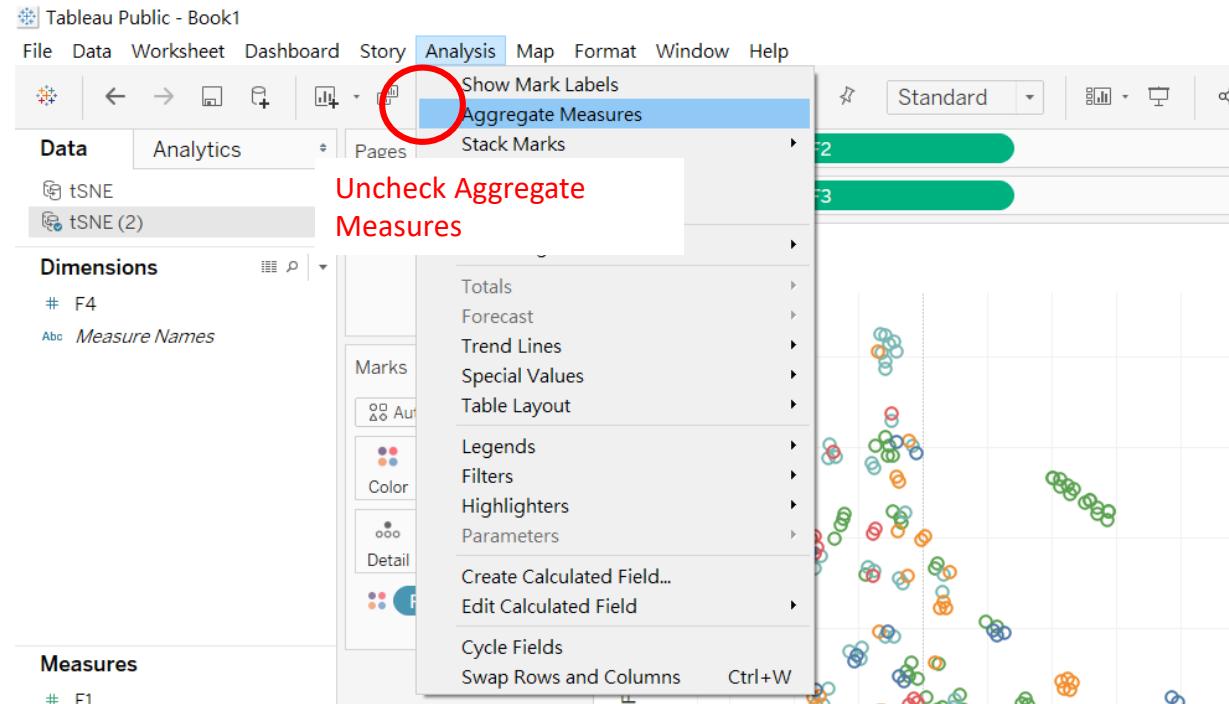


Visualize the downloaded file in Tableau public

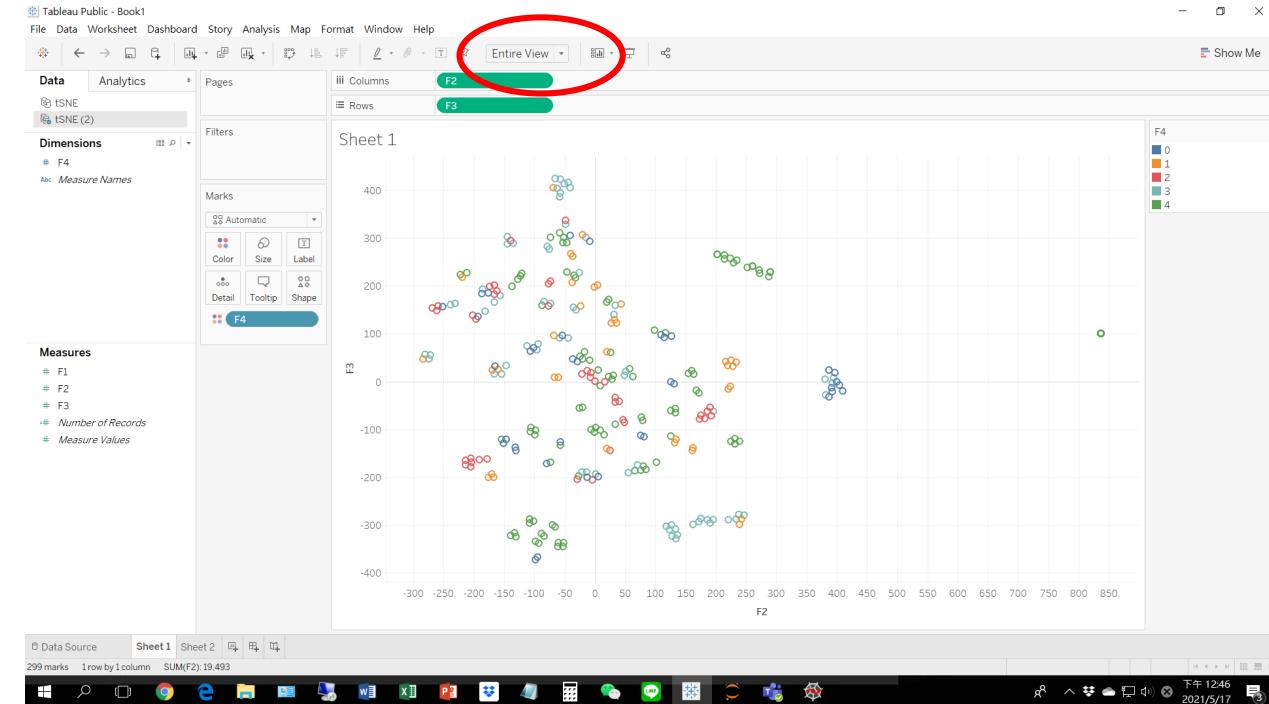


Visualize the downloaded file in Tableau public

disable Aggregate Measures



Select entire view

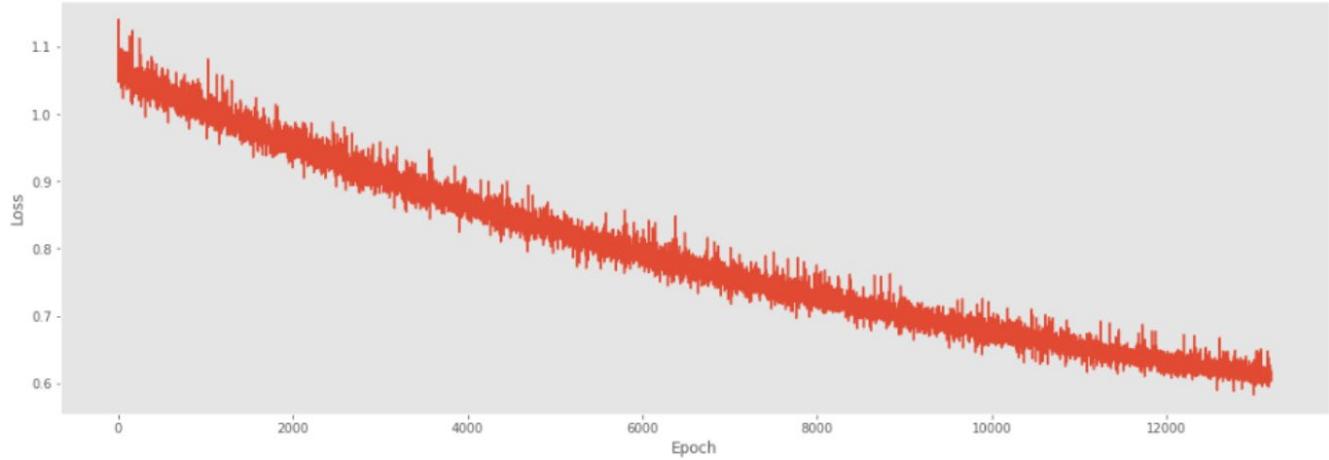


HW6 (1)

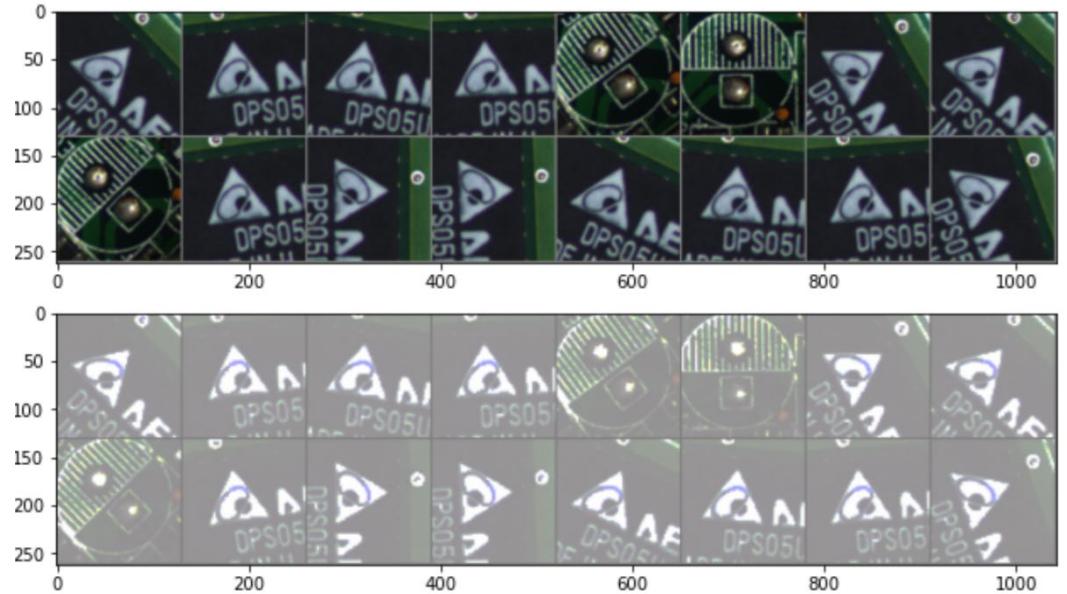
- Train an AE to learn a compact representation (try latent vector of size 20, 30, or 50) of your own images, e.g., facial expression. Test with 10 happy and 10 angry faces.
- Show the recovered image.
- Send the latent vectors to t -SNE to see whether they form clusters.



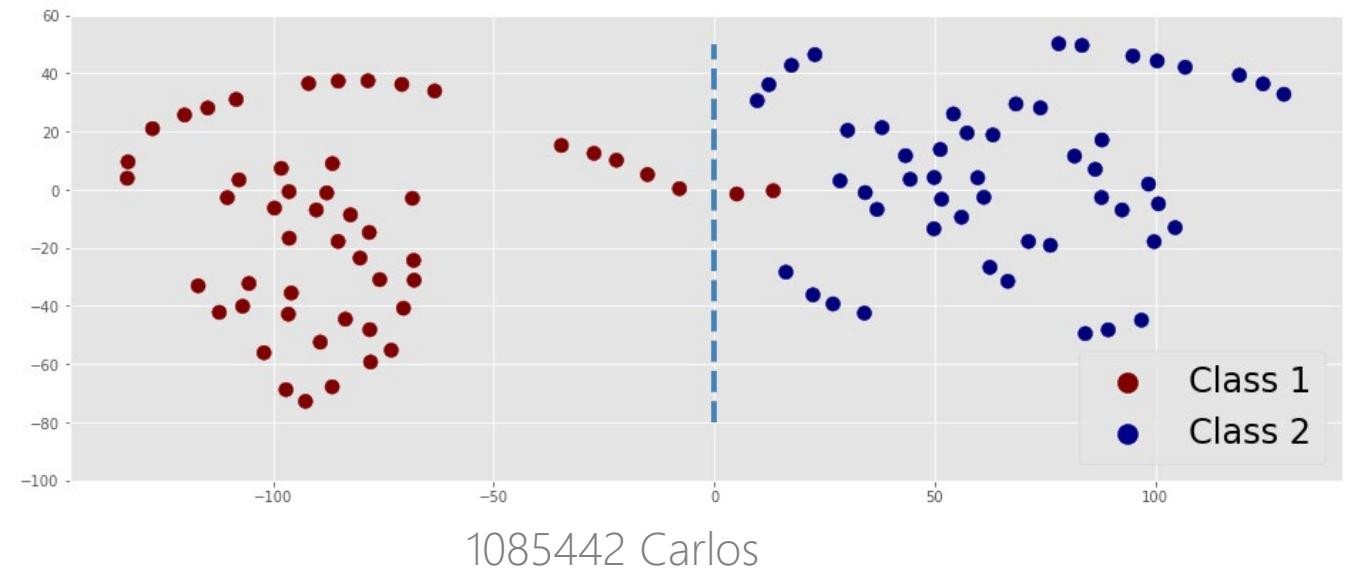
Class 1 = 100, Class 2 = 100, Latent vector size = 64, 2000 epochs, batch=32, Image size = 180x180x3



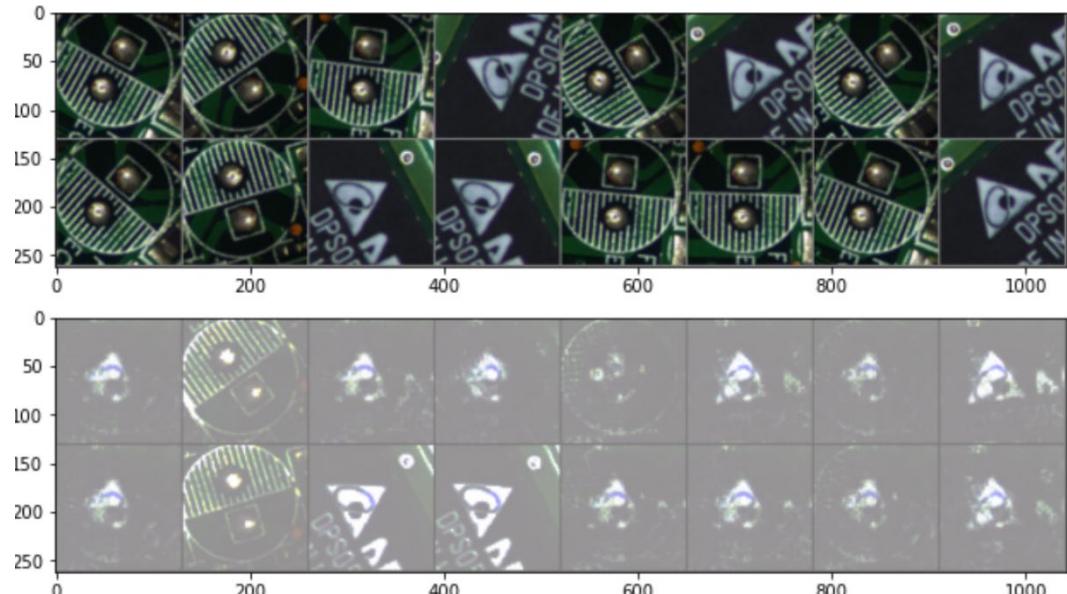
Recovered training images



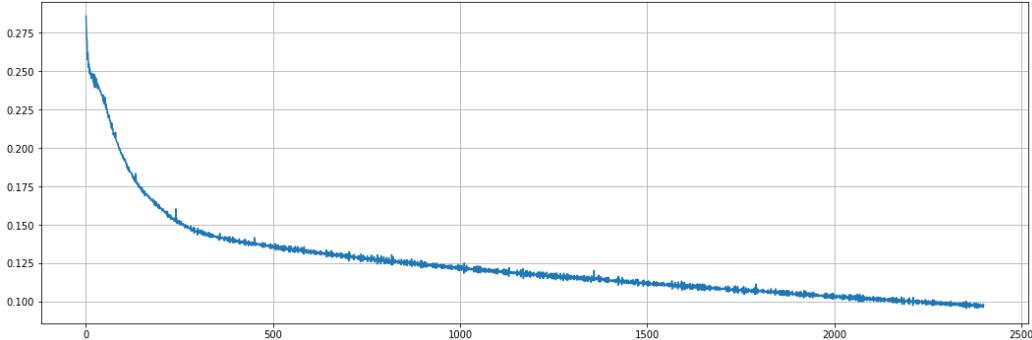
t-SNE (perplexity=10) results of training images



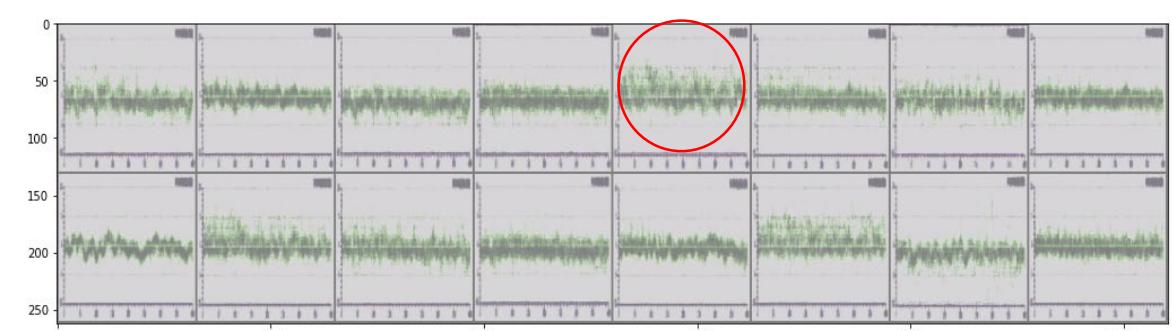
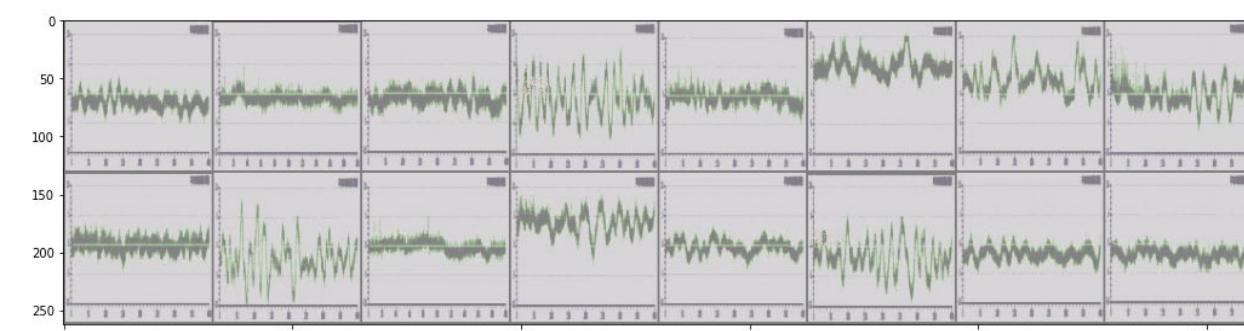
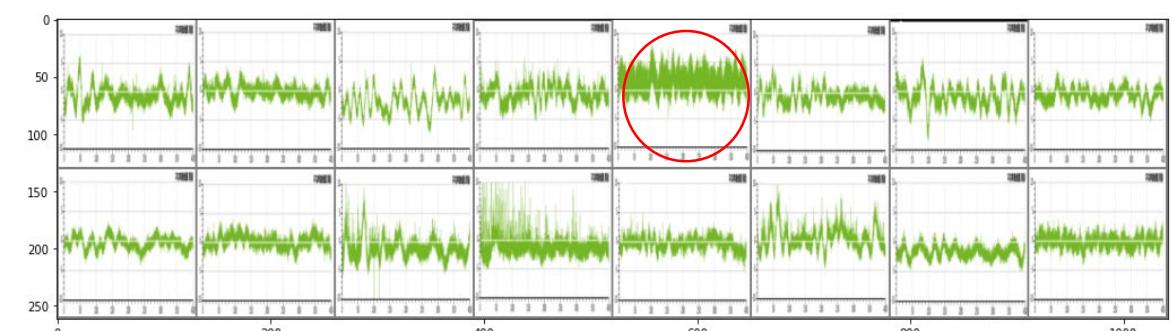
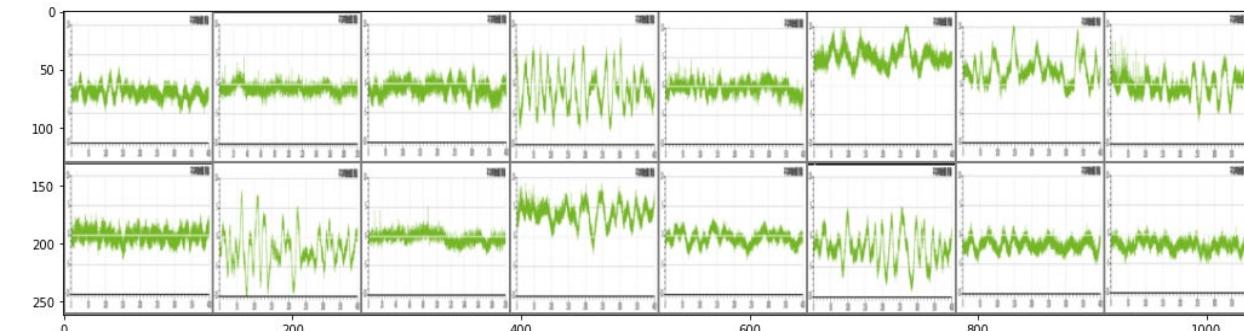
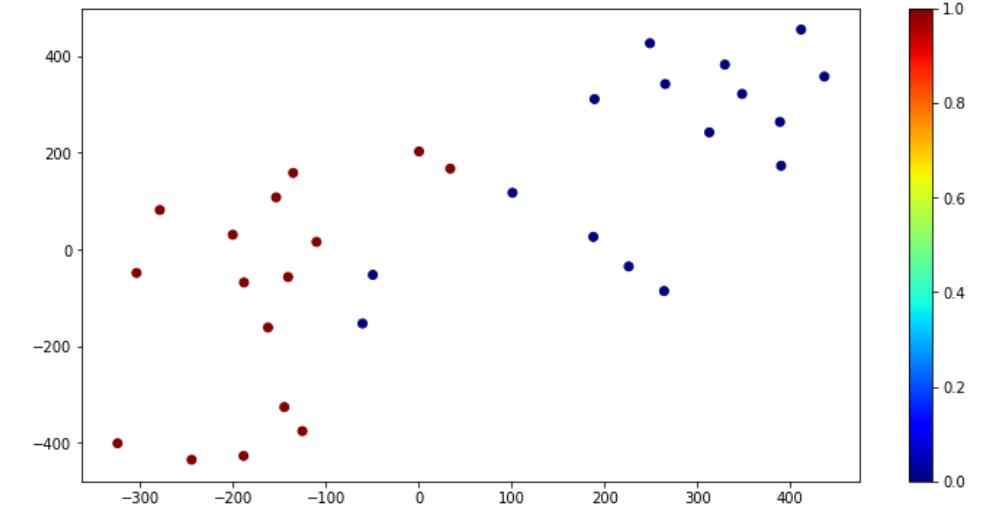
Recovered test images



Normal = 16, Abnormal = 16, Latent vector size =
32, 1200 epochs



t-SNE (perplexity=?) results of training images

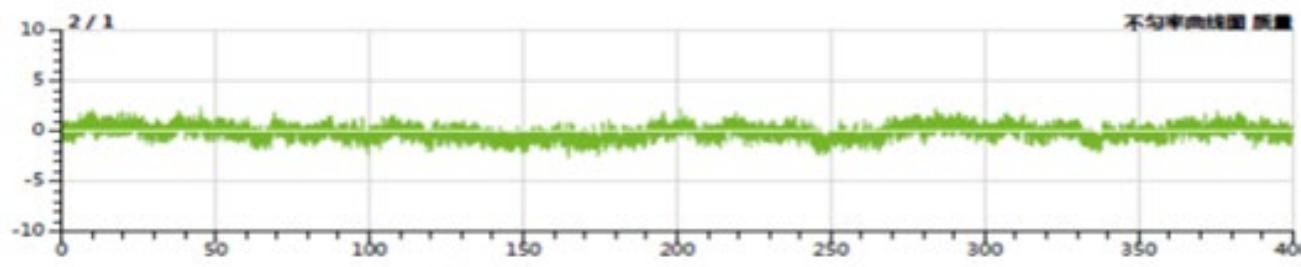
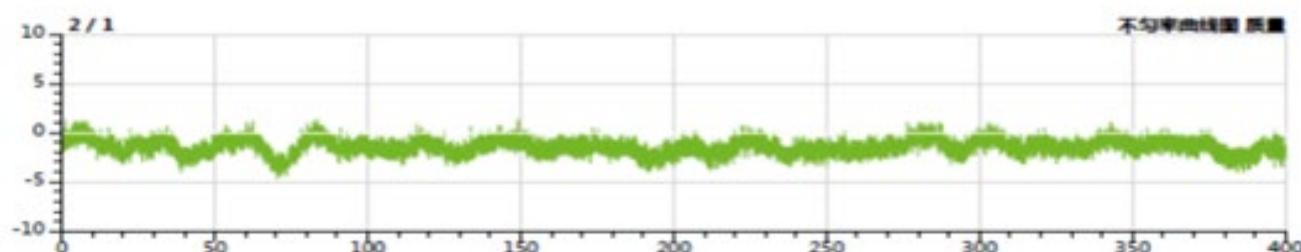
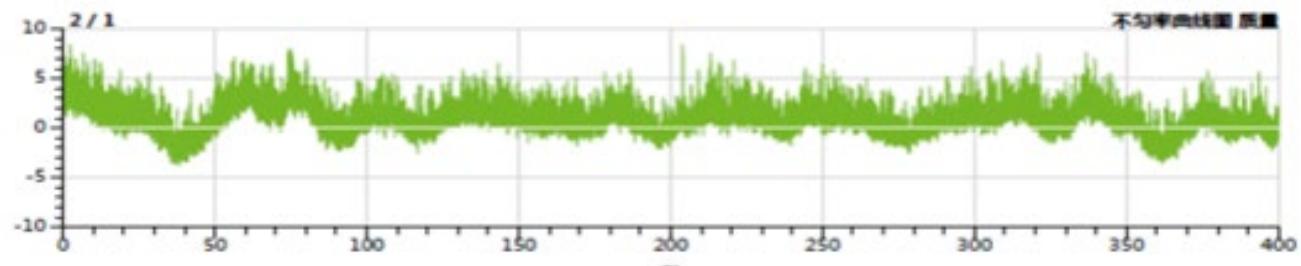
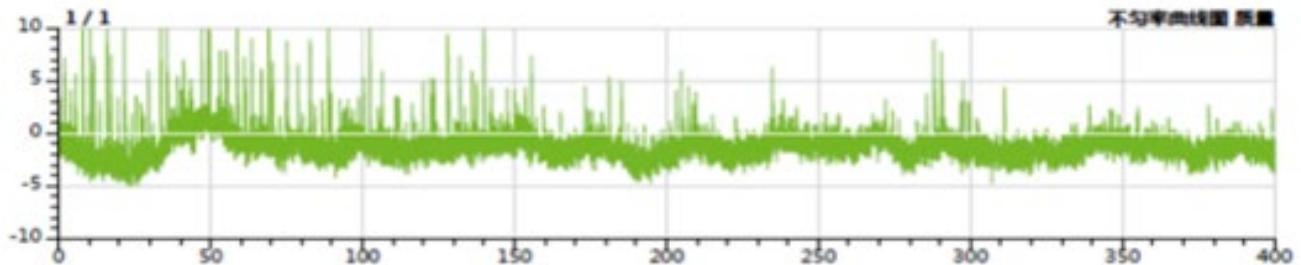


Recovered training images

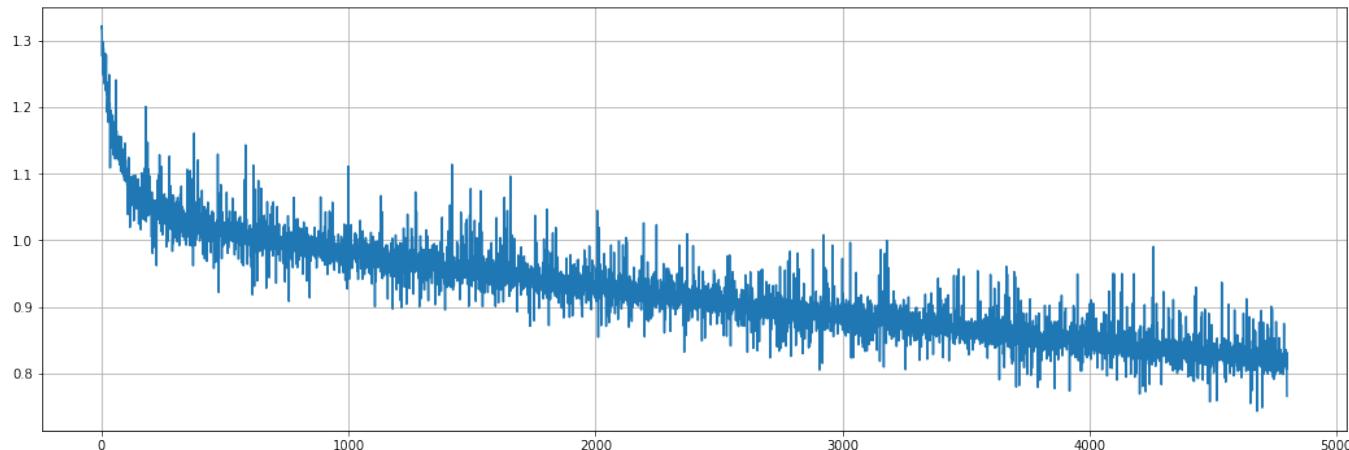
Recovered un-seen test images

1071229 詹心好

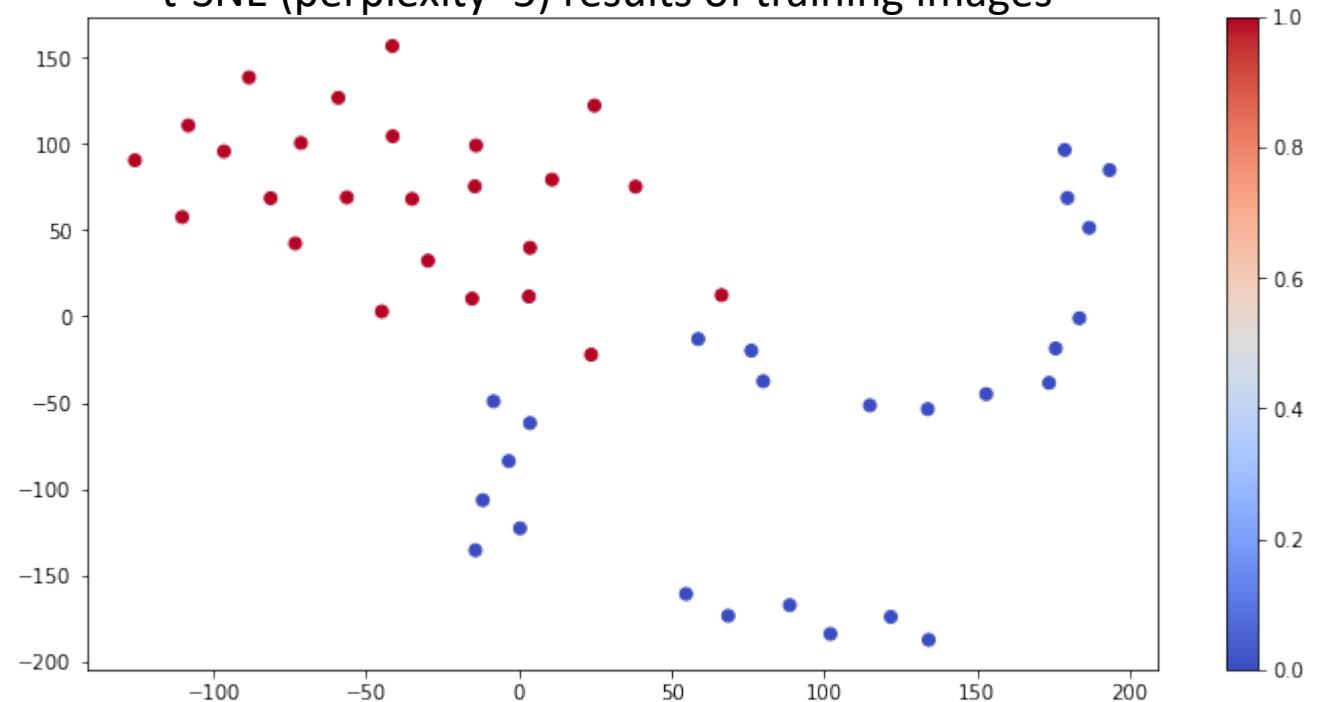
Image size = 128*128*3
Class1 = 16 , Class2 = 16
Latent vector size = 128
Batch = 128
epoch : 20000



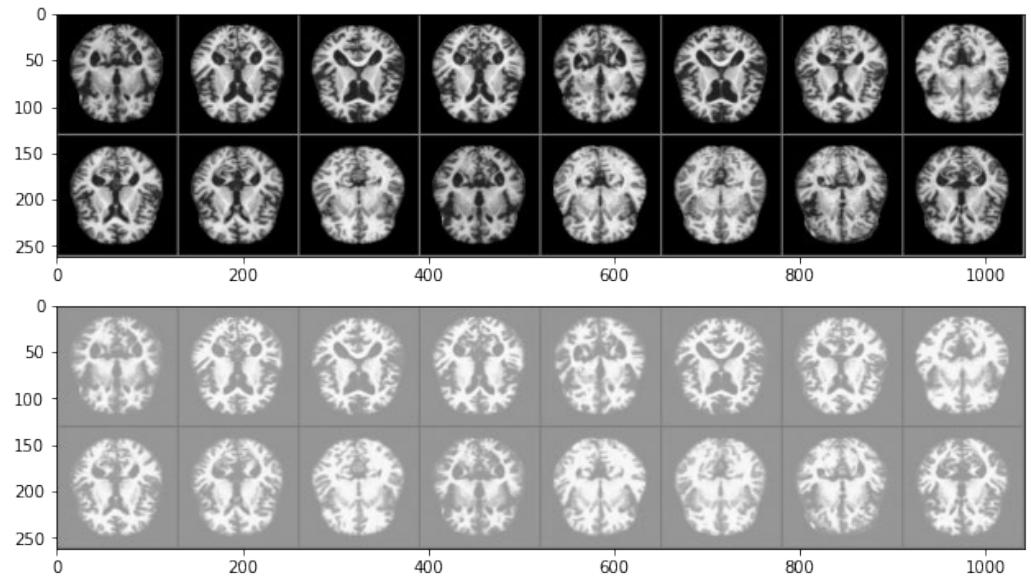
Dementia = 25, health = 25, Latent vector size = 64, 1200 epochs, batch= , Image size = 128x128x3



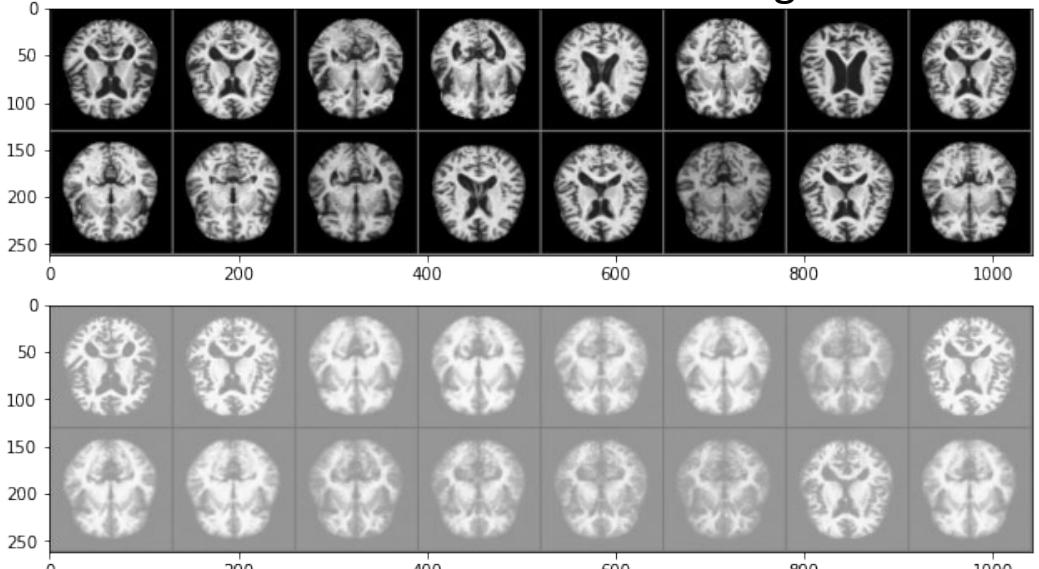
t-SNE (perplexity=5) results of training images



Recovered training images



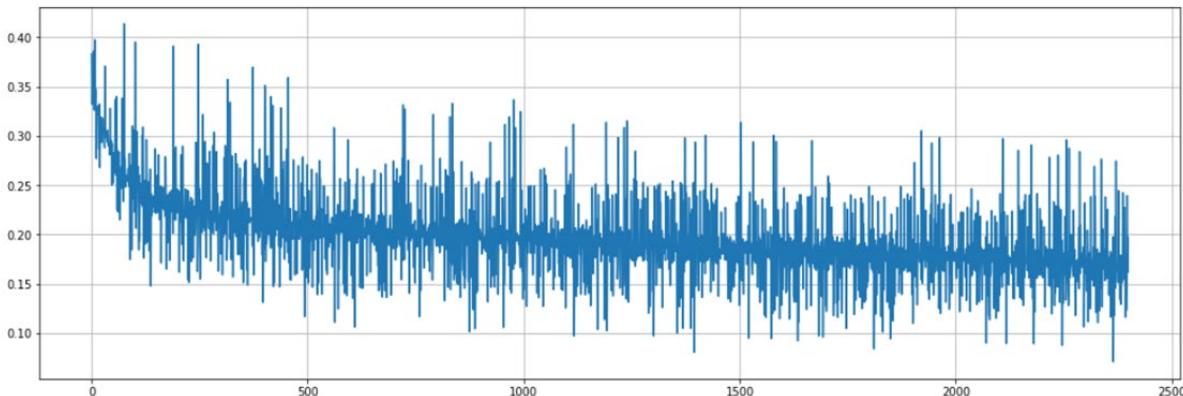
Recovered un-seen test images



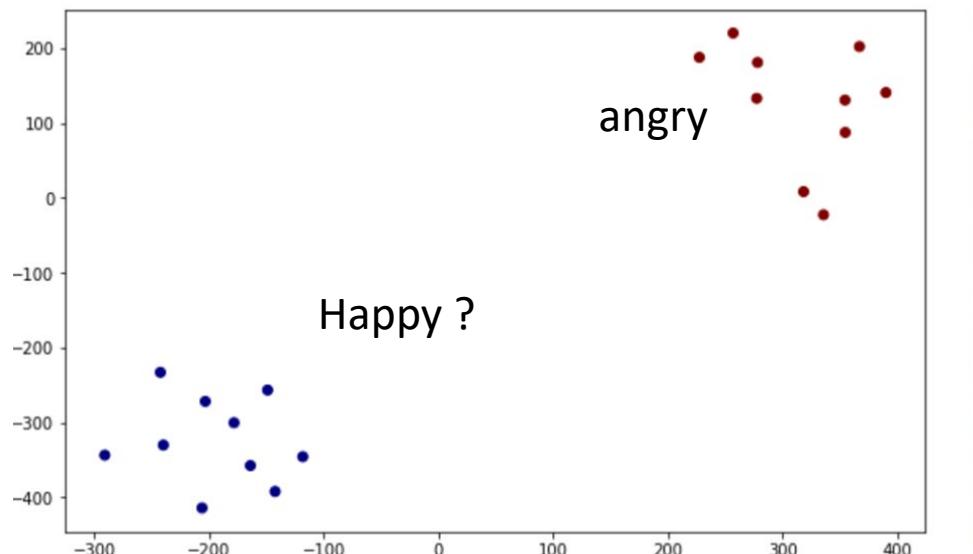
1095434 Sara

27

Happy = 10, Angry = 10, Latent vector size = 64, 1160 epochs,
batch=16, Image size = 128x128x3

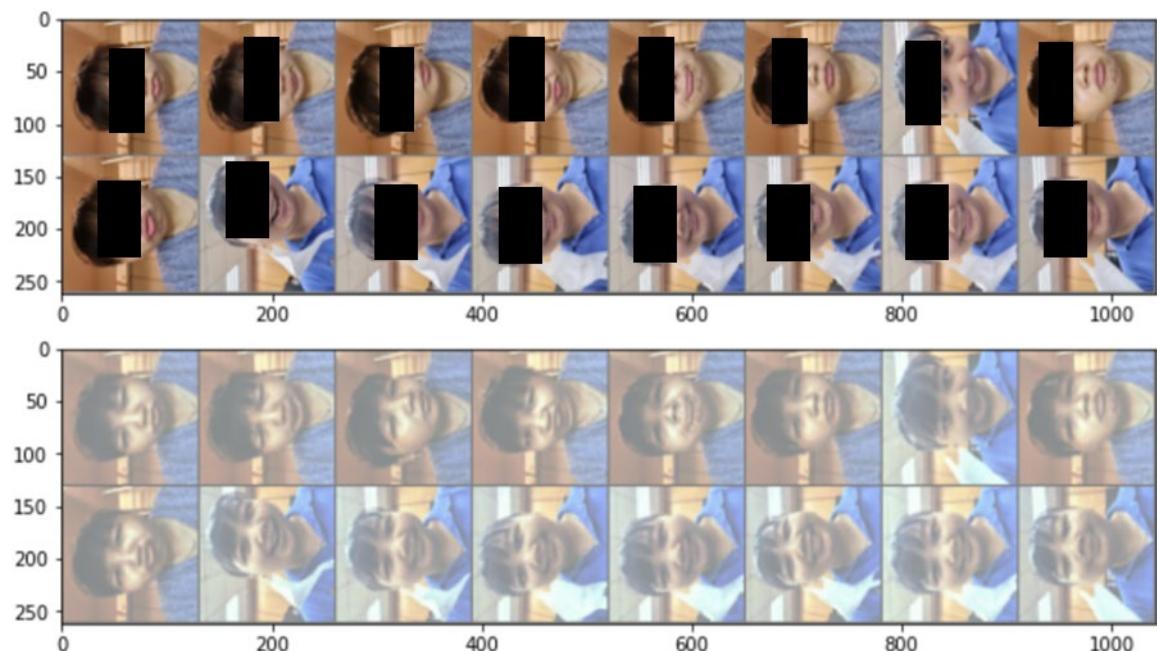


t-SNE (perplexity=5) results of training images

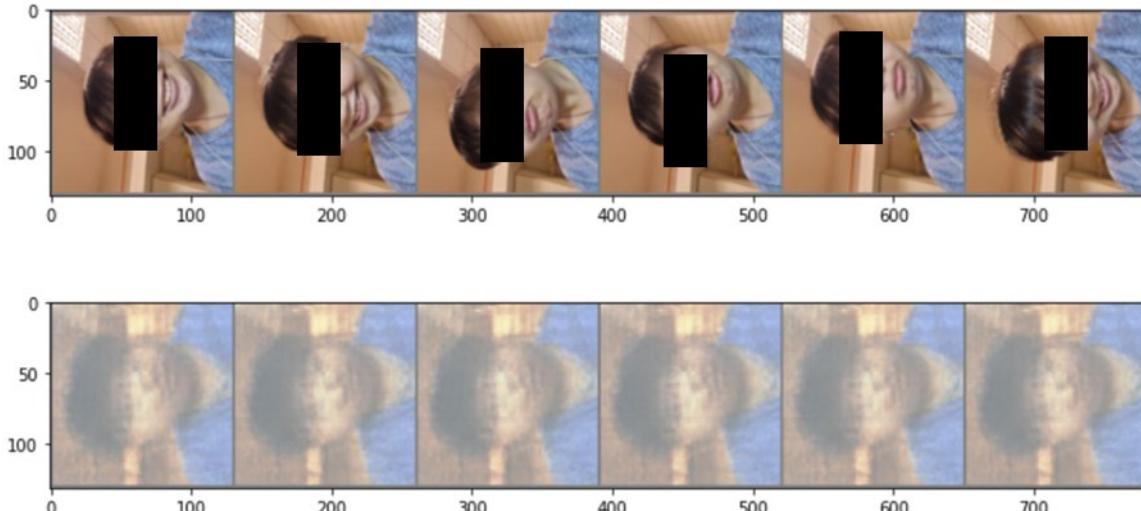


1061259 Keren

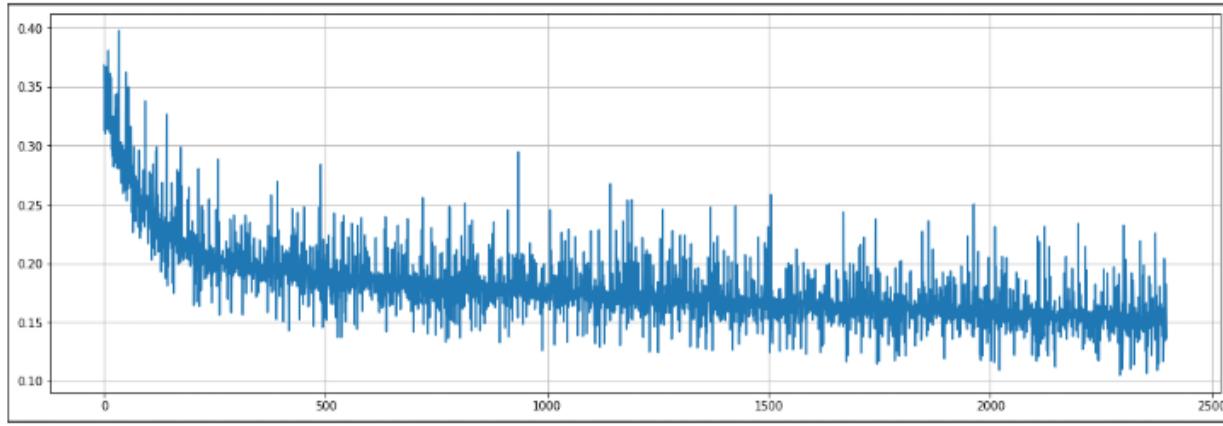
Recovered training images



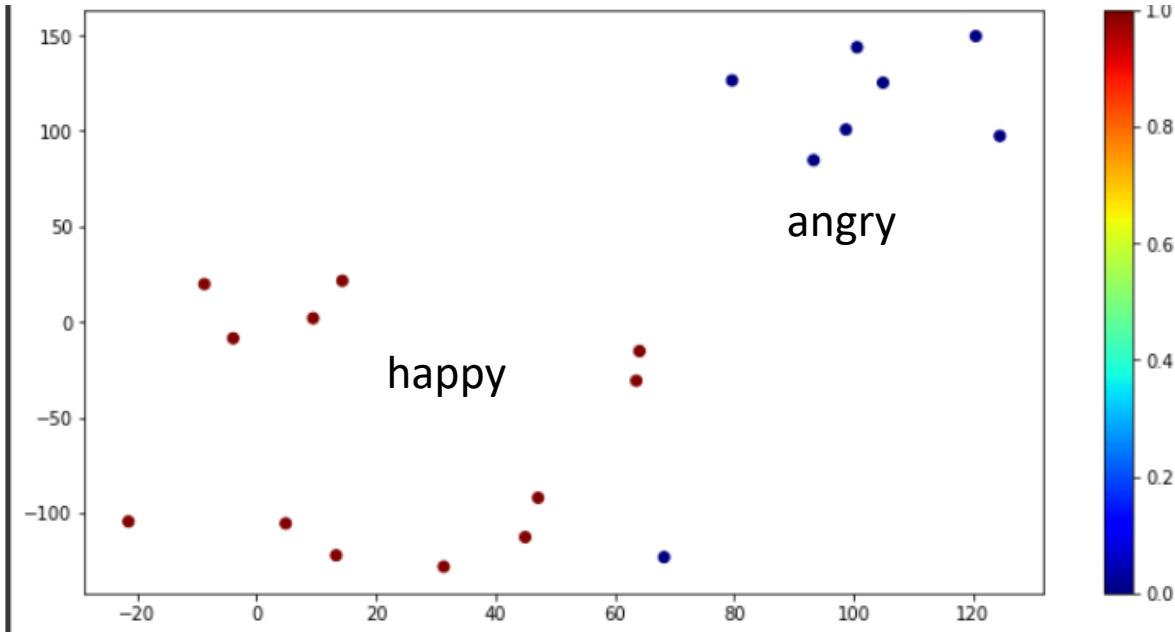
Recovered un-seen test images



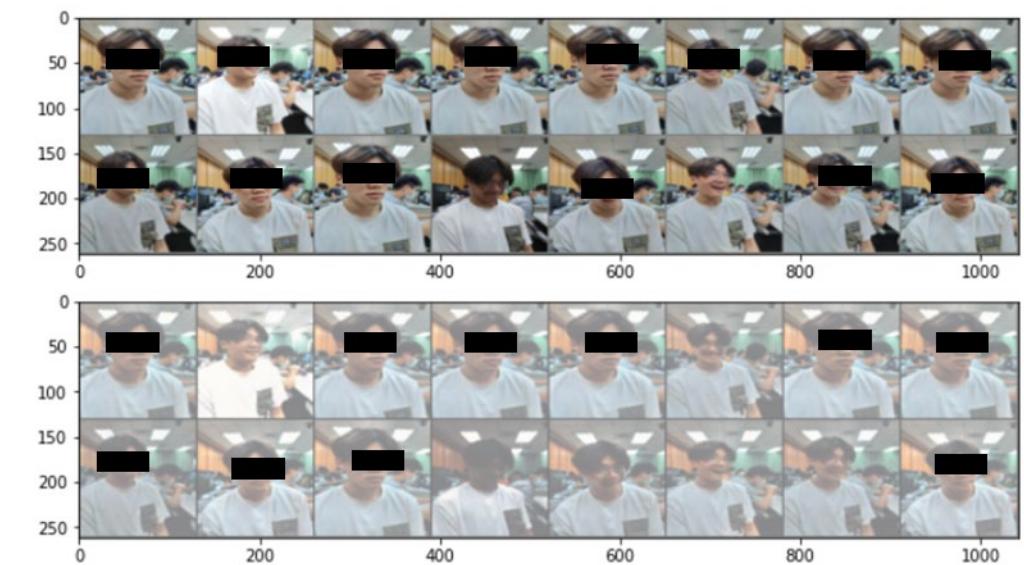
Happy = 12, Angry = 8, Latent vector size = 20, 1200 epochs,
batch=?, Image size = ?x?x3



t-SNE (perplexity=5) results of training images



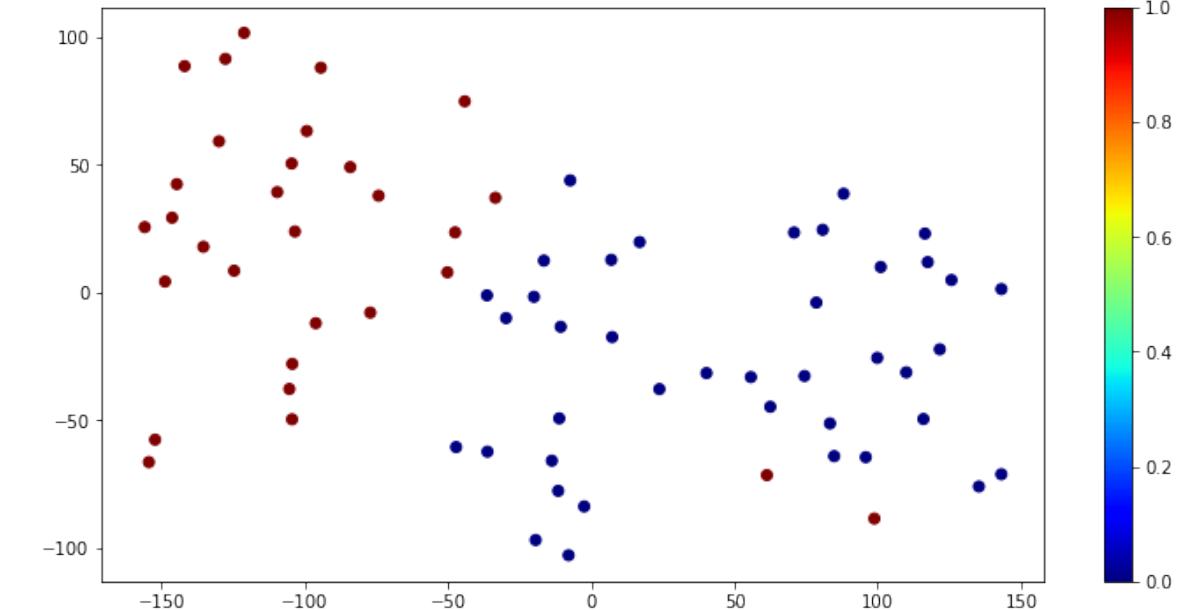
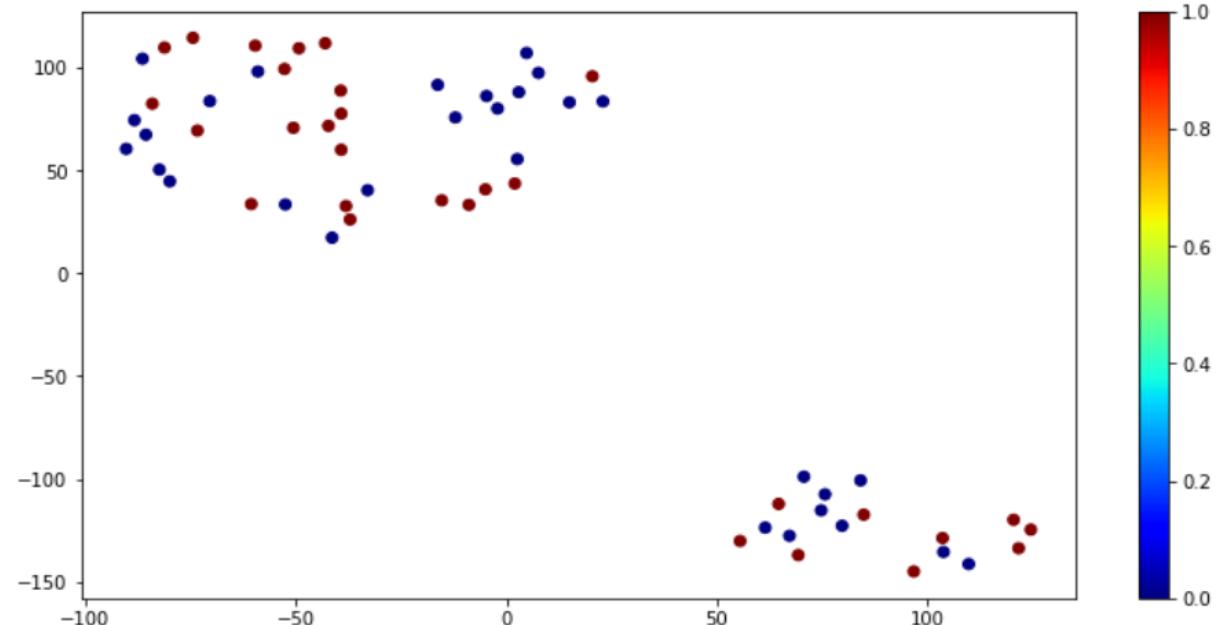
Recovered training images



Recovered un-seen test images



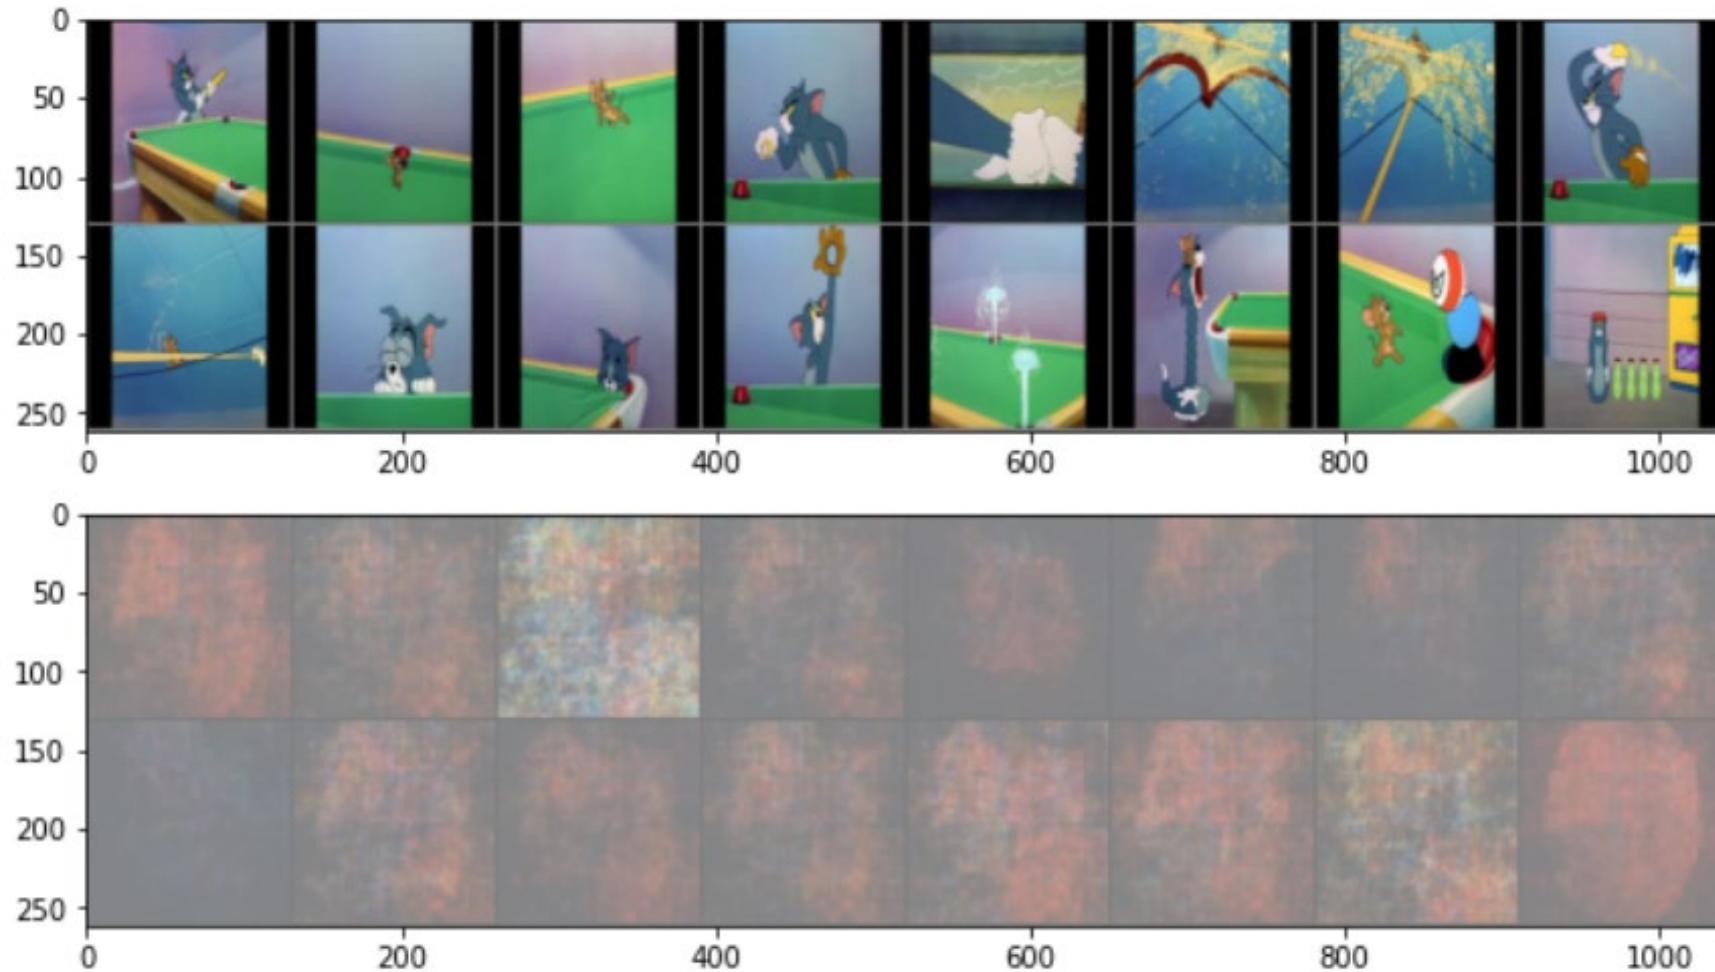
AE's classification idea is different from human's



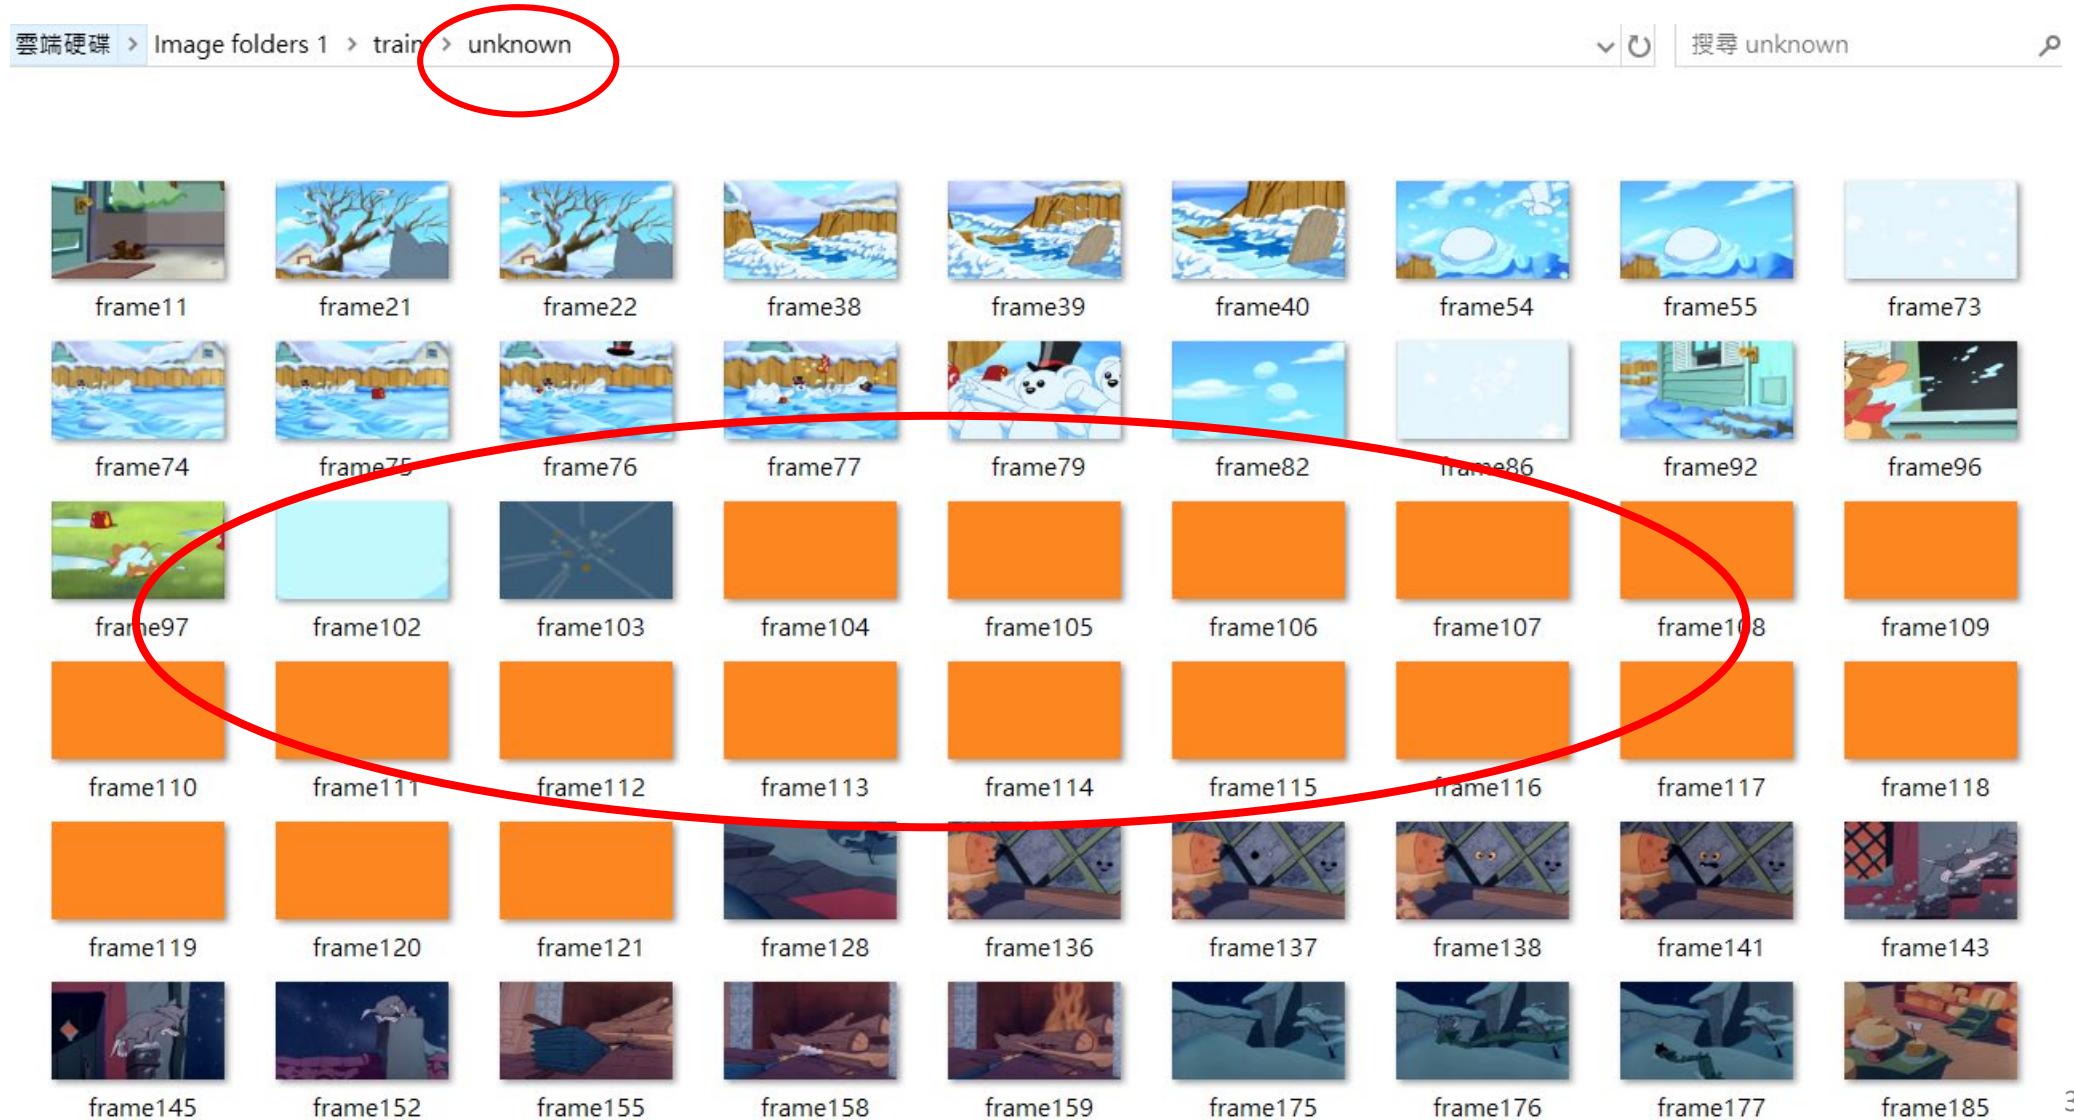
More on using AE on Tom & Jerry images

Results are still not good after 1200 epochs

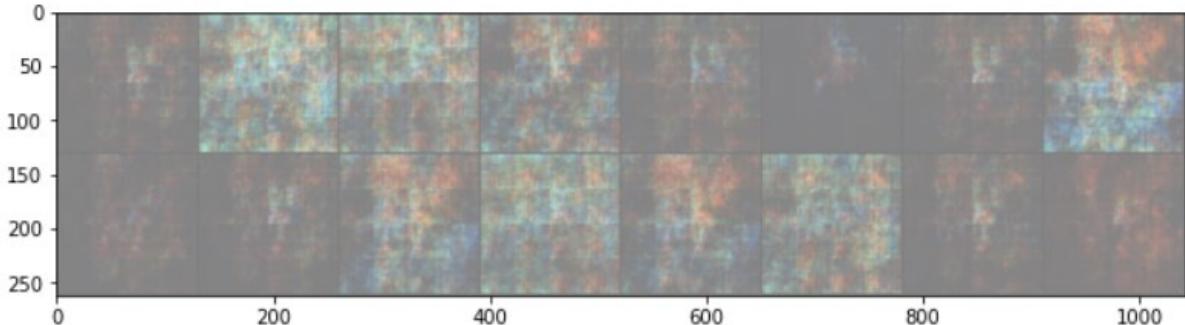
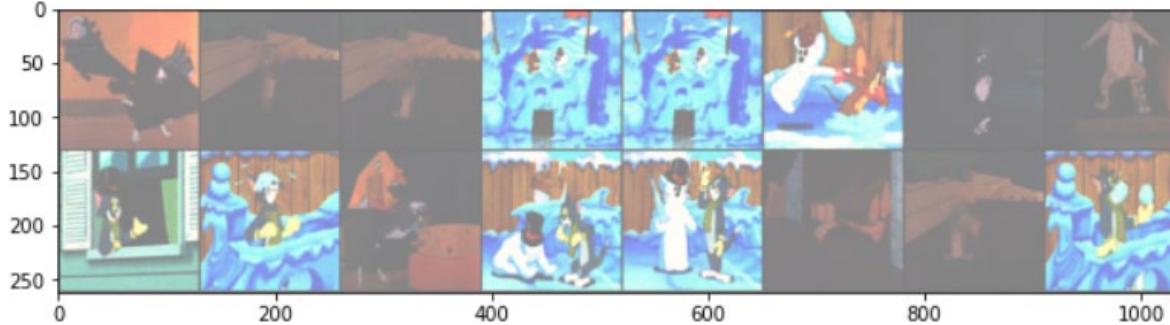
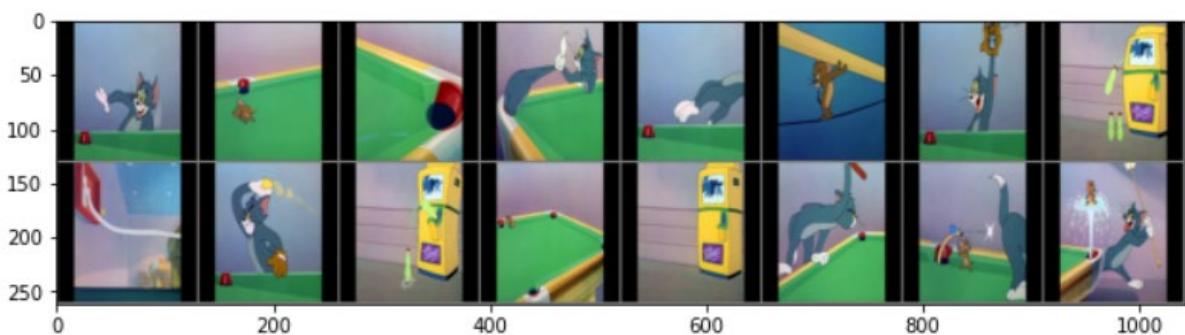
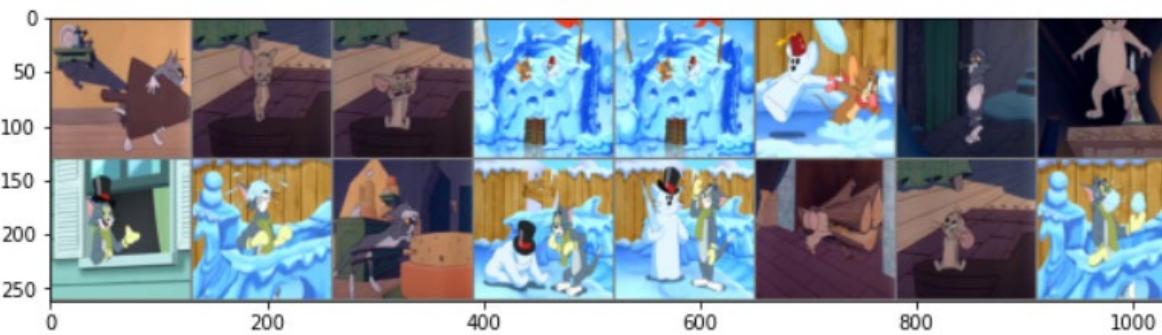
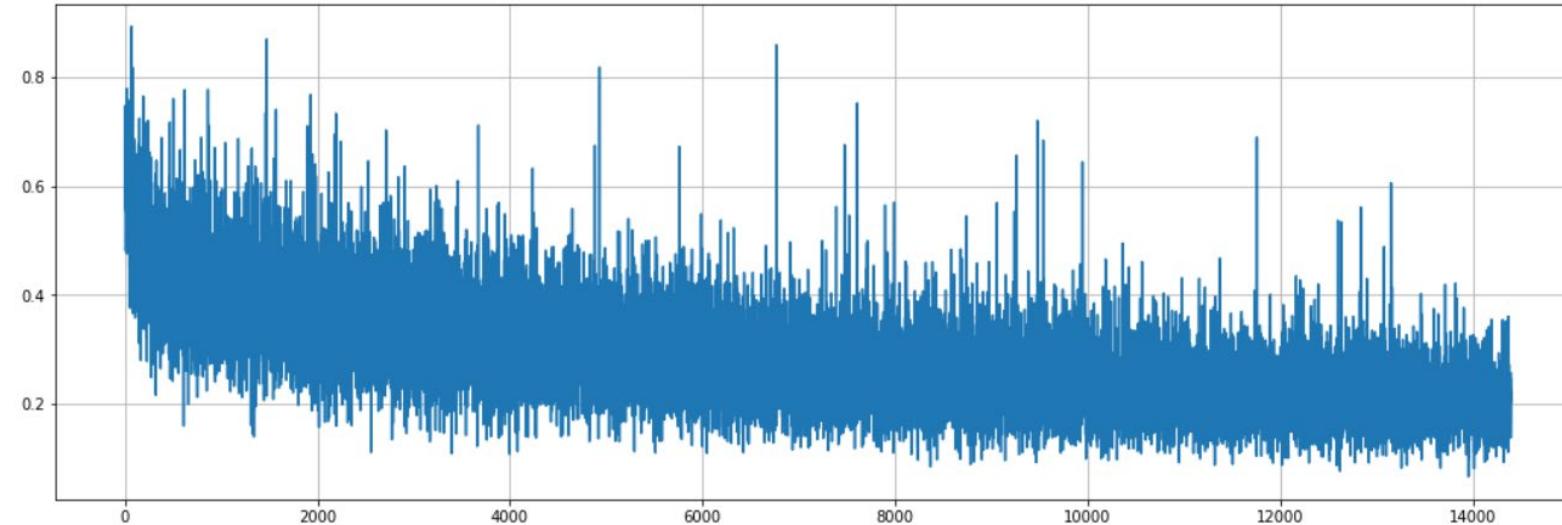
Test on un-seen images – fails to reconstruct the input images



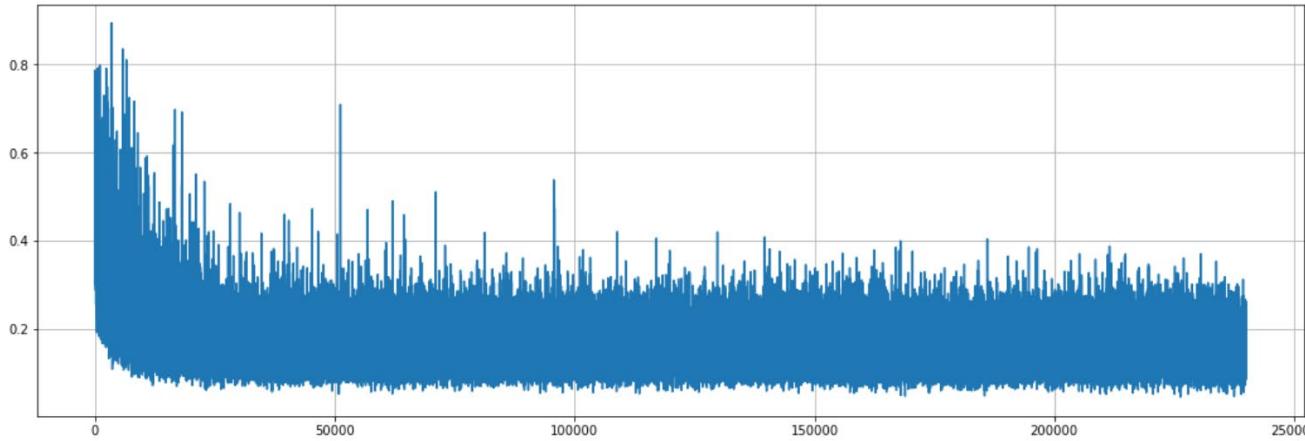
Remove the "unknown" folder?



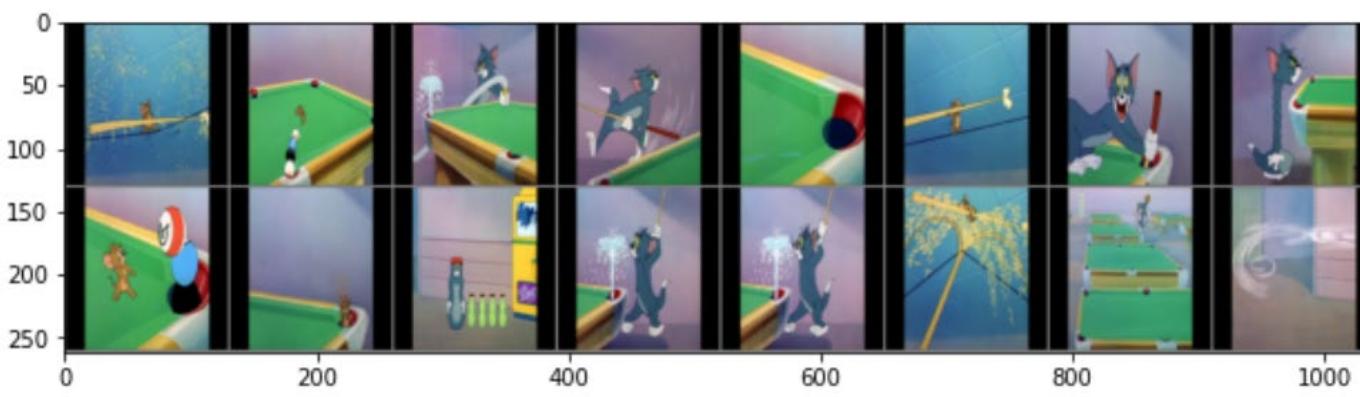
Train 1200 epochs after removing the "unknown" folder



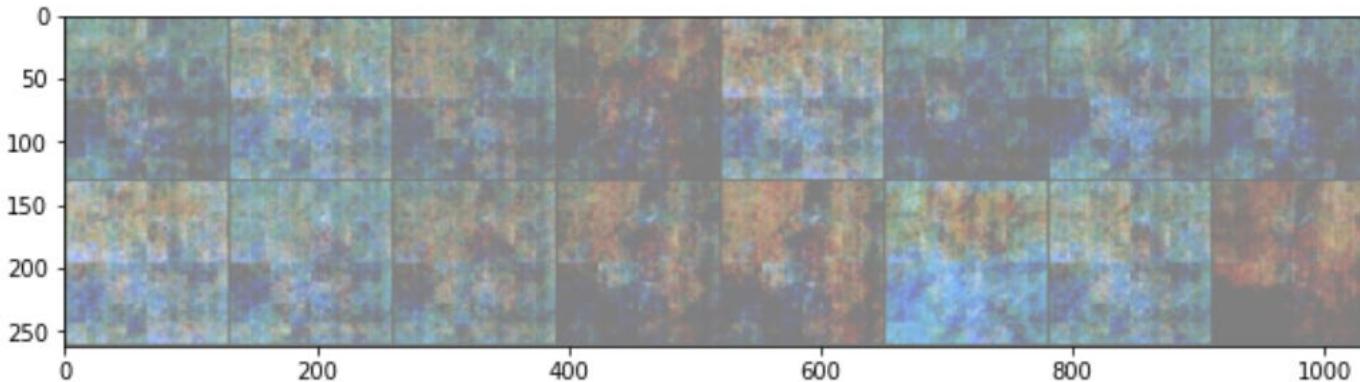
Train 20,000 epochs



1071226 陳仁慶



Still failed to recover
un-seen images

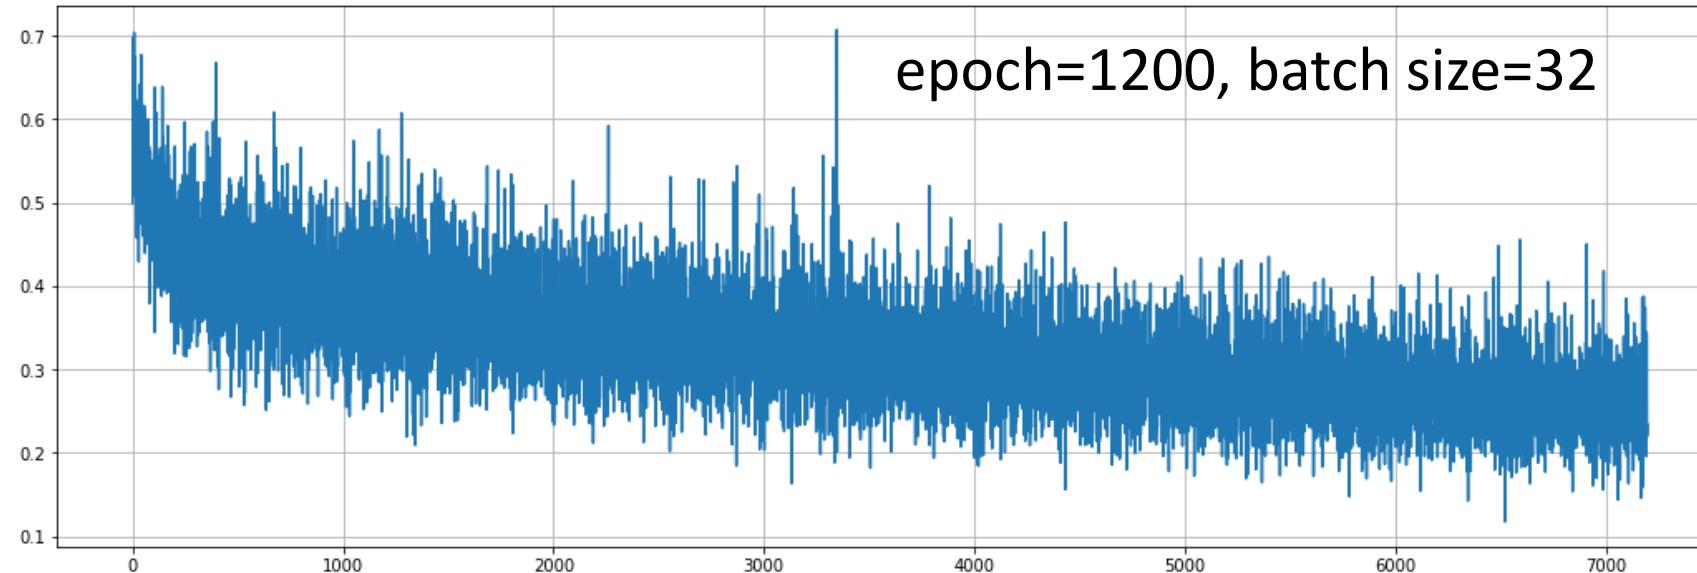
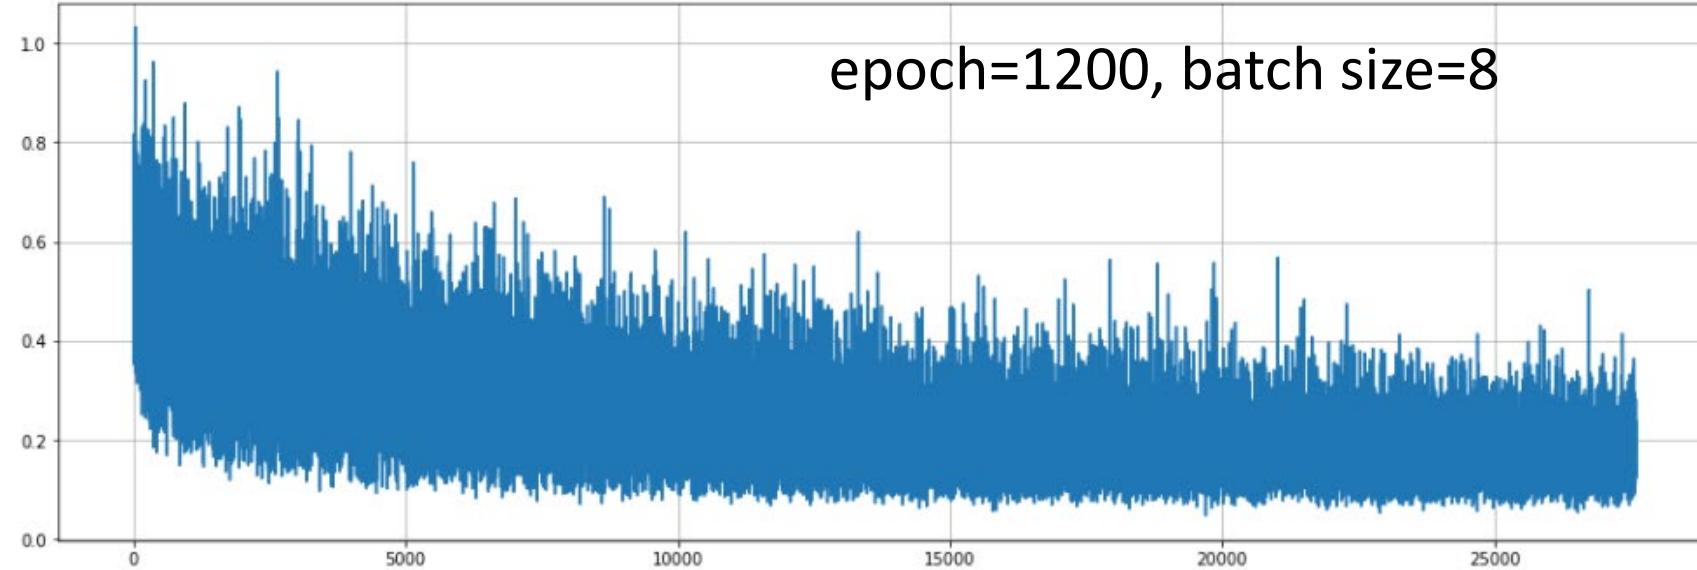


How about adjusting batch size?

```
self.encoder = nn.Sequential(  
    nn.Conv2d(3, 32, kernel_size=2, stride=2),  
    nn.BatchNorm2d(32, eps=1e-05, momentum=0.1, af  
    nn.ReLU(),  
    nn.Conv2d(32, 64, kernel_size=2, stride=2),  
    nn.BatchNorm2d(64, eps=1e-05, momentum=0.1, af  
    nn.ReLU(),  
    nn.Conv2d(64, 128, kernel_size=2, stride=2),  
    nn.BatchNorm2d(128, eps=1e-05, momentum=0.1, a  
    nn.ReLU(),  
    nn.Conv2d(128, 256, kernel_size=2, stride=2),  
    nn.BatchNorm2d(256, eps=1e-05, momentum=0.1, a  
    nn.ReLU(),  
    nn.Conv2d(256, 512, kernel_size=2, stride=2),  
    nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, a  
    nn.ReLU(),  
    nn.Conv2d(512, 1024, kernel_size=2, stride=2),  
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1,  
    nn.ReLU(),  
    nn.Conv2d(1024, 1024, kernel_size=2, stride=2)  
    nn.BatchNorm2d(1024, eps=1e-05, momentum=0.1,  
    nn.ReLU(),  
    Flatten(),  
    nn.Linear(in_features=i, out_features=o),  
)
```

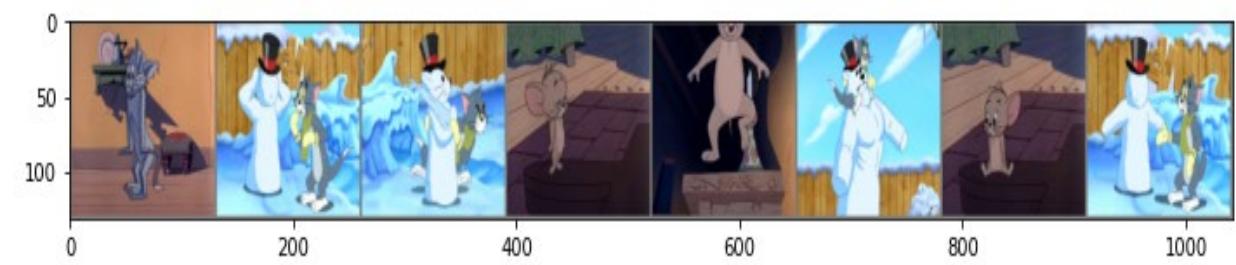
[12]:

```
import torch.utils.data as Data  
loader = Data.DataLoader(  
    dataset=train_dataset,  
    batch_size=16,  
    shuffle=True)
```

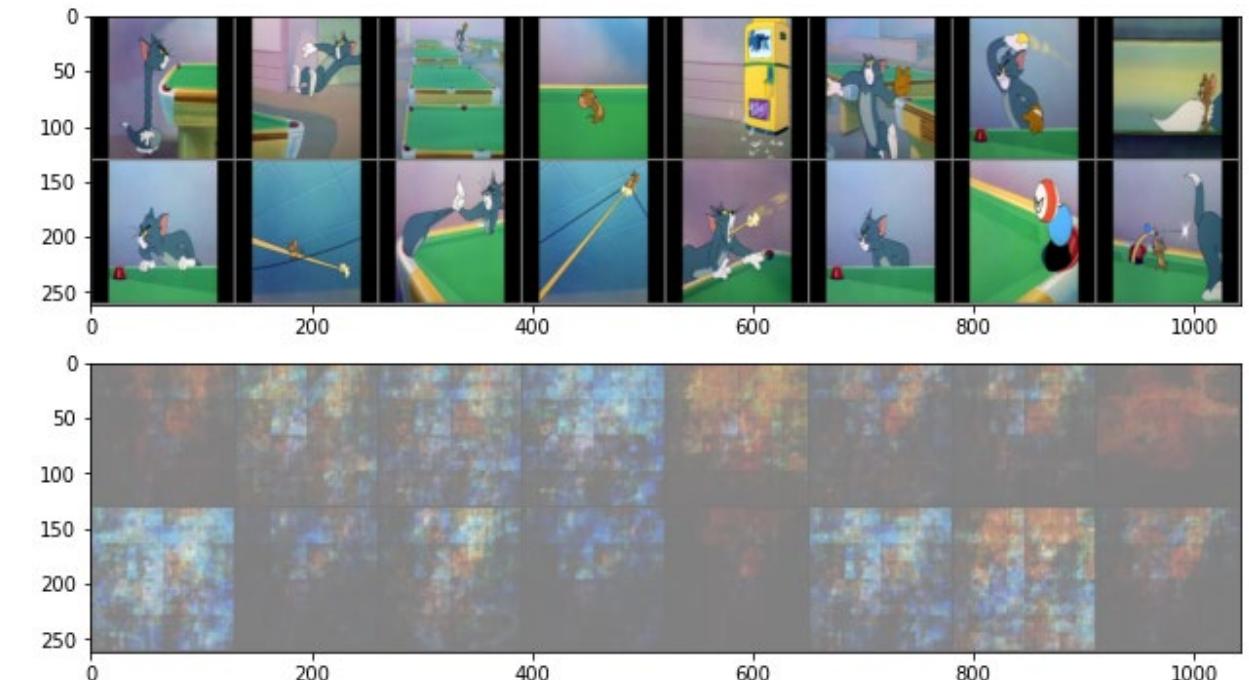
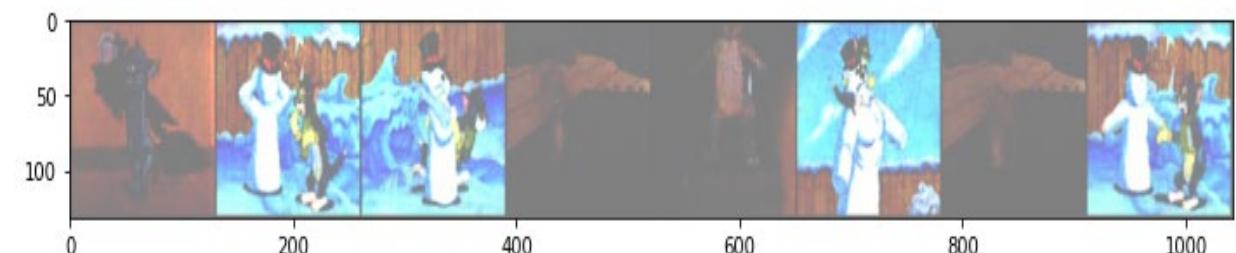
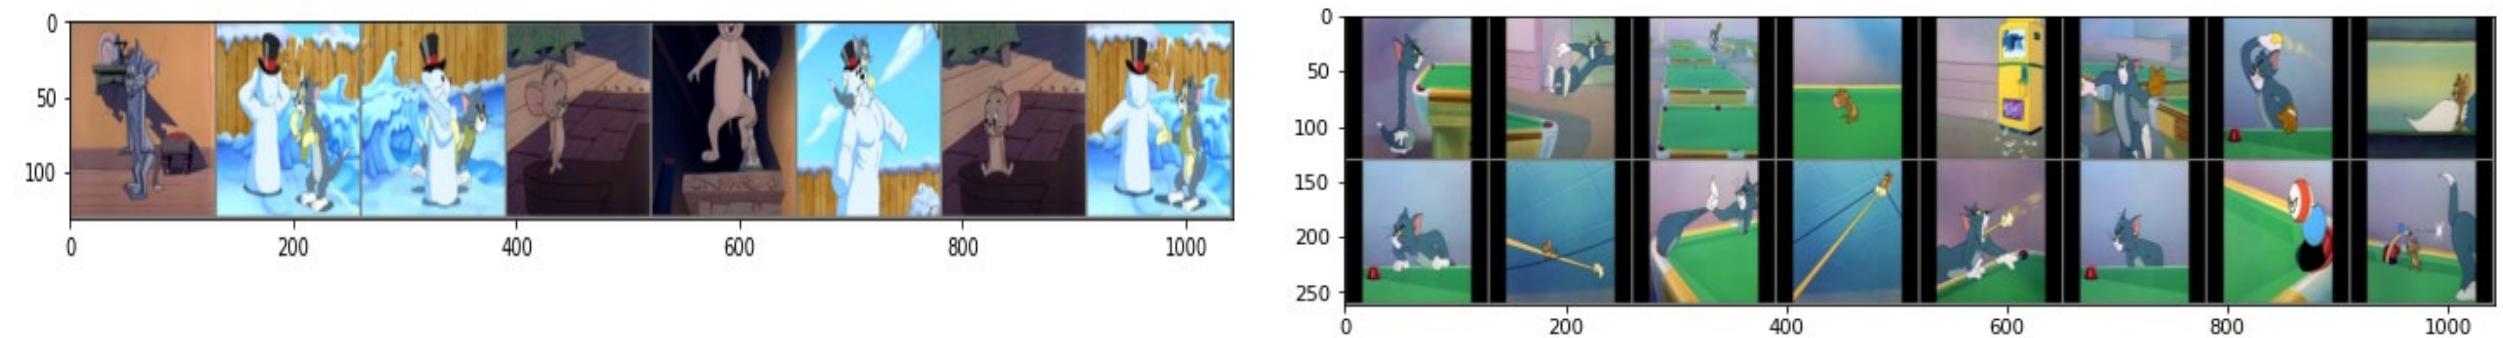


epoch=1200, batch size=8

Train:

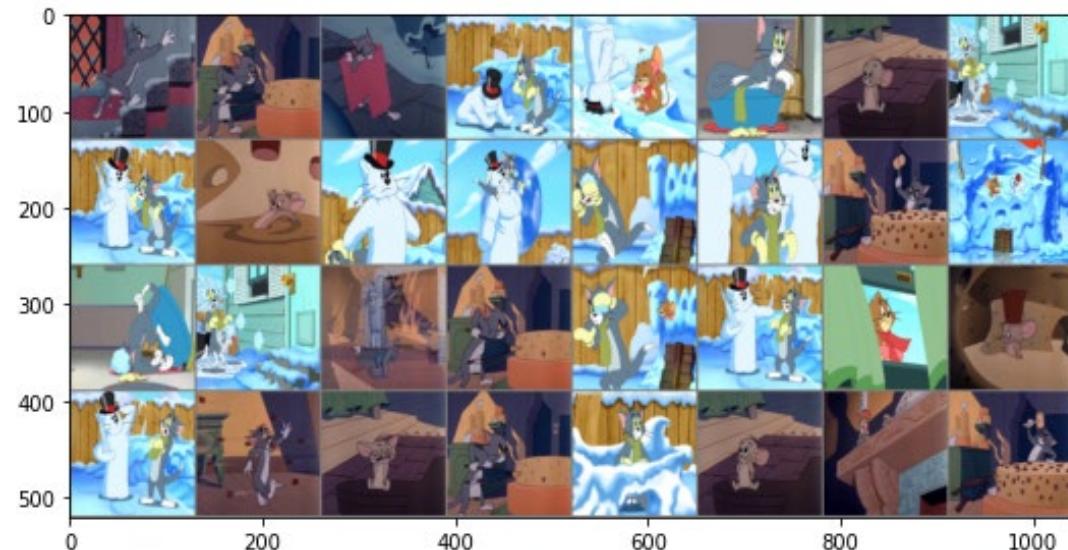


Test:



epoch=1200, batch size=32

Train:



Test:

