


Vibrational Auto-Encoder (VAE)

Why VAE?

Assume 1-d code




Practice


- Run "7.2.Conv_VAE.ipynb"



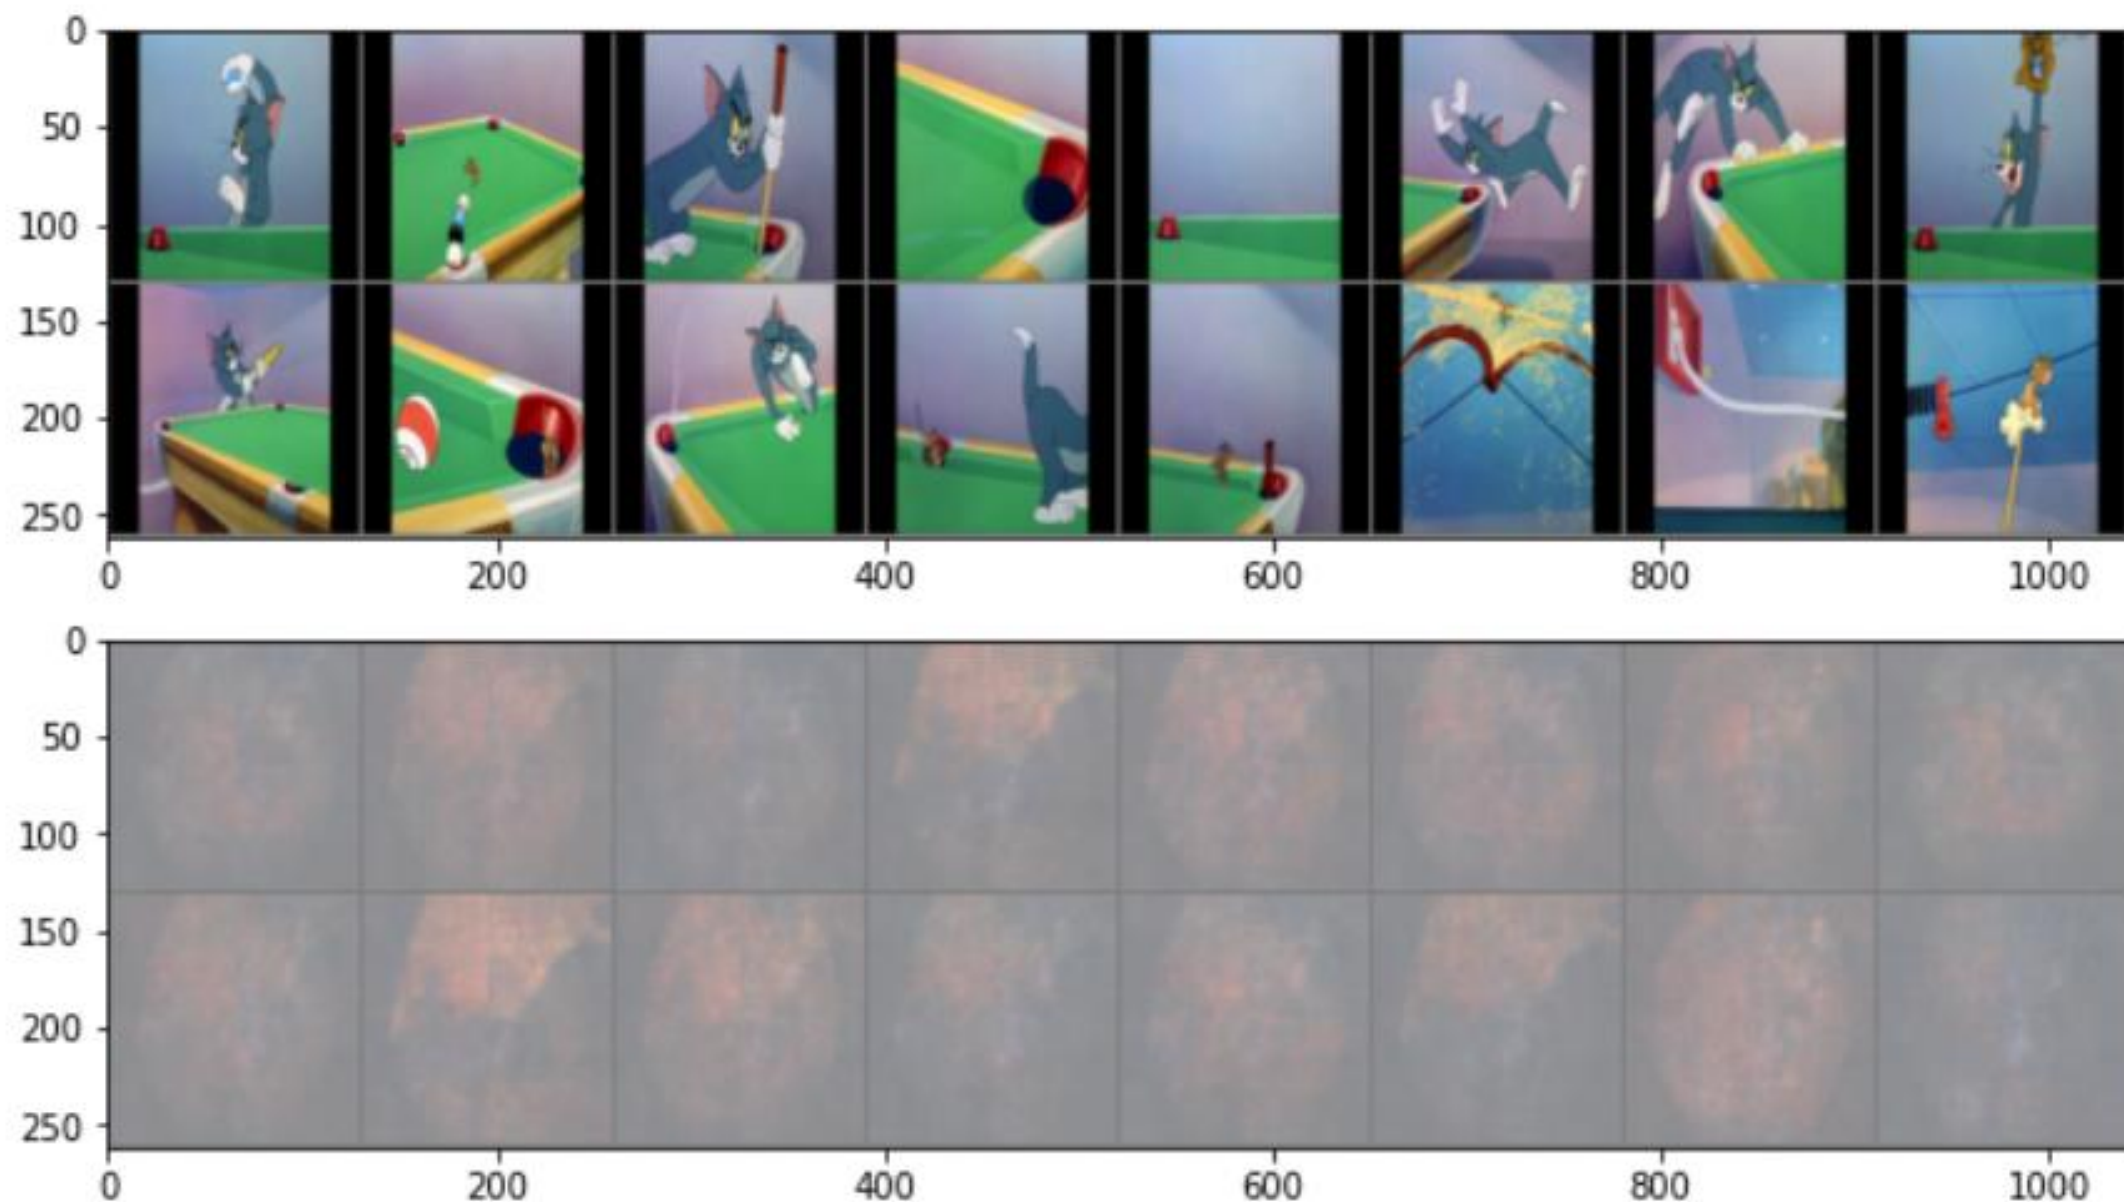
Loss plot of epoch 0-400




Training images recovered after training for 400 epochs




NN fails to recover un-seen test images if trained for 400 epochs




Loss plot of epoch 400-800




Training images recovered after training for 800 epochs



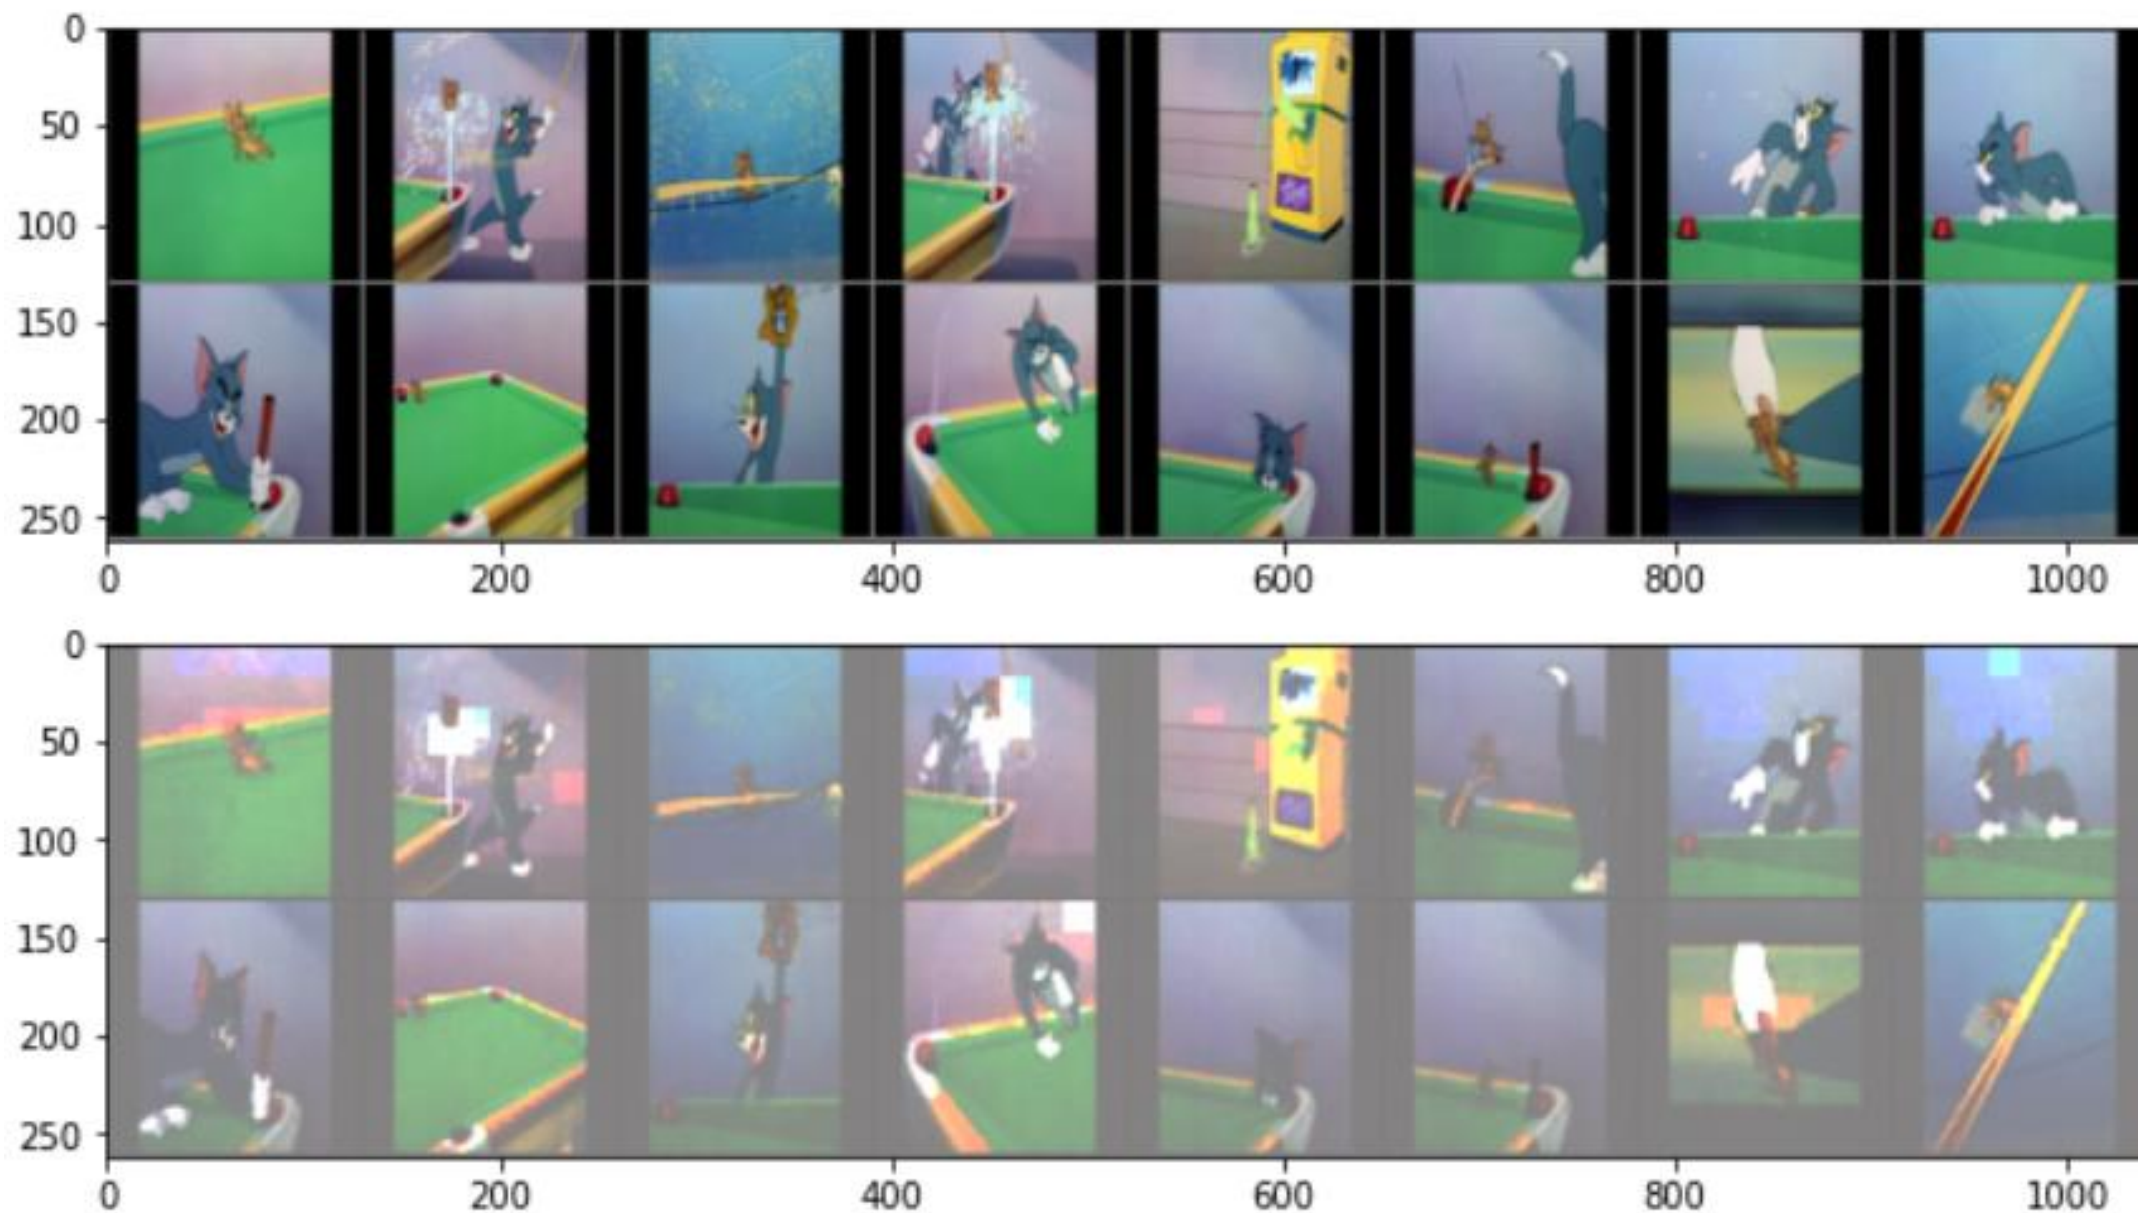
After training for 800 epochs, the NN can recover un-seen test images




Loss plot of epoch 800-1200








Training images recovered after training for 1200 epochs



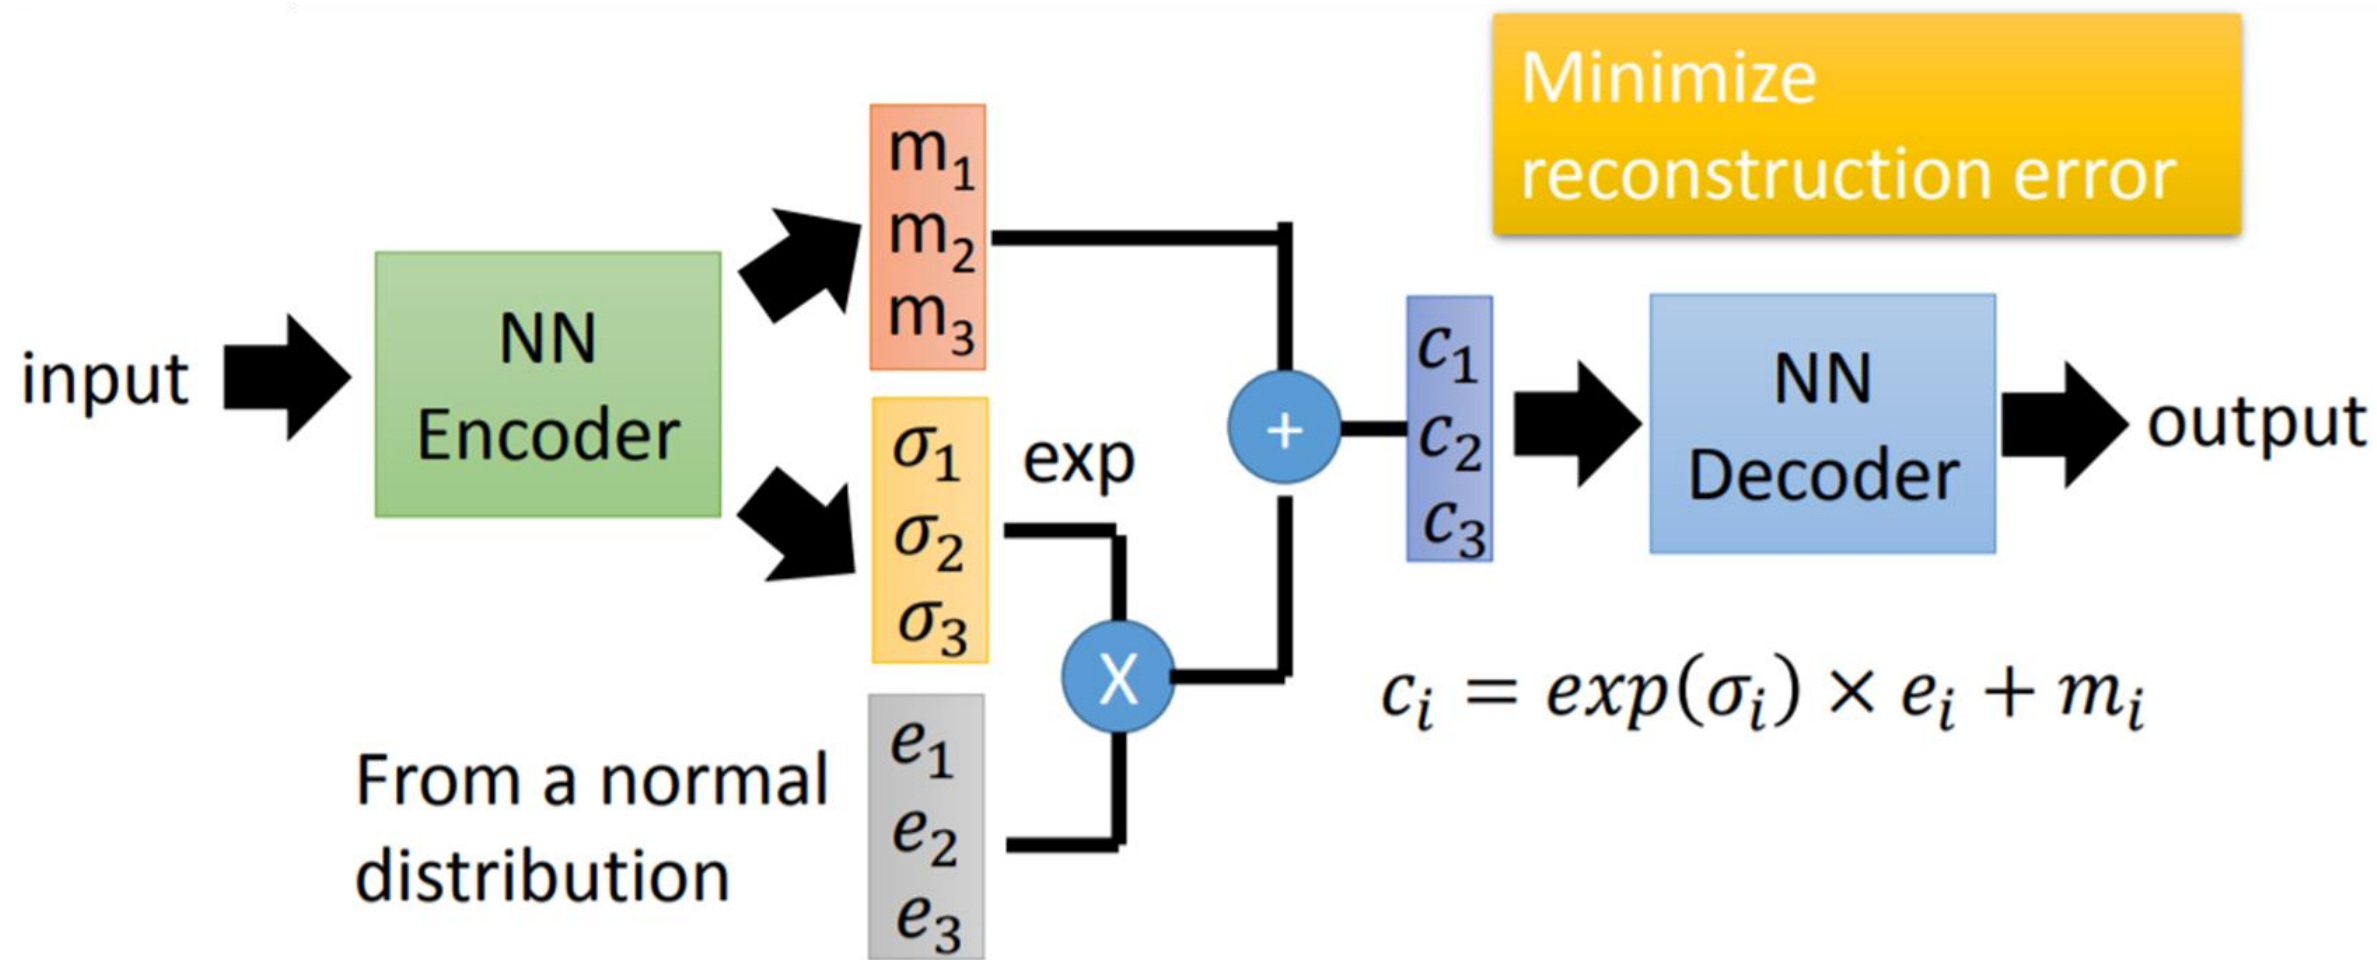
After training for 1200 epochs, the NN can recover un-seen test images



Other interesting VAE results



VAE

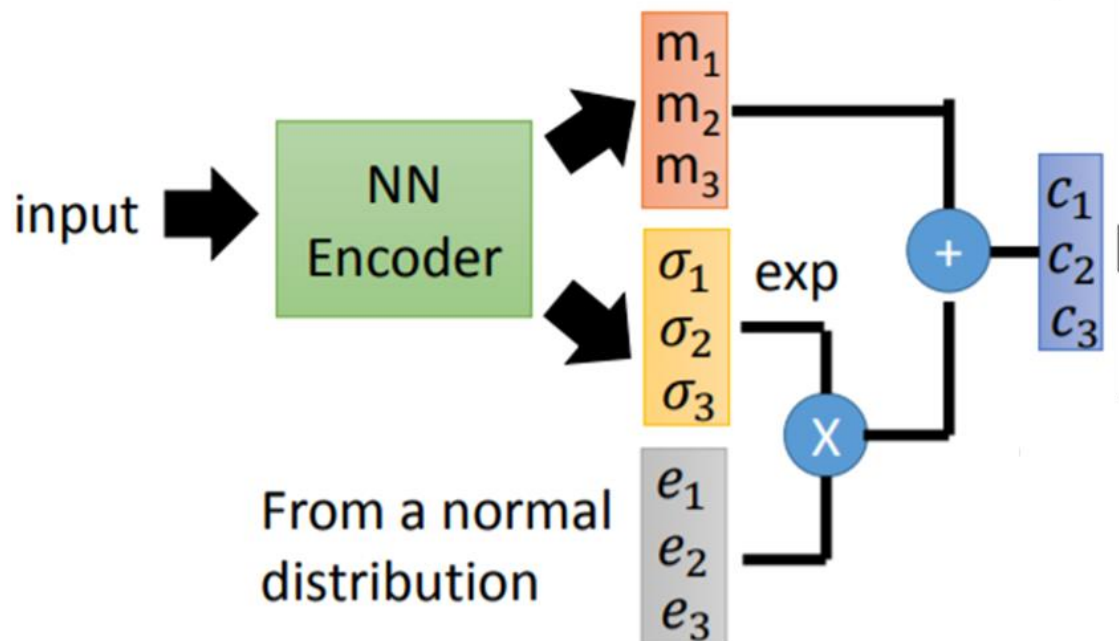


Encoder

```
[15]: for batchX, _ in loader:  
      break;  
      print(batchX.shape)
```

```
torch.Size([16, 3, 128, 128])
```

```
(fc1): Linear(in_features=1024, out_features=64,  
(fc2): Linear(in_features=1024, out_features=64,  
(fc3): Linear(in_features=64, out_features=1024,
```



```
16]: h = model.encoder(batchX.to(device))  
      print(h.shape)
```

```
torch.Size([16, 1024])
```

```
[17]: mu=model.fc1(h)  
      print(mu.shape)
```

```
torch.Size([16, 64])
```

```
[18]: logvar=model.fc2(h)  
      print(logvar.shape)
```

```
torch.Size([16, 64])
```

```
[19]: std = logvar.mul(0.5).exp_()  
      print(std.shape)
```

```
torch.Size([16, 64])
```

```
[20]: esp=torch.randn(*mu.size())  
      print(esp.shape)
```

```
torch.Size([16, 64])
```

```
[21]: z=mu+std*esp.to(device)  
      print(z.shape)
```

```
torch.Size([16, 64])
```

m_1
 m_2
 m_3

σ_1
 σ_2
 σ_3

σ_1
 σ_2
 σ_3

exp

e_1
 e_2
 e_3

c_1
 c_2
 c_3

Decoder

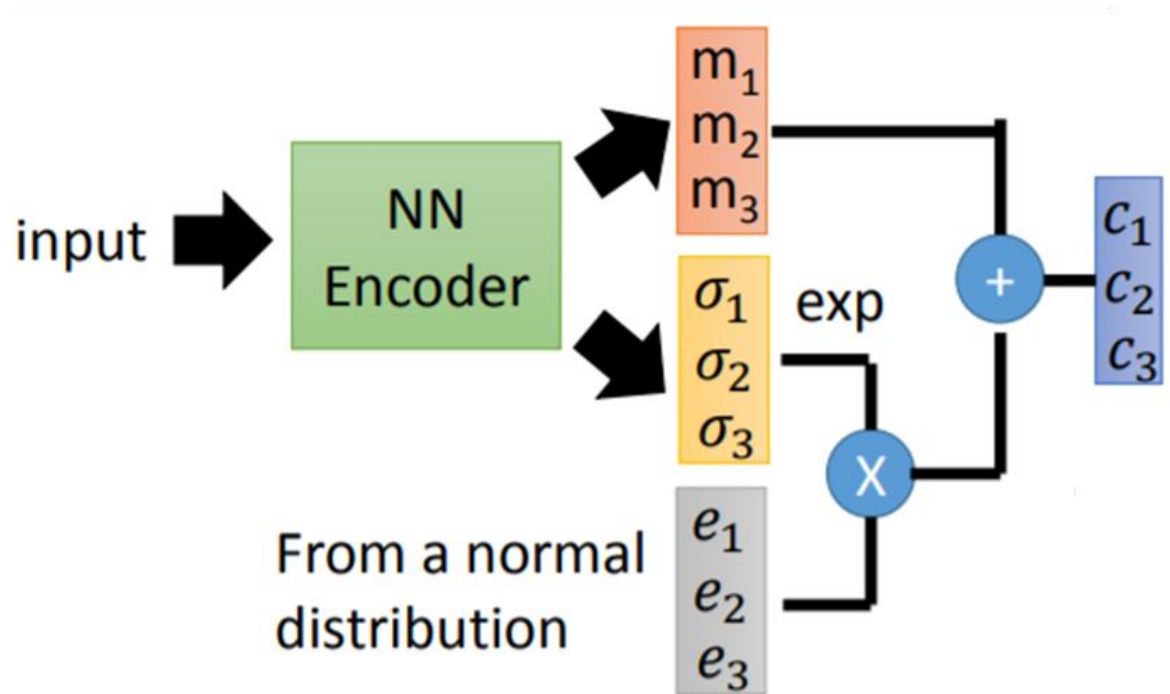
```
[22] z = model.fc3(z)  
      print(z.shape)
```

```
torch.Size([16, 1024])
```

```
[23] z = model.decoder(z)  
      print(z.shape)
```

```
torch.Size([16, 3, 128, 128])
```

Loss function



We want σ_i close to 0
(variance close to 1)

Minimize

$$\sum_{i=1}^3 (\underbrace{\exp(\sigma_i)}_{\text{blue}} - \underbrace{(1 + \sigma_i)}_{\text{red}} + \underbrace{(m_i)^2}_{\text{purple}})$$

L2 regularization

Loss function

```
[9]: def loss_fn(recon_x, x, mu, logvar):  
    #BCE = F.binary_cross_entropy(recon_x, x, size_average=False).to(device)  
    MSE = F.mse_loss(recon_x, x, reduction='sum')  
    # see Appendix B from VAE paper:  
    # Kingma and Welling. Auto-Encoding Variational Bayes. ICLR, 2014  
    # 0.5 * sum(1 + log(sigma^2) - mu^2 - sigma^2)  
    KLD = -0.5*torch.mean(1+logvar-mu.pow(2)-logvar.exp()).to(device)  
    return MSE+KLD, MSE, KLD
```

Minimize

$$\sum_{i=1}^3 (\underbrace{\exp(\sigma_i)}_{\text{L2 regularization}} - \underbrace{(1 + \sigma_i)}_{\text{L2 regularization}} + \underbrace{(m_i)^2}_{\text{L2 regularization}})$$

L2 regularization

```
[23]: tensorY,mu,logvar = model(batchX.to(device))  
print(tensorY.shape)
```


```
torch.Size([16, 3, 128, 128])
```

```
[24]: loss, mse,kld = loss_fn(tensorY, batchX.to(device), mu, logvar)  
print(loss)
```

```
tensor(627375.3750, device='cuda:0', grad_fn=<AddBackward0>)
```

$$\text{loss} = \text{MSE}(x, \hat{x}) + \text{KL}(q(z|x) || P(z)), \text{ why?}$$

Source: <https://towardsdatascience.com/understanding-variational-autoencoders-vaes-f70510919f73>



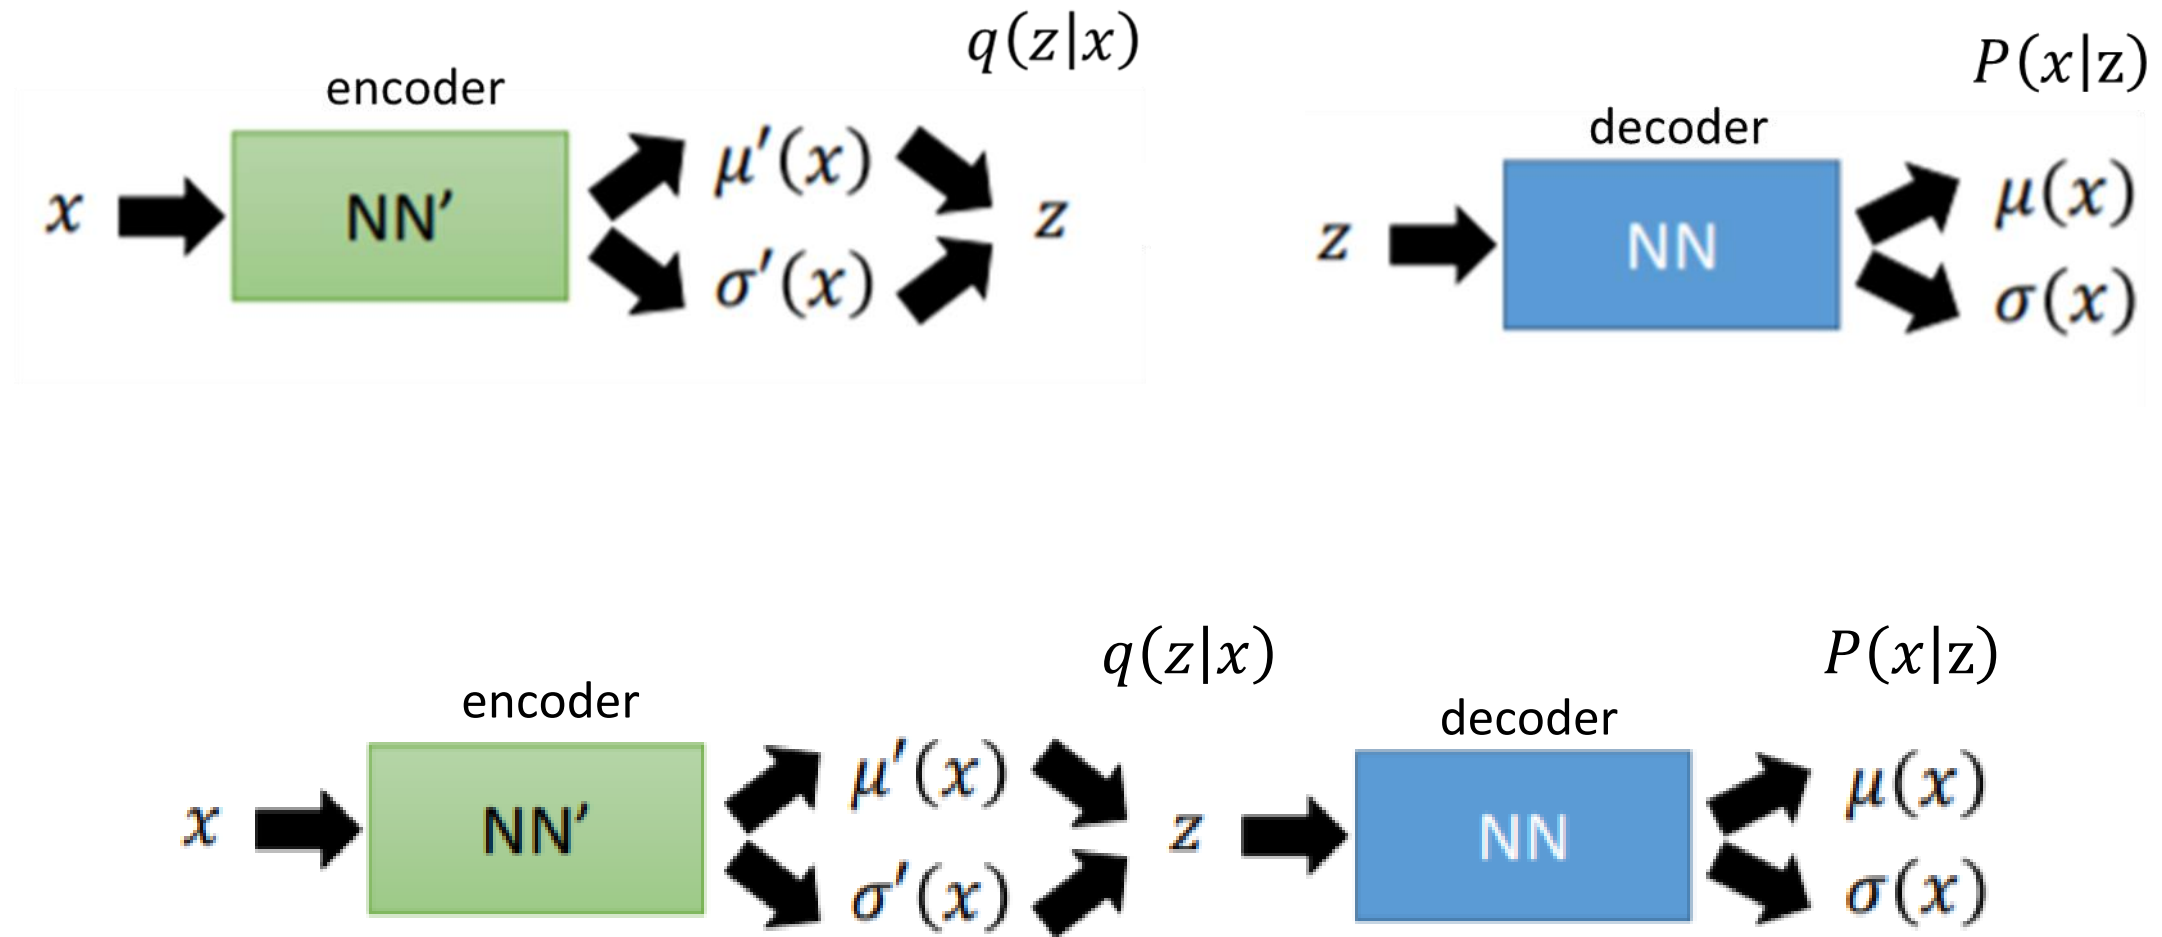
$$\text{loss} = ||x - \hat{x}||^2 + \text{KL}[N(\mu_x, \sigma_x), N(0, I)] = ||x - d(z)||^2 + \text{KL}[N(\mu_x, \sigma_x), N(0, I)]$$

$$D_{KL}(p||q) = \sum_{i=1}^N p(x_i) \log\left(\frac{p(x_i)}{q(x_i)}\right)$$

Minimize

$$\sum_{i=1}^3 (\exp(\sigma_i) - (1 + \sigma_i) + (m_i)^2)$$

The decoder and encoder of VAE model two conditional probabilities



Gaussian mixture model

Gaussian Mixture Model

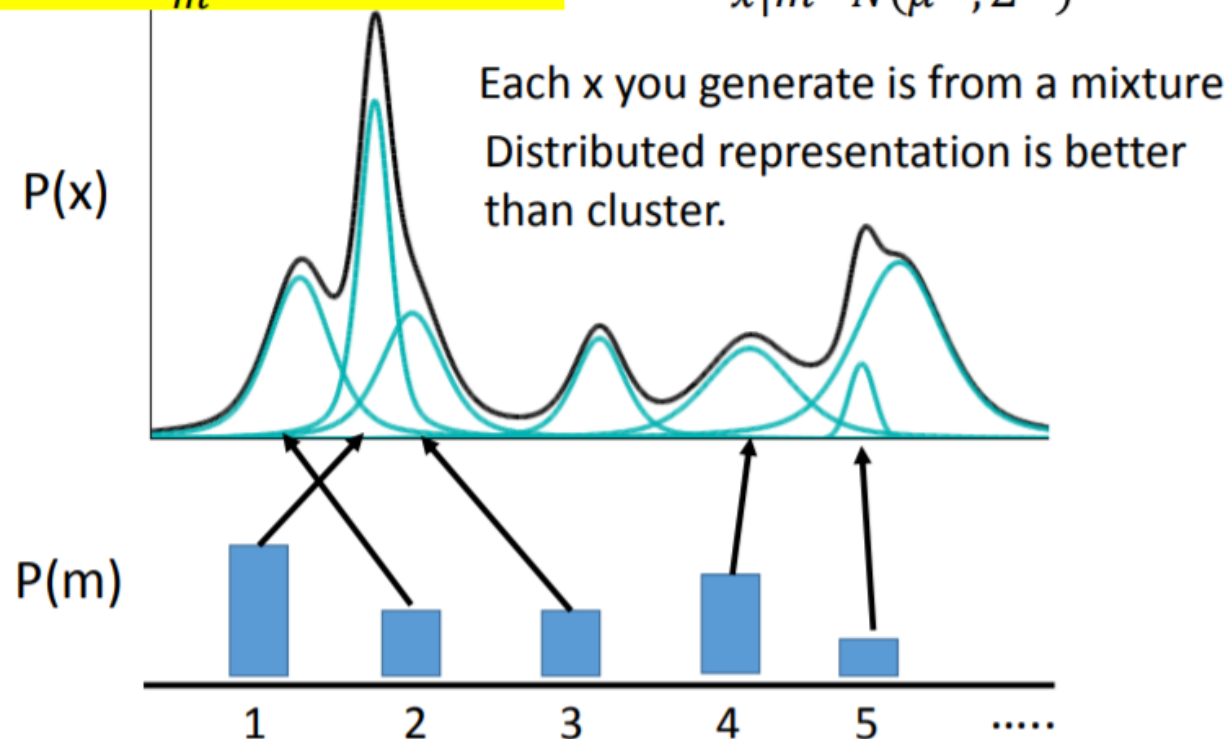
$$P(x) = \sum_m P(m)P(x|m)$$

How to sample?


$m \sim P(m)$ (multinomial)

m is an integer

$x|m \sim N(\mu^m, \Sigma^m)$




The probability of sampling an output image x from latent vector space z can be modelled as a Gaussian mixture model



Gaussian Mixture Model

$$P(x) = \int_z P(z) P(x|z) dz$$



We want to train a decoder NN that can maximize the likelihood of observing the training images

Maximizing Likelihood

$$P(x) = \int_z P(z)P(x|z)dz$$

$P(z)$ is normal distribution

$$x|z \sim N(\mu(z), \sigma(z))$$

$\mu(z), \sigma(z)$ is going to be estimated

$$L = \sum_x \log P(x)$$

Maximizing the likelihood of the observed x

Recap: maximize the likelihood of observing the classification of the training data

Training Data	x^1	x^2	x^3	$\dots \dots$	x^N
	C_1	C_1	C_2		C_1

$$\max \quad L(w, b) = f_{w,b}(x^1) f_{w,b}(x^2) (1 - f_{w,b}(x^3)) \cdots f_{w,b}(x^N)$$


$$\min \quad -\ln L(w, b) = -\ln f_{w,b}(x^1) - \ln f_{w,b}(x^2) - \ln (1 - f_{w,b}(x^3)) \cdots$$

\hat{y}^n : 1 for class 1, 0 for class 2

$$= \sum_n \underbrace{-\left[\hat{y}^n \ln f_{w,b}(x^n) + (1 - \hat{y}^n) \ln (1 - f_{w,b}(x^n)) \right]}_{\text{Cross entropy between two Bernoulli distribution}}$$


Rewrite the maximum likelihood item $\log P(x)$ as the summation of a lower bound L_b and KL divergence

$$\begin{aligned}
 \log P(x) &= \int_z q(z|x) \log P(x) dz && \text{q(z|x) can be any distribution} \\
 &= \int_z q(z|x) \log \left(\frac{P(z, x)}{P(z|x)} \right) dz = \int_z q(z|x) \log \left(\frac{P(z, x) q(z|x)}{q(z|x) P(z|x)} \right) dz \\
 &= \int_z q(z|x) \log \left(\frac{P(z, x)}{q(z|x)} \right) dz + \underbrace{\int_z q(z|x) \log \left(\frac{q(z|x)}{P(z|x)} \right) dz}_{KL(q(z|x) || P(z|x))} \\
 &\geq \int_z q(z|x) \log \left(\frac{P(x|z) P(z)}{q(z|x)} \right) dz && \text{lower bound } L_b \quad \geq 0
 \end{aligned}$$



$$D_{KL}(q||p) = \sum_{i=1}^N q(x_i) \log \left(\frac{q(x_i)}{p(x_i)} \right)$$

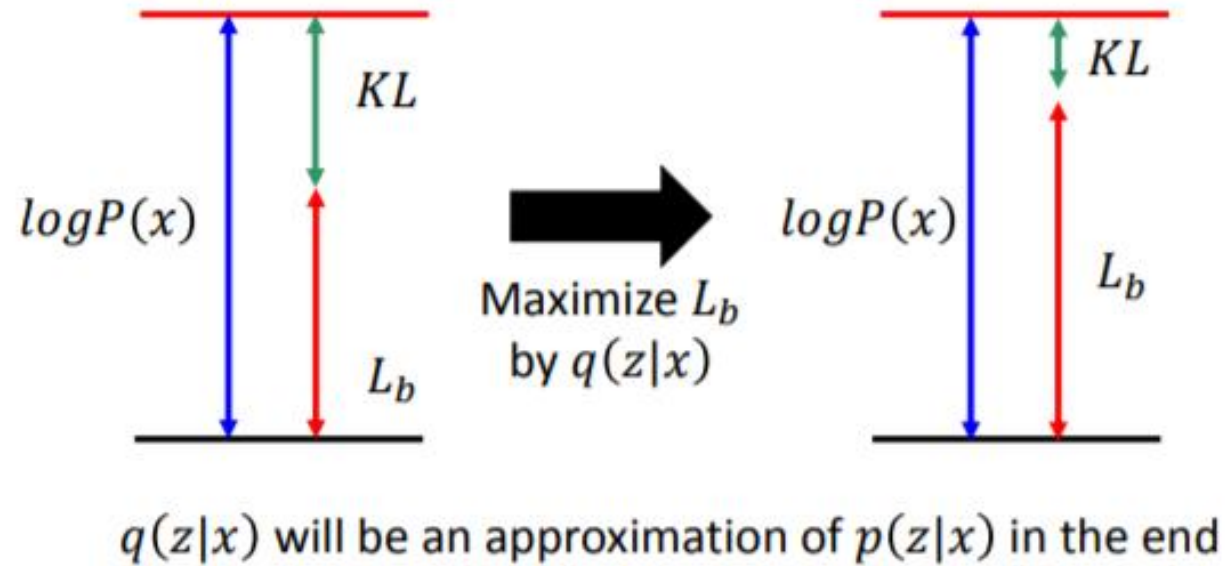
If we maximum L_b by adjusting $P(x|z)$ and $q(z|x)$ simultaneously, then we can maximum L_b and at the same time minimize the KL distance



$$\log P(x) = L_b + KL(q(z|x) || P(z|x))$$

$$L_b = \int_z q(z|x) \log \left(\frac{P(x|z)P(z)}{q(z|x)} \right) dz$$


Find $P(x|z)$ and $q(z|x)$
maximizing L_b



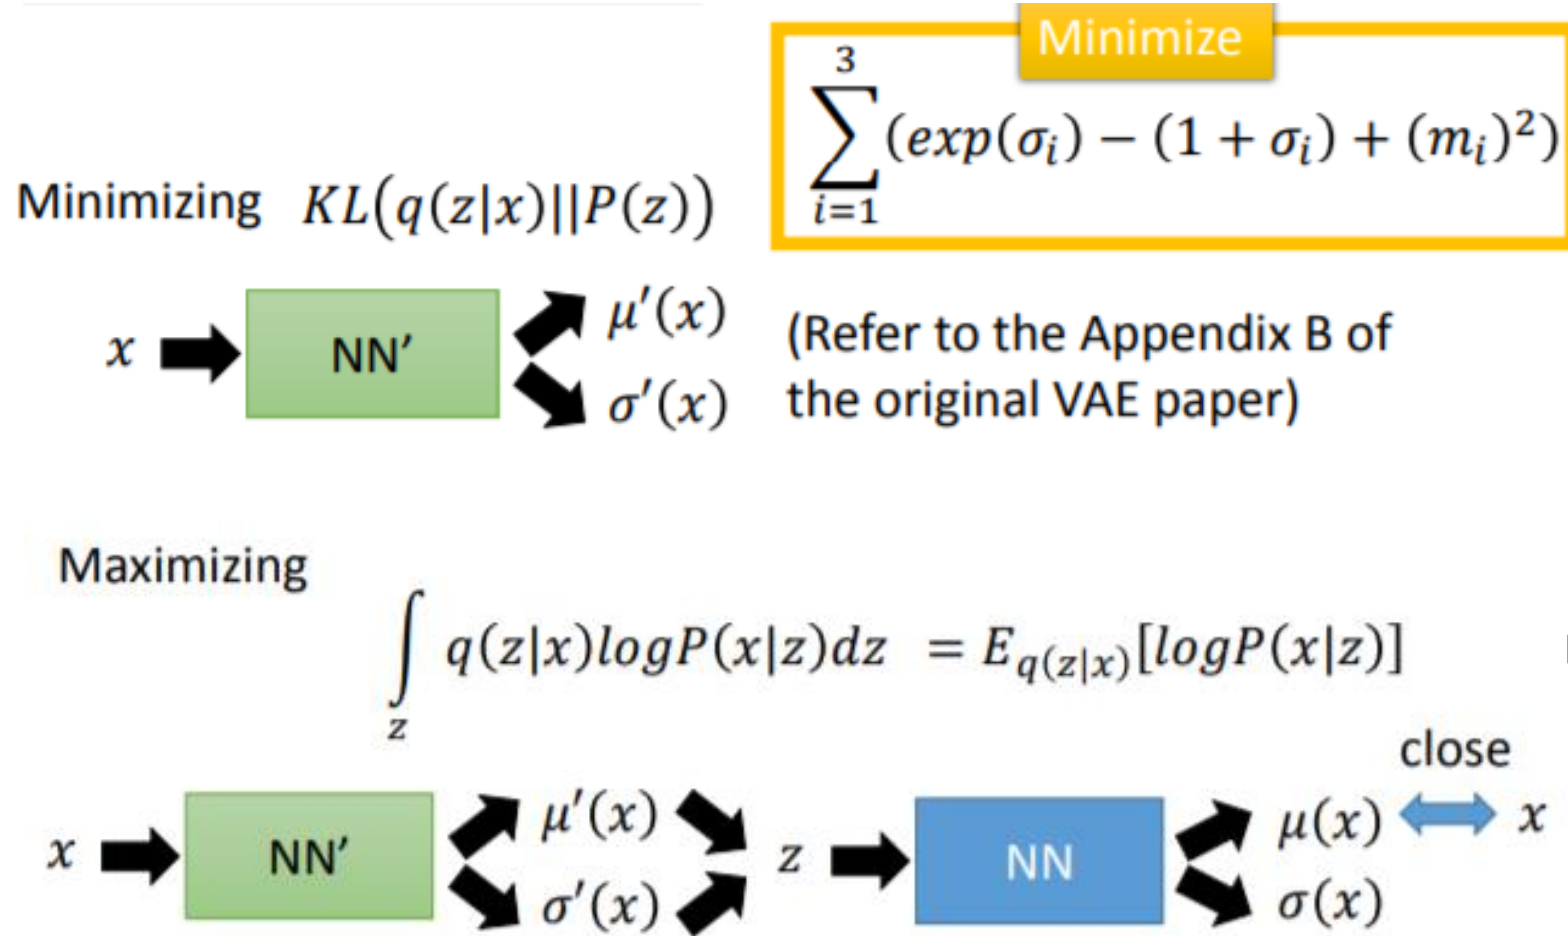
Rewrite the lower bound L_b as the summation of two terms: $-KL(q(z|x)||P(z))$, and

$$\int_z q(z|x) \log P(x|z) dz$$

$$\begin{aligned} \max \quad L_b &= \int_z q(z|x) \log \left(\frac{P(z, x)}{q(z|x)} \right) dz = \int_z q(z|x) \log \left(\frac{P(x|z)P(z)}{q(z|x)} \right) dz \\ &= \underbrace{\int_z \boxed{q(z|x)} \log \left(\boxed{\frac{P(z)}{q(z|x)}} \right) dz}_{-KL(q(z|x)||P(z))} + \int_z q(z|x) \log P(x|z) dz \end{aligned}$$



max. L_b can be done by min. $KL(q(z|x)||P(z))$ and max. $\int_z q(z|x) \log P(x|z) dz$. That is why loss = KLD + MSE (x, \hat{x})



HW6 (2)

- Train an VAE to learn a compact representation (try latent vector of size 20, 30, 50) of your facial expression. Test with 10 happy and 10 angry faces.
- Show the recovered image.
- Send the latent vectors to t -SNE to see whether they form clusters.

