
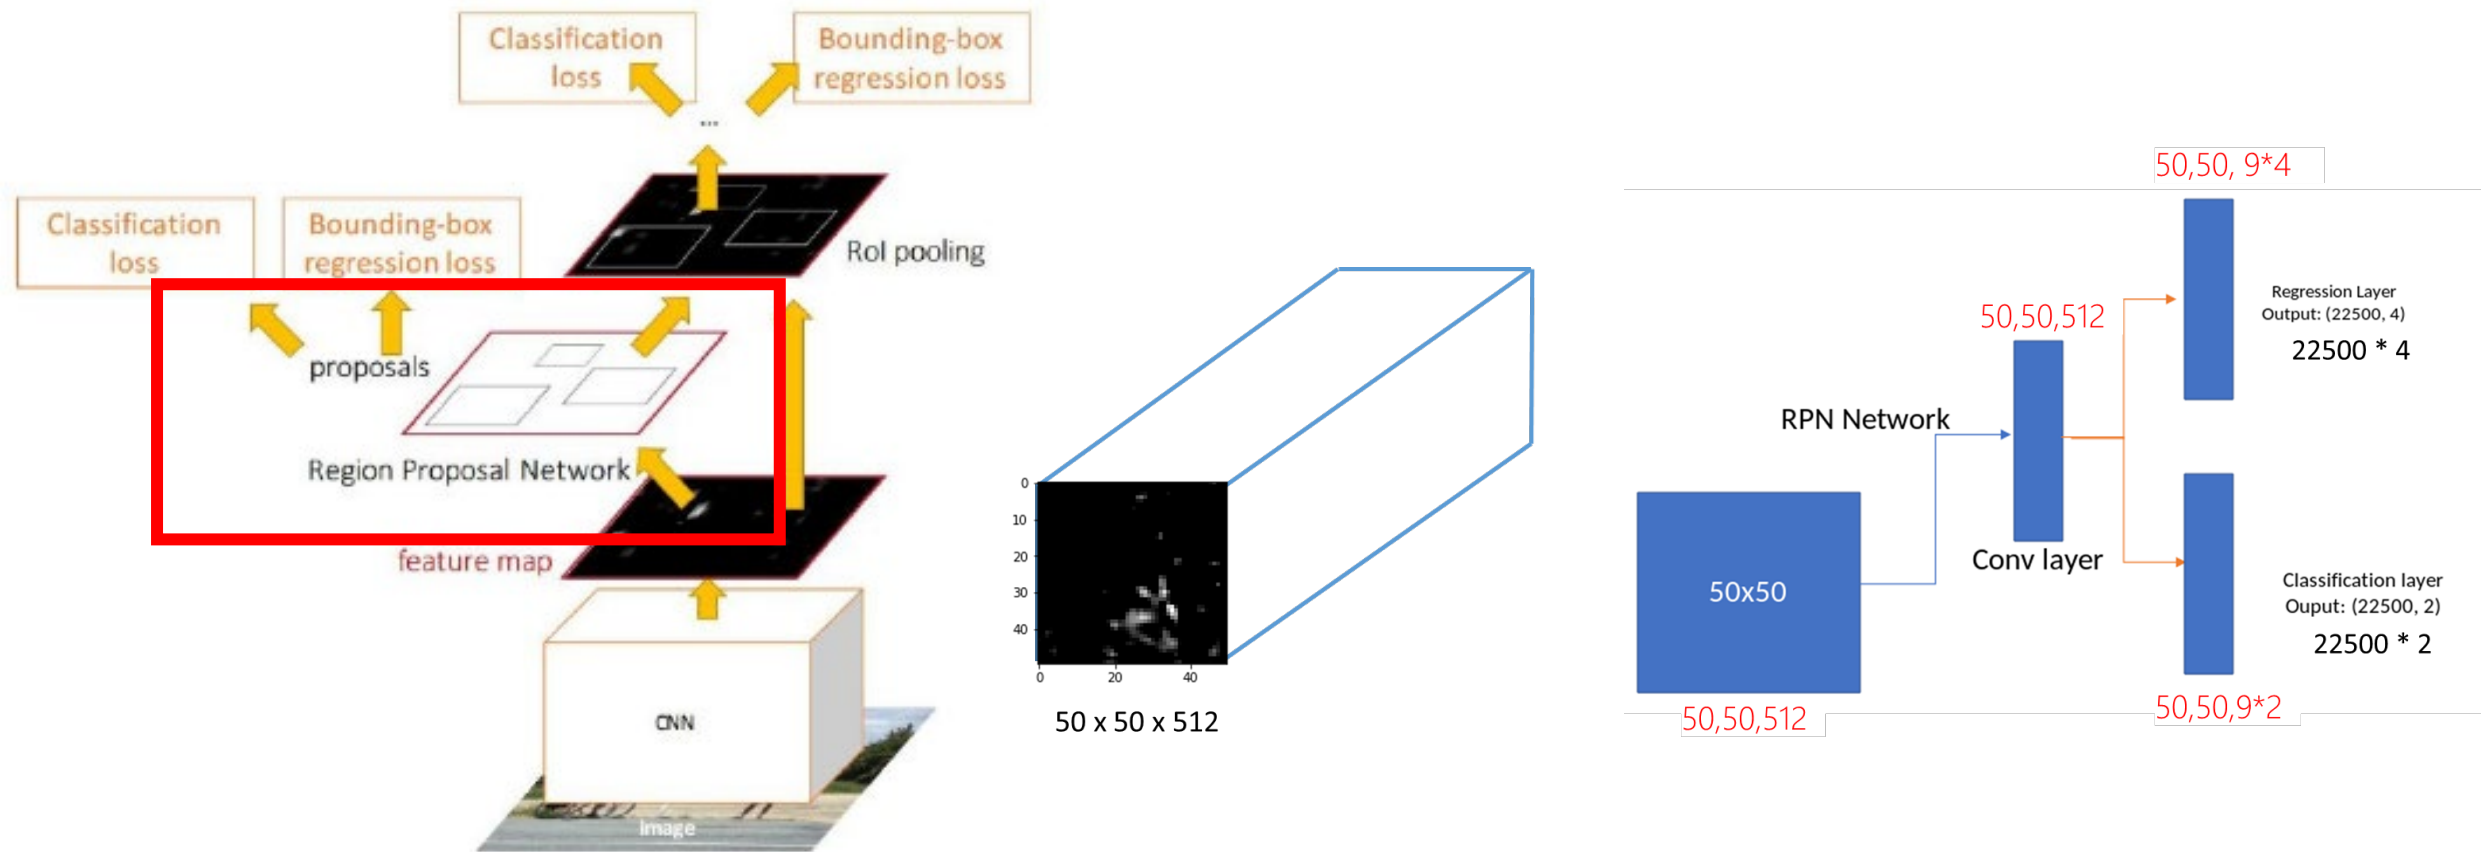


Recap: FasterRCNN



RPN (region proposal network)

Takes feature map as input and predict 22,500 ROIs (region of interests)



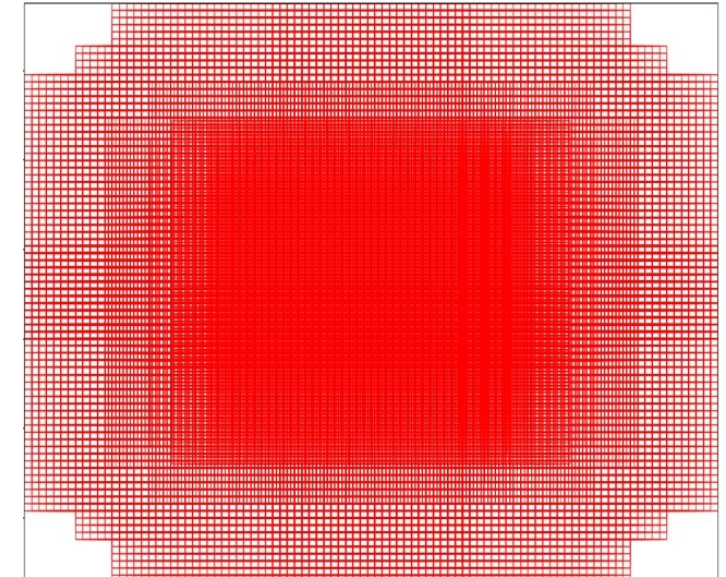
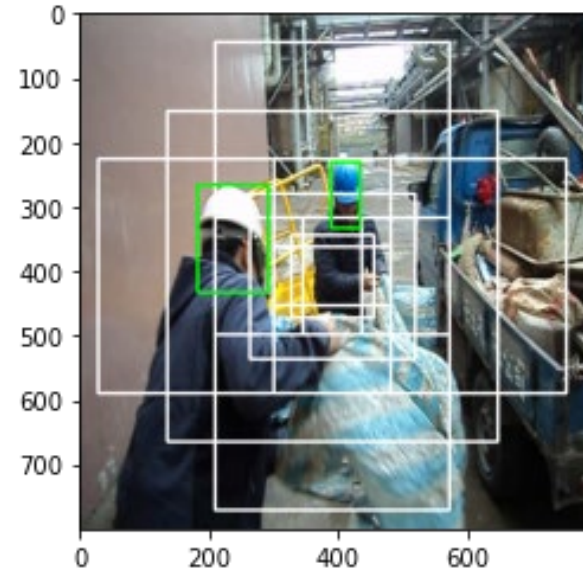
Ground-truth labels to calculate RPN classification loss

16*16 anchor points

9 anchor boxes for each anchor pt

$$16*16*9 = 22,500$$

8940 valid anchor boxes



Label the 8,940 valid anchor boxes

1: IOU > 0.7 (may contain object)

0: IOU < 0.3 (background)

-1: ignore

Sample a batch of anchor boxes to train RPN: 128 positive examples and 128 negative examples (label 0). Change the labels of all other valid anchor boxes to -1(ignore) at this mini-batch training.

RPN classification loss

$$L(p_i, t_i) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, \hat{p}_i) + \lambda \frac{1}{N_{reg}} \sum_i p_i L_{reg}(t_i, \hat{t}_i)$$

$$N_{cls} = 256$$

\hat{p}_i Probability for class 0 and 1 predicted by RPN

$$p_i = \begin{cases} 0, & \text{negative label} \\ 1, & \text{positive label} \end{cases}$$

$$L_{cls}(p_i, \hat{p}_i) = CE(p_i, \hat{p}_i) = - \sum_{k=1}^2 p_k \ln(\hat{p}_k)$$

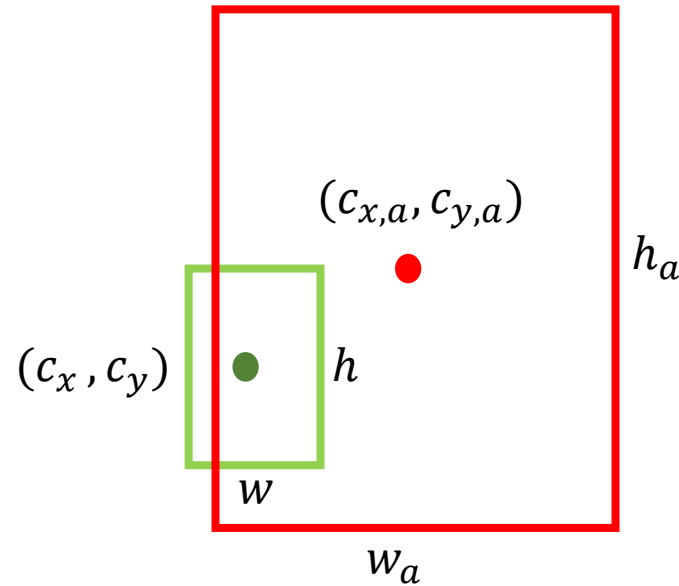
```
# For classification we use cross-entropy loss
rpn_cls_loss = F.cross_entropy(rpn_score, gt_rpn_score)
print(rpn_cls_loss)
```

```
rpn_loc = pred_anchor_locs[0]
rpn_score = pred_cls_scores[0]

gt_rpn_loc = torch.from_numpy(anchor_locations)
gt_rpn_score = torch.from_numpy(anchor_labels)
```

Ground-truth values to calculate RPN bounding box regression loss

For each valid anchor box, use the ground truth bbox with maximum IOU to calculate a normalized location representation



$$d_x = \frac{c_x - c_{x,a}}{w_a} \quad d_y = \frac{c_y - c_{y,a}}{h_a}$$

$$d_w = \log\left(\frac{w}{w_a}\right) \quad d_h = \log\left(\frac{h}{h_a}\right)$$

Smooth L1 loss

$$Loss_2 = \frac{1}{N} \sum_{i=1}^N (y^i - \hat{y}^i)^2$$

$$Loss_1 = \frac{1}{N} \sum_{i=1}^N |y^i - \hat{y}^i|$$

$$smooth\ Loss_1 = \begin{cases} 0.5x^2 \times 1/\sigma^2 & \text{if } |x| < 1/\sigma^2 \\ |x| - 0.5 & \text{otherwise} \end{cases}$$

RPN bounding box regression loss

$$L(p_i, t_i) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, \hat{p}_i) + \lambda \frac{1}{N_{reg}} \sum_i p_i L_{reg}(t_i, \hat{t}_i)$$

$$N_{reg} = 128$$

$$t_i = [d_x, d_y, d_w, d_h]$$

$$\hat{t}_i = [\hat{d}_x, \hat{d}_y, \hat{d}_w, \hat{d}_h]$$

$$p_i = \begin{cases} 0, & \text{negative label} \\ 1, & \text{positive label} \end{cases}$$

$$L_{reg} = \begin{cases} 0.5(t_i - \hat{t}_i)^2 \times 1/\sigma^2 & \text{if } |t_i - \hat{t}_i| < 1/\sigma^2 \\ |t_i - \hat{t}_i| - 0.5 & \text{otherwise} \end{cases} \quad \sigma = 3 \text{ for RPN training}$$

RPN bounding box regression loss

$$L(p_i, t_i) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, \hat{p}_i) + \lambda \frac{1}{N_{reg}} \sum_i L_{reg}(t_i, \hat{t}_i)$$

```
# For Regression we use smooth L1 loss as defined in the Fast RCNN paper
```

```
pos = gt_rpn_score > 0
```

```
mask = pos.unsqueeze(1).expand_as(rpn_loc)
```

```
print(mask.shape)
```

```
# take those bounding boxes which have positive labels
```

```
mask_loc_preds = rpn_loc[mask].view(-1, 4)
```

```
mask_loc_targets = gt_rpn_loc[mask].view(-1, 4)
```

```
print(mask_loc_preds.shape, mask_loc_targets.shape)
```


```
x = torch.abs(mask_loc_targets.cpu() - mask_loc_preds.cpu())
```

```
rpn_loc_loss = ((x < 1).float() * 0.5 * x**2) + ((x >= 1).float() * (x-0.5))
```

```
print(rpn_loc_loss.sum())
```

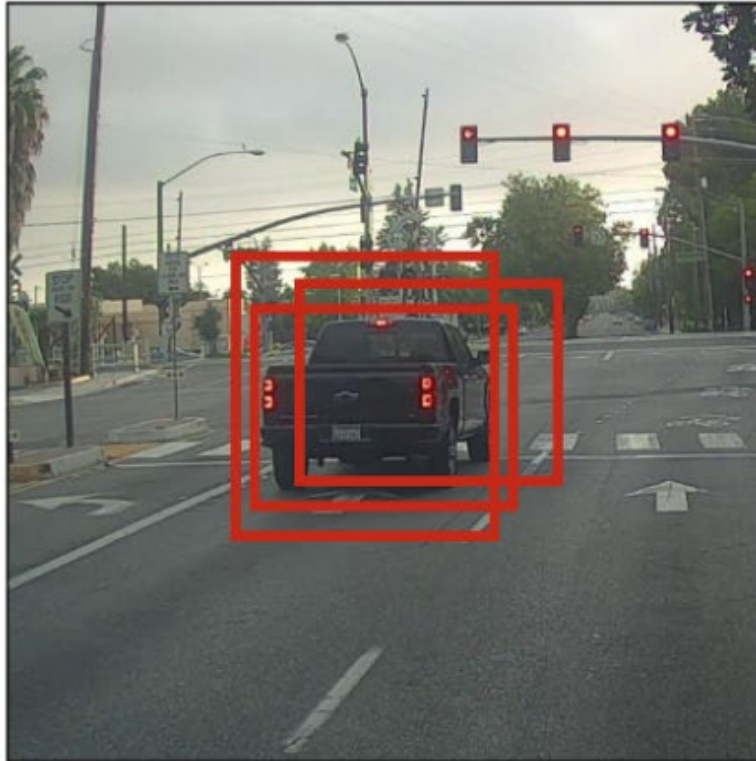

Pass ROIs to FastRCNN detector

Use NMS, IOU to reduce the number of ROI from 22500 \rightarrow 2000 \rightarrow 128



Non-maximum Suppression (NMS)

Before non-max suppression



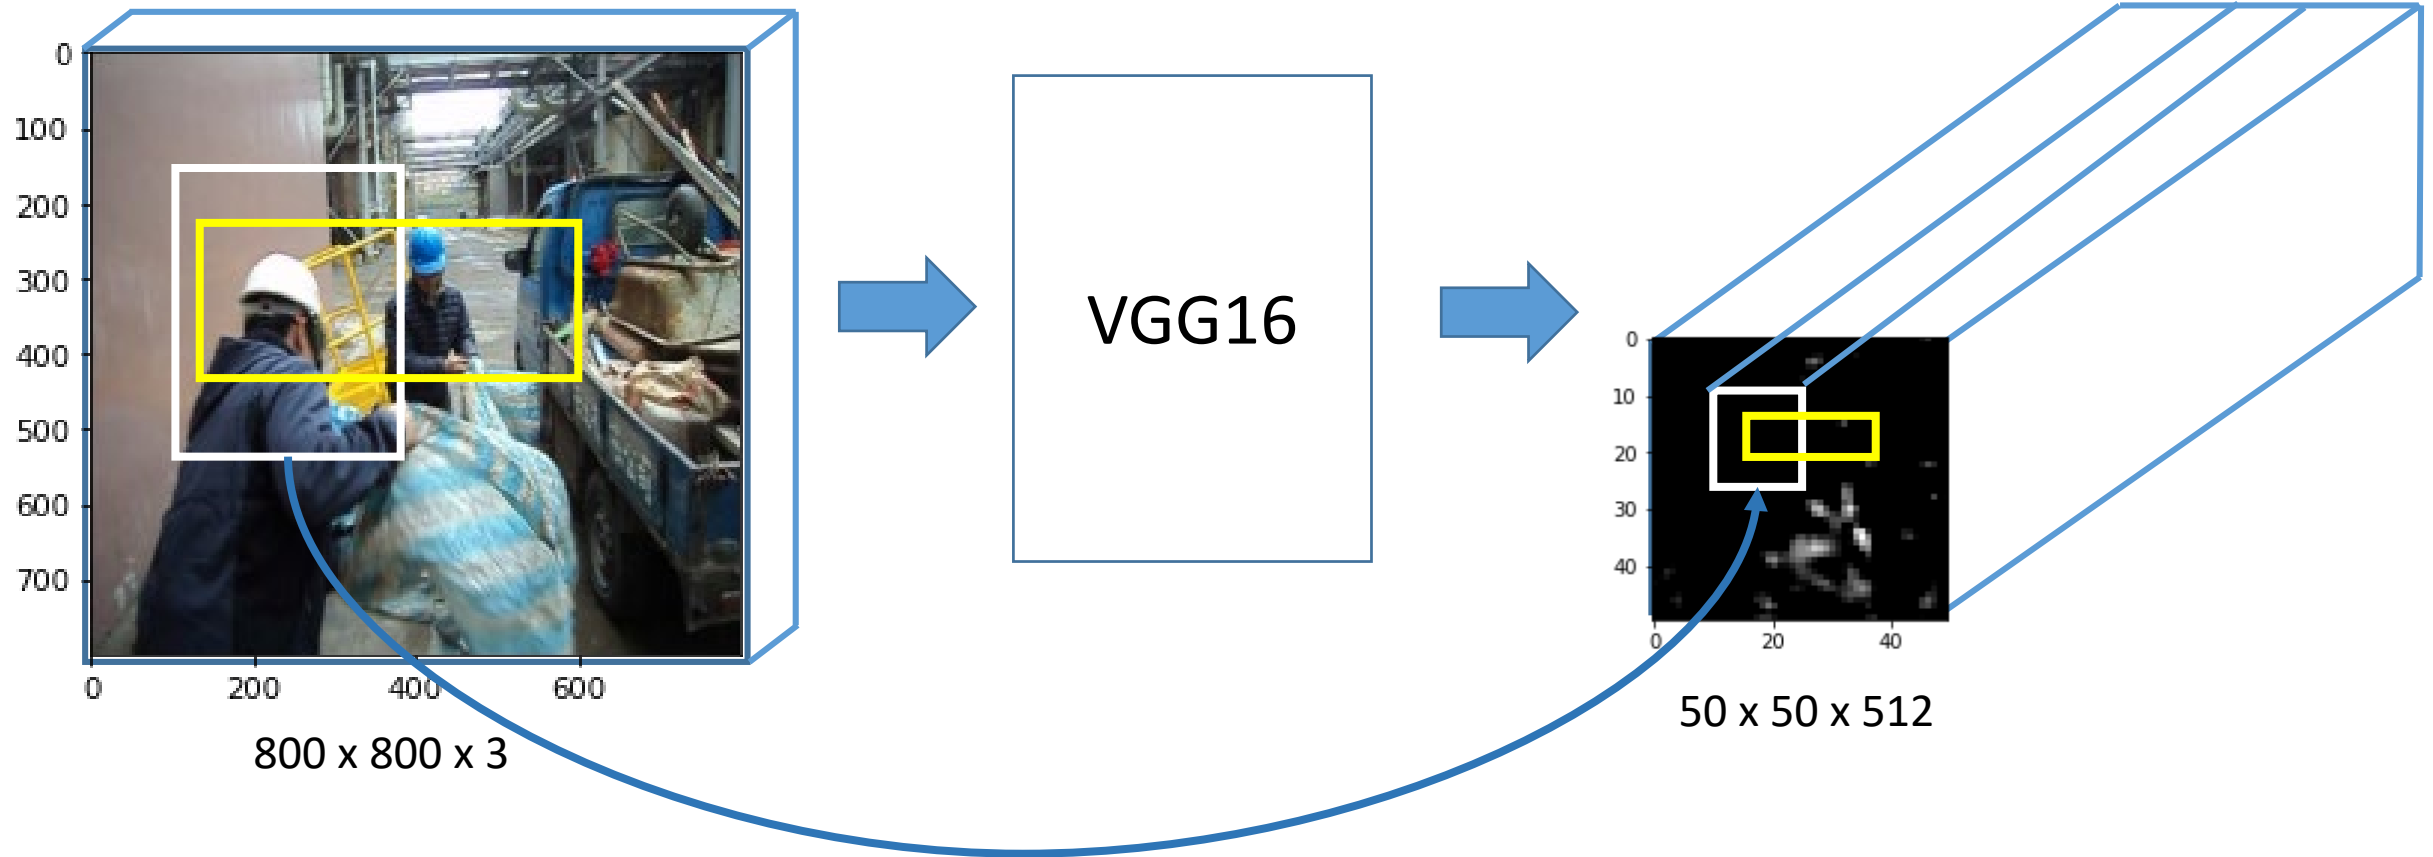
Non-Max
Suppression



After non-max suppression


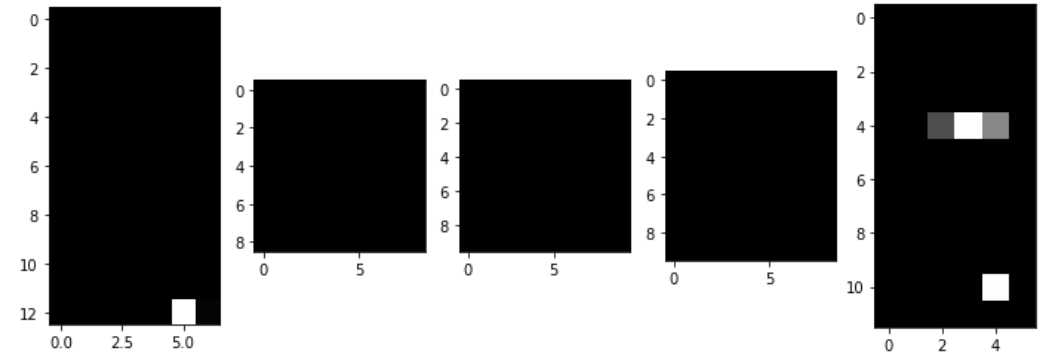

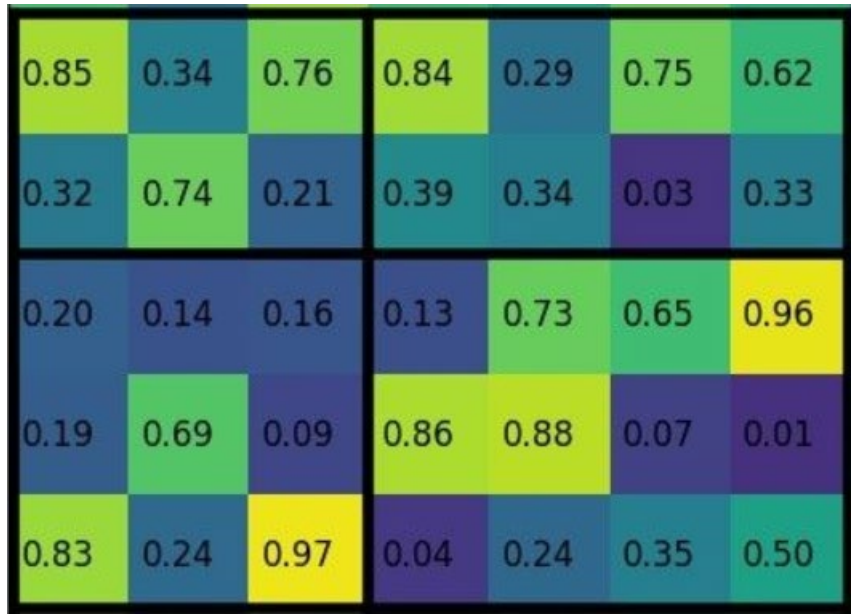


Extract the feature maps of the 128 ROI samples





ROI Pooling

Extract the feature maps of the 128 ROI samples, adjust to the same size $H=7$, $W=7$ using max pooling (ROI Pooling)



https://blog.csdn.net/qq_35586657/article/details/97885290

FastRCNN detector



FastRCNN detector

```
# Reshape the tensor so that we can pool
k = output.view(output.size(0), -1)
print(k.shape) # 25088 = 7*7*512
```

```
roi_head_classifier = nn.Sequential(*[nn.Linear(25088, 4096), nn.Linear(4096, 4096)])
cls_loc = nn.Linear(4096, 2 * 4).to(device) # (1 classes 安全帽 + 1 background. Each i
cls_loc.weight.data.normal_(0, 0.01)
cls_loc.bias.data.zero_()

score = nn.Linear(4096, 2).to(device) # (1 classes, 安全帽 + 1 background)
```

No of object classes you want to predict + 1 (background)

```
# passing the output of roi-pooling to ROI head
k = roi_head_classifier(k.to(device))
roi_cls_loc = cls_loc(k)
roi_cls_score = score(k)
print(roi_cls_loc.shape, roi_cls_score.shape)
```

```
torch.Size([128, 8]) torch.Size([128, 2])
```

Class practice

- Prepare a training image that has at least two classes of objects to be recognized. Mark the 2 bounding boxes that represent 2 classes of objects. Pass the image + bbox through FasterRCNN to calculate training loss.

FastRCNN detector classification loss

$$L(p_i, t_i) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, \hat{p}_i) + \lambda \frac{1}{N_{reg}} \sum_i p_i L_{reg}(t_i, \hat{t}_i)$$

$$N_{cls} = 128$$

\hat{p}_i Probabilities for each class

p_i 0 or 1

$$L_{cls}(p_i, \hat{p}_i) = CE(p_i, \hat{p}_i) = - \sum_{k=1}^c p_k \ln(\hat{p}_k)$$

FastRCNN detector bounding box regression loss

$$L(p_i, t_i) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, \hat{p}_i) + \lambda \frac{1}{N_{reg}} \sum_i p_i L_{reg}(t_i, \hat{t}_i)$$

$$N_{reg} = 128$$

$$t_i = [x_1, y_1, x_2, y_2]$$

$$\hat{t}_i = [\hat{x}_1, \hat{y}_1, \hat{x}_2, \hat{y}_2]$$

$$p_i = \begin{cases} 0, & \text{negative label} \\ 1, & \text{positive label} \end{cases}$$

$$L_{reg} = \begin{cases} 0.5(t_i - \hat{t}_i)^2 \times 1/\sigma^2 & \text{if } |t_i - \hat{t}_i| < 1/\sigma^2 \\ |t_i - \hat{t}_i| - 0.5 & \text{otherwise} \end{cases}$$

$\sigma = 1$ for FastRCNN training

Fine tune FasterRCNN to detect our own objects

PyTorch tutorial

TorchVision Object Detection Finetuning Tutorial

Finetune a pre-trained Mask R-CNN model.

Image/Video



<https://pytorch.org/tutorials/>

PyTorch tutorial

TORCHVISION OBJECT DETECTION FINETUNING TUTORIAL

- TIP

To get the most of this tutorial, we suggest using this [Colab Version](#). This will allow you to experiment with the information presented below.

For this tutorial, we will be finetuning a pre-trained [Mask R-CNN](#) model in the [Penn-Fudan Database for Pedestrian Detection and Segmentation](#). It contains 170 images with 345 instances of pedestrians, and we will use it to illustrate how to use the new features in torchvision in order to train an instance segmentation model on a custom dataset.

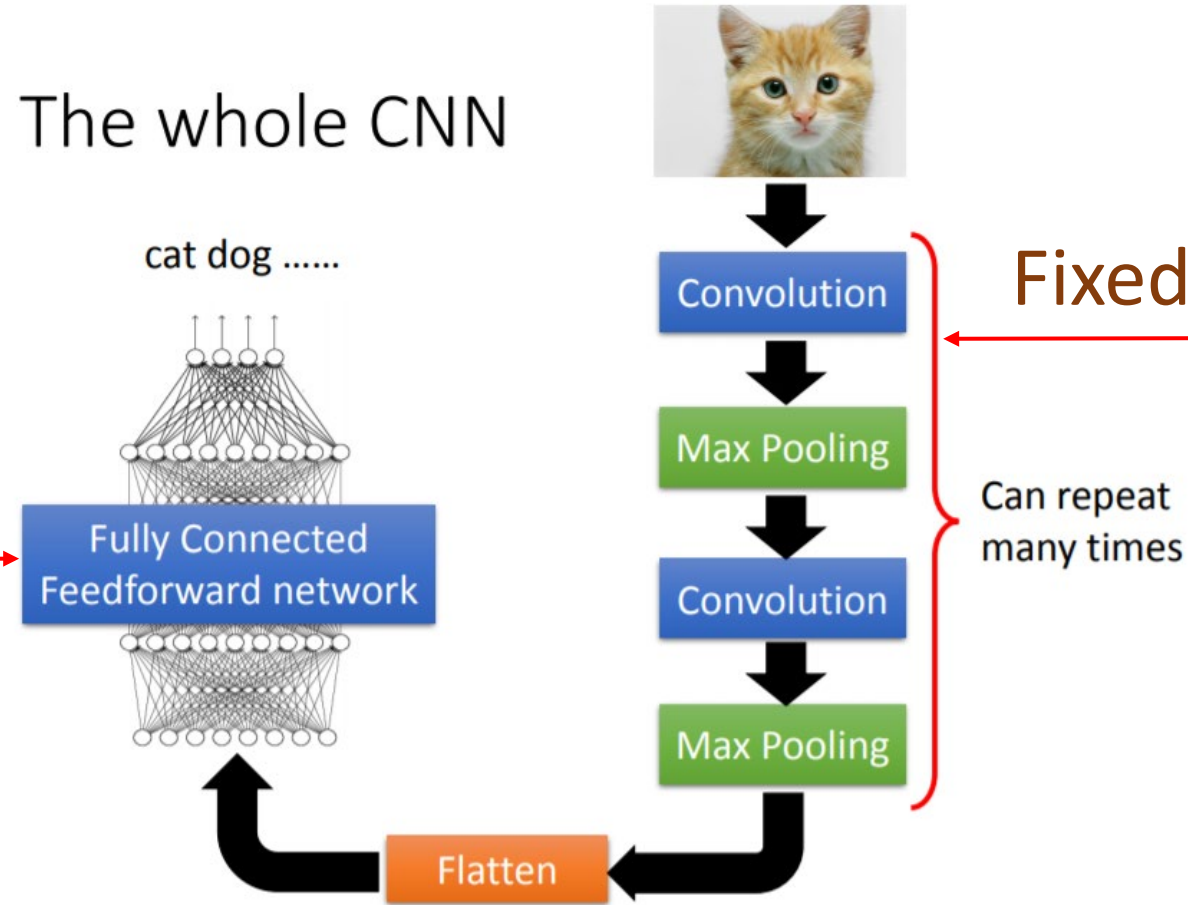
<https://pytorch.org/tutorials/>

Class practice

FasterRCNN(3) Fine_tune.ipynb

Recap – Fine-tune a pre-trained image classifier

The whole CNN



Reference: 李弘毅 ML Lecture 10
<https://youtu.be/FrKWiRv254g>

Train

```
AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(4, 4))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(64, 128, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=12288, out_features=1000, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=1000, out_features=1000, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=1000, out_features=1000, bias=True)
  )
)
```


Recap – Fine-tune a pre-trained image classifier

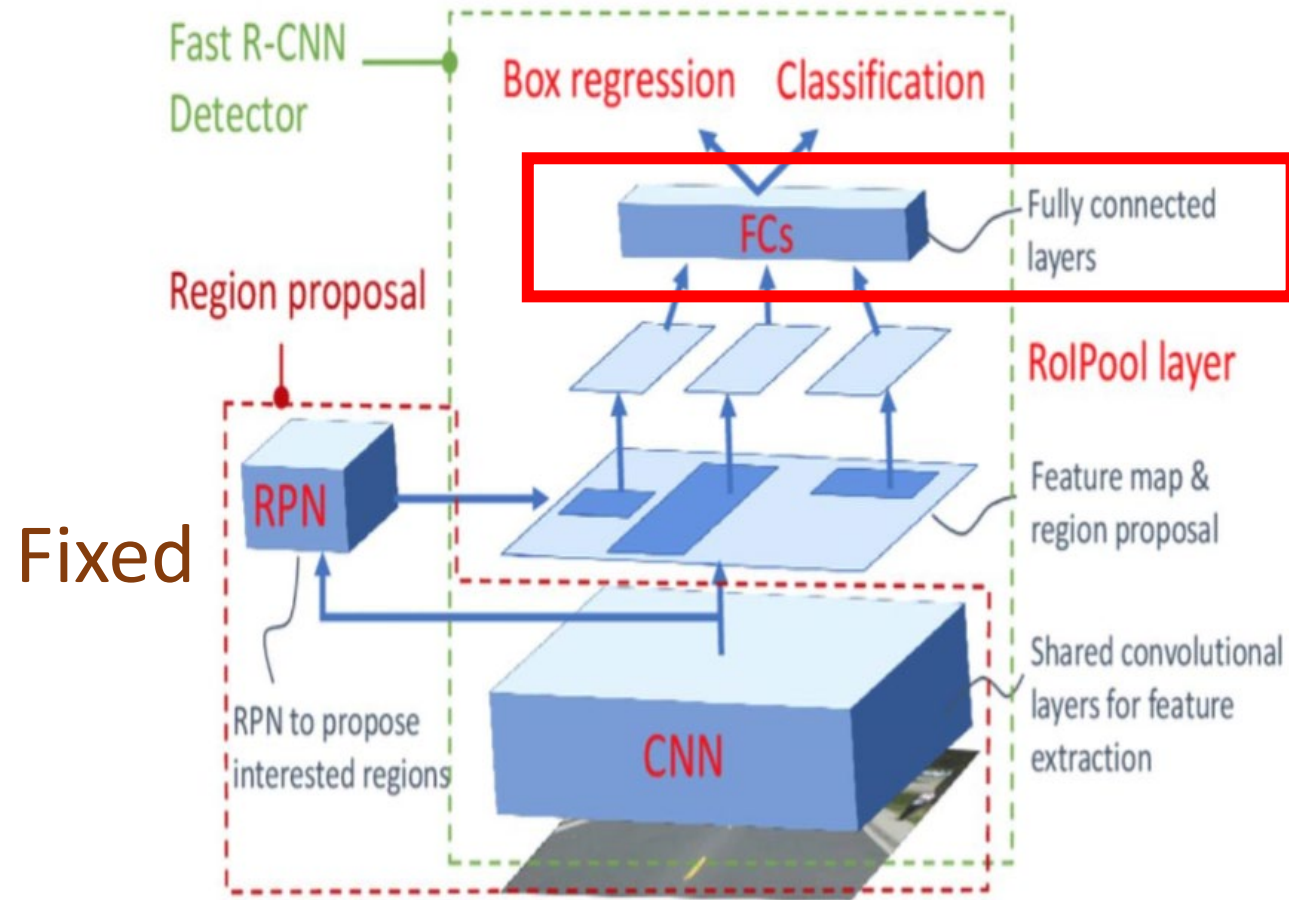
```
In [3]: import torch.nn as nn
# fix the weight of convolution layers
model.features.eval()

# modify classifier
model.classifier = torch.nn.Sequential(
    nn.Linear(25088, 4096),
    nn.ReLU(inplace=True),
    nn.Dropout(p=0.5, inplace=False),
    nn.Linear(4096, 4096),
    nn.ReLU(inplace=True),
    nn.Dropout(p=0.5, inplace=False),
    torch.nn.Linear(4096, 5))
```

Fixed

Train

Fine-tune FasterRCNN



Train

Fine-tune FasterRCNN

```
import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor

# Load a model pre-trained pre-trained on COCO
model = torchvision.models.detection.fasterrcnn_resnet50_fpn(pretrained=True)

# replace the classifier
num_classes = 2 # 1 class (person) + background

# get number of input features for the classifier
in_features = model.roi_heads.box_predictor.cls_score.in_features

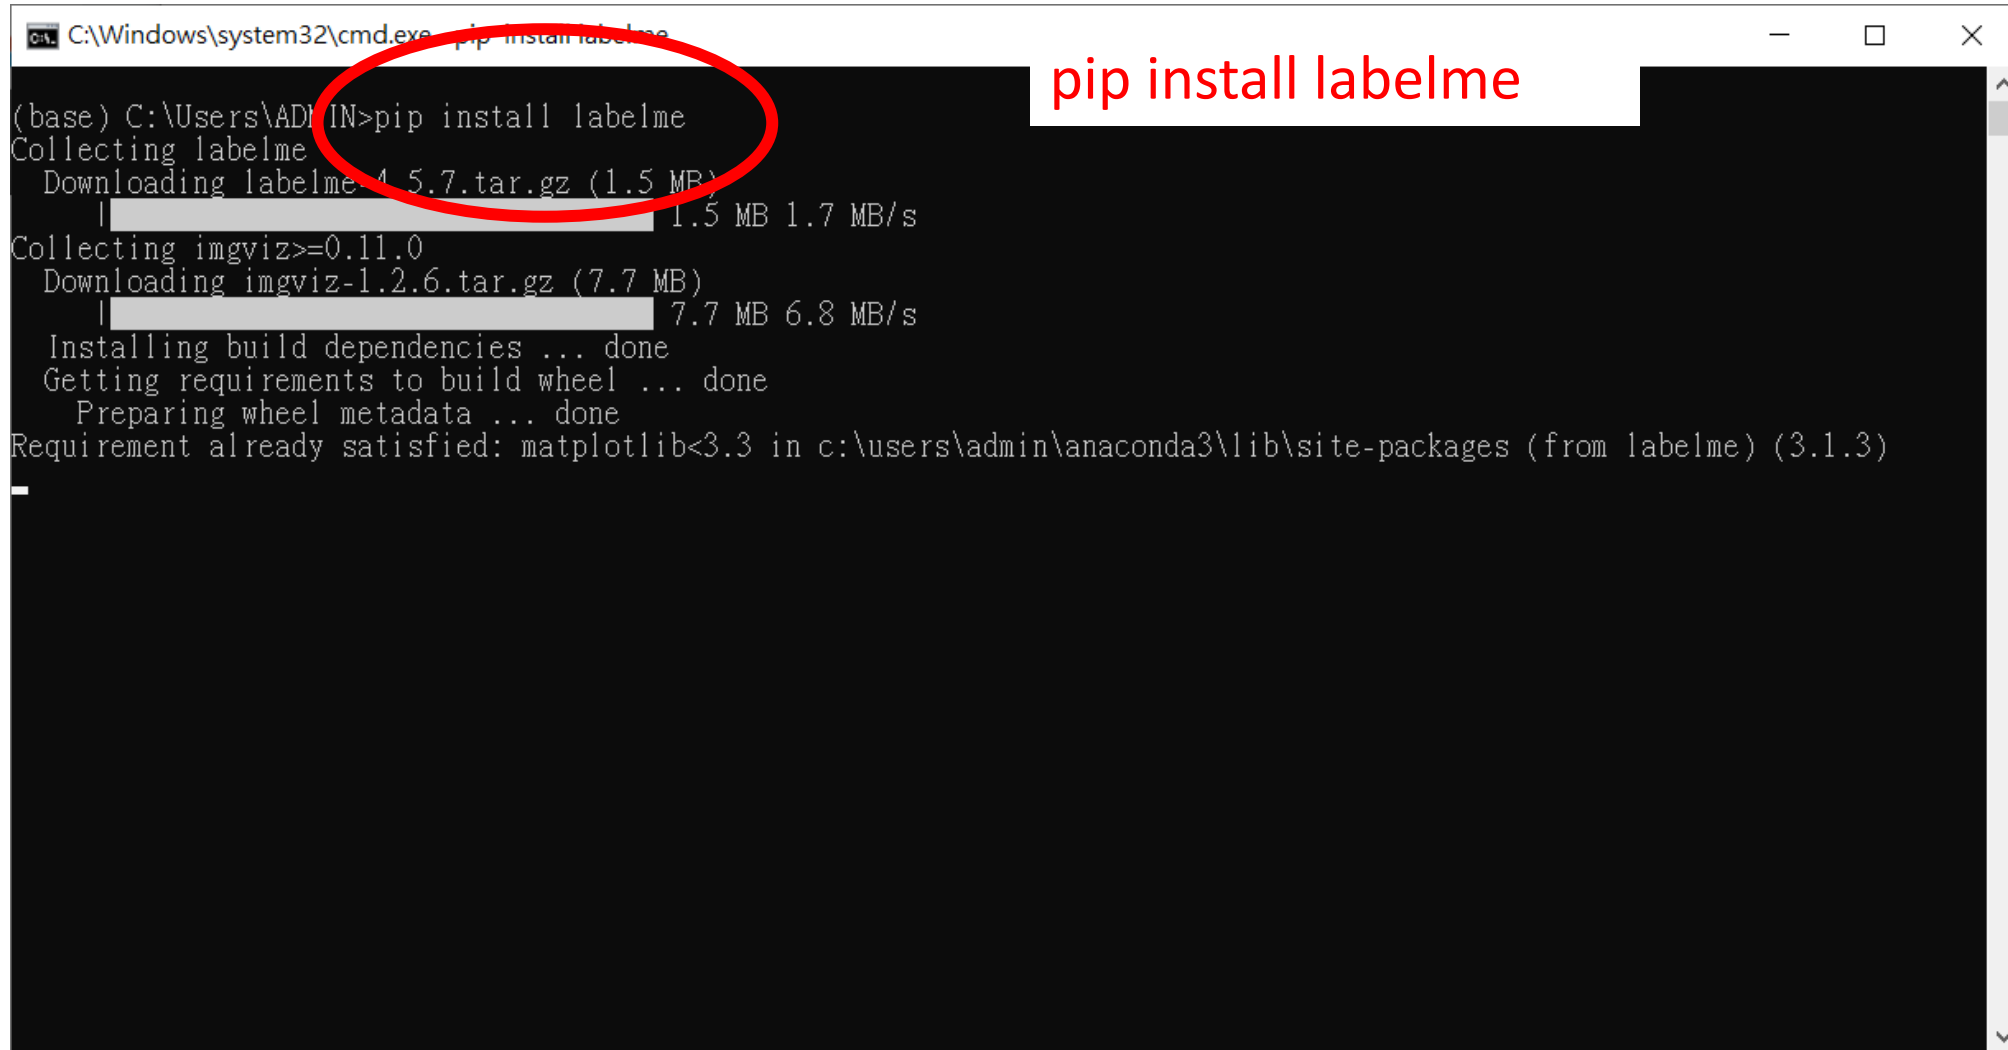
# replace the pre-trained head with a new one
model.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)
```

Train

Class practice – Train your own object detector

1. Select an object that is not in the COCO dataset and collect 10 pictures that contain this object.
2. Use LabelMe to label the object and save to json. Then convert json files to mask images.
3. Run "FasterRCNN(3) Fine_tune.ipynb" to fine tune the pre-trained FasterRCNN to train a customized NN that can recognize your own object.

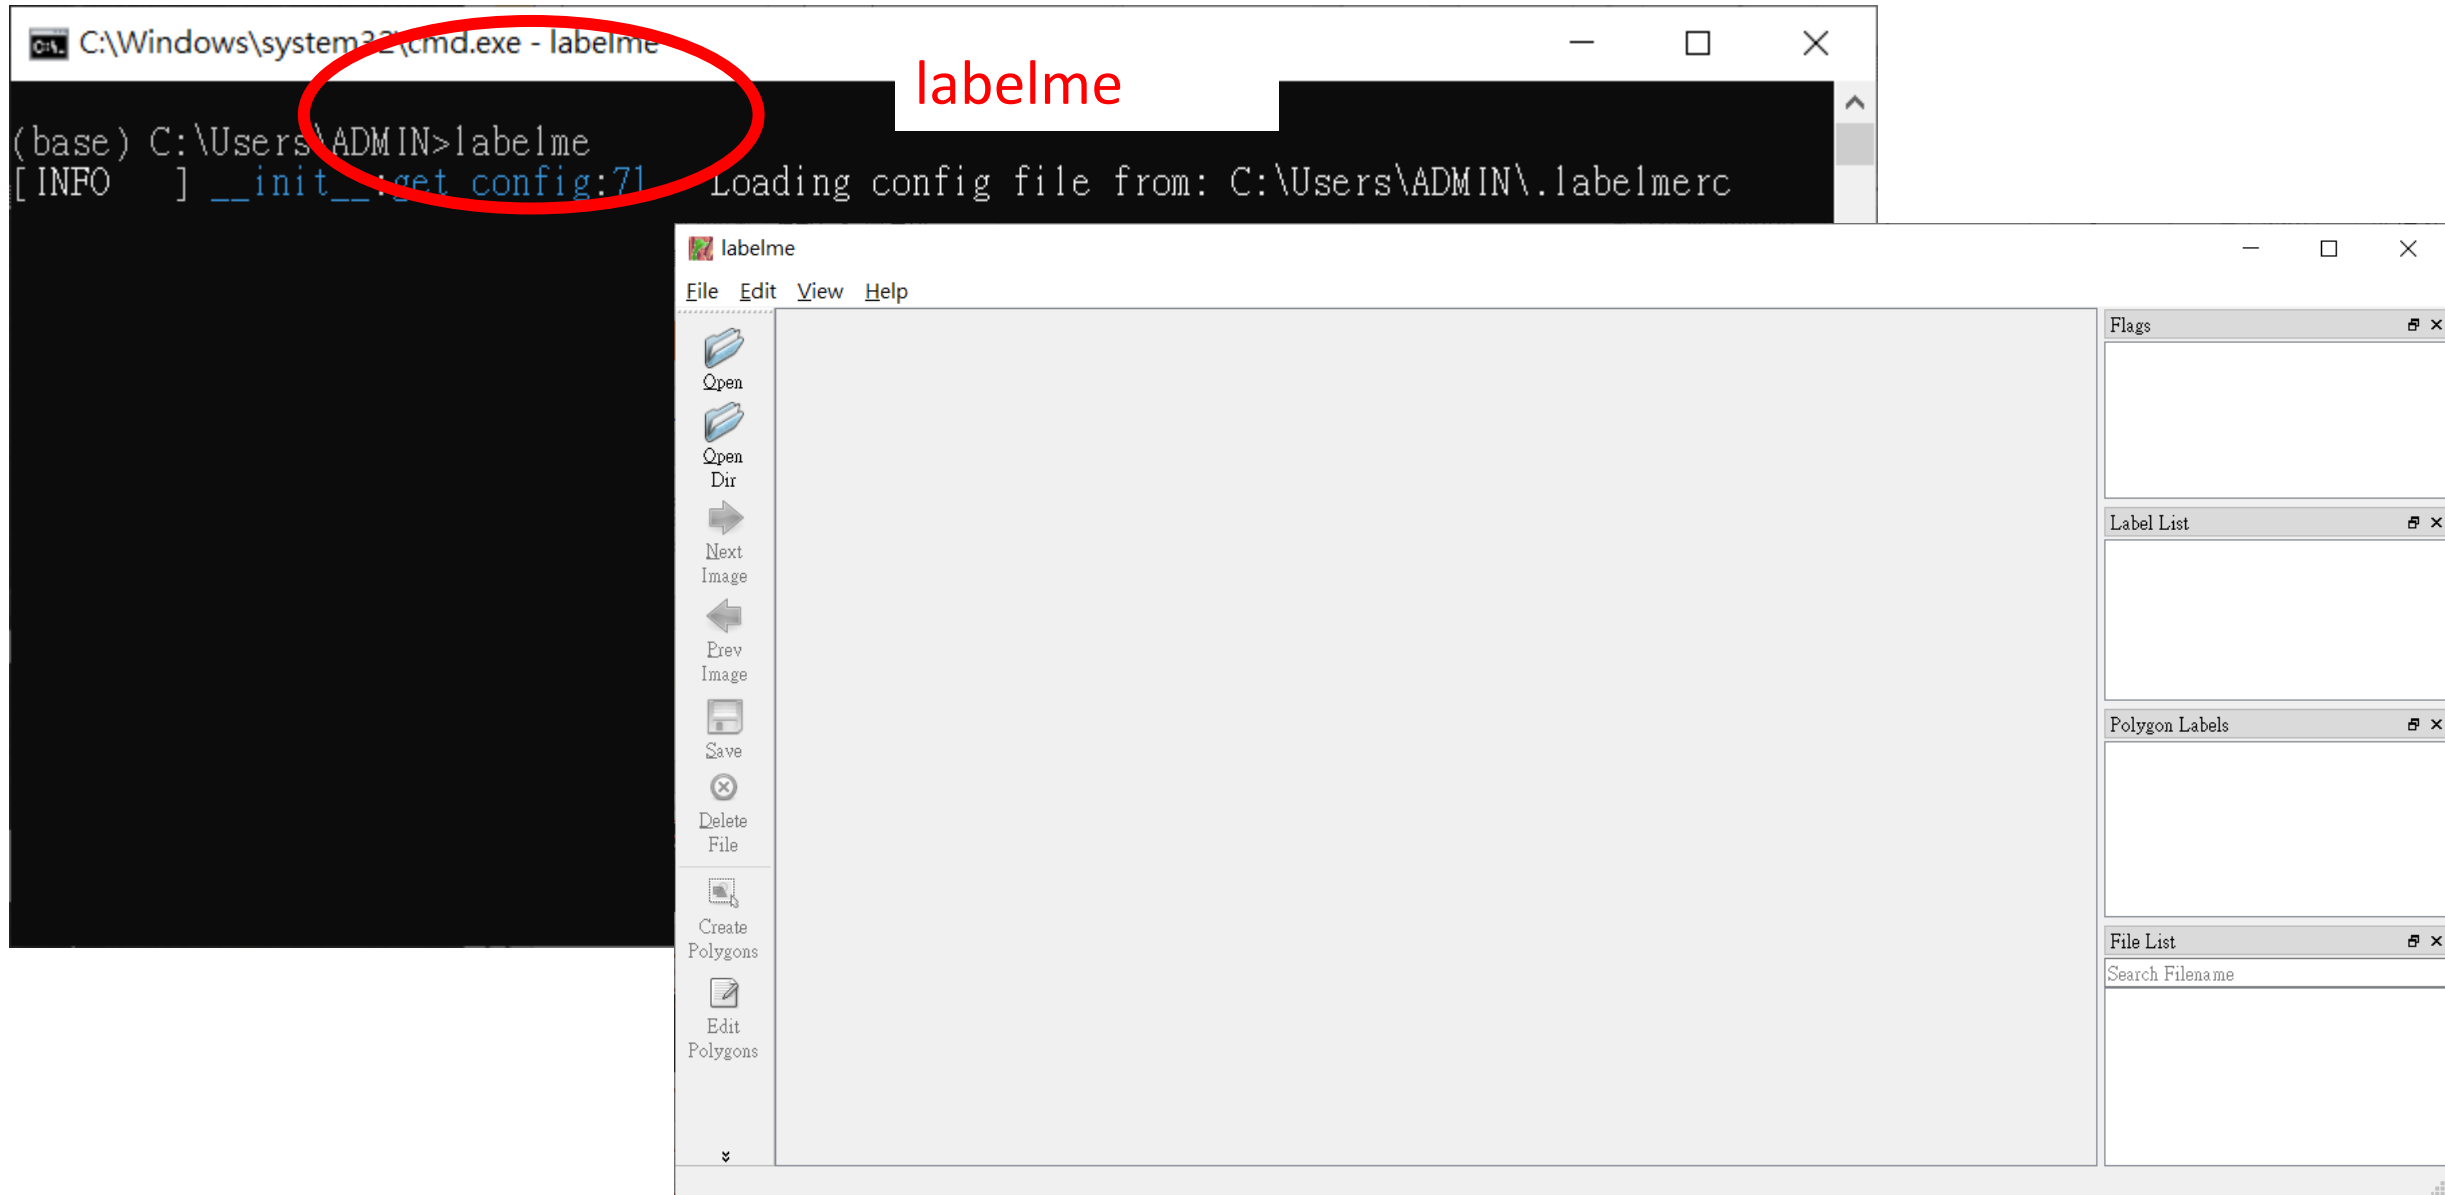
pip install labelme in your Anaconda environment



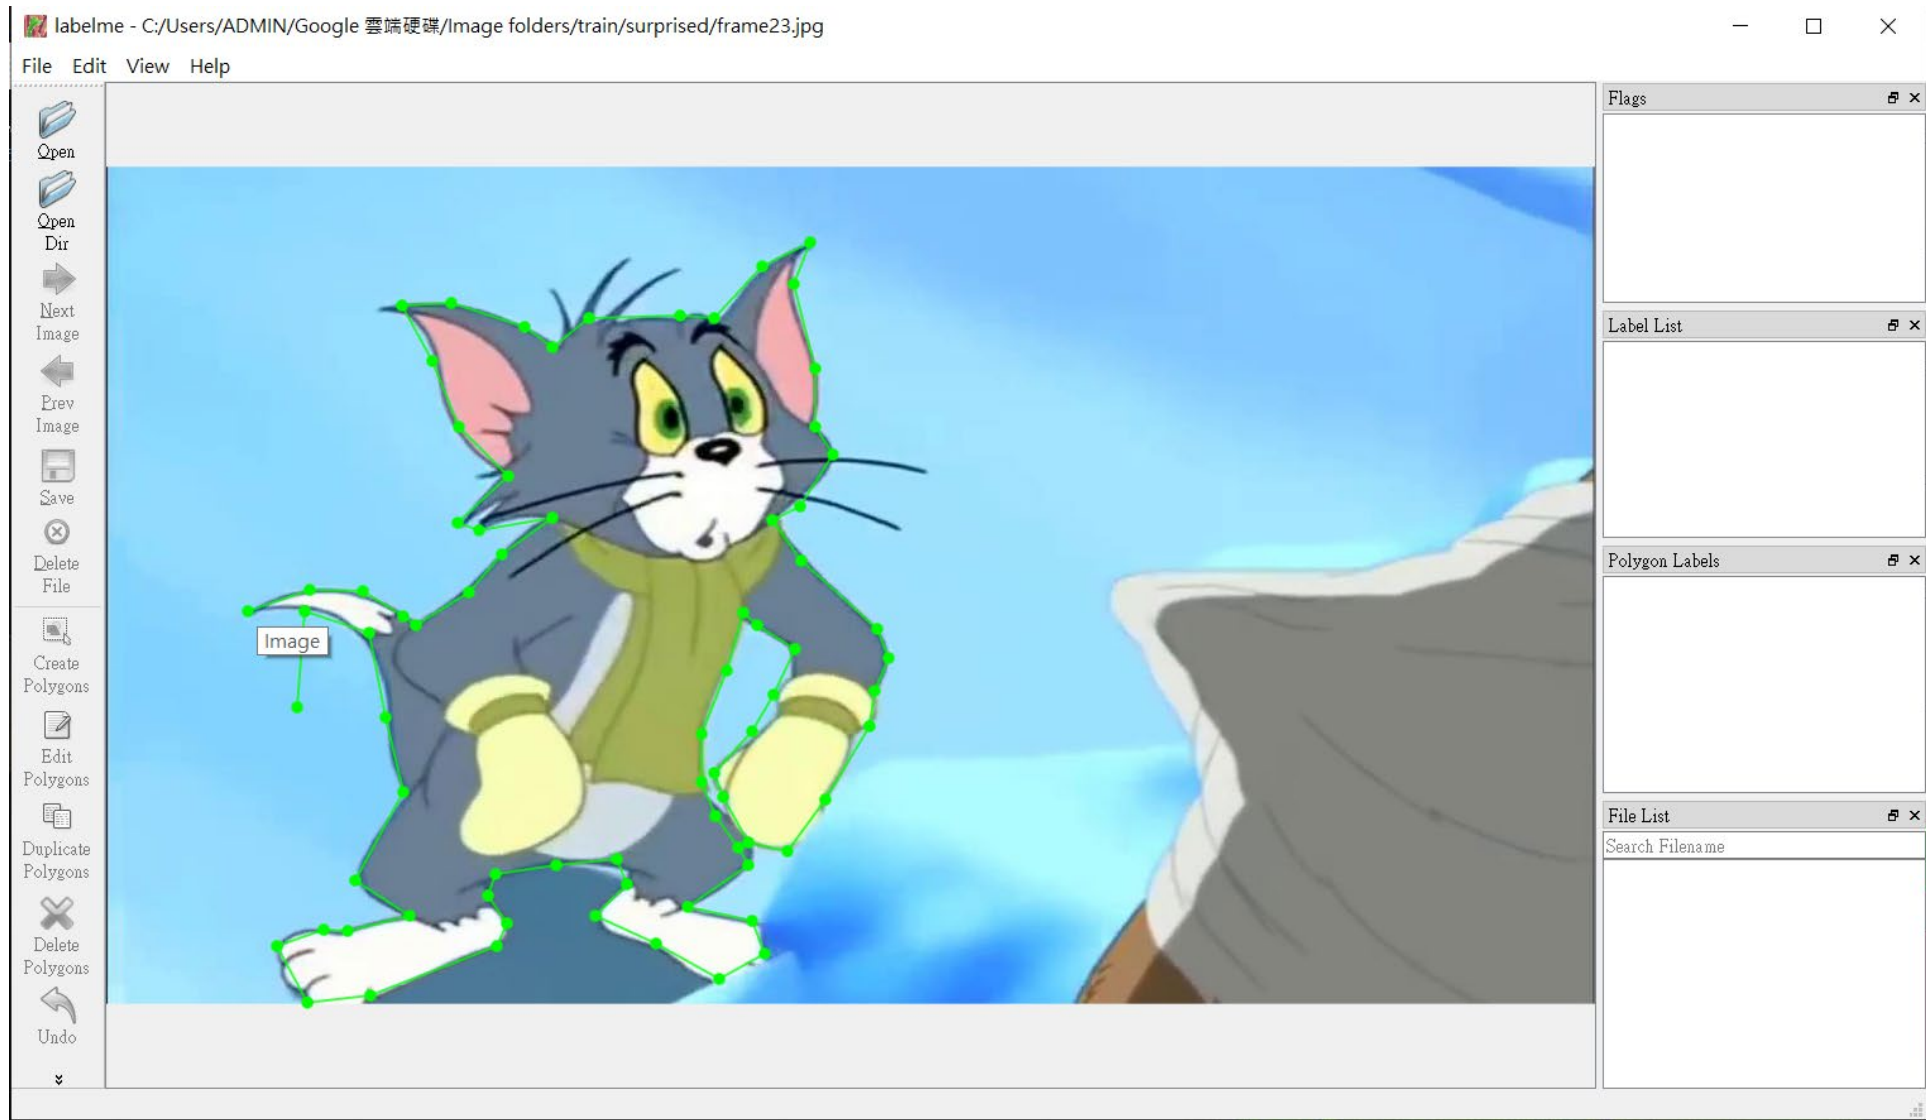
```
C:\Windows\system32\cmd.exe - pip install labelme

(base) C:\Users\ADMIN>pip install labelme
Collecting labelme
  Downloading labelme-4.5.7.tar.gz (1.5 MB)
    | 1.5 MB 1.7 MB/s
Collecting imgviz>=0.11.0
  Downloading imgviz-1.2.6.tar.gz (7.7 MB)
    | 7.7 MB 6.8 MB/s
Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing wheel metadata ... done
Requirement already satisfied: matplotlib<3.3 in c:\users\admin\anaconda3\lib\site-packages (from labelme) (3.1.3)
```

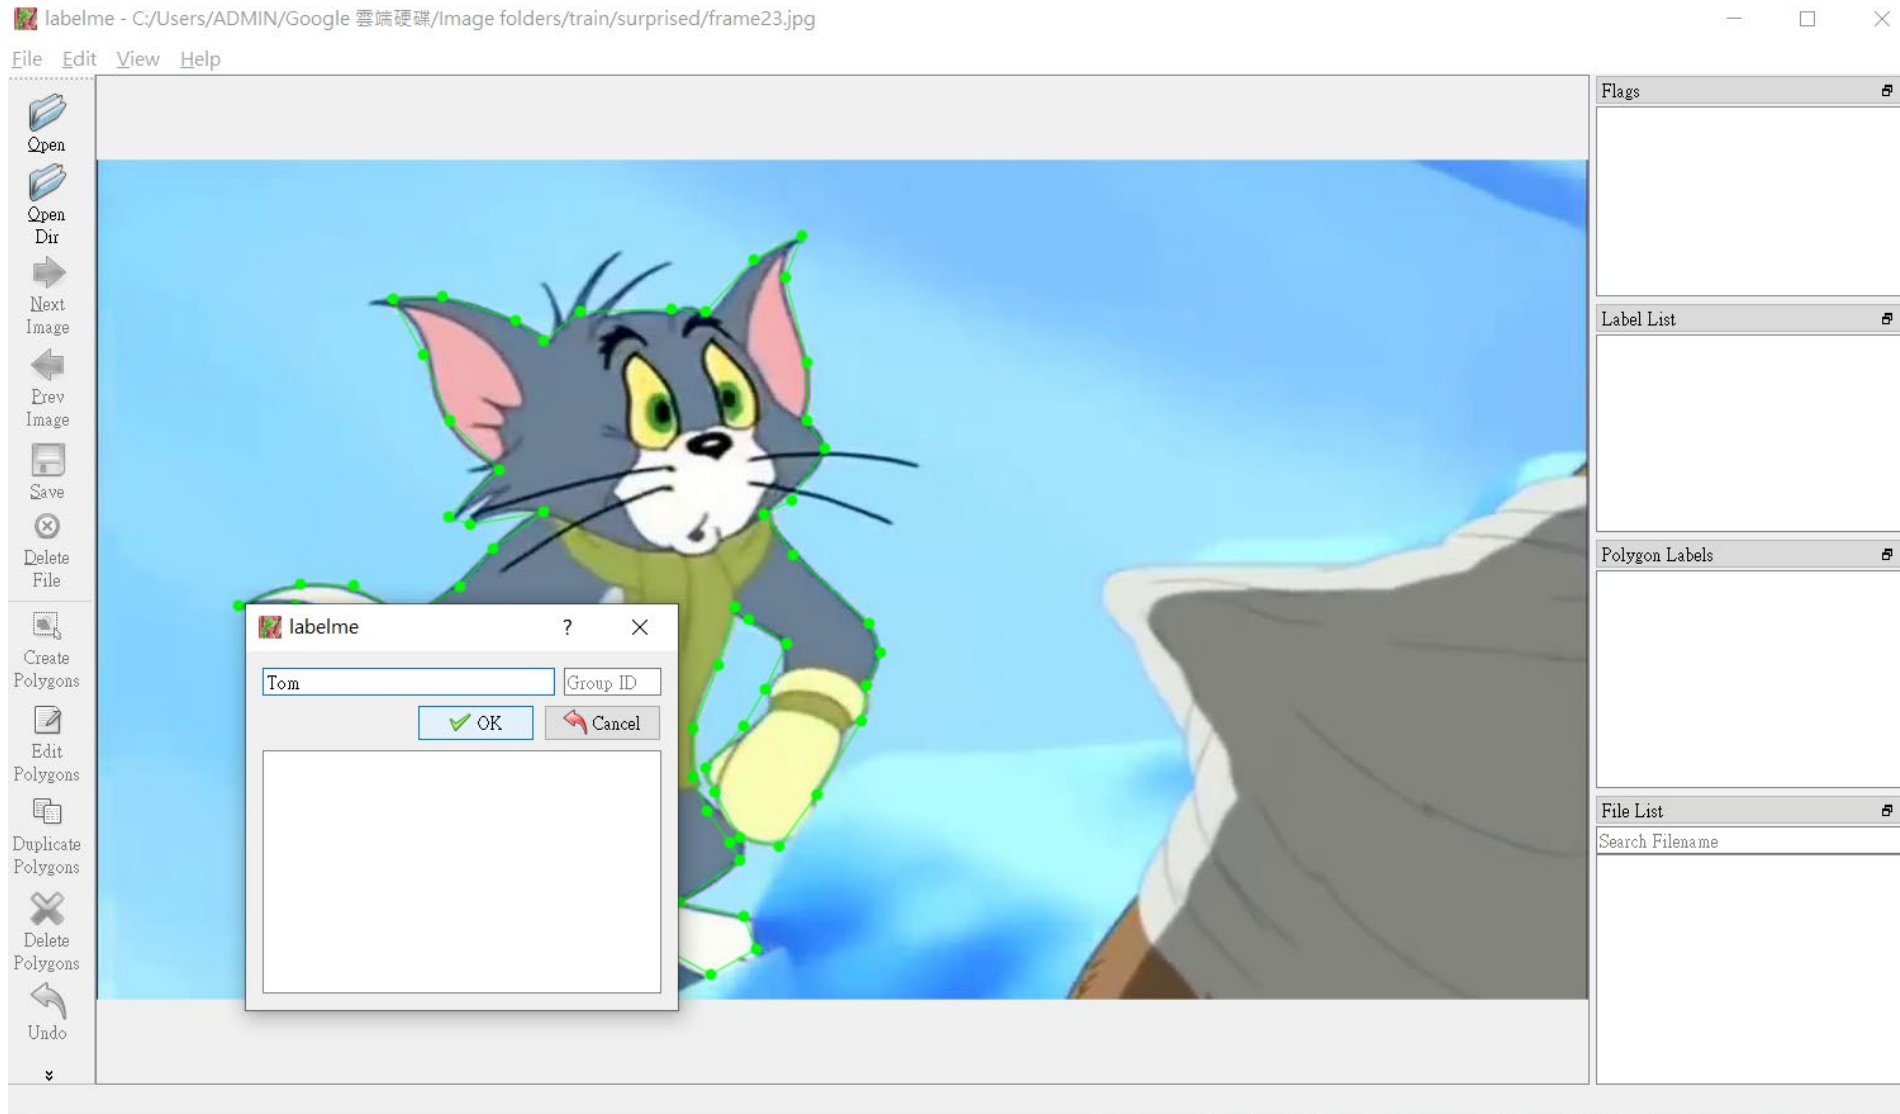
Run labelme



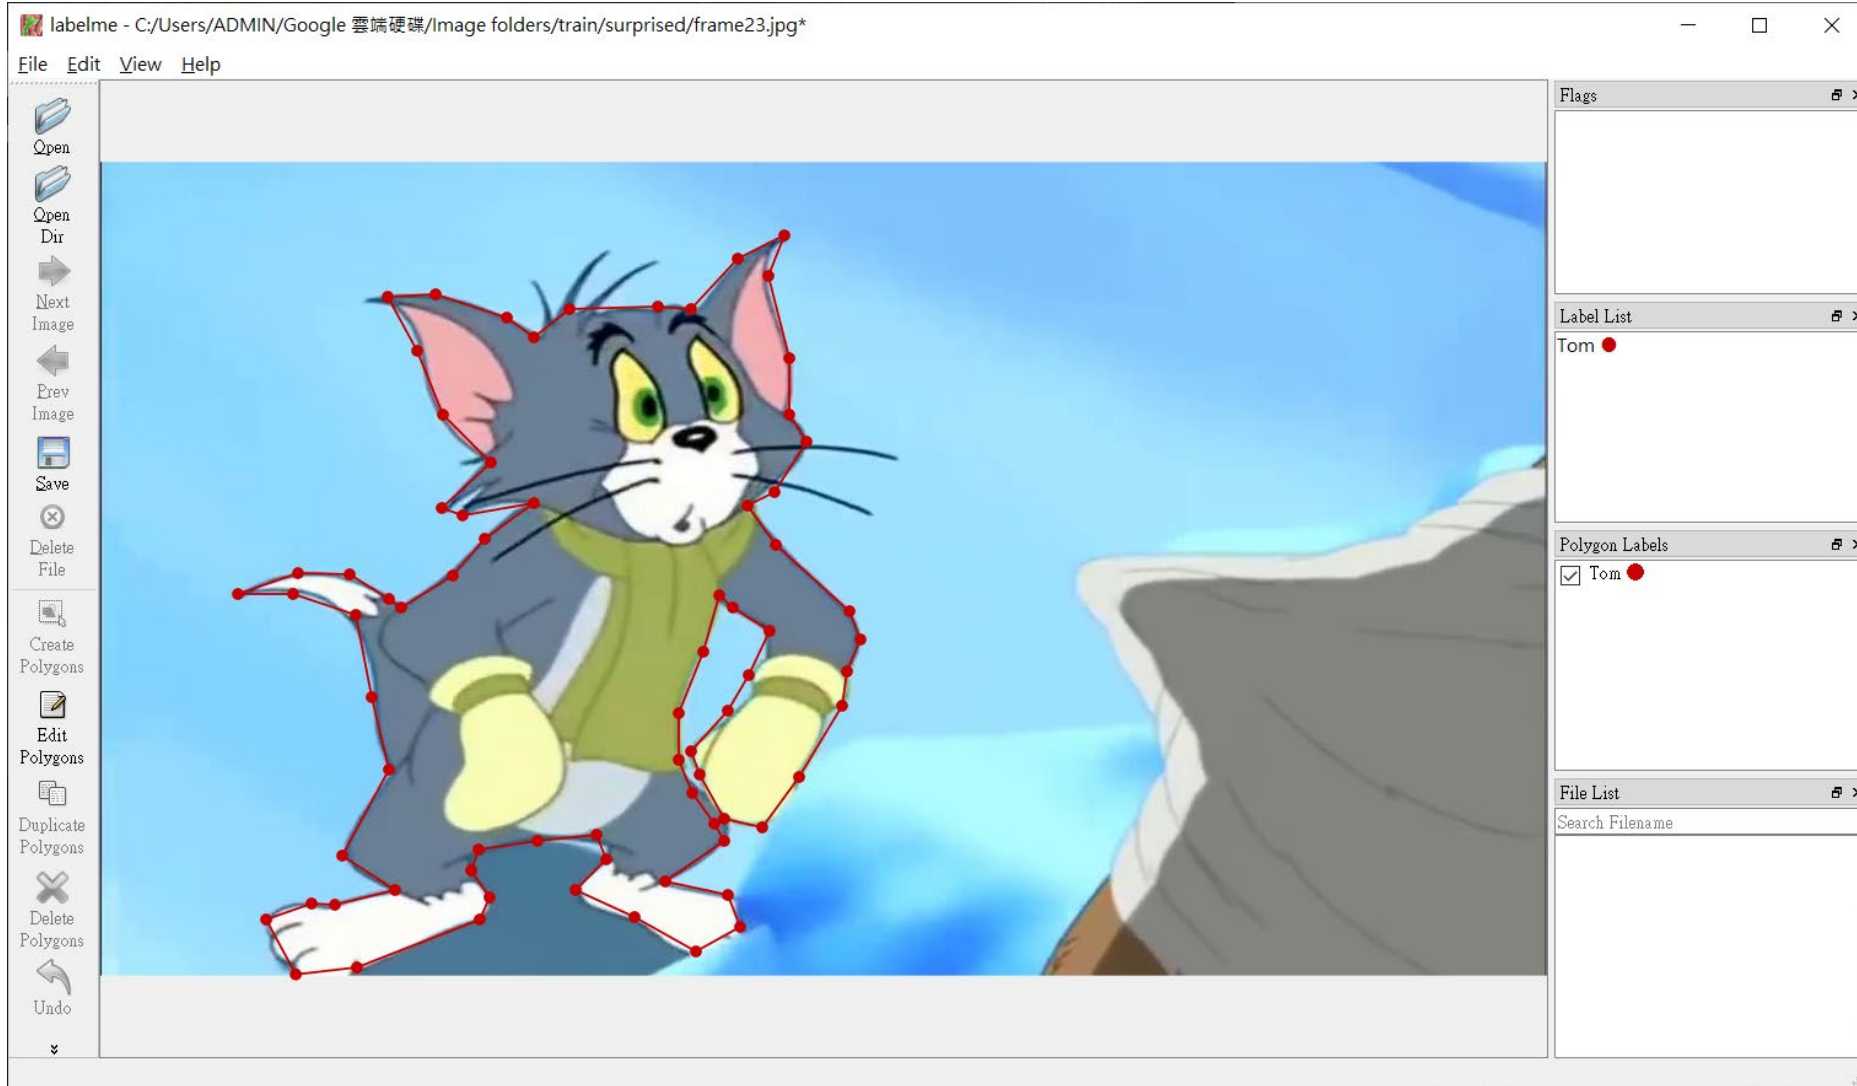
Load an image and draw boundary



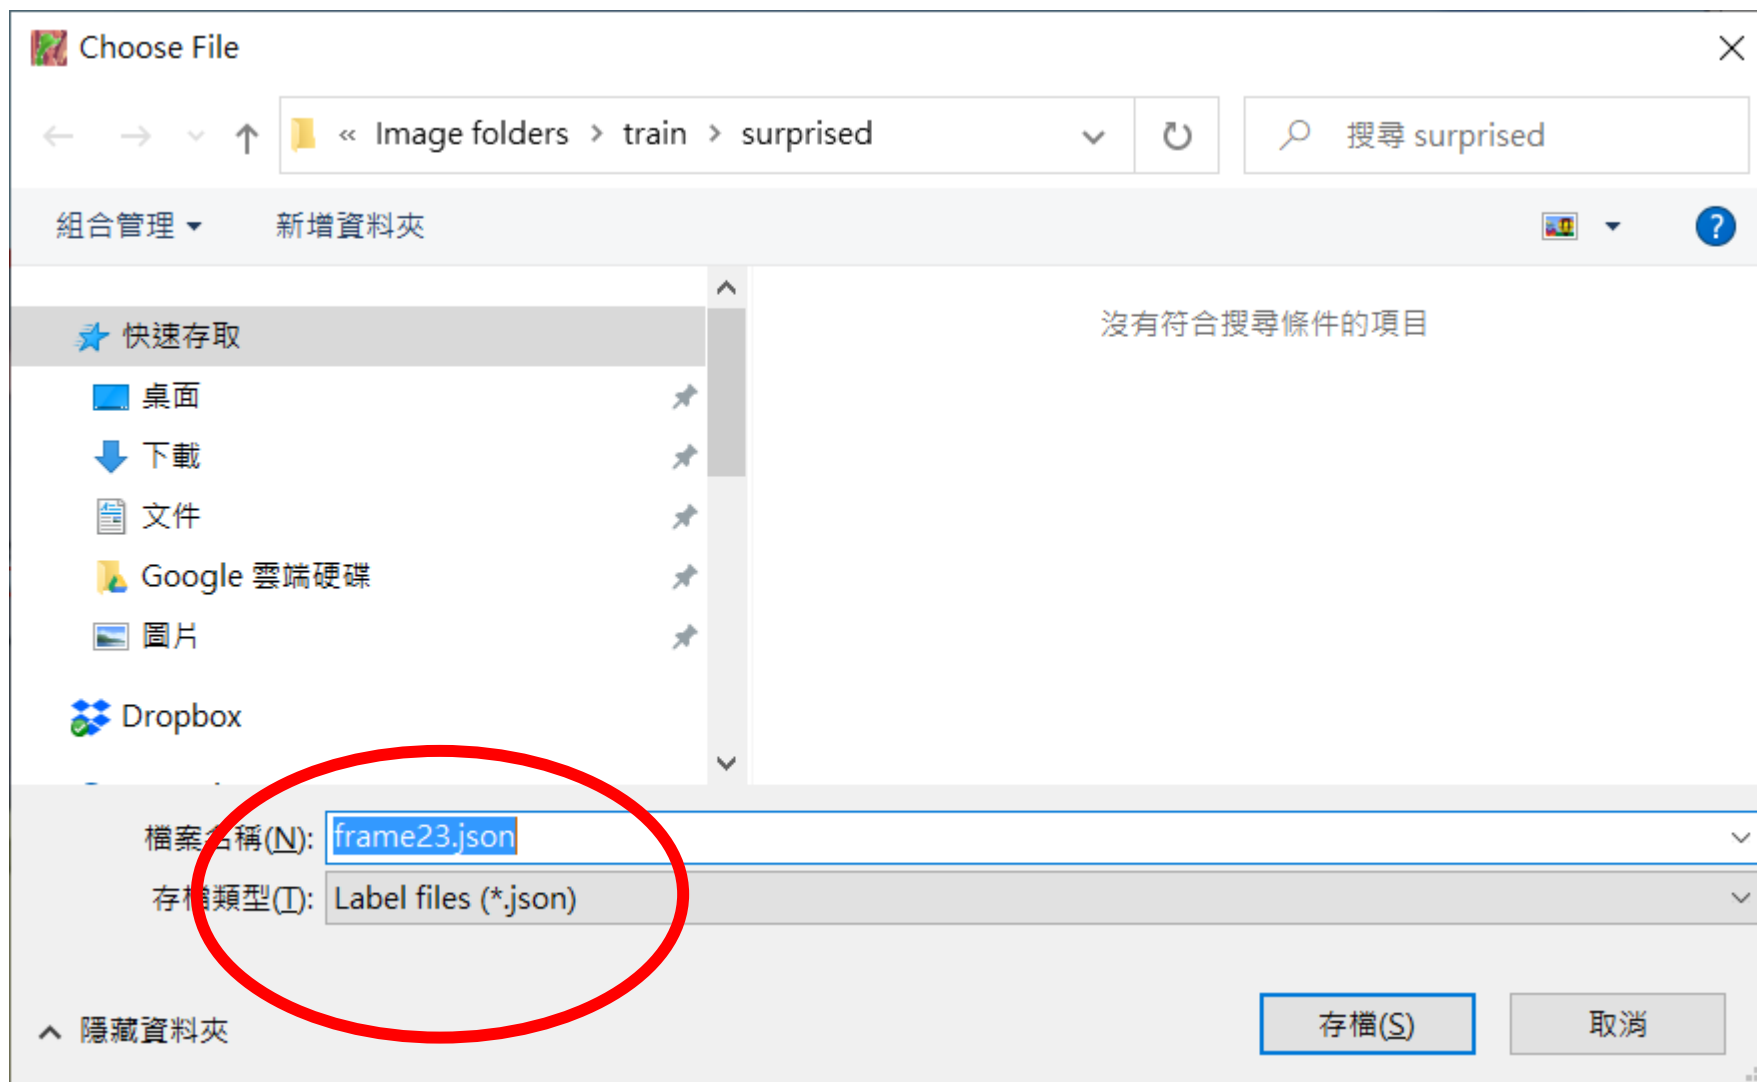
Save label



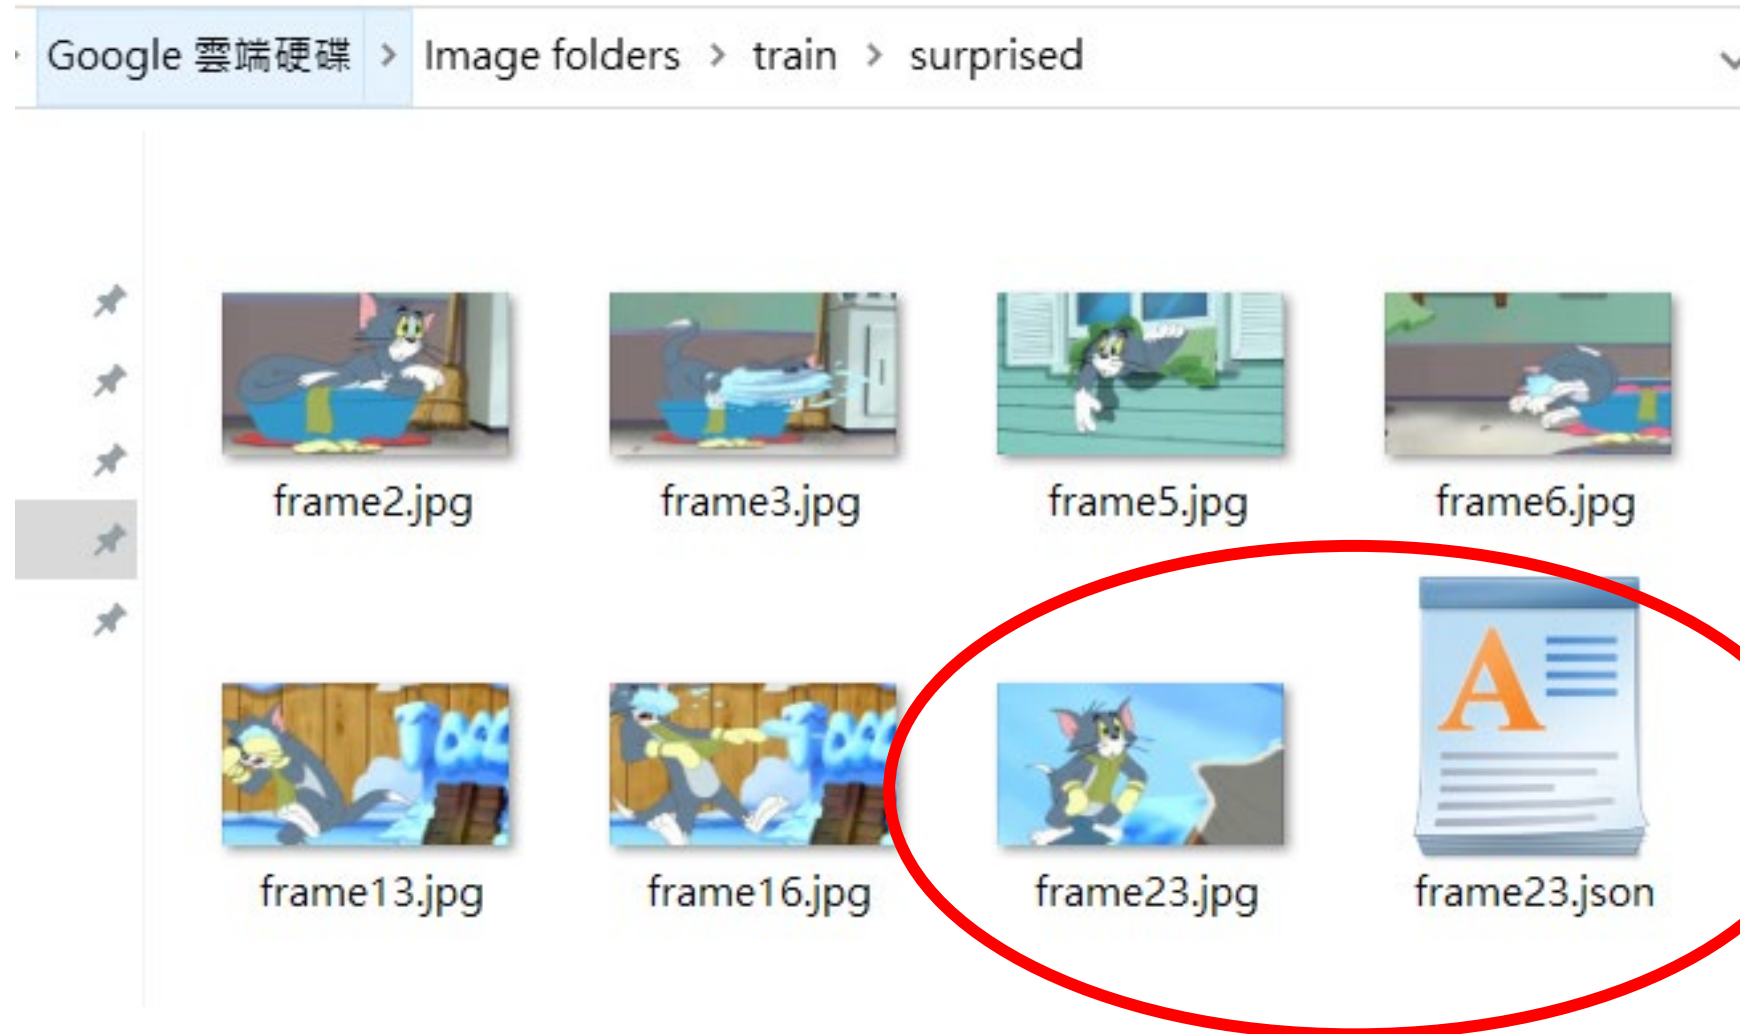
Saved label



Save boundary to json file



Saved json file

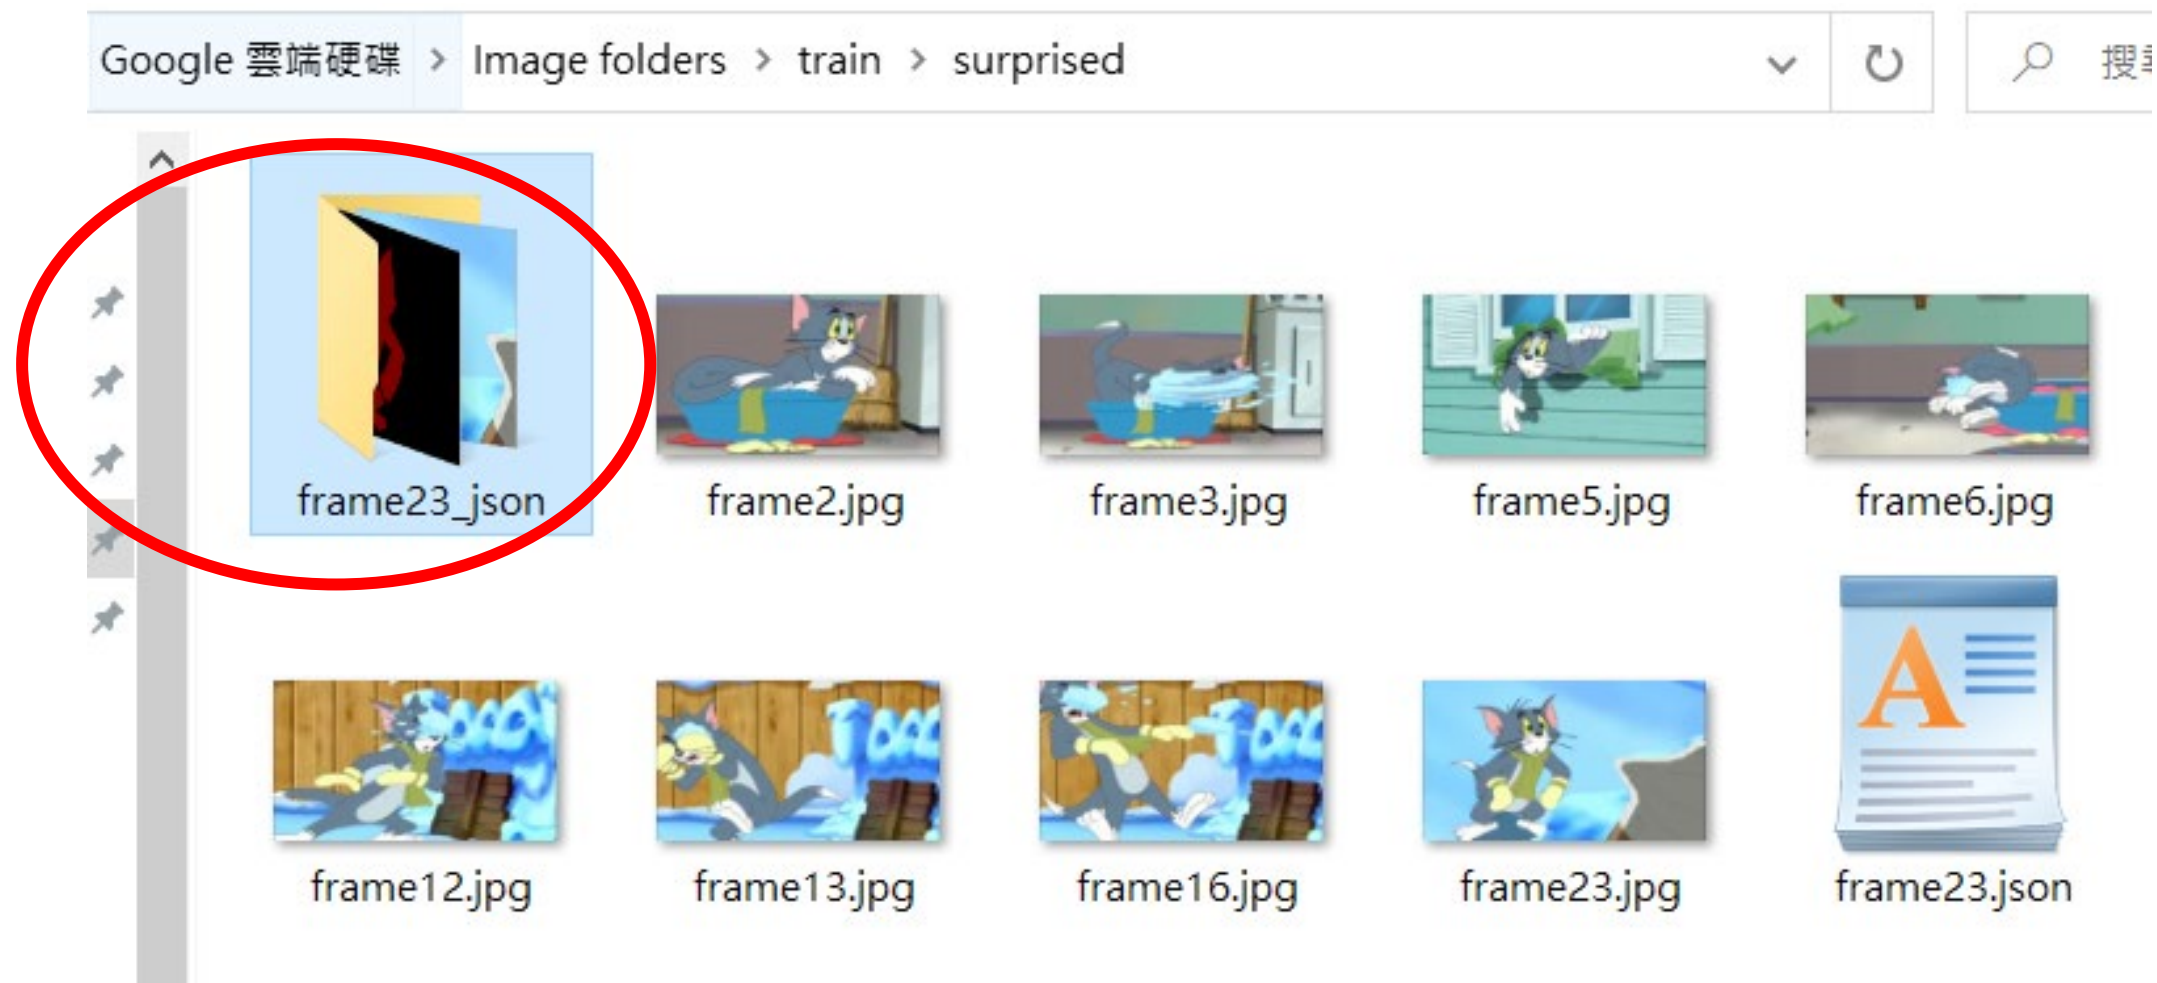


Convert json file to mask image

cd to the folder where you save the *.json file
Labelme_json_to_dataset *.json

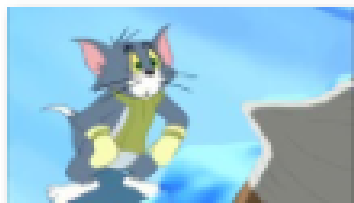
```
(base) C:\Users\ADMIN>cd C:\Users\ADMIN\Google 雲端硬碟\Image folders\train\surprised  
(base) C:\Users\ADMIN\Google 雲端硬碟\Image folders\train\surprised>labelme_json_to_dataset frame23.json  
[WARNING] json_to_dataset:main:16 - This script is aimed to demonstrate how to convert the JSON file to a single image dataset.  
[WARNING] json_to_dataset:main:20 - It won't handle multiple JSON files to generate a real-use dataset.  
[INFO] json_to_dataset:main:77 - Saved to: frame23.json  
(base) C:\Users\ADMIN\Google 雲端硬碟\Image folders\train\surprised>
```

Mask images are saved in a folder



Mask image

雲端硬碟 > Image folders > train > surprised > frame23_json



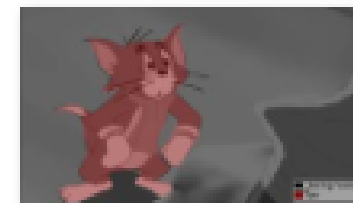
img.png



label.png



label_names.txt



label_viz.png

label_names.txt - 記事本
檔案(F) 編輯(E) 格式(O) 檢視(V) 說明
background
Tom

Save RGB and mask images on your Google drive

My Drive > Object Detection Image Folder ▾

Name



mask



pic

My Drive > Object Detection Image Folder > pic ▾

Files



0001.jpg



0002.jpg



0003.jpg



0005.jpg



0006.jpg



0007.jpg

Save RGB and mask images on your Google drive

My Drive > Object Detection Image Folder > mask ▾

Files



 0001.png



 0002.png



 0003.png



 0005.png



 0006.png



 0007.png

Split training and test set

```
# split the dataset in train and test set
```

```
torch.manual_seed(1)
```

```
indices = torch.randperm(len(dataset)).tolist()
```

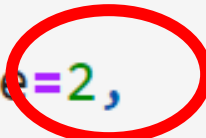
```
dataset = torch.utils.data.Subset(dataset, indices[:-50])
```

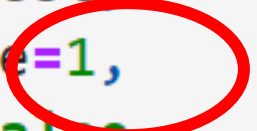
```
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])
```

I have 89 labelled images, of which 50 were used as training images and the remaining 39 used as test.

Batch size

```
# define training and validation data loaders
```

```
data_loader = torch.utils.data.DataLoader(  
    dataset,  
    batch_size=2,  Try larger batch size  
    shuffle=True,  
    num_workers=4,  
    collate_fn=utils.collate_fn)
```

```
data_loader_test = torch.utils.data.DataLoader(  
    dataset_test,  
    batch_size=1,   
    shuffle=False,  
    num_workers=4,  
    collate_fn=utils.collate_fn)
```

Training performance evaluation

$$L(p_i, t_i) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, \hat{p}_i) + \lambda \frac{1}{N_{reg}} \sum_i p_i L_{reg}(t_i, \hat{t}_i)$$

```
Epoch: [0] [ 0/20] eta: 0:01:04 lr: 0.000268 loss: 3.3367 (3.3367) loss_classifier: 0.7079 (0.7079)
loss_box_reg: 0.1177 (0.1177) loss_mask: 2.4933 (2.4933) loss_objectness: 0.0119 (0.0119) loss_rpn_bo
x_reg: 0.0060 (0.0060) time: 3.2353 data: 2.5175 max mem: 2483
Epoch: [0] [10/20] eta: 0:00:07 lr: 0.002897 loss: 1.3361 (1.7297) loss_classifier: 0.1779 (0.2886)
loss_box_reg: 0.0578 (0.0657) loss_mask: 1.0733 (1.2952) loss_objectness: 0.0762 (0.0733) loss_rpn_bo
x_reg: 0.0060 (0.0068) time: 0.7238 data: 0.3247 max mem: 2759
Epoch: [0] [19/20] eta: 0:00:00 lr: 0.005000 loss: 0.8906 (1.2921) loss_classifier: 0.1090 (0.2024)
loss_box_reg: 0.0579 (0.0652) loss_mask: 0.6534 (0.9624) loss_objectness: 0.0416 (0.0560) loss_rpn_bo
x_reg: 0.0045 (0.0061) time: 0.6082 data: 0.2312 max mem: 2759
Epoch: [0] Total time: 0:00:12 (0.6111 s / it)
```

Testing performance evaluation

- To evaluate object detection models like FasterRCNN and YOLO, the mean average precision (mAP) is used. The mAP compares the ground-truth bounding box to the detected box and returns a score. The higher the score, the more accurate the model is in its detections.


Precision and recall

Table 3 Confusion matrix

	Reference (high-risk)	Reference (low-risk)
Predicted (high-risk)	True positive (TP)	False positive (FP)
Predicted (low-risk)	False negative (FN)	True negative (TN)

$$Recall = \frac{TP}{TP + FN}$$

“ When a model has high recall but low precision, then the model classifies most of the positive samples correctly but it has many false positives (i.e. classifies many Negative samples as Positive). When a model has high precision but low recall, then the model is accurate when it classifies a sample as Positive but it may classify only some of the positive samples.




<https://blog.paperspace.com/mean-average-precision/>

Average precision

The **average precision (AP)** is a way to summarize the precision-recall curve into a single value representing the average of all precisions.

$$AP = \sum_{k=0}^{k=n-1} [Recalls(k) - Recalls(k+1)] * Precisions(k)$$



$Recalls(n) = 0, Precisions(n) = 1$
 $n = \text{Number of thresholds.}$



AP is the weighted sum of precisions at each threshold where the weight is the increase in recall.

mAP (mean average precision)

$$mAP = \frac{1}{n} \sum_{k=1}^n AP_k \quad n = \text{number of classes}$$





HW4 – Object detector

- Fine-tune pre-trained FasterRCNN to detect your own objects.
- Show training loss plot.
- Show detection results.


Automatically labelled photo-realistic images

Accelerate computer vision model training with the synthetic image data generated using Unity's perception package



2D bounding boxes




3D bounding boxes



Class segmentation



Instance segmentation



<https://unity.com/products/computer-vision>

