

Generative Adversarial Network (GAN)

Practice

- Open "8.1. GAN.ipynb"



Generator takes a feature vector and generates an output image (using deconvolution to do up-sampling)

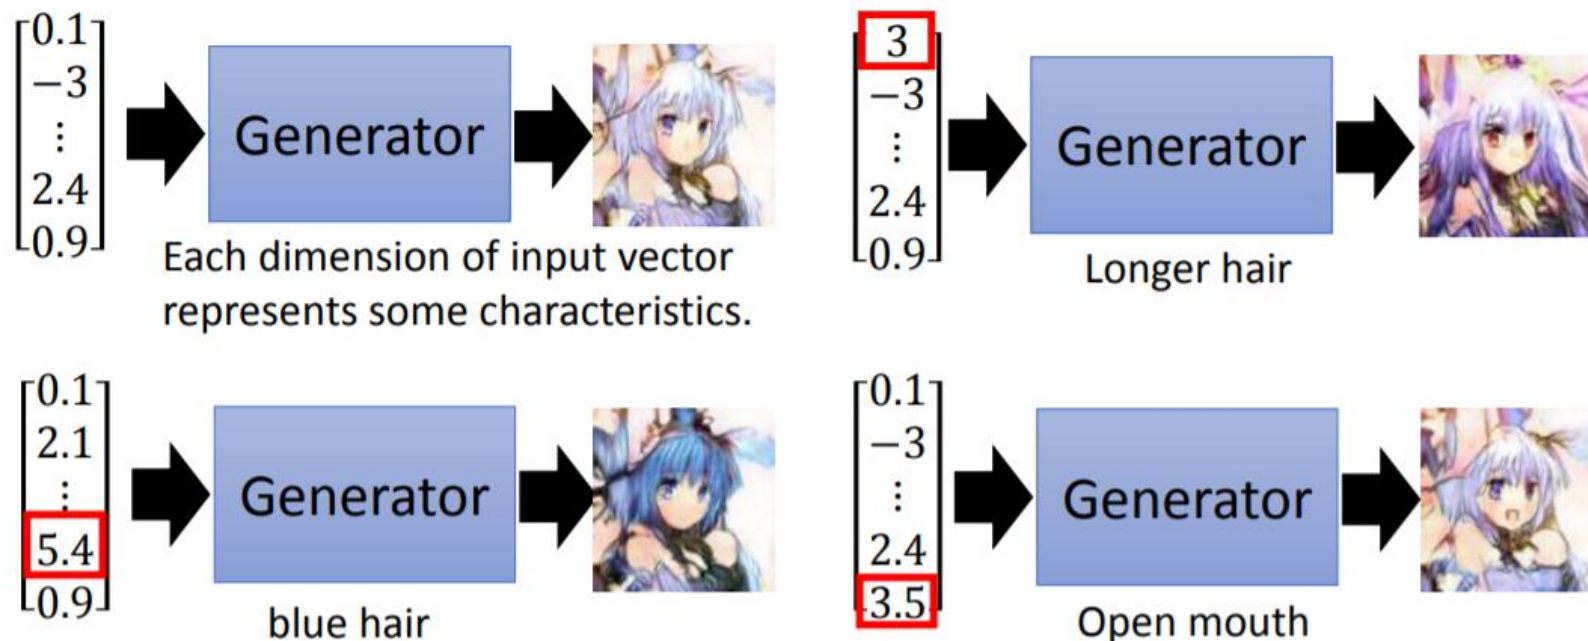
```
[10]: latent_size= 128
```

```
[11]: generator = nn.Sequential(  
    # in: latent_size x 1 x 1  
  
    nn.ConvTranspose2d(latent_size, 512, kernel_size=4, stride=1, padding=0, bias=False),  
    nn.BatchNorm2d(512),  
    nn.ReLU(True),  
    # out: 512 x 4 x 4  
  
    nn.ConvTranspose2d(512, 256, kernel_size=4, stride=2, padding=1, bias=False),  
    nn.BatchNorm2d(256),  
    nn.ReLU(True),  
    # out: 256 x 8 x 8  
  
    nn.ConvTranspose2d(256, 128, kernel_size=4, stride=2, padding=1, bias=False),  
    nn.BatchNorm2d(128),  
    nn.ReLU(True),  
    # out: 128 x 16 x 16  
  
    nn.ConvTranspose2d(128, 64, kernel_size=4, stride=2, padding=1, bias=False),  
    nn.BatchNorm2d(64),  
    nn.ReLU(True),  
    # out: 64 x 32 x 32
```

Each dimension in the feature vector represent a property of the output image

```
[10]: latent_size= 128
```

```
[11]: generator = nn.Sequential(  
    # in: latent_size x 1 x 1  
  
    nn.ConvTranspose2d(latent_size, 512, kernel_size=4, stride=1, padding=0, bias=False),  
    nn.BatchNorm2d(512),
```



Discriminator takes an input image and tells whether it is a true image or not

Larger value means real, smaller value means fake

```
[15]: discriminator = nn.Sequential(  
    # in: 3 x 128 x 128  
    nn.Conv2d(3, 64, kernel_size=4, stride=2, padding=1, bias=False),  
    nn.BatchNorm2d(64),  
    nn.LeakyReLU(0.2, inplace=True),  
    # out: 64 x 64 x 64  
  
    nn.Conv2d(64, 128, kernel_size=4, stride=2, padding=1, bias=False),  
    nn.BatchNorm2d(128),  
    nn.LeakyReLU(0.2, inplace=True),  
    # out: 128 x 32 x 32  
  
    nn.Conv2d(128, 256, kernel_size=4, stride=2, padding=1, bias=False),  
    nn.BatchNorm2d(256),  
    nn.LeakyReLU(0.2, inplace=True),  
    # out: 256 x 16 x 16  
  
    nn.Conv2d(256, 512, kernel_size=4, stride=2, padding=1, bias=False),  
    nn.BatchNorm2d(512),  
    nn.LeakyReLU(0.2, inplace=True),  
    # out: 512 x 8 x 8  
  
    nn.Conv2d(512, 1024, kernel_size=4, stride=2, padding=1, bias=False),  
    nn.BatchNorm2d(1024),  
    nn.LeakyReLU(0.2, inplace=True),  
    # out: 1024 x 4 x 4  
  
    nn.Conv2d(1024, 1, kernel_size=4, stride=1, padding=0, bias=False)  
    # out: 1 x 1 x 1
```



Reference: 李弘毅 GAN Lecture 1 (2018)

Step1 – Fix G and train D

- Initialize generator and discriminator

G

D

[12]: `generator.to(device)`

- In each training iteration:

[16]: `discriminator.to(device)`

Step 1: Fix generator G, and update discriminator D



Database

samp

gene

ok

randomly
sampled

vector

[35]:

```
for epoch in range(epochs):
    if(epoch % 10 == 0):
        print(epoch, end=",")
    for real_images, _ in train_d1:
        # Train discriminator
        loss_d, real_score, fake_score = train_discriminator(real_images)
        # Train generator
        loss_g = train_generator(opt_g)
```


Train discriminator

- Sample m examples $\{x^1, x^2, \dots, x^m\}$ from database
- Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from a distribution
- Obtaining generated data $\{\tilde{x}^1, \tilde{x}^2, \dots, \tilde{x}^m\}$, $\tilde{x}^i = G(z^i)$

Update discriminator parameters θ_d to maximize


- $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log D(x^i) + \frac{1}{m} \sum_{i=1}^m \log(1 - D(\tilde{x}^i))$
- $\theta_d \leftarrow \theta_d + \eta \nabla \tilde{V}(\theta_d)$

```
[19]: def train_discriminator(real_images, opt_d):  
    # Clear discriminator gradients  
    opt_d.zero_grad()  
  
    # Pass real images through discriminator  
    real_preds = discriminator(real_images)  
    real_targets = torch.ones(real_images.size(0))  
    real_loss = F.binary_cross_entropy(real_preds, real_targets)  
    real_score = torch.mean(real_preds).item()  
  
    # Generate fake images  
    latent = torch.randn(batch_size, latent_size, device=device)  
    fake_images = generator(latent.to(device))  
  
    # Pass fake images through discriminator  
    fake_targets = torch.zeros(fake_images.size(0))  
    fake_preds = discriminator(fake_images)  
    fake_loss = F.binary_cross_entropy(fake_preds, fake_targets)  
    fake_score = torch.mean(fake_preds).item()  
  
    # Update discriminator weights  
    loss = real_loss + fake_loss  
    loss.backward()  
    opt_d.step()  
    return loss.item(), real_score, fake_score
```

Step2 – Fix D and train G

Step 2: Fix discriminator D, and update generator G

Generator learns to “fool” the discriminator



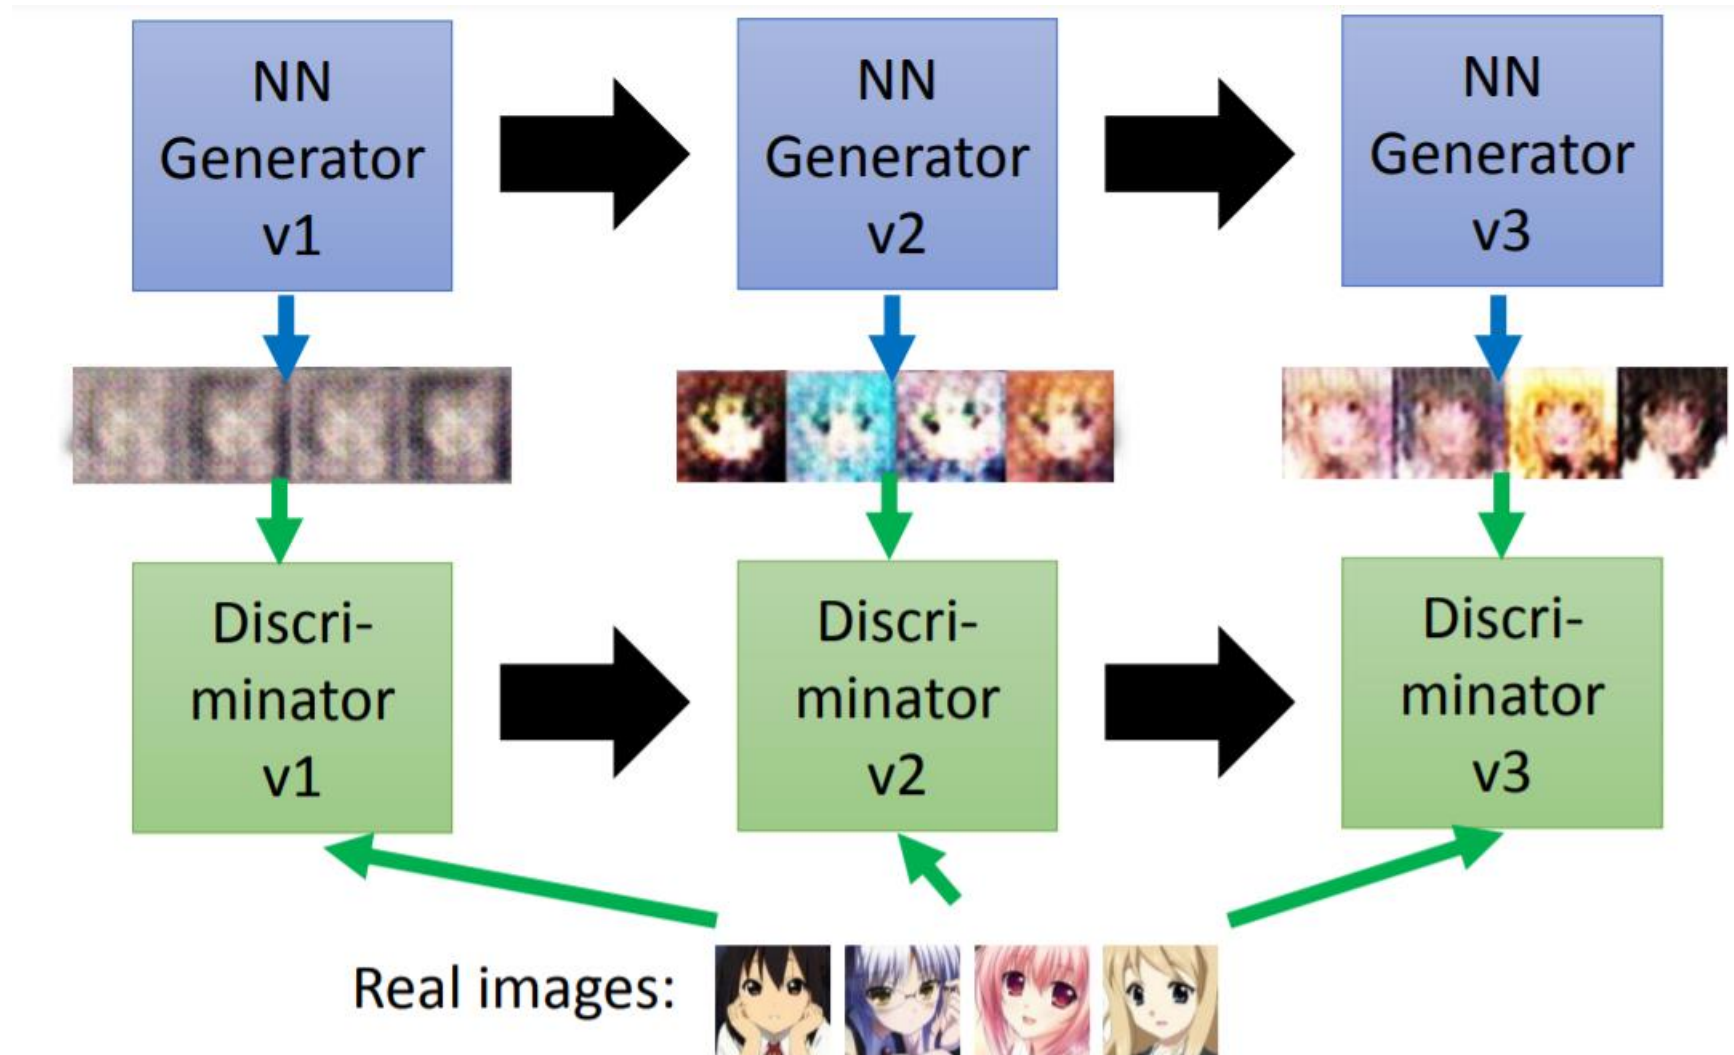
```
[35]: for epoch in range(epochs):
        if(epoch % 10 ==0):
            print(epoch, end=",")
        for real_images, _ in train_d1:
            # Train discriminator
            loss_d, real_score, fake_score = train_discriminator(real_images)
            # Train generator
            loss_g = train_generator(opt_g)
```


Train generator

-
- Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from a distribution
 - Update generator parameters θ_g to maximize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log \left(D \left(G(z^i) \right) \right)$
 - $\theta_g \leftarrow \theta_g - \eta \nabla \tilde{V}(\theta_g)$

```
[28]: def train_generator(opt_g):  
    # Clear generator gradients  
    opt_g.zero_grad()  
  
    # Generate fake images  
    latent = torch.randn(batch_size, latent_size,  
        fake_images = generator(latent)  
  
    # Try to fool the discriminator  
    preds = discriminator(fake_images)  
    targets = torch.ones(batch_size, 1, device=device)  
    loss = F.binary_cross_entropy(preds, targets)  
  
    # Update generator weights  
    loss.backward()  
    opt_g.step()  
  
    return loss.item()
```

An evolution process



Visualize the fake images from G during the evolution process

```
[17]: sample_dir = 'generated'  
os.makedirs(sample_dir, exist_ok=True)
```

```
[34]: fixed_latent = torch.randn(64, latent_size,  
      # used to generate saved images
```

```
if(epoch % 50 ==0):  
    # Log losses & scores (last batch)  
    print("Epoch [{}/{}], loss_g: {:.4f}, loss_d:  
          epoch+1, epochs, loss_g, loss_d, real_score)  
    # Save generated images  
    save_samples(epoch+start_idx, fixed_latent, s
```

```
[18]: def save_samples(index, latent_tensors, show=True):  
    fake_images = generator(latent_tensors)  
    fake_fname = 'generated-images-{:0=4d}.png'.format(index)  
    save_image(denorm(fake_images), os.path.join(sample_dir, f  
    print('Saving', fake_fname)  
    if show:  
        fig, ax = plt.subplots(figsize=(8, 8))  
        ax.set_xticks([]); ax.set_yticks([])  
        ax.imshow(make_grid(fake_images.cpu().detach(), nrow=8).
```




8.1. GAN.ipynb

檔案 編輯 檢視畫面 插入 執行階段 工具 說明 無法儲存變更

共用



檔案



gdrive
generated

generated-images-0001.png
generated-images-0051.png
generated-images-0101.png
generated-images-0151.png
generated-images-0201.png

sample_data

+ 程式碼 + 文字 複製到雲端硬碟




```
# Log losses & scores (l
print("Epoch [{} / {}], loss_
    epoch+1, epochs, l
# Save generated images
latent = torch.randn(batch_
save_samples(epoch+start_idx,
#save_samples(epoch+start_idx,
```

```
0, Epoch [1/1200], loss_g: 5.6803, los
Saving generated-images-0001.png
10, 20, 30, 40, 50, Epoch [51/1200], loss_
Saving generated-images-0051.png
60, 70, 80, 90, 100, Epoch [101/1200], los
Saving generated-images-0101.png
110, 120, 130, 140, 150, Epoch [151/1200],
Saving generated-images-0151.png
160, 170, 180, 190, 200, Epoch [201/1200],
Saving generated-images-0201.png
210, 220,
```

RAM
磁碟

編輯

generated-images-0201.png X generated-images-0001.png X



檔案



<>

gdrive

generated

generated-images-0001.png

generated-images-0051.png

generated-images-0101.png

generated-images-0151.png

generated-images-0201.png

generated-images-0251.png


sample_data

29.48 GB 可用

+ 程式碼 + 文字 複製到雲端硬碟

```
Saving generated-images-0001.png  
10, 20, 30, 40, 50, Epoch [51/1200], loss_  
Saving generated-images-0051.png  
60, 70, 80, 90, 100, Epoch [101/1200], los  
Saving generated-images-0101.png  
110, 120, 130, 140, 150, Epoch [151/1200],  
Saving generated-images-0151.png  
160, 170, 180, 190, 200, Epoch [201/1200],  
Saving generated-images-0201.png  
210, 220, 230, 240, 250, Epoch [251/1200],  
Saving generated-images-0251.png  
260, 270, 280,
```

generated-images-0251.png



檔案



gdrive
generated


generated-images-0001.png
generated-images-0051.png
generated-images-0101.png
generated-images-0151.png
generated-images-0201.png
generated-images-0251.png
generated-images-0301.png

sample_data

+ 程式碼 + 文字 複製到雲端硬碟

```
Saving generated-images-0001.png  
10, 20, 30, 40, 50, Epoch [51/1200], loss_  
Saving generated-images-0051.png  
60, 70, 80, 90, 100, Epoch [101/1200], los  
Saving generated-images-0101.png  
110, 120, 130, 140, 150, Epoch [151/1200],  
Saving generated-images-0151.png  
160, 170, 180, 190, 200, Epoch [201/1200],  
Saving generated-images-0201.png  
210, 220, 230, 240, 250, Epoch [251/1200],  
Saving generated-images-0251.png  
260, 270, 280, 290, 300, Epoch [301/1200],  
Saving generated-images-0301.png  
310, 320, 330, 340,
```

generated-images-0301.png X



Practice


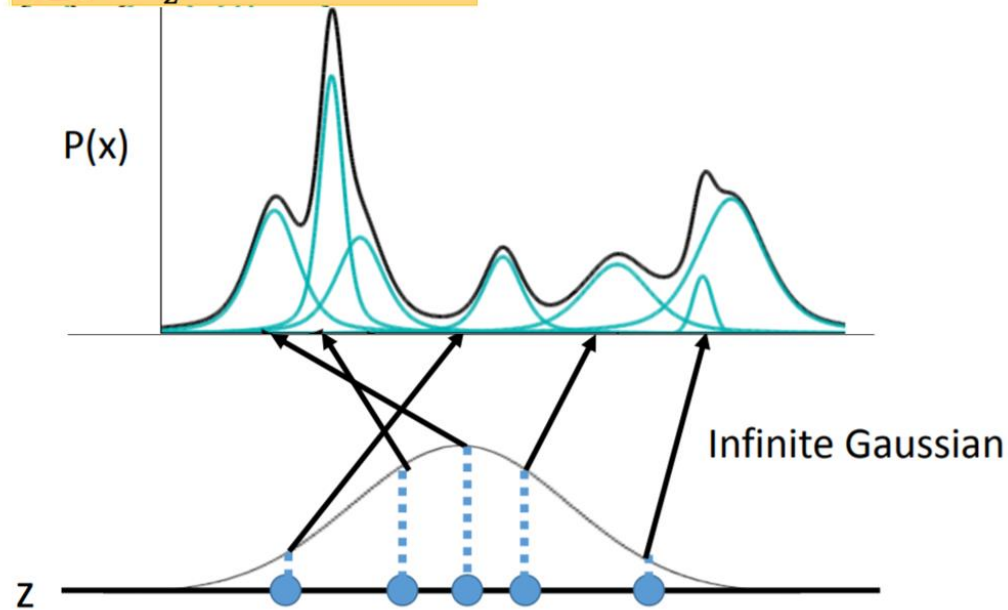
- Run "8.1. GAN.ipynb" in Colab. Try to generate Tom&Jerry



The generator can be modeled as a probability distribution model P_G

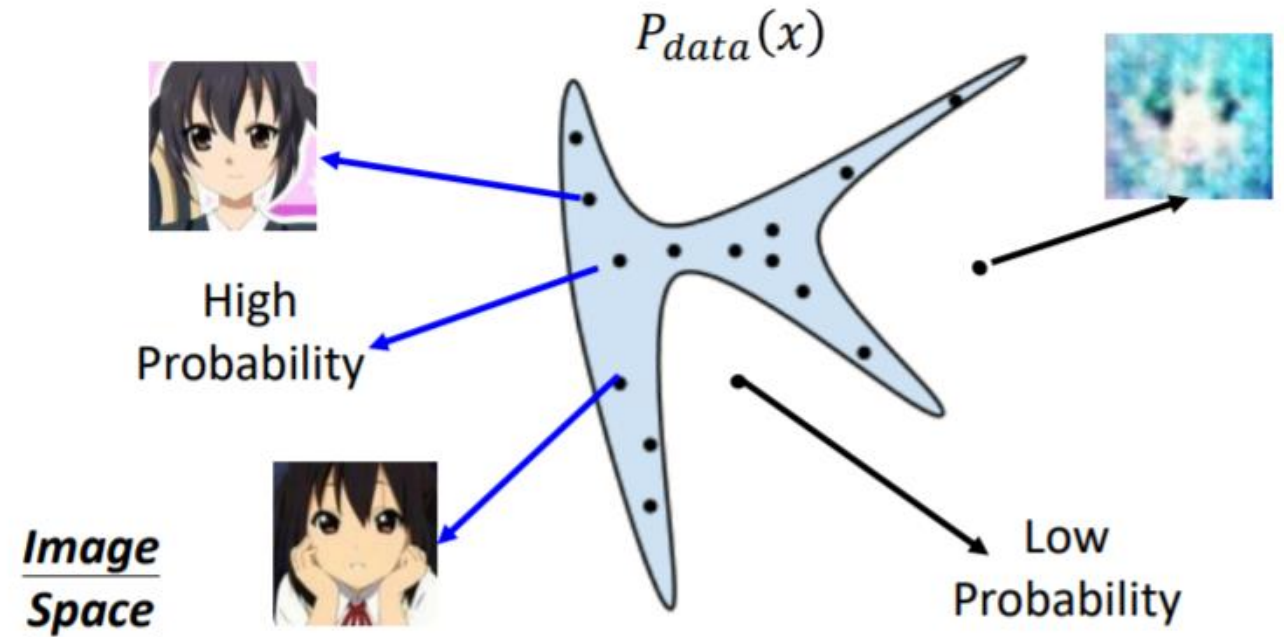
Gaussian Mixture Model

$$P(x) = \int_z P(z)P(x|z)dz$$



X : an image (a high-dimensional vector)

- We want to find data distribution $P_{data}(x)$



Based on the distribution model P_G , we want to maximize the likelihood of observing the training images x_1, x_2, \dots, x_m

- Given a data distribution $P_{data}(x)$ (We can sample from it.)
- We have a distribution $P_G(x; \theta)$ parameterized by θ
 - We want to find θ such that $P_G(x; \theta)$ close to $P_{data}(x)$
 - E.g. $P_G(x; \theta)$ is a Gaussian Mixture Model, θ are means and variances of the Gaussians


Sample $\{x^1, x^2, \dots, x^m\}$ from $P_{data}(x)$

We can compute $P_G(x^i; \theta)$

Likelihood of generating the samples

$$L = \prod_{i=1}^m P_G(x^i; \theta)$$

Find θ^* maximizing the likelihood



Review: Based on the distribution model P_G , we want to maximize the likelihood of observing the training images $x_1, x_2, x_3, \dots, x_m$

HW6(2) – VAE

Maximizing Likelihood

$$P(x) = \int_z P(z)P(x|z)dz$$

$P(z)$ is normal distribution
 $x|z \sim N(\mu(z), \sigma(z))$
 $\mu(z), \sigma(z)$ is going to be estimated

$$L = \sum_x \log P(x)$$

Maximizing the likelihood of the observed x

$$L = p(x^1) \times p(x^2) \times p(x^3) \times \dots p(x^m) = \prod_{i=1, \dots, m} P(x^i)$$

HW4, 5 – MLP, CNN classifier

Training
Data

x^1	x^2	x^3	...	x^N
C_1	C_1	C_2		C_1

$$L(w, b) = f_{w,b}(x^1) f_{w,b}(x^2) (1 - f_{w,b}(x^3)) \dots f_{w,b}(x^N)$$

$$-\ln L(w, b) = -\ln f_{w,b}(x^1) - \ln f_{w,b}(x^2) - \ln(1 - f_{w,b}(x^3)) \dots$$

\hat{y}^n : 1 for class 1, 0 for class 2

$$= \sum_n - \left[\hat{y}^n \ln f_{w,b}(x^n) + (1 - \hat{y}^n) \ln (1 - f_{w,b}(x^n)) \right]$$

Cross entropy between two Bernoulli distribution

Compare with RL PPO: Based on the distribution model $P_{\theta}(\tau)$,
We want to maximize the likelihood of obtaining best rewards

$\tau = (s_1, a_1, r_1, s_2, a_2, r_2, \dots s_T, a_T)$ See RL training video in my Github/Intelligent-Robot

$$p_{\theta}(\tau) = p(s_1)p_{\theta}(a_1|s_1)p(s_2|s_1, a_1)p_{\theta}(a_2|s_2)p(s_3|s_2, a_2) \dots$$

$$R(\tau) = \sum_{t=1}^T r_t \quad \text{The reward of doing this trajectory under current policy}$$

$$\bar{R}_{\theta} = \sum R(\tau) p_{\theta}(\tau) = E_{\tau \sim p_{\theta}(\tau)}[R(\tau)]$$

$$\max_{\theta} E[\bar{R}_{\theta}]$$

Maximum likelihood estimation = minimum KL divergence

$$\theta^* = \arg \max_{\theta} \prod_{i=1}^m P_G(x^i; \theta) = \arg \min_{\theta} KL(P_{data} || P_G)$$

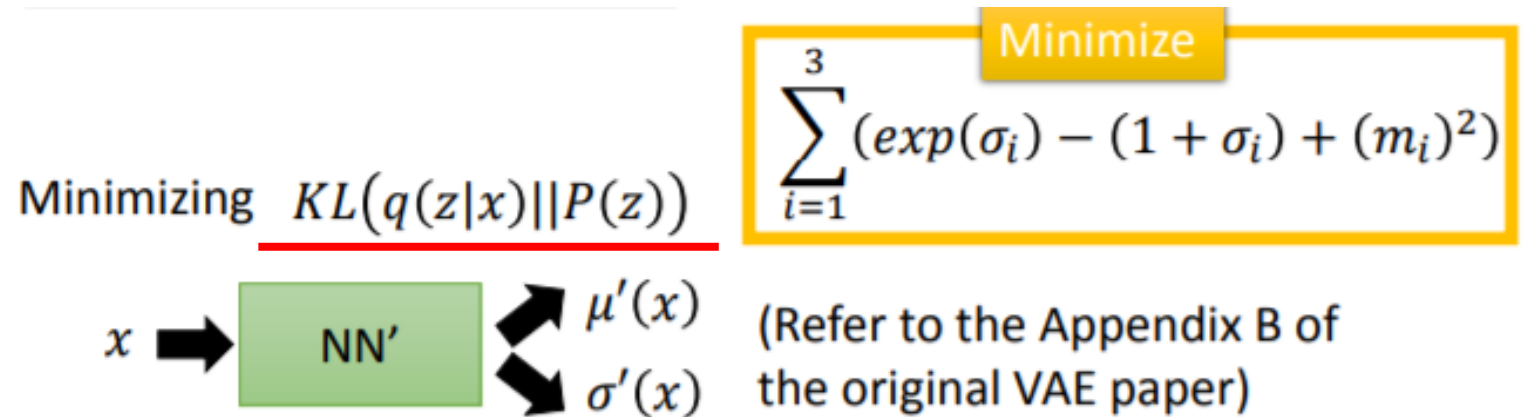
$$\begin{aligned}\theta^* &= \arg \max_{\theta} \prod_{i=1}^m P_G(x^i; \theta) = \arg \max_{\theta} \log \prod_{i=1}^m P_G(x^i; \theta) \\ &= \arg \max_{\theta} \sum_{i=1}^m \log P_G(x^i; \theta) \quad \{x^1, x^2, \dots, x^m\} \text{ from } P_{data}(x) \\ &\approx \arg \max_{\theta} E_{x \sim P_{data}} [\log P_G(x; \theta)] \\ &= \arg \max_{\theta} \int_x P_{data}(x) \log P_G(x; \theta) dx - \int_x P_{data}(x) \log P_{data}(x) dx \\ &= \arg \min_{\theta} \underline{KL(P_{data} || P_G)} \quad \text{How to define a general } P_G?\end{aligned}$$

$$\begin{aligned}\int_x P_{data}(x) \log P_{data}(x) dx &= \log P_{data}(x) \\ D_{KL}(q || p) &= \sum_{i=1}^N q(x_i) \log \left(\frac{q(x_i)}{p(x_i)} \right)\end{aligned}$$

Compare with VAE and PPO: : Maximum likelihood estimation = minimum KL divergence

$$\theta^* = \arg \max_{\theta} \prod_{i=1}^m P_G(x^i; \theta) = \arg \min_{\theta} KL(P_{data} || P_G)$$

Variational AE

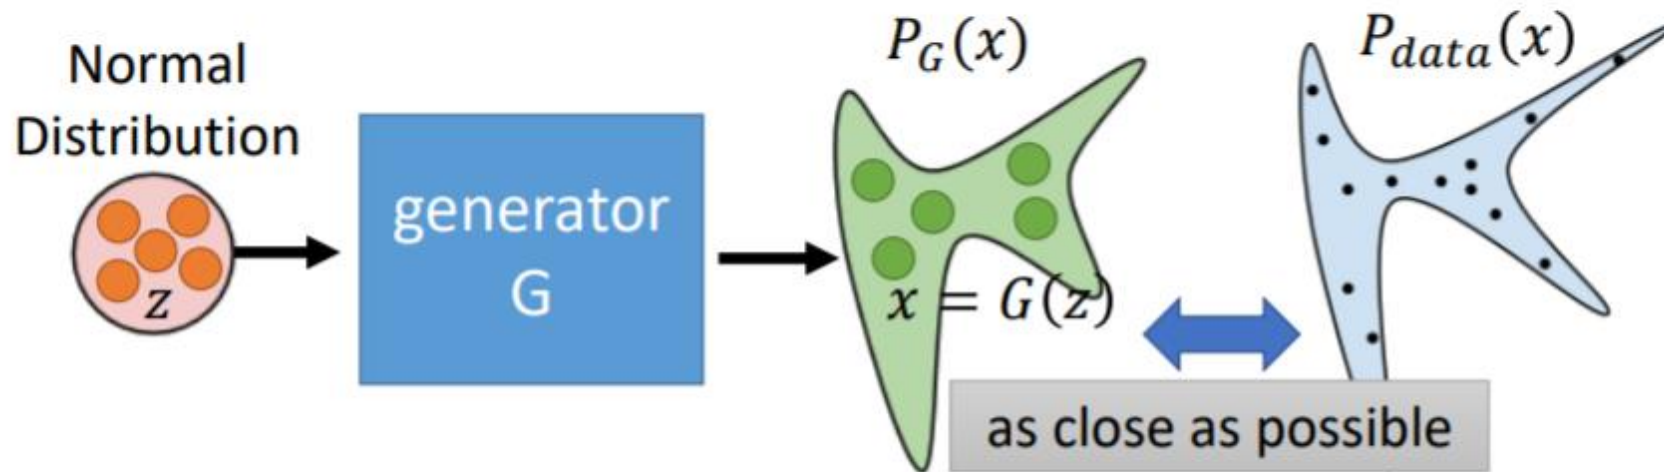


Proximal policy optimization (PPO)

$$J_{PPO}^{\theta'}(\theta) = J^{\theta'}(\theta) - \beta KL(\theta, \theta')$$

How to define a general P_G ? – use a generator NN

- A generator G is a network. The network defines a probability distribution P_G



$$G^* = \arg \min_G \underline{Div}(P_G, P_{data})$$

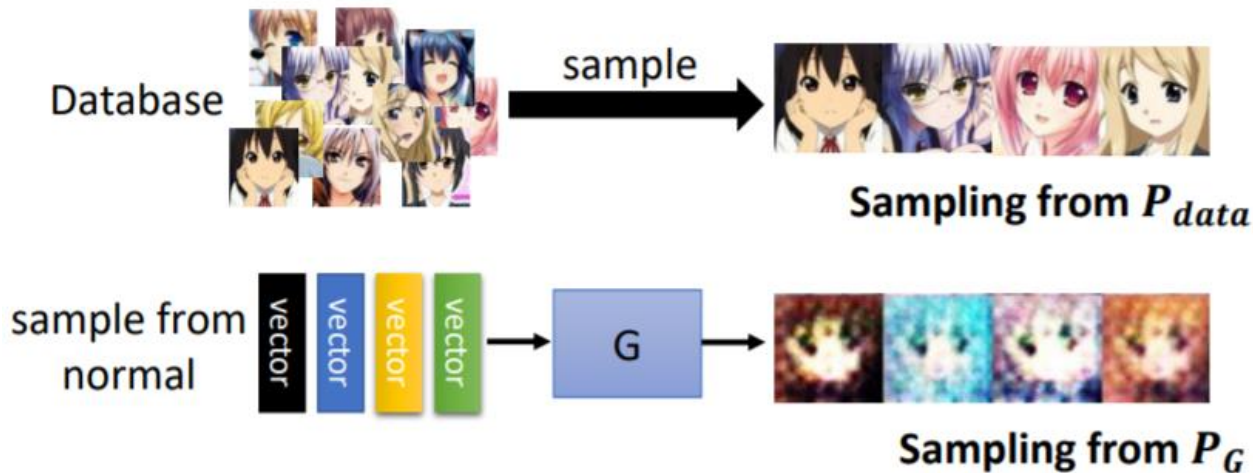
Divergence between distributions P_G and P_{data}

How to compute the divergence?

How to compute the divergence? – train a discriminator NN to compute the divergence

$$G^* = \arg \min_G \text{Div}(P_G, P_{data})$$

Although we do not know the distributions of P_G and P_{data} , we can sample from them.




```
[35]: for epoch in range(epochs):  
    if(epoch % 10 ==0):  
        print(epoch, end=",")  
    for real_images, _ in train_dl:  
        # Train discriminator  
        loss_d, real_score, fake_score = train_discriminator(real_images.  
        # Train generator  
        loss_g = train_generator(opt_g)
```

```
[19]: def train_discriminator(real_images, opt_d):  
    # Clear discriminator gradients  
    opt_d.zero_grad()  
  
    # Pass real images through discriminator  
    real_preds = discriminator(real_images)  
    real_targets = torch.ones(real_images.size(0))  
    real_loss = F.binary_cross_entropy(real_preds, real_targets)  
    real_score = torch.mean(real_preds).item()  
  
    # Generate fake images  
    latent = torch.randn(batch_size, latent_size,  
    fake_images = generator(latent.to(device))
```

Train a discriminator D to compute the divergence

★ : data sampled from P_{data}
★ : data sampled from P_G



Example Objective Function for D

$$V(G, D) = E_{x \sim P_{data}} [\log D(x)] + E_{x \sim P_G} [\log(1 - D(x))]$$

(G is fixed)

Training: $D^* = \arg \max_D V(D, G)$

[Goodfellow, et al., NIPS, 2014]

```
[19]: def train_discriminator(real_images,
# Clear discriminator gradients
opt_d.zero_grad()

# Pass real images through discriminator
real_preds = discriminator(real_images)
real_targets = torch.ones(real_images.size(0))
real_loss = F.binary_cross_entropy(real_preds, real_targets)
real_score = torch.mean(real_preds)

# Generate fake images
latent = torch.randn(batch_size, 1)
fake_images = generator(latent.to(device))

# Pass fake images through discriminator
fake_targets = torch.zeros(fake_images.size(0))
fake_preds = discriminator(fake_images)
fake_loss = F.binary_cross_entropy(fake_preds, fake_targets)
fake_score = torch.mean(fake_preds)

# Update discriminator weights
loss = real_loss + fake_loss
loss.backward()
opt_d.step()
```

Review: Loss function for binary classifier

Using the example objective function is exactly the same as training a binary classifier.

The discriminator is a binary classifier (logistic regression) to classify real vs fake

Cross entropy:

$$C(f(x^n), \hat{y}^n) = -[\hat{y}^n \ln f(x^n) + (1 - \hat{y}^n) \ln(1 - f(x^n))]$$

$$V(G, D) = E_{x \sim P_{data}}[\log D(x)] + E_{x \sim P_G}[\log(1 - D(x))]$$

(G is fixed)

Adversarial neural networks G and D

Difficult to train!!

(1) Train generator – max. likelihood = min. KL divergence between P_G and P_{data}

$$\theta^* = \arg \max_{\theta} \prod_{i=1}^m P_G(x^i; \theta) = \arg \min_{\theta} KL(P_{data} || P_G)$$

(2) Train discriminator – max. divergence between P_G and P_{data}



small divergence

★ : data sampled from P_{data}
★ : data sampled from P_G



large divergence

$$V(G, D) = E_{x \sim P_{data}} [\log D(x)] + E_{x \sim P_G} [\log(1 - D(x))]$$

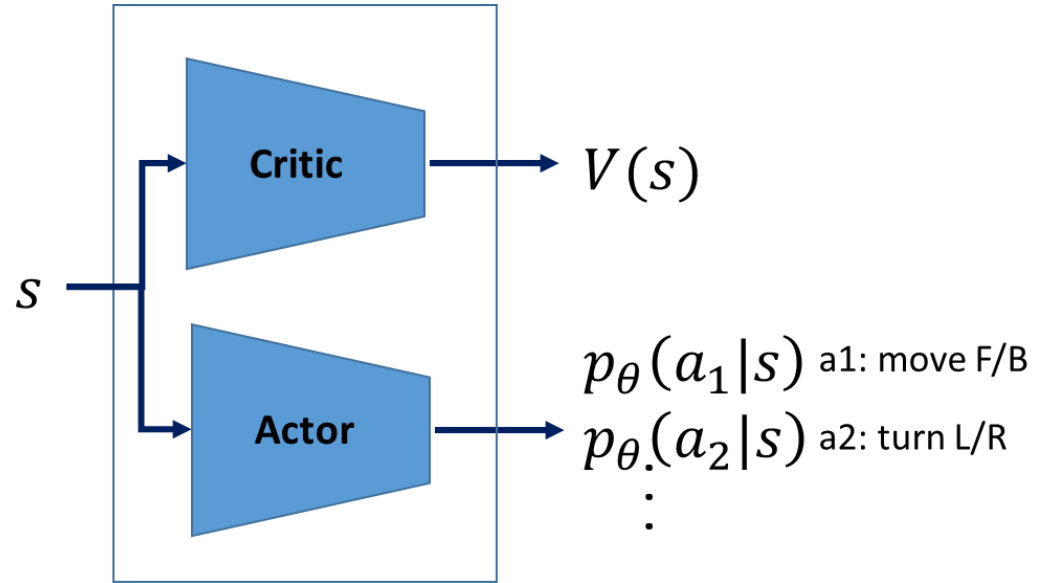
(G is fixed)

Training:

$$D^* = \arg \max_D V(D, G)$$

Compare with RL PPO: Train two neural networks actor and critic together

$$L = c_v L_v + L_\pi - \beta L_{reg}$$



(1) Actor – Learns the best actions (that can have maximum long-term rewards)

$$L_\pi = \sum_{(s_t, a_t)} \min \left(\frac{p_\theta(a_t|s_t)}{p_{\theta'}(a_t|s_t)} A^{\theta'}(s_t, a_t), \text{clip} \left(\frac{p_\theta(a_t|s_t)}{p_{\theta'}(a_t|s_t)}, 1 - \varepsilon, 1 + \varepsilon \right) A^{\theta'}(s_t, a_t) \right)$$

(2) Critic – Learns the expected value of the long-term reward.

$$L_v = \text{MSE of } (\text{return} - v)$$

The objective function of \mathcal{D} is related to JS divergence

GAN is sensitive to hyper-parameter tuning and its performance range is large. Different GANs' performances are similar

