

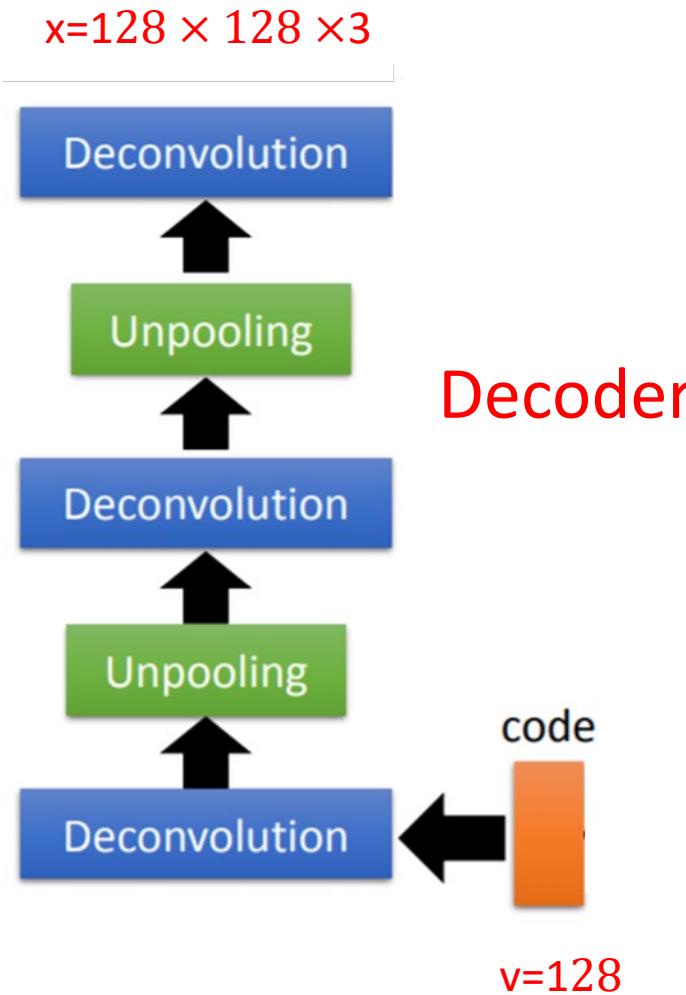
Generative Adversarial Network (GAN)

Practice

- Open "8.1. GAN.ipynb"



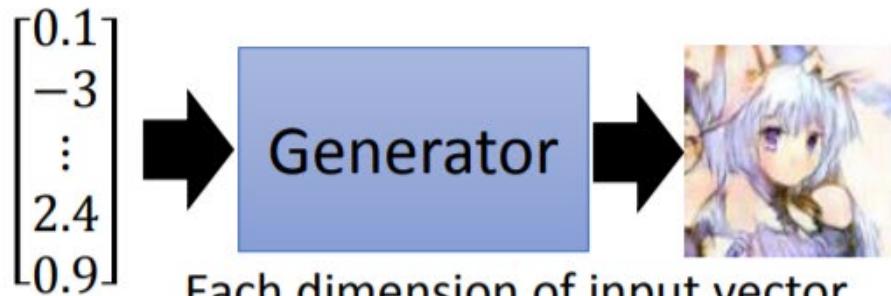
Generator



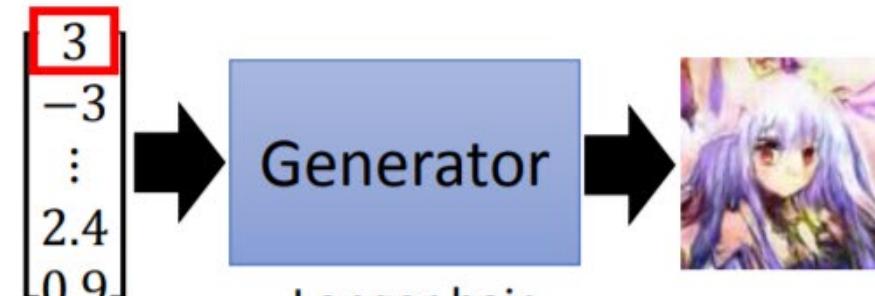
```
[10]: latent_size= 128
```

```
[11]: generator = nn.Sequential(  
# in: latent_size x 1 x 1  
  
nn.ConvTranspose2d(latent_size, 512, kernel_size=4, stride=1,  
nn.BatchNorm2d(512),  
nn.ReLU(True),  
# out: 512 x 4 x 4  
  
nn.ConvTranspose2d(512, 256, kernel_size=4, stride=2, padding=  
nn.BatchNorm2d(256),  
nn.ReLU(True),  
# out: 256 x 8 x 8  
  
nn.ConvTranspose2d(256, 128, kernel_size=4, stride=2, padding=  
nn.BatchNorm2d(128),  
nn.ReLU(True),  
# out: 128 x 16 x 16  
  
nn.ConvTranspose2d(128, 64, kernel_size=4, stride=2, padding=1  
nn.BatchNorm2d(64),  
nn.ReLU(True),  
# out: 64 x 32 x 32  
  
nn.ConvTranspose2d(64, 32, kernel_size=4, stride=2, padding=1,  
nn.BatchNorm2d(32),  
nn.ReLU(True),  
# out: 32 x 64 x 64  
  
nn.ConvTranspose2d(32, 3, kernel_size=4, stride=2, padding=1,  
nn.Tanh()  
# out: 3 x 128 x 128  
)
```

Each dimension in the feature vector represent a property of the output image



Each dimension of input vector represents some characteristics.



Longer hair

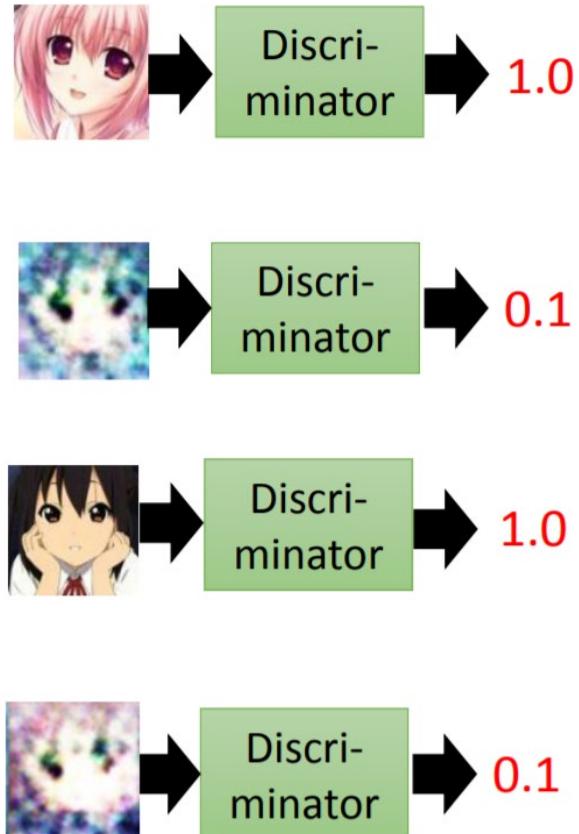


blue hair

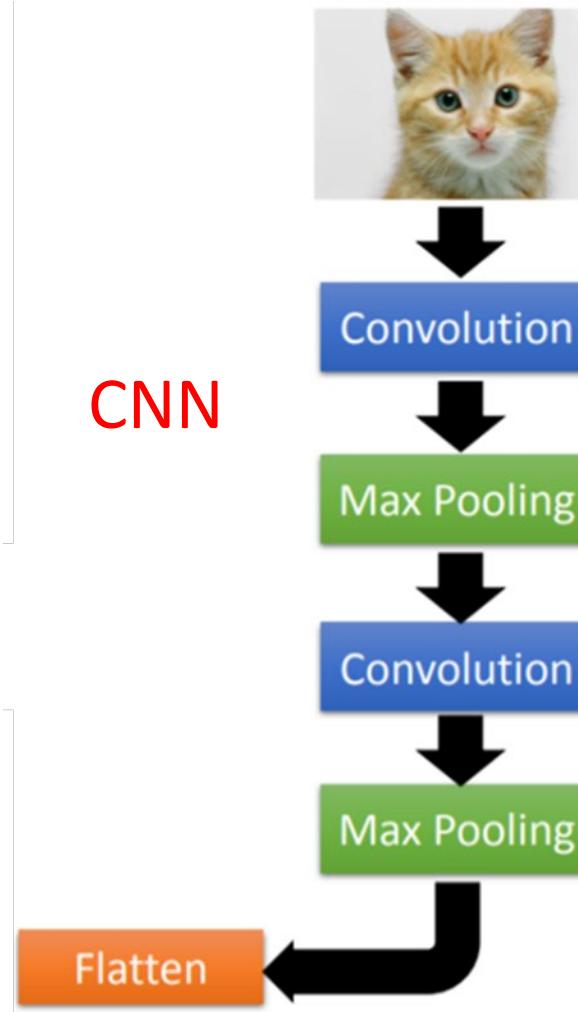


Open mouth

Discriminator



CNN



[15]: discriminator = nn.Sequential(

in: 3 x 128 x 128

nn.Conv2d(3, 64, kernel_size=4, stride=2, p:
nn.BatchNorm2d(64),
nn.LeakyReLU(0.2, inplace=True),
out: 64 x 64 x 64

nn.Conv2d(64, 128, kernel_size=4, stride=2,
nn.BatchNorm2d(128),
nn.LeakyReLU(0.2, inplace=True),
out: 128 x 32 x 32

nn.Conv2d(128, 256, kernel_size=4, stride=2,
nn.BatchNorm2d(256),
nn.LeakyReLU(0.2, inplace=True),
out: 256 x 16 x 16

nn.Conv2d(256, 512, kernel_size=4, stride=2,
nn.BatchNorm2d(512),
nn.LeakyReLU(0.2, inplace=True),
out: 512 x 8 x 8

nn.Conv2d(512, 1024, kernel_size=4, stride=2,
nn.BatchNorm2d(1024),
nn.LeakyReLU(0.2, inplace=True),
out: 1024 x 4 x 4

nn.Conv2d(1024, 1, kernel_size=4, stride=1,
out: 1 x 1 x 1

nn.Flatten(),
nn.Sigmoid()

Reference: 李弘毅 GAN Lecture 1 (2018)

Step1 – Fix G and train D

- Initialize generator and discriminator

- In each training iteration:



[12]: `generator.to(device)`

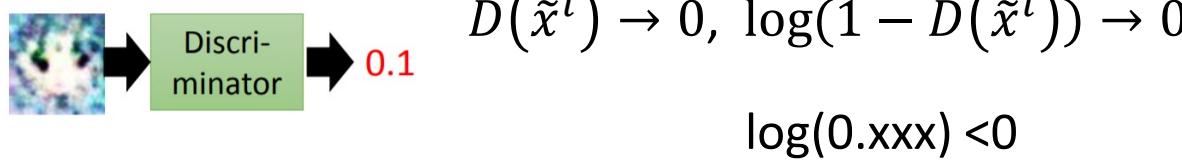
[16]: `discriminator.to(device)`

Step 1: Fix generator G, and update discriminator D

```
[35]: for epoch in range(epochs):
    if(epoch % 10 ==0):
        print(epoch, end=",")
    for real_images, _ in train_dl:
        # Train discriminator
        loss_d, real_score, fake_score = train_discriminator(real_images.
        # Train generator
        loss_g = train_generator(opt_g)
```

Train discriminator D

- Sample m examples $\{x^1, x^2, \dots, x^m\}$ from database
- Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from a distribution
- Obtaining generated data $\{\tilde{x}^1, \tilde{x}^2, \dots, \tilde{x}^m\}, \tilde{x}^i = G(z^i)$



Update discriminator parameters θ_d to maximize

- $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \underline{\log D(x^i)} + \frac{1}{m} \sum_{i=1}^m \underline{\log (1 - D(\tilde{x}^i))}$
- $\theta_d \leftarrow \theta_d + \eta \nabla \tilde{V}(\theta_d)$

[19]: `def train_discriminator(real_images, opt_d):`

```
# Clear discriminator gradients  
opt_d.zero_grad()
```

```
# Pass real images through discriminator
```

```
real_preds = discriminator(real_images)  
real_targets = torch.ones(real_images.size(0))  
real_loss = F.binary_cross_entropy(real_preds,  
real_score = torch.mean(real_preds).item()
```

```
# Generate fake images
```

```
latent = torch.randn(batch_size, latent_size)  
fake_images = generator(latent.to(device))
```

```
# Pass fake images through discriminator
```

```
fake_targets = torch.zeros(fake_images.size(0))  
fake_preds = discriminator(fake_images)  
fake_loss = F.binary_cross_entropy(fake_preds,  
fake_score = torch.mean(fake_preds).item()
```

```
# Update discriminator weights
```

```
loss = real_loss + fake_loss  
loss.backward()  
opt_d.step()  
return loss.item(), real_score, fake_score
```

Review – Cross entropy

Measures the differences between the true probability p_i and the predicted probability q_i

$$H(p, q) = - \sum_i p_i \ln(q_i)$$

A	B	C	D	E	F	G	H	I	J	K
z1	z2	y-hat	p1	p2	EXP(z1)	EXP(z2)	F+G	q1	q2	-(P1*LN(Q1)+P2*LN(Q2))
-0.018	0.0855	0	1	0	0.982	1.089	2.071	0.474	0.526	0.74624
-0.0244	0.0741	0	1	0	0.976	1.077	2.053	0.475	0.525	0.74361
-0.0187	0.085	0	1	0	0.981	1.089	2.070	0.474	0.526	0.74634
-0.0258	0.0687	1	0	1	0.975	1.071	2.046	0.476	0.524	0.64701
-0.0267	0.0617	1	0	1	0.974	1.064	2.037	0.478	0.522	0.64992
										0.70662

Step2 – Fix D and train G

Step 2: Fix discriminator D, and update generator G

Generator learns to “fool” the discriminator

```
[35]: for epoch in range(epochs):
    if(epoch % 10 ==0):
        print(epoch, end=",")
    for real_images, _ in train_dl:
        # Train discriminator
        loss_d, real_score, fake_score =
        # Train generator
        loss_g = train_generator(opt_g)
```

- Sample m noise samples{ z^1, z^2, \dots, z^m } from a distribution
- Update generator parameters θ_g to maximize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log(D(G(z^i)))$
 - $\theta_g \leftarrow \theta_g - \eta \nabla \tilde{V}(\theta_g)$

```
[28]: def train_generator(opt_g):
    # Clear generator gradients
    opt_g.zero_grad()

    # Generate fake images
    latent = torch.randn(batch_size, latent_size,
    fake_images = generator(latent)

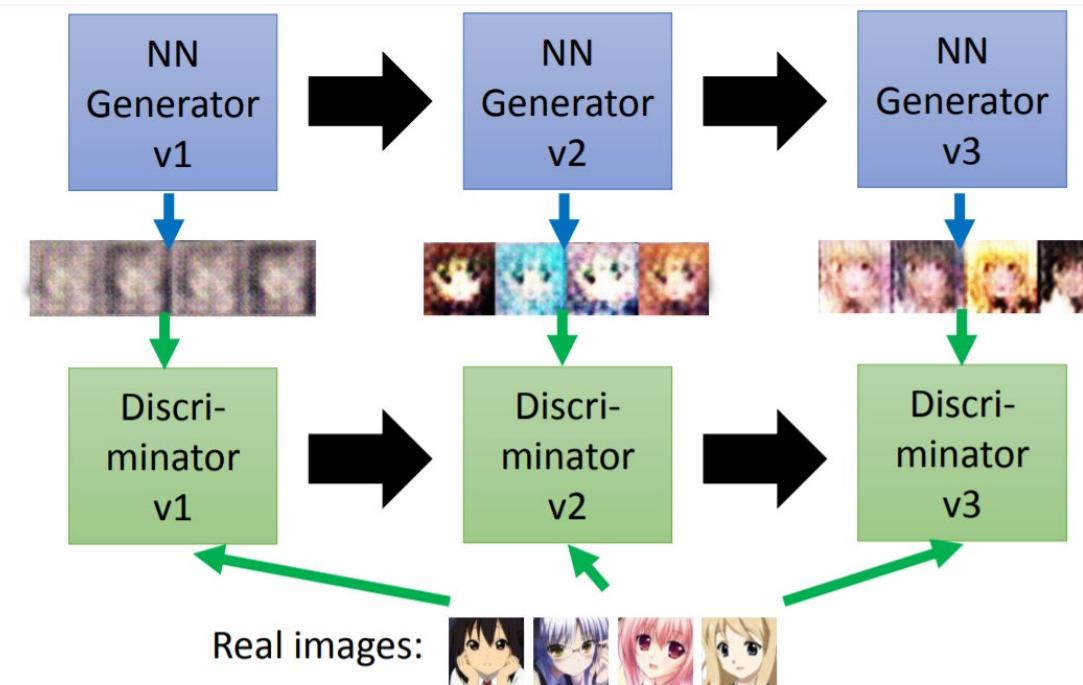
    # Try to fool the discriminator
    preds = discriminator(fake_images)
    targets = torch.ones(batch_size, 1, device=device)
    loss = F.binary_cross_entropy(preds, targets)

    # Update generator weights
    loss.backward()
    opt_g.step()

    return loss.item()
```

An evolution process

```
[35]: for epoch in range(epochs):  
    if(epoch % 10 ==0):  
        print(epoch, end=",")  
    for real_images, _ in train_dl:  
        # Train discriminator  
        loss_d, real_score, fake_score = train_discriminator(real_images)  
        # Train generator  
        loss_g = train_generator(opt_g)
```



Visualize fake images generated by G during the evolution process

```
[17]: sample_dir = 'generated'  
os.makedirs(sample_dir, exist_ok=True)
```

```
[34]: fixed_latent = torch.randn(64, latent_size,  
# used to generate saved images
```

```
if(epoch % 50 ==0):  
    # Log Losses & scores (Last batch)  
    print("Epoch [{}/{}], loss_g: {:.4f}, loss_d: {:.4f}, real_scores: {:.4f}, fake_scores: {:.4f} ".format(epoch+1, epochs, loss_g, loss_d, real_scores, fake_scores))  
    # Save generated images  
save_samples(epoch+start_idx, fixed_latent, sample_dir)
```

```
[18]: def save_samples(index, latent_tensors, show=True):  
    fake_images = generator(latent_tensors)  
    fake_fname = 'generated-images-{:0=4d}.png'.format(index)  
    save_image(denorm(fake_images), os.path.join(sample_dir, fake_fname))  
    print('Saving', fake_fname)  
    if show:  
        fig, ax = plt.subplots(figsize=(8, 8))  
        ax.set_xticks([]); ax.set_yticks([])  
        ax.imshow(make_grid(fake_images.cpu().detach(), nrow=8).permute(1, 2, 0))
```



8.1. GAN.ipynb

檔案 編輯 檢視畫面 插入 執行階段 工具 說明 無法儲存變更

共用



檔案



<>
..
gdrive

generated

generated-images-0001.png
generated-images-0051.png
generated-images-0101.png
generated-images-0151.png
generated-images-0201.png

sample_data

+ 程式碼 + 文字

複製到雲端硬碟

```
# Log losses & scores ()  
print("Epoch [{}/{}], loss_ epoch+1, epochs, 1c  
# Save generated images  
latent = torch.randn(batch_ save_samples(epoch+start_idx,  
#save_samples(epoch+start_id
```

```
0, Epoch [1/1200], loss_g: 5.6803, los Saving generated-images-0001.png  
10, 20, 30, 40, 50, Epoch [51/1200], loss_ Saving generated-images-0051.png  
60, 70, 80, 90, 100, Epoch [101/1200], los Saving generated-images-0101.png  
110, 120, 130, 140, 150, Epoch [151/1200], Saving generated-images-0151.png  
160, 170, 180, 190, 200, Epoch [201/1200], Saving generated-images-0201.png  
210, 220,
```

RAM
磁碟

編輯

generated-images-0201.png X generated-images-0001.png X





8.1. GAN.ipynb

共用



檔案 編輯 檢視畫面 插入 執行階段 工具 說明 無法儲存變更

檔案



..

gdrive

generated

generated-images-0001.png

generated-images-0051.png

generated-images-0101.png

generated-images-0151.png

generated-images-0201.png

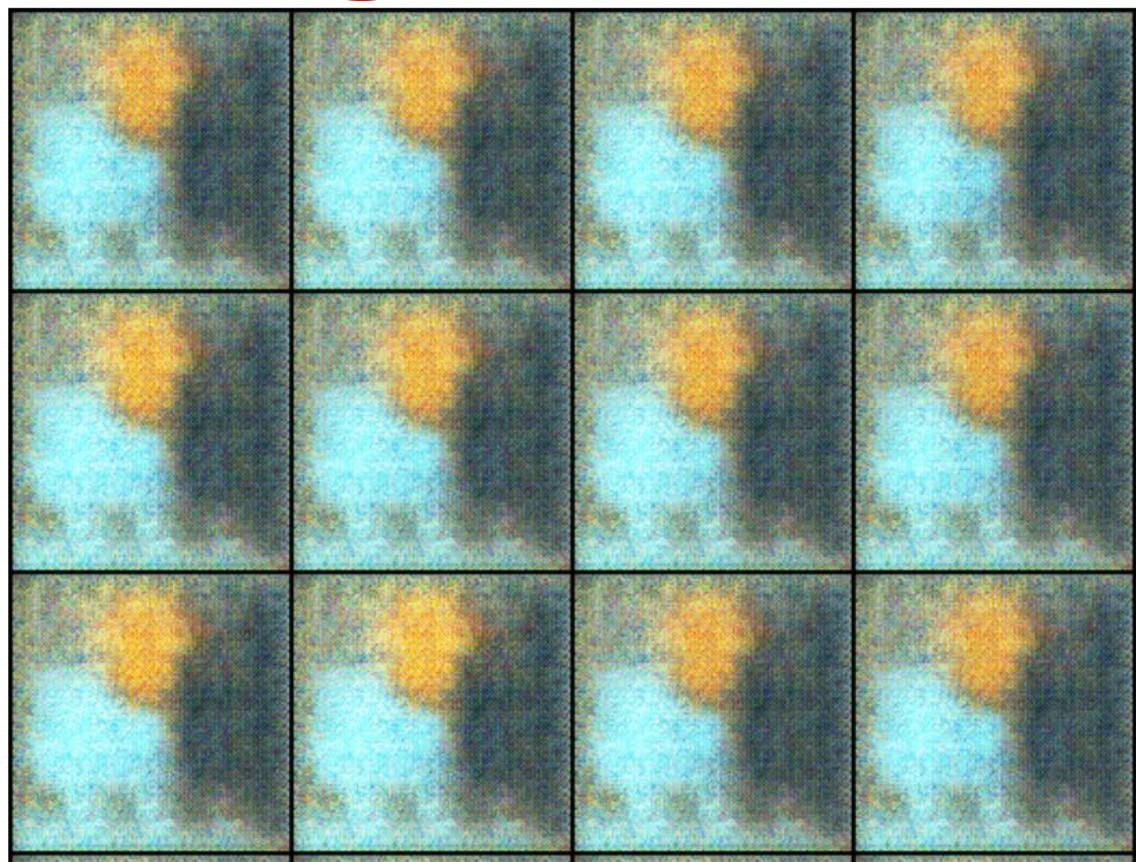
generated-images-0251.png

sample_data

+ 程式碼 + 文字 | 複製到雲端硬碟

```
Saving generated-images-0001.png  
10, 20, 30, 40, 50, Epoch [51/1200], loss_  
Saving generated-images-0051.png  
60, 70, 80, 90, 100, Epoch [101/1200], los  
Saving generated-images-0101.png  
110, 120, 130, 140, 150, Epoch [151/1200],  
Saving generated-images-0151.png  
160, 170, 180, 190, 200, Epoch [201/1200],  
Saving generated-images-0201.png  
210, 220, 230, 240, 250, Epoch [251/1200],  
Saving generated-images-0251.png  
260, 270, 280,
```

generated-images-0251.png X



磁碟 29.48 GB 可用



8.1. GAN.ipynb

共用



T

檔案 編輯 檢視畫面 插入 執行階段 工具 說明 無法儲存變更

檔案



..

gdrive

generated

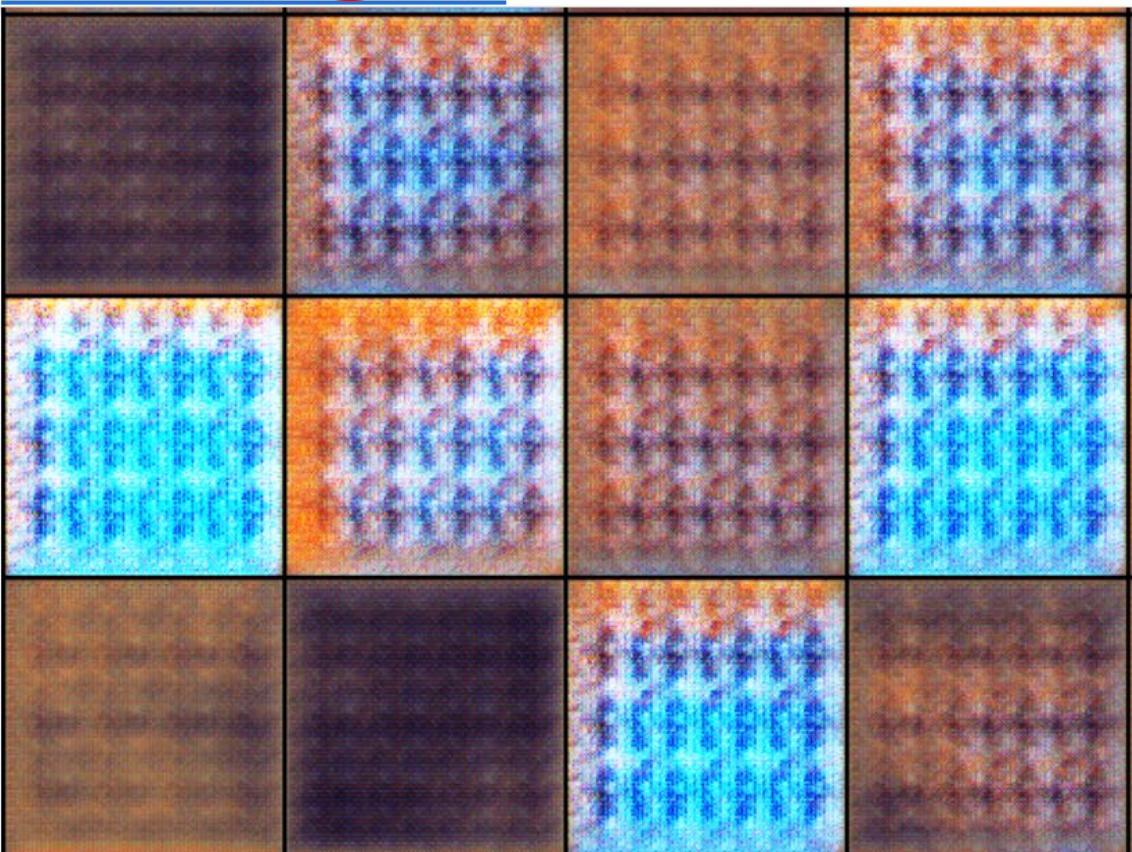
- generated-images-0001.png
- generated-images-0051.png
- generated-images-0101.png
- generated-images-0151.png
- generated-images-0201.png
- generated-images-0251.png
- generated-images-0301.png

sample_data

+ 程式碼 + 文字 複製到雲端硬碟

```
Saving generated-images-0001.png  
10, 20, 30, 40, 50, Epoch [51/1200], loss_  
Saving generated-images-0051.png  
60, 70, 80, 90, 100, Epoch [101/1200], los  
Saving generated-images-0101.png  
110, 120, 130, 140, 150, Epoch [151/1200],  
Saving generated-images-0151.png  
160, 170, 180, 190, 200, Epoch [201/1200],  
Saving generated-images-0201.png  
210, 220, 230, 240, 250, Epoch [251/1200],  
Saving generated-images-0251.png  
260, 270, 280, 290, 300, Epoch [301/1200],  
Saving generated-images-0301.png  
310, 320, 330, 340,
```

generated-images-0301.png X



Epoch =41201



34151



28651



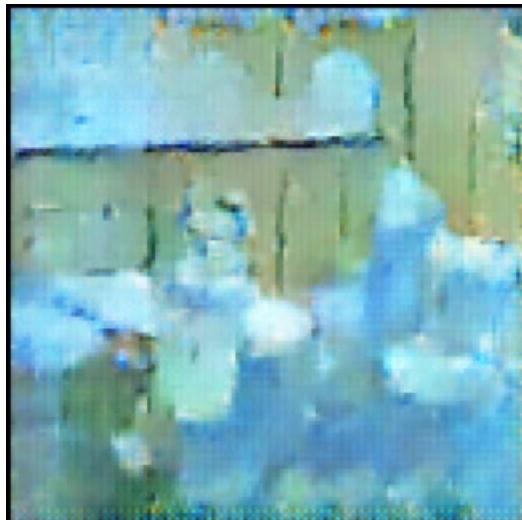
25901



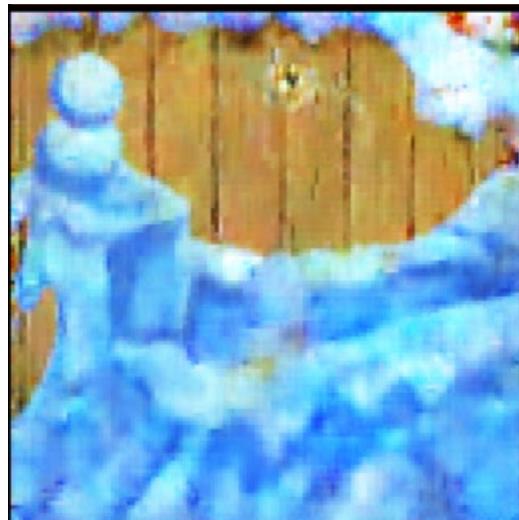
23201



16401



15551



14301



Save the generated images as a video

```
[37]: img_array = []
for filename in files:
    img = cv2.imread(filename)
height, width, layers = img.shape
size = (width,height)
img_array.append(img)

out = cv2.VideoWriter('GANTrainingVideo.avi',cv2.VideoWriter_fourcc(*'DIVX'), 15, size)

for i in range(len(img_array)):
    out.write(img_array[i])
out.release()
```

Tom & Jerry video <https://youtu.be/uDEGITFmhiQ>

HW7 (1)

- Use your own images, e.g., facial expression to train a GAN.
- Show the generated images.
- Try latent vectors like (all 1), (all 0.5), (all 0.3), (0, 0, 1, 0,), (one dimension goes from 0 to 1 and other dimensions fixed), ... to see the generated images.



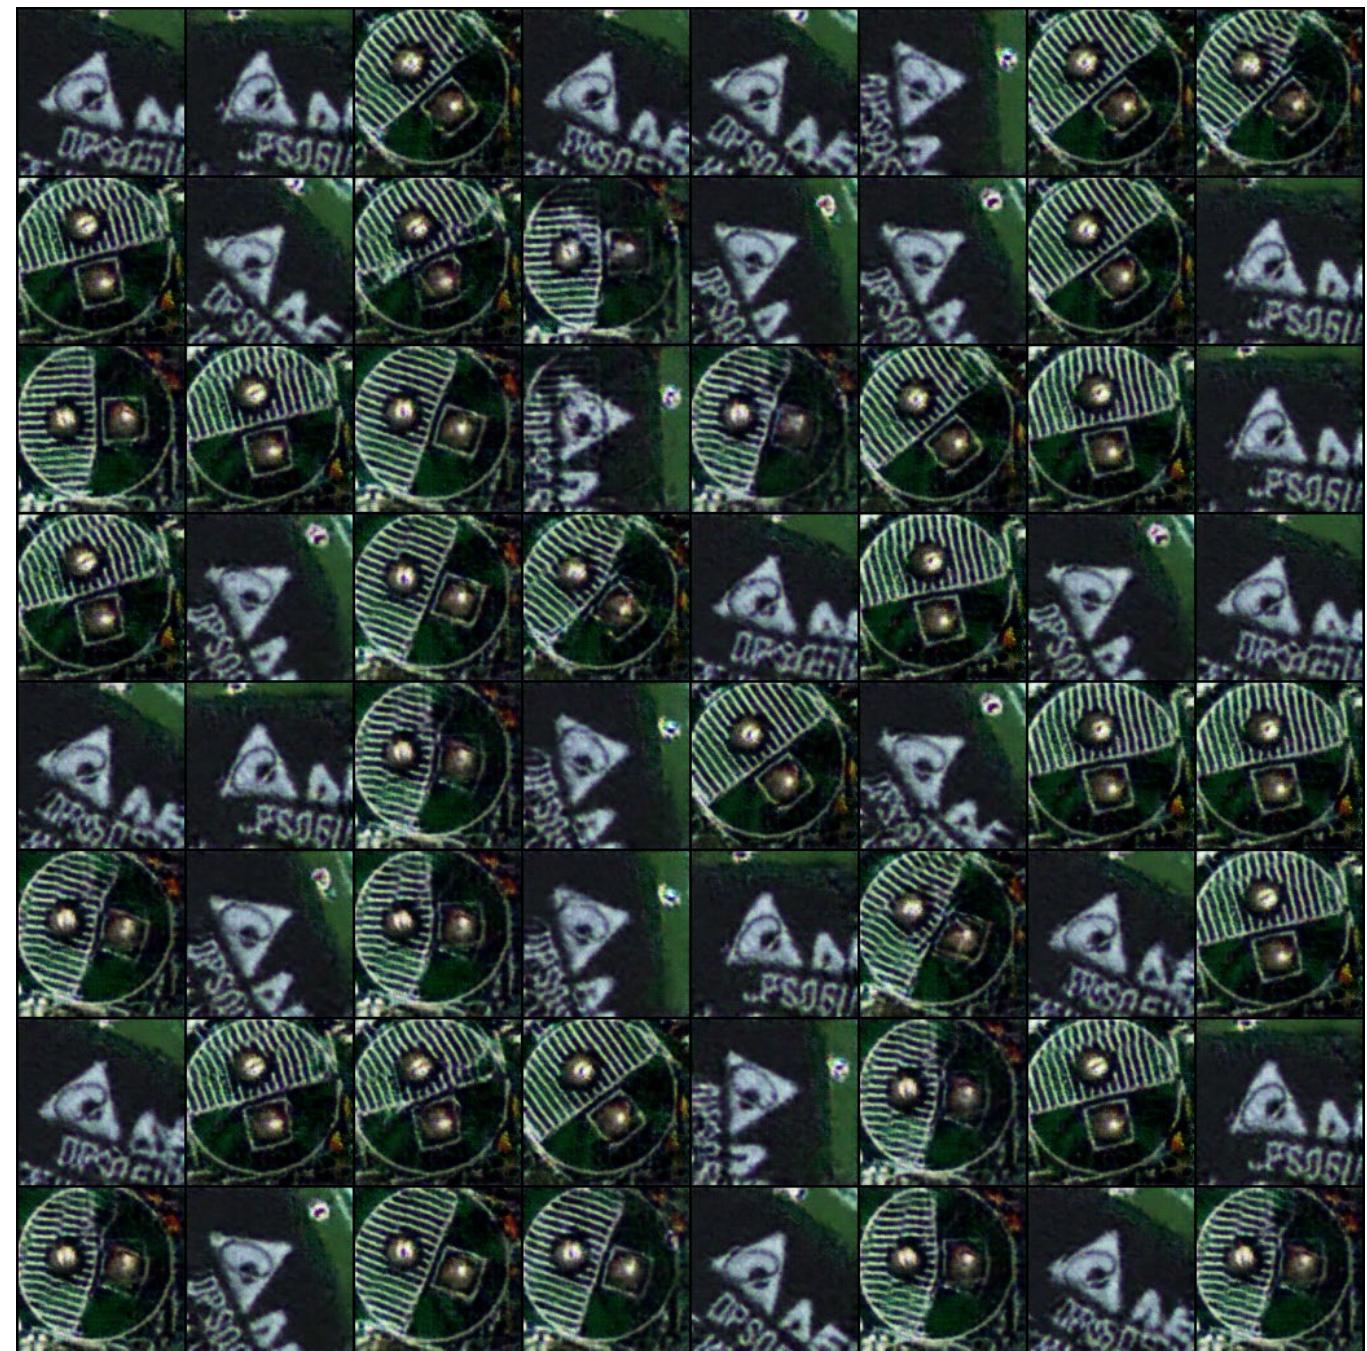
Images generated from G after
5651 epochs

Image size = 180x180x3

Class 1 = 100, Class 2 = 100

Latent vector size = 128
batch=32

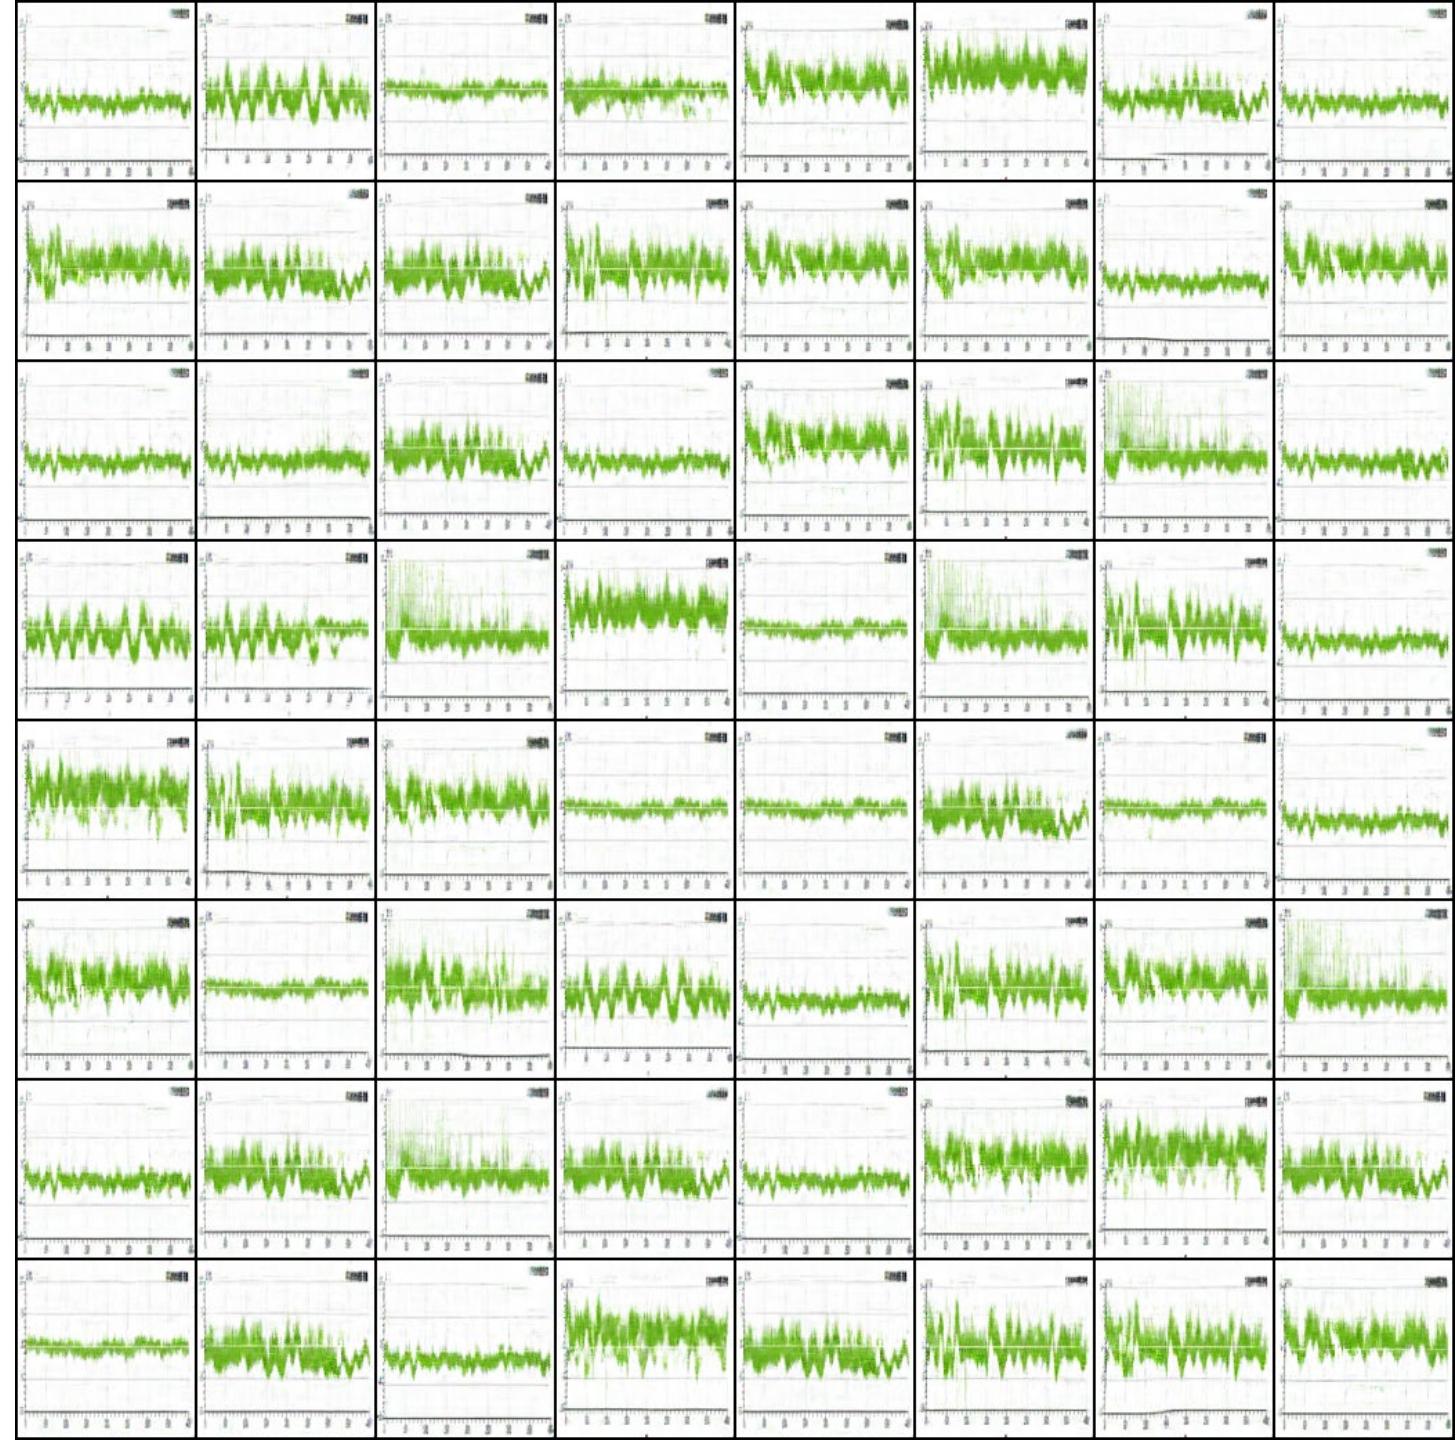
Run on 2080Ti, takes about 15 min.



1085442 Carlos

Images-19951

1071228 陳仁慶



Images generated from G after
2100 epochs

Image size = 128x128x3

Class 1 = 10, Class 2 = 10

Latent vector size = 128
batch=128

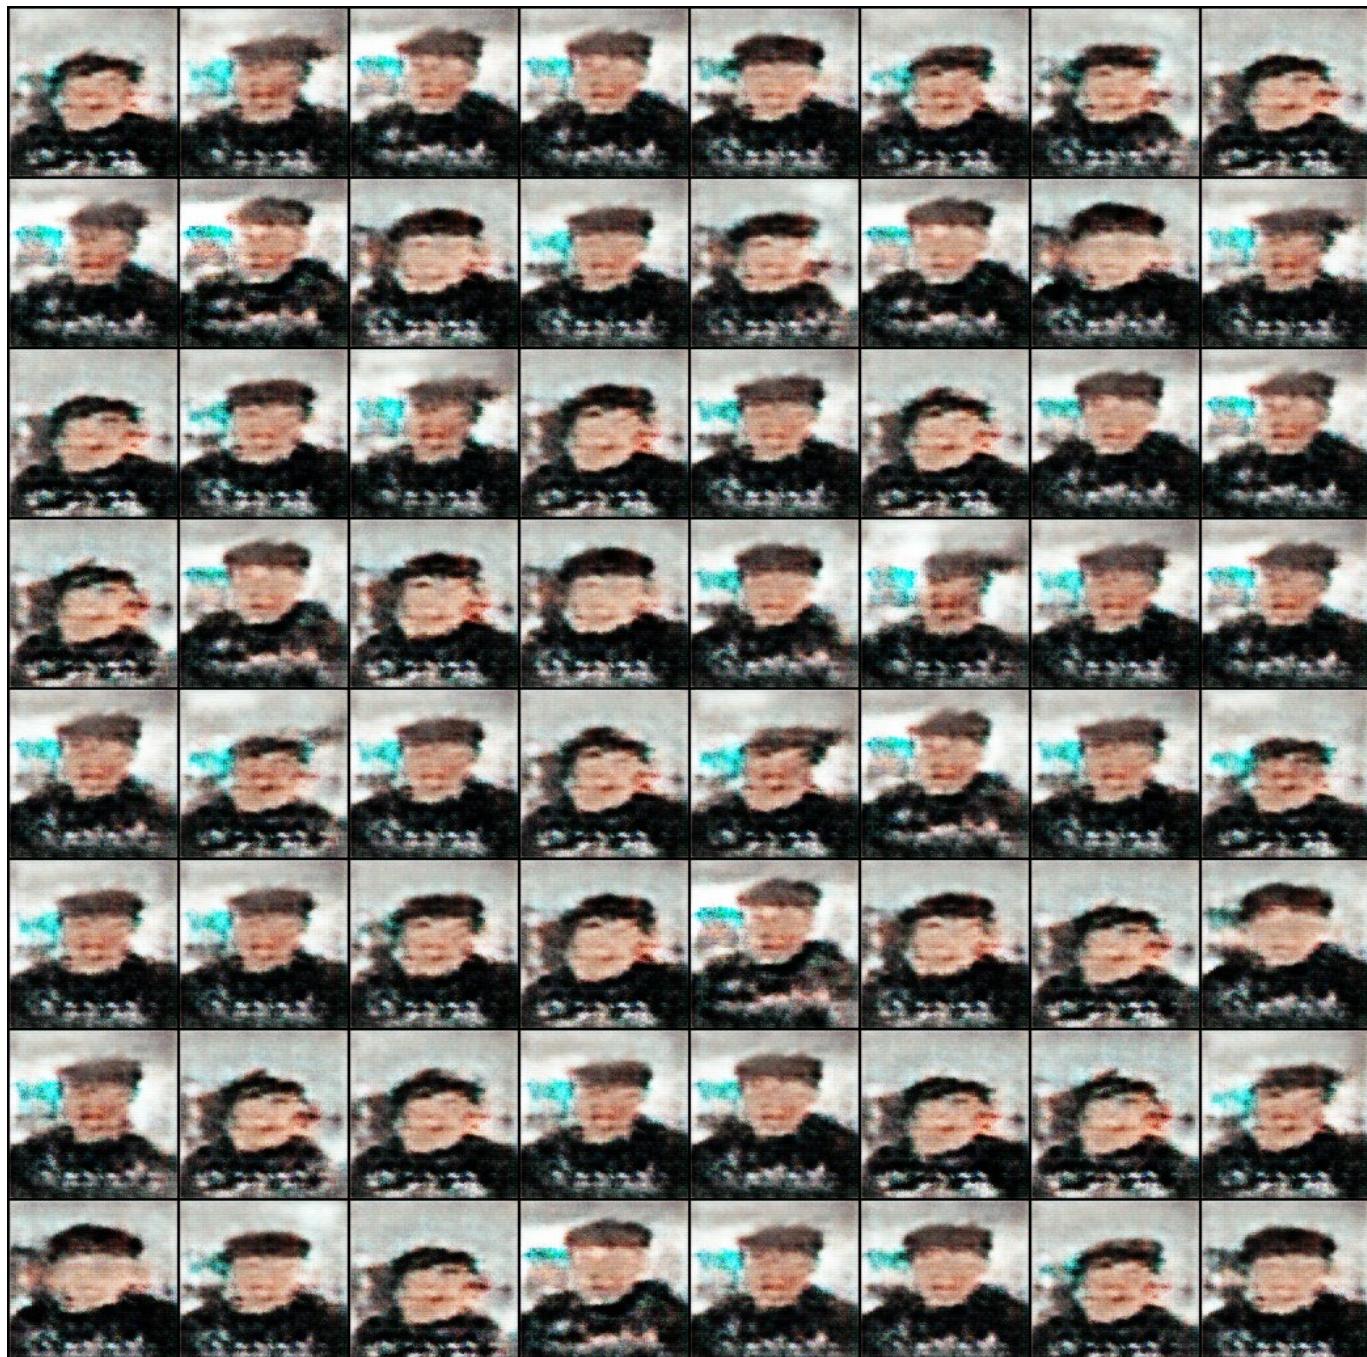
Run on Colab Tesla K80, takes about 120
min.



Images generated from G after
3051 epochs

Image size = 128x128x3
Class 1 = 10, Class 2 = 10
Latent vector size = 128
batch=128

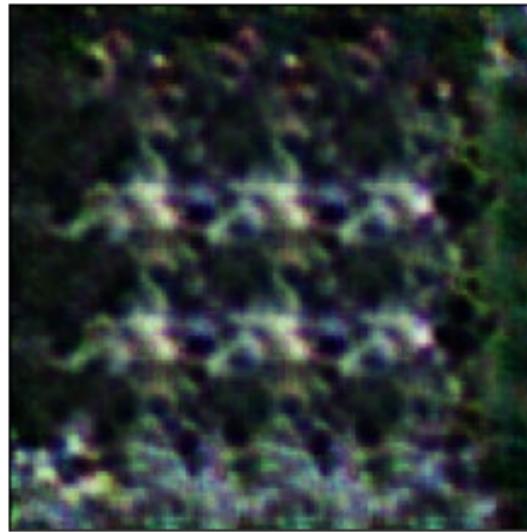
Run on Colab Tesla T4, takes about 60 min.



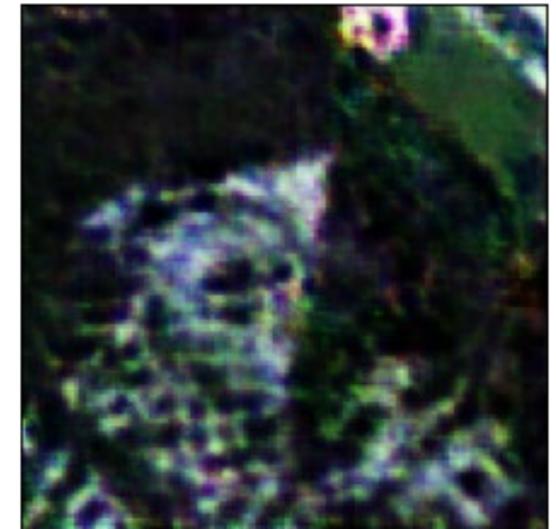
Input vector: [1, 1, 1...1]



[0, 0, 0...0]



[-1, -1, -1...-1]



```
x_ones = torch.ones(1, latent_size, 1, 1)
x_negs = -torch.ones(1, latent_size, 1, 1)
x_zeros = torch.zeros(1, latent_size, 1, 1)
```

```
fake_image = denorm(generator(x_ones.to(device)).detach().cpu()[0].permute(1, 2, 0))
print('Image size:', fake_image.shape)
plt.figure(figsize=(4, 4))
plt.imshow(fake_image.numpy().astype(float))
plt.xticks([])
plt.yticks([])
```

Image size: torch.Size([128, 128, 3])

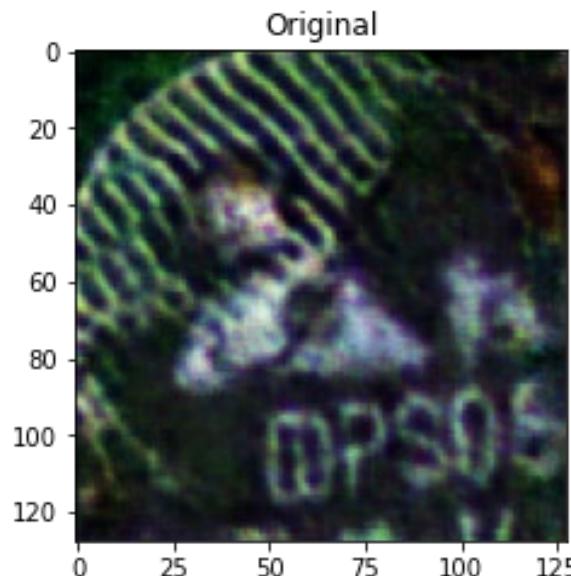
```

# Change a feature from the feature vector
feature_index = 3
new_ones = torch.ones(1, latent_size, 1, 1)
fake1 = denorm(generator(new_ones.to(device)).detach().cpu()[0].permute(1, 2, 0)).numpy().astype(float)
new_ones[:, feature_index, :, :] = 50
fake2 = denorm(generator(new_ones.to(device)).detach().cpu()[0].permute(1, 2, 0)).numpy().astype(float)
print('Fake image size: ', fake1.shape)
fig, ax = plt.subplots(1, 2, figsize=(8, 4))
ax[0].imshow(fake1)
ax[1].imshow(fake2)
ax[0].set_title('Original')
ax[1].set_title('Feature changed')

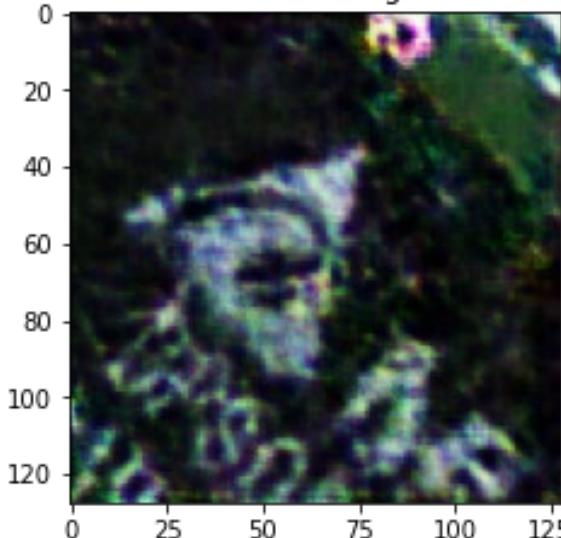
```

Fake image size: (128, 128, 3)

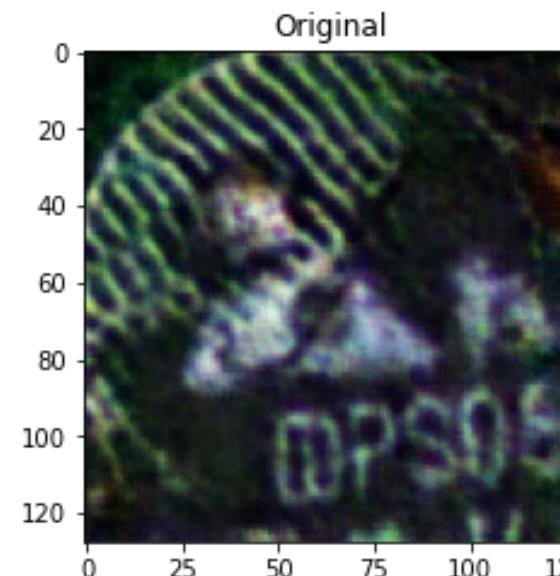
Feature 3=50



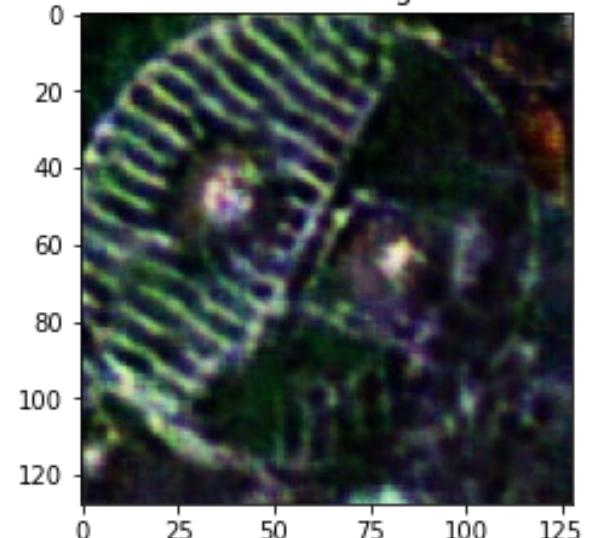
Feature changed



Feature 123=-50



Feature changed

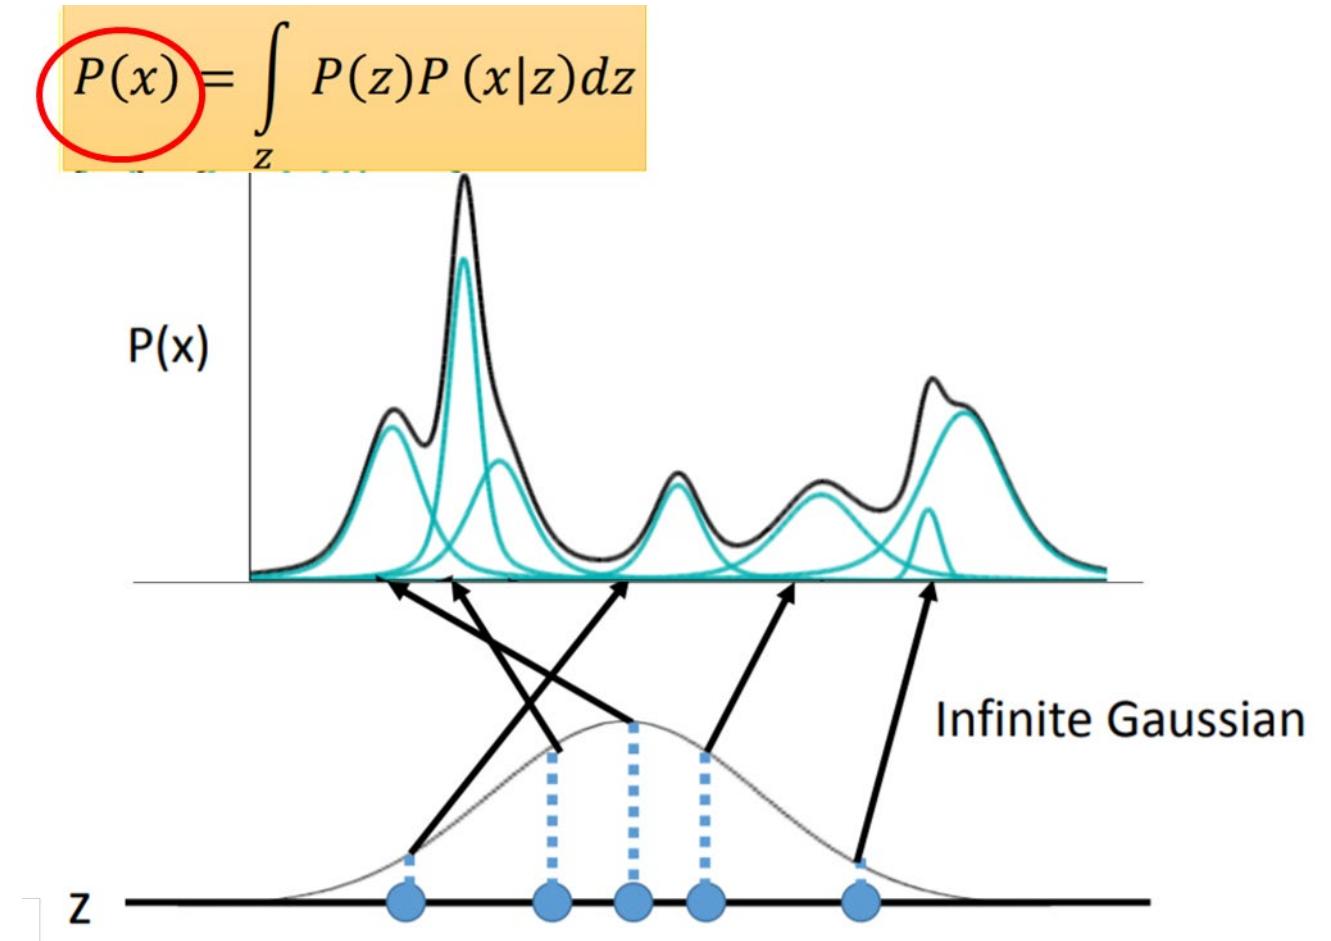
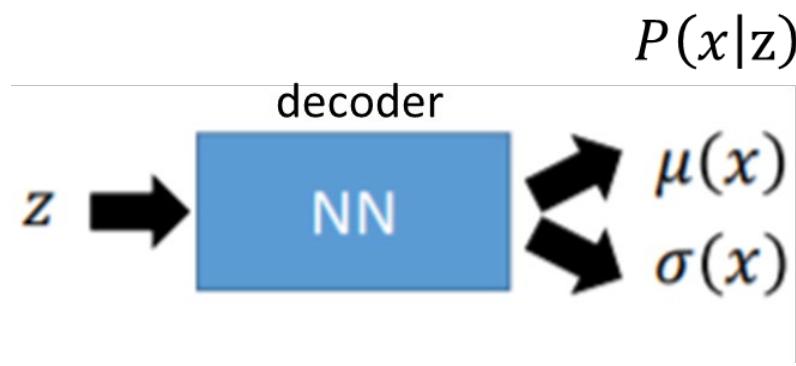


Theory behind GAN, VAE, PPO

The concept of deep generative model

1. Use probability distribution model to model the image generation process or the state-action generation process.
2. With this probability model, we want to maximize the likelihood to observe the training data or observe the trajectories with higher long-term rewards.
3. Maximize the likelihood can be achieved by minimizing the KL divergence.
3. Use DNN to represent a more general probability model.
4. Use DNN to represent a more general divergence measure.

(VAE) 1. Modelling the image generation process as a mixed Gaussian distribution model



(VAE) 2. Given the image generation model, we want to maximize the likelihood to observe the training images

Maximizing Likelihood

$$P(x) = \int_z P(z)P(x|z)dz$$

$$L = \sum_x \log P(x)$$

$P(z)$ is normal distribution

$x|z \sim N(\mu(z), \sigma(z))$

$\mu(z), \sigma(z)$ is going to be estimated

Maximizing the likelihood of the observed x

$$L = p(x^1) \times p(x^2) \times p(x^3) \times \cdots p(x^m) = \prod_{i=1, \dots, m} P(x^i)$$

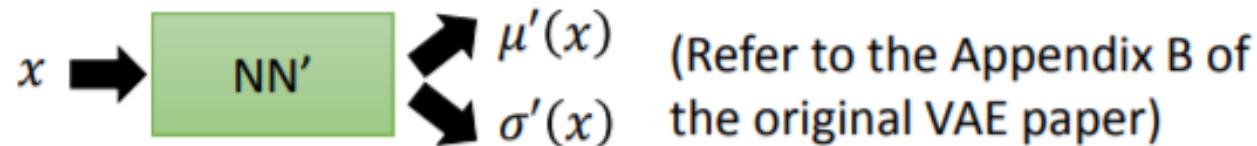
(VAE) 3. Maximize the likelihood can be achieved by minimizing KL divergence

Max. L_b can be done by min. $KL(q(z|x)||P(z))$ and max. $\int_z q(z|x) \log P(x|z) dz$. That is why loss = KLD + MSE (x, \hat{x})

Minimizing $KL(q(z|x)||P(z))$

$$\sum_{i=1}^3 (\exp(\sigma_i) - (1 + \sigma_i) + (m_i)^2)$$

Minimize



Maximizing

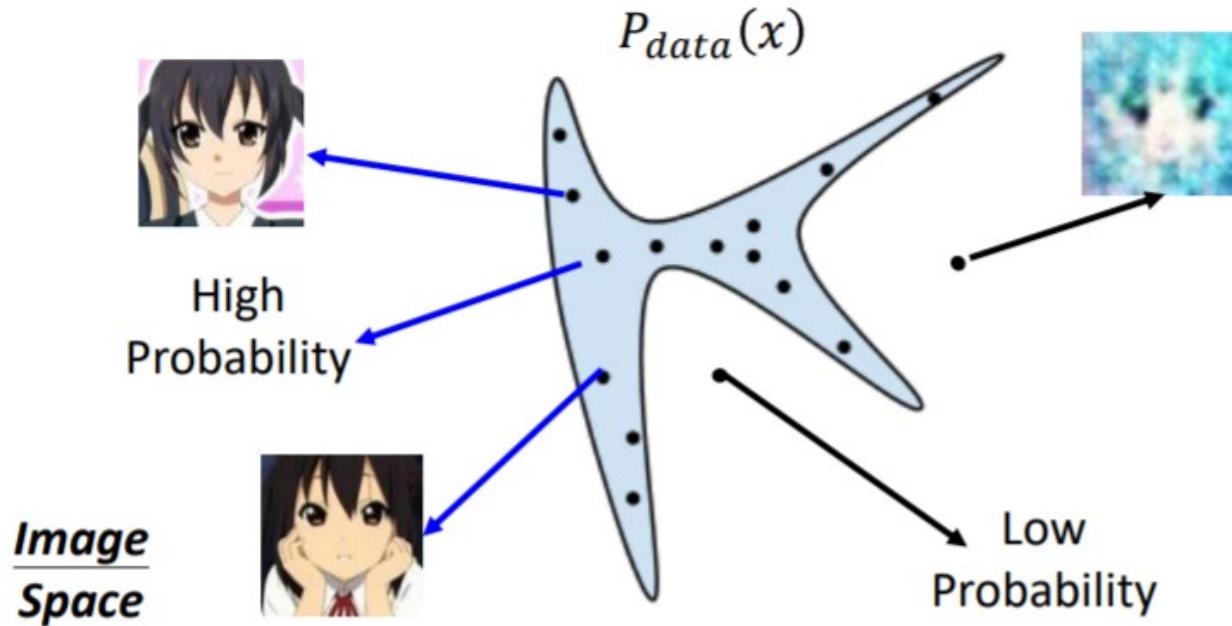
$$\int_z q(z|x) \log P(x|z) dz = E_{q(z|x)}[\log P(x|z)]$$



(GAN) 1. Modelling the generator as a probability distribution model P_G

X: an image (a high-dimensional vector)

- We want to find data distribution $P_{data}(x)$



(GAN) 2. Given the image generation model, we want to maximize the likelihood to observe the training images

- Given a data distribution $P_{data}(x)$ (We can sample from it.)
- We have a distribution $P_G(x; \theta)$ parameterized by θ
 - We want to find θ such that $P_G(x; \theta)$ close to $P_{data}(x)$
 - E.g. $P_G(x; \theta)$ is a Gaussian Mixture Model, θ are means and variances of the Gaussians

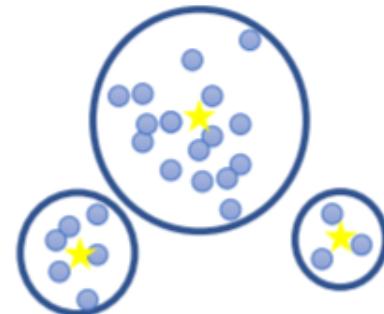
Sample $\{x^1, x^2, \dots, x^m\}$ from $P_{data}(x)$

We can compute $P_G(x^i; \theta)$

Likelihood of generating the samples

$$L = \prod_{i=1}^m P_G(x^i; \theta)$$

Find θ^* maximizing the likelihood



(GAN) 3. Maximize the likelihood can be achieved by minimizing the KL divergence

Maximum likelihood estimation = minimum
KL divergence

$$\theta^* = \arg \max_{\theta} \prod_{i=1}^m P_G(x^i; \theta) = \arg \min_{\theta} KL(P_{data} || P_G)$$

$$\begin{aligned}\theta^* &= \arg \max_{\theta} \prod_{i=1}^m P_G(x^i; \theta) = \arg \max_{\theta} \log \prod_{i=1}^m P_G(x^i; \theta) \\ &= \arg \max_{\theta} \sum_{i=1}^m \log P_G(x^i; \theta) \quad \{x^1, x^2, \dots, x^m\} \text{ from } P_{data}(x) \\ &\approx \arg \max_{\theta} E_{x \sim P_{data}} [\log P_G(x; \theta)] \\ &= \arg \max_{\theta} \int_x P_{data}(x) \log P_G(x; \theta) dx - \int_x P_{data}(x) \log P_{data}(x) dx \\ &= \arg \min_{\theta} KL(P_{data} || P_G) \quad \text{How to define a general } P_G?\end{aligned}$$

$$\int_x P_{data}(x) \log P_{data}(x) dx = \log P_{data}(x)$$

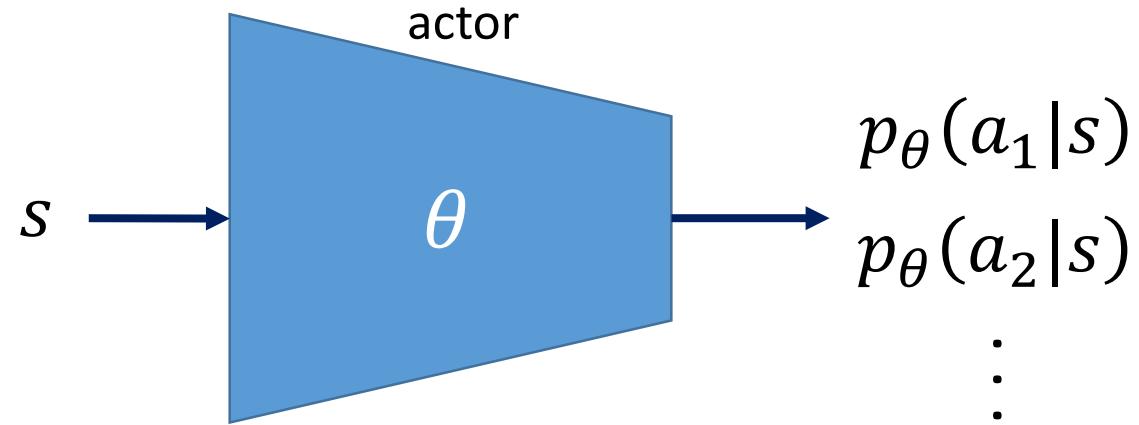
$$D_{KL}(q || p) = \sum_{i=1}^N q(x_i) \log\left(\frac{q(x_i)}{p(x_i)}\right)$$

(Classification) 1. Assuming the training data is generated from $y_1 = P_{w,b}(C_1 | x) = \sigma(w \cdot x + b)$ 2. Maximize the likelihood to observe the training data

Training Data	<table style="width: 100%; border-collapse: collapse;"> <tr> <td style="text-align: center;">x^1</td><td style="text-align: center;">x^2</td><td style="text-align: center;">x^3</td><td style="text-align: center;">.....</td><td style="text-align: center;">x^N</td></tr> <tr> <td style="text-align: center;">C_1</td><td style="text-align: center;">C_1</td><td style="text-align: center;">C_2</td><td style="text-align: center;">.....</td><td style="text-align: center;">C_1</td></tr> </table>	x^1	x^2	x^3	x^N	C_1	C_1	C_2	C_1
x^1	x^2	x^3	x^N							
C_1	C_1	C_2	C_1							

$$\begin{aligned}
 L(w, b) &= f_{w,b}(x^1)f_{w,b}(x^2)\left(1 - f_{w,b}(x^3)\right)\cdots f_{w,b}(x^N) \\
 -\ln L(w, b) &= \ln f_{w,b}(x^1) + \ln f_{w,b}(x^2) + \ln \left(1 - f_{w,b}(x^3)\right)\cdots \\
 &\quad \hat{y}^n: 1 \text{ for class 1, } 0 \text{ for class 2} \\
 &= \sum_n -\left[\hat{y}^n \ln f_{w,b}(x^n) + (1 - \hat{y}^n) \ln \left(1 - f_{w,b}(x^n)\right)\right] \\
 &\quad \text{Cross entropy between two Bernoulli distribution}
 \end{aligned}$$

(PPO) 1. Modelling the state-action mapping network as $p_\theta(a|s)$



(PPO) 2. Given the probability model, we want to maximize the likelihood to observe trajectories with higher long-term rewards

$$\tau = (s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_T, a_T)$$

$$p_\theta(\tau) = p(s_1)p_\theta(a_1|s_1)p(s_2|s_1, a_1)p_\theta(a_2|s_2)p(s_3|s_2, a_2)\dots$$

$$\bar{R}_\theta = \sum R(\tau) p_\theta(\tau) = E_{\tau \sim p_\theta(\tau)}[R(\tau)] \quad R(\tau) = \sum_{t=1}^T r_t$$

$$\max_\theta E[\bar{R}_\theta]$$

$$\begin{aligned} \nabla \bar{R}_\theta &= \sum R(\tau) \nabla p_\theta(\tau) = E_{\tau \sim p_\theta(\tau)}[R(\tau) \nabla \log p_\theta(\tau)] \approx \frac{1}{N} \sum_{n=1}^N R(\tau^n) \nabla \log p_\theta(\tau^n) \\ &= \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} R(\tau^n) \nabla \log p_\theta(a_t^n | s_t^n) \end{aligned}$$

(PPO) 3. To facilitate off-policy sampling, we minimize the KL divergence

Off-policy

$$\nabla \bar{R}_\theta = E_{(s_t, a_t) \sim \pi_{\theta'}} \left[\frac{p_\theta(a_t | s_t)}{p_{\theta'}(a_t | s_t)} A^{\theta'}(s_t, a_t) \nabla \log p_\theta(a_t^n | s_t^n) \right]$$

Loss function

$$J^{\theta'}(\theta) = E_{(s_t, a_t) \sim \pi_{\theta'}} \left[\frac{p_\theta(a_t | s_t)}{p_{\theta'}(a_t | s_t)} A^{\theta'}(s_t, a_t) \right]$$

Proximal policy optimization (PPO)

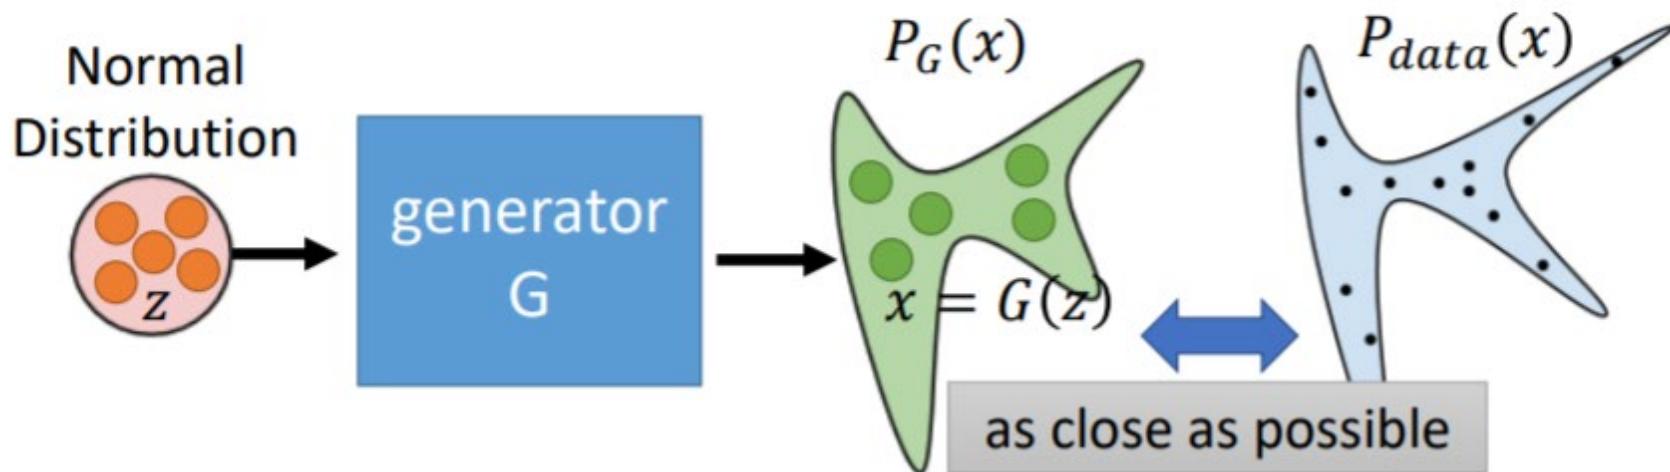
$$J_{PPO}^{\theta'}(\theta) = J^{\theta'}(\theta) - \beta KL(\theta, \theta')$$

$$J_{PPO2}^{\theta'}(\theta) = \sum_{(s_t, a_t)} \min \left(\frac{p_\theta(a_t | s_t)}{p_{\theta'}(a_t | s_t)} A^{\theta'}(s_t, a_t), \text{clip} \left(\frac{p_\theta(a_t | s_t)}{p_{\theta'}(a_t | s_t)}, 1 - \varepsilon, 1 + \varepsilon \right) A^{\theta'}(s_t, a_t) \right)$$

Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms. arXiv preprint arXiv:1707.06347.

4. Use DNN to represent a more general probability model

- A generator G is a network. The network defines a probability distribution P_G



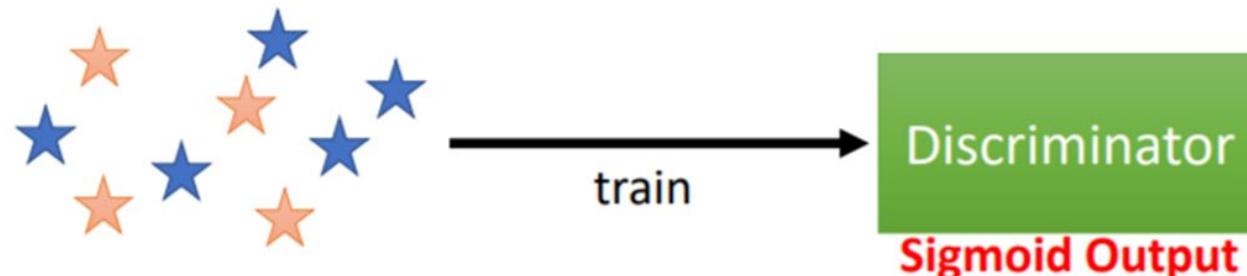
$$G^* = \arg \min_G \underline{Div}(P_G, P_{data})$$

Divergence between distributions P_G and P_{data}
How to compute the divergence?

5. Use DNN to represent a more general divergence measure

How to compute the divergence when a NN is used to represent a more general probability model? – train a discriminator NN to compute the divergence

- ★ : data sampled from P_{data}
- ★ : data sampled from P_G



Example Objective Function for D

$$V(G, D) = E_{x \sim P_{data}}[\log D(x)] + E_{x \sim P_G}[\log(1 - D(x))]$$

(G is fixed)

Training: $D^* = \arg \max_D V(D, G)$

[Goodfellow, et al., NIPS, 2014]

The objective function for D is related to JS divergence

$$V(D, G) = E_{x \sim P_{data}} [\log D(x)] + E_{x \sim P_G} [\log(1 - D(x))]$$

- Given G , what is the optimal D^* maximizing

$$\begin{aligned} V &= E_{x \sim P_{data}} [\log D(x)] + E_{x \sim P_G} [\log(1 - D(x))] \\ &= \int_x P_{data}(x) \log D(x) dx + \int_x P_G(x) \log(1 - D(x)) dx \\ &= \int_x [P_{data}(x) \log D(x) + P_G(x) \log(1 - D(x))] dx \end{aligned}$$

Assume that $D(x)$ can be any function

- Given x , the optimal D^* maximizing

$$P_{data}(x) \log D(x) + P_G(x) \log(1 - D(x))$$

The objective function for D is related to JS divergence

$$V(D, G) = E_{x \sim P_{data}} [\log D(x)] + E_{x \sim P_G} [\log(1 - D(x))]$$

$$\begin{aligned}\max_D V(G, D) &= V(G, D^*) & D^*(x) &= \frac{P_{data}(x)}{P_{data}(x) + P_G(x)} \\&= -2\log 2 + \int_x P_{data}(x) \log \frac{P_{data}(x)}{(P_{data}(x) + P_G(x))/2} dx \\&\quad + \int_x P_G(x) \log \frac{P_G(x)}{(P_{data}(x) + P_G(x))/2} dx \\&= -2\log 2 + \text{KL}\left(P_{data} \parallel \frac{P_{data} + P_G}{2}\right) + \text{KL}\left(P_G \parallel \frac{P_{data} + P_G}{2}\right) \\&= -2\log 2 + 2JSD(P_{data} \parallel P_G) \quad \text{Jensen-Shannon divergence}\end{aligned}$$

The objective function for D is implemented as binary cross entropy

$$V(D, G) = E_{x \sim P_{data}}[\log D(x)] + E_{x \sim P_G}[\log(1 - D(x))]$$

Binary Classifier

D is a binary classifier with sigmoid output (can be deep)

$\{x^1, x^2, \dots, x^m\}$ from $P_{data}(x)$  Positive examples

$\{\tilde{x}^1, \tilde{x}^2, \dots, \tilde{x}^m\}$ from $P_G(x)$  Negative examples

Minimize Cross-entropy

Cross entropy:

$$C(f(x^n), \hat{y}^n) = -[\hat{y}^n \ln f(x^n) + (1 - \hat{y}^n) \ln(1 - f(x^n))]$$

(GAN) Two adversarial neural networks G and D

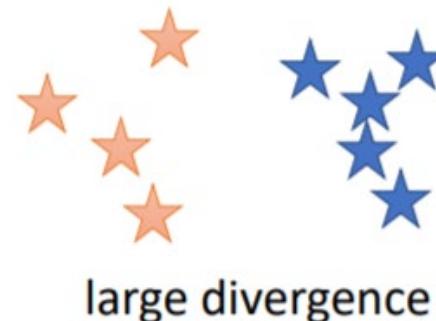
(1) Train generator

$$G^* = \arg \min_G \text{Div}(P_G, P_{\text{data}})$$



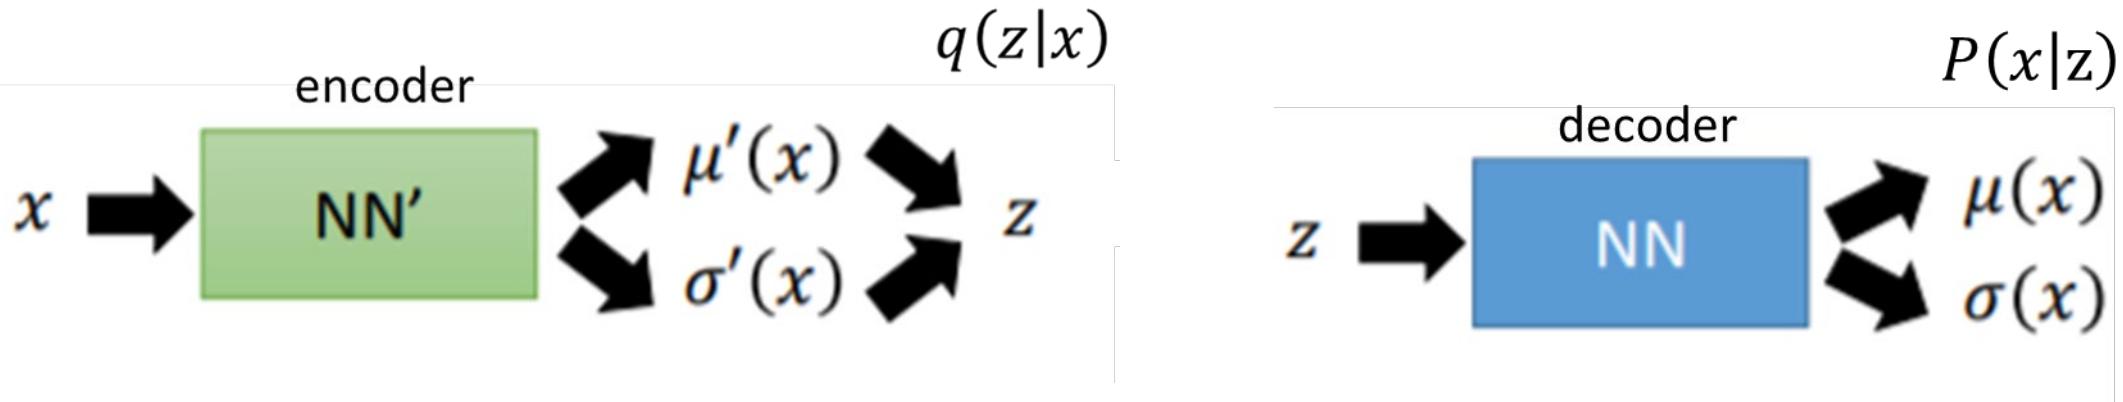
(2) Train discriminator

$$D^* = \arg \max_D V(D, G)$$

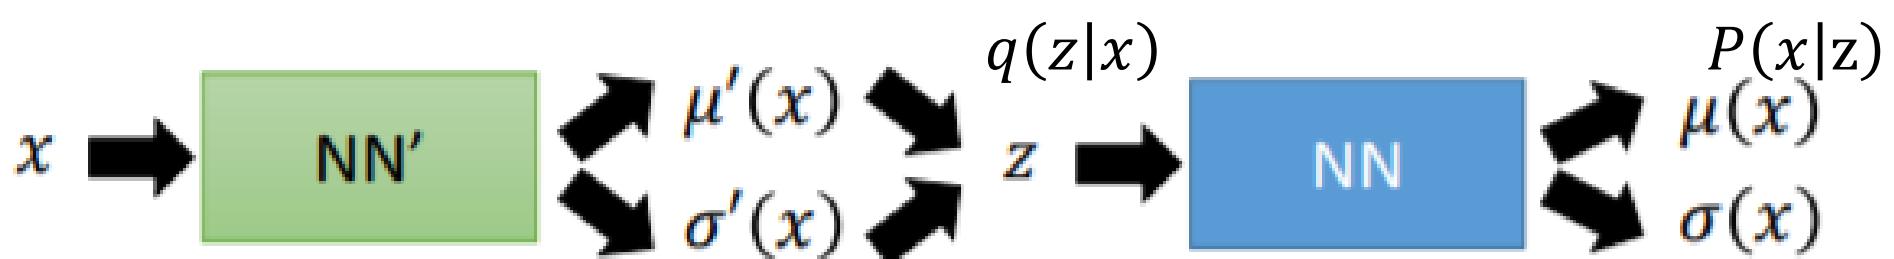


$$V(D, G) = E_{x \sim P_{\text{data}}}[\log D(x)] + E_{x \sim P_G}[\log(1 - D(x))]$$

(VAE) Two sequential neural networks, one encode and one decode

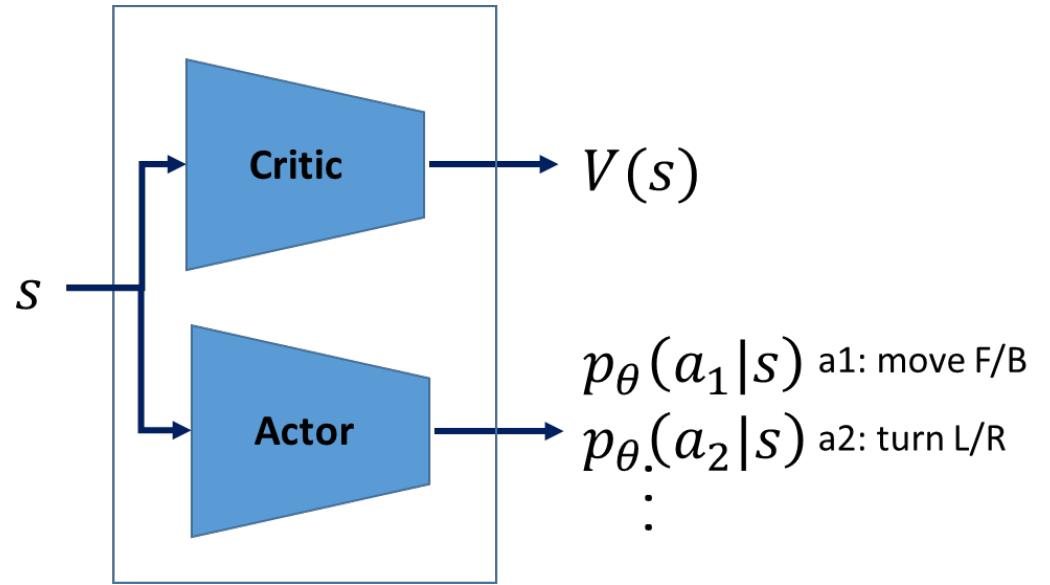


$$L = MSE(x, \hat{x}) + KL(q(z|x) || P(z))$$



(PPO-AC) Two parallel neural networks, one actor and one critic

$$L = c_v L_v + L_\pi - \beta L_{reg}$$



(1) Actor – Learns the best actions (that can have maximum long-term rewards)

$$L_\pi = \sum_{(s_t, a_t)} \min \left(\frac{p_\theta(a_t|s_t)}{p_{\theta'}(a_t|s_t)} A^{\theta'}(s_t, a_t), \text{clip} \left(\frac{p_\theta(a_t|s_t)}{p_{\theta'}(a_t|s_t)}, 1 - \varepsilon, 1 + \varepsilon \right) A^{\theta'}(s_t, a_t) \right)$$

(2) Critic – Learns the expected value of the long-term reward.

$$L_v = \text{MSE of (return} - \nu)$$

Algorithm

- Given G_0
- Find D_0^* maximizing $V(G_0, D)$

$V(G_0, D_0^*)$ is the JS divergence between $P_{data}(x)$ and $P_{G_0}(x)$

- $\theta_G \leftarrow \theta_G - \eta \partial V(G, D_0^*) / \partial \theta_G \rightarrow$ Obtain G_1
- Find D_1^* maximizing $V(G_1, D)$

Decrease JS divergence(?)

$V(G_1, D_1^*)$ is the JS divergence between $P_{data}(x)$ and $P_{G_1}(x)$

- $\theta_G \leftarrow \theta_G - \eta \partial V(G, D_1^*) / \partial \theta_G \rightarrow$ Obtain G_2
-

Decrease JS divergence(?)

$$G^* = \arg \min_G \max_D V(D, G)$$

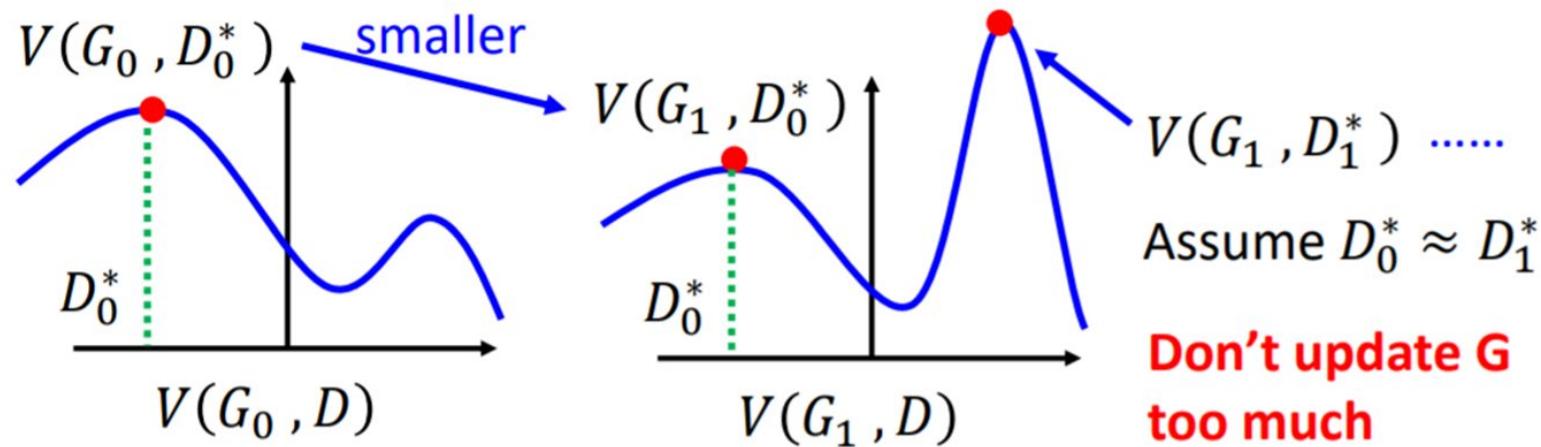
$$D^* = \arg \max_D V(D, G)$$

$$\begin{aligned} & V(D, G) \\ &= E_{x \sim P_{data}} [\log D(x)] + E_{x \sim P_G} [\log(1 - D(x))] \end{aligned}$$

$$G^* = \arg \min_G \text{Div}(P_G, P_{data})$$

Do not update G too much...

When we update G from G_0 to G_1 , the divergence measurement function $V(G_0, D^*)$ will no longer fully represent the JS divergence between P_G and P_{data}



(PPO): We do not update $p_\theta(a_t|s_t)$ too much so that we can calculate $p_{\theta'}(a_t|s_t)$ off-policy

$$L_\pi = \sum_{(s_t, a_t)} \min \left(\frac{p_\theta(a_t|s_t)}{p_{\theta'}(a_t|s_t)} A^{\theta'}(s_t, a_t), \text{clip} \left(\frac{p_\theta(a_t|s_t)}{p_{\theta'}(a_t|s_t)}, 1 - \varepsilon, 1 + \varepsilon \right) A^{\theta'}(s_t, a_t) \right)$$

But we should update D as much as possible

- In each training iteration:

- Sample m examples $\{x^1, x^2, \dots, x^m\}$ from data distribution $P_{data}(x)$
- Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from the prior $P_{prior}(z)$
- Obtaining generated data $\{\tilde{x}^1, \tilde{x}^2, \dots, \tilde{x}^m\}, \tilde{x}^i = G(z^i)$
- Update discriminator parameters θ_d to maximize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log D(x^i) + \frac{1}{m} \sum_{i=1}^m \log (1 - D(\tilde{x}^i))$
 - $\theta_d \leftarrow \theta_d + \eta \nabla \tilde{V}(\theta_d)$
- Sample another m noise samples $\{z^1, z^2, \dots, z^m\}$ from the prior $P_{prior}(z)$
- Update generator parameters θ_g to minimize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log D(x^i) + \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^i)))$
 - $\theta_g \leftarrow \theta_g - \eta \nabla \tilde{V}(\theta_g)$

Learning
D

Repeat
k times

Learning
G

Only
Once

In [19]: `def train_discriminator(real_images, opt_d):`

`# Clear discriminator gradients`
`opt_d.zero_grad()`

For loop for k times?

`# Pass real images through discriminator`
`real_preds = discriminator(real_images)`
`real_targets = torch.ones(real_images.size(0), 1, device)`
`real_loss = F.binary_cross_entropy(real_preds, real_targets)`
`real_score = torch.mean(real_preds).item()`

`# Generate fake images`

`latent = torch.randn(batch_size, latent_size, 1, 1, device)`
`fake_images = generator(latent.to(device))`

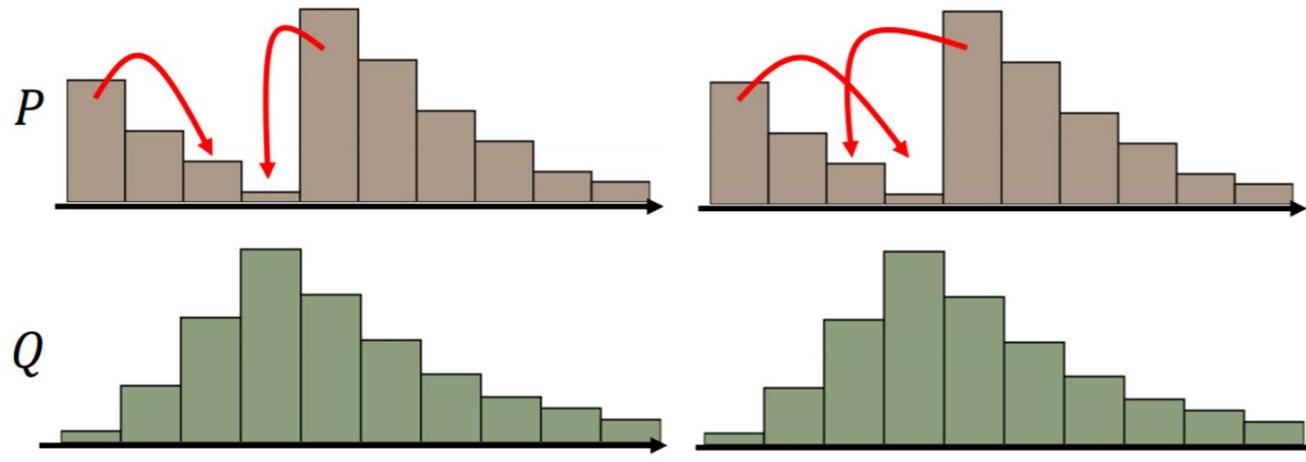
`# Pass fake images through discriminator`
`fake_targets = torch.zeros(fake_images.size(0), 1, device)`
`fake_preds = discriminator(fake_images)`
`fake_loss = F.binary_cross_entropy(fake_preds, fake_targets)`
`fake_score = torch.mean(fake_preds).item()`

`# Update discriminator weights`
`loss = real_loss + fake_loss`
`loss.backward()`
`opt_d.step()`

`return loss.item(), real_score, fake_score`

Wasserstein GAN (WGAN)

W-GAN: Earth move's distance



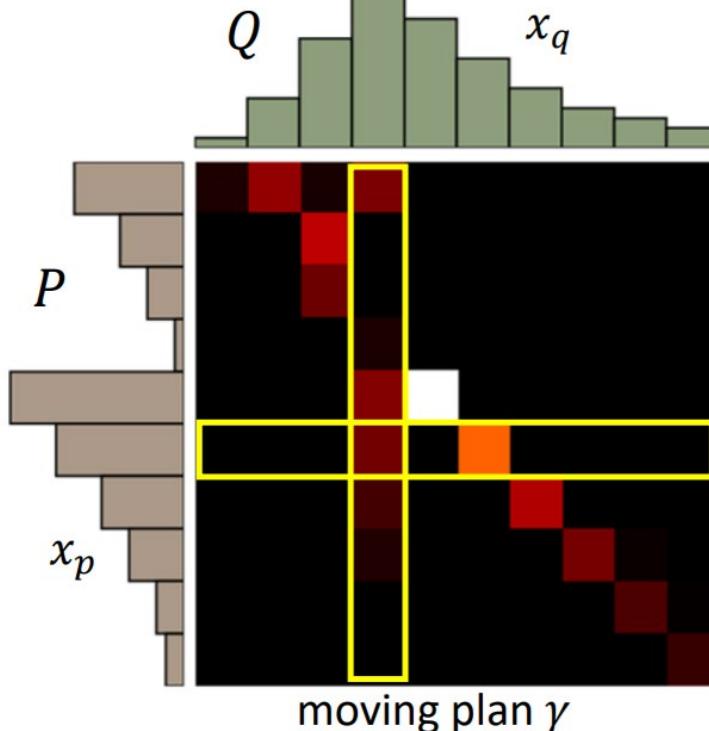
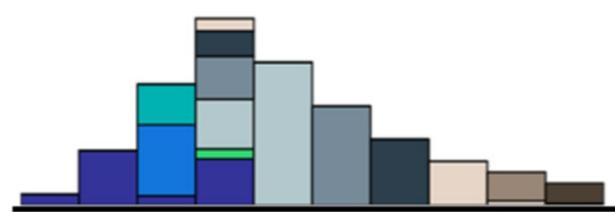
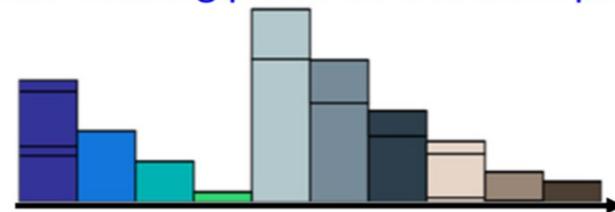
Average distance of a plan γ :

$$B(\gamma) = \sum_{x_p, x_q} \gamma(x_p, x_q) \|x_p - x_q\|$$

Earth Mover's Distance:

$$W(P, Q) = \min_{\gamma \in \Pi} B(\gamma)$$

Best “moving plans” of this example



W-GAN

Evaluate wasserstein distance between P_{data} and P_G

$$V(G, D) = \max_{D \in \underline{1-Lipschitz}} \left\{ E_{x \sim P_{data}}[D(x)] - E_{x \sim P_G}[D(x)] \right\}$$

D has to be smooth enough.

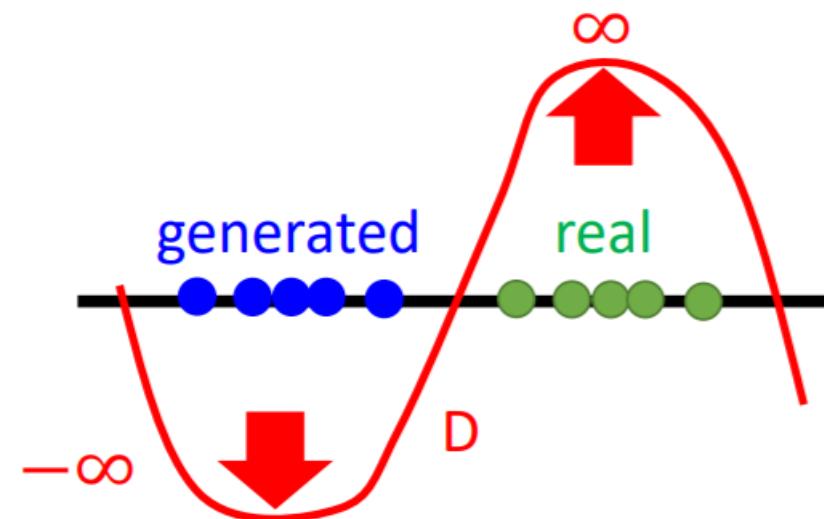
Lipschitz Function

$$\|f(x_1) - f(x_2)\| \leq K \|x_1 - x_2\|$$

Output
change

Input
change

K=1 for "1 – Lipschitz"



GAN is sensitive to hyper-parameter tuning and its performance range is large. Different GANs' performances are similar

