

(1) Reach goal

# Training setting

$$s = (\Delta x, \Delta y, \Delta z, \theta_B, \theta_U, \theta_L, \theta_W)$$

$$a = (\Delta\theta_B, \Delta\theta_U, \Delta\theta_L, \Delta\theta_W)$$

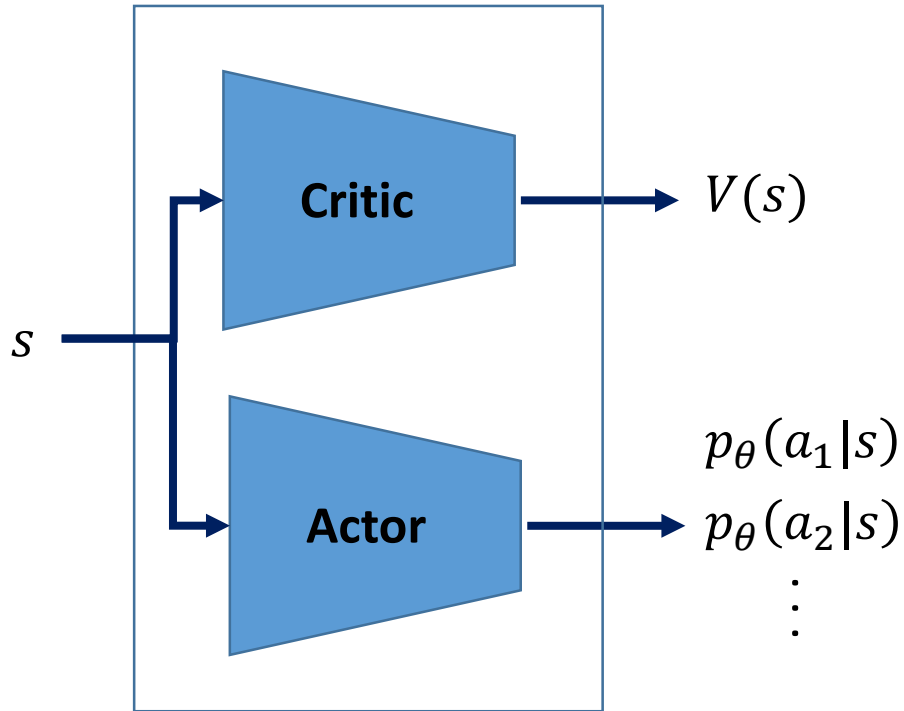
$$r = \begin{cases} -0.005 & \text{per step} \\ -5 & \text{collision, out of range} \\ +20 & \text{goal, } d \leq 0.5 \end{cases}$$

No. of training environment = 9

Goal initialize = randomly positioned in polar system  $\theta = -80 \sim 80$ ,  $r = 0.8 \sim 1.5$

Arm initialize:  $(\theta_B = 0, \theta_U = 45, \theta_L = 45, \theta_W = 45)$

# Summary of training



NN: 7-512-512-512-4

Time horizon = 2000

Buffer size = 20480

Batch size = 2048

$$\tau = (s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_T, a_{T=2000})$$

$$\nabla \bar{R}_\theta \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \left( \sum_{t'=t}^{T_n} \gamma^{t'-t} r_{t'}^n - b \right) \nabla \log p_\theta(a_t^n | s_t^n)$$

$V^{\pi_\theta}(s_t^i)$  Expected value of  $b$   
 $E[G_t^i] = Q^{\pi_\theta}(s_t^i, a_t^i)$  Expected value of  $G_t^i$

$$A^\theta(s_t, a_t) = (r_t^n + V^{\pi_\theta}(s_{t+1}^n) - V^{\pi_\theta}(s_t^n))$$

$$L_v = (G_t^n - V^{\pi_\theta}(s_t^n))^2 = (r_t^n + \gamma V^{\pi_\theta}(s_{t+1}^n) - V^{\pi_\theta}(s_t^n))^2$$

$$L_\pi = \sum_{(s_t, a_t)} \min \left( \frac{p_\theta(a_t | s_t)}{p_{\theta'}(a_t | s_t)} A^{\theta'}(s_t, a_t), \text{clip} \left( \frac{p_\theta(a_t | s_t)}{p_{\theta'}(a_t | s_t)}, 1 - \varepsilon, 1 + \varepsilon \right) A^{\theta'}(s_t, a_t) \right)$$

$$L = L_\pi + c_v L_v + c_{reg} L_{reg}$$


# HW4(1)

- Describe the training setting
- Show tensor board plots and discuss your training performance
- Describe your test performance




# Typical successful training experiences, goal reach $d \leq 0.25$


Cumulative Reward  
tag: Environment/Cumulative Reward




Episode Length  
tag: Environment/Episode Length




Cumulative Reward  
tag: Environment/Cumulative Reward




Episode Length  
tag: Environment/Episode Length



Policy Loss  
tag: Losses/Policy Loss




Value Loss  
tag: Losses/Value Loss




# Failed training, goal reach $d \leq 0.25$

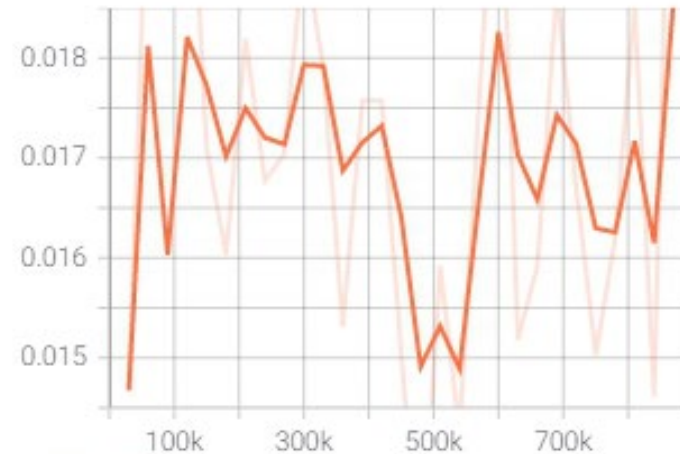
Cumulative Reward  
tag: Environment/Cumulative Reward




Episode Length  
tag: Environment/Episode Length



Policy Loss  
tag: Losses/Policy Loss



Value Loss  
tag: Losses/Value Loss

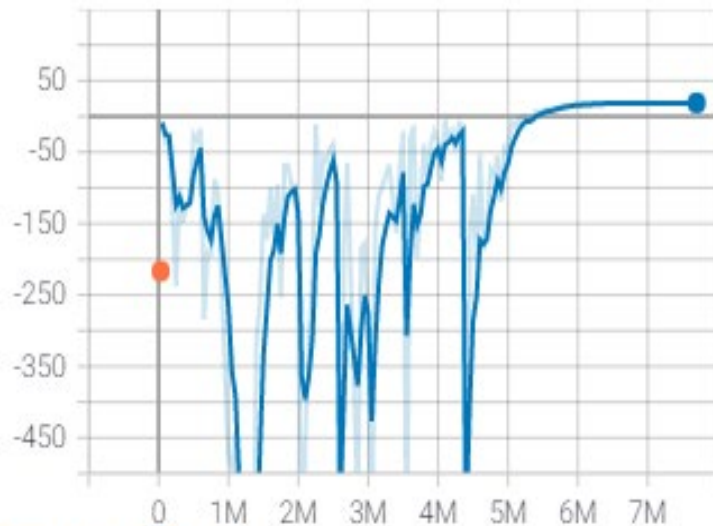


# Successful training, goal reach $d \leq 0.1$

In this training, the point reach threshold is set to be 0.1, which makes it very difficult for AI to learn to reach goal. For the first 5M, the reward performance looks un-hope and most people will give up. But luckily the reward goes to maximum and become steady after 6M.

Cumulative Reward


tag: Environment/Cumulative Reward



run to download ▼

Episode Length


tag: Environment/Episode Length



run to download ▼


Losses

Policy Loss  
tag: Losses/Policy Loss



run to download ▼

Value Loss  
tag: Losses/Value Loss




run to download ▼


(2) Pick up goal and place it to white area



# Open "s3 – Pick up red and put to white"



# Training setting



$$s = (sNo, \Delta_1 x, \Delta_1 y, \Delta_1 z, \theta_B, \theta_U, \theta_L, \theta_W) \quad s = (sNo, \Delta_2 x, \Delta_2 y, \Delta_2 z, \theta_B, \theta_U, \theta_L, \theta_W)$$

$$r = \begin{cases} -0.005 \\ -5 \\ +20 \end{cases} \quad d_1 \leq 0.5$$

$$r = \begin{cases} -0.005 & \text{per step} \\ -5 & \text{collision, out of range} \\ +20 & \text{goal, } d_2 \leq 0.5 \end{cases}$$

NN: 8-512-512-512-4

Time horizon = 2000

Buffer size = 20480

Batch size = 2048

$$a = (\Delta\theta_B, \Delta\theta_U, \Delta\theta_L, \Delta\theta_W)$$


No. of training environment = 9

Goal initialize = randomly positioned in polar system  $\theta = -80 \sim 80$ ,  $r = 0.8 \sim 1.5$

Goal2 initialize = same as goal 1

Arm initialize:  $(\theta_B = 0, \theta_U = 45, \theta_L = 45, \theta_W = 45)$

# State




$$s = (sNo, \Delta_1 x, \Delta_1 y, \Delta_1 z, \theta_B, \theta_U, \theta_L, \theta_W) \quad s = (sNo, \Delta_2 x, \Delta_2 y, \Delta_2 z, \theta_B, \theta_U, \theta_L, \theta_W)$$

```
sensor.AddObservation(stage);
```


```
if (stage == 1)  
    sensor.AddObservation(EndTouchPlane.position - goalUpTouchPt.posi  
else //stage =2  
    sensor.AddObservation(goalDownTouchPt.position - goal2UpTouchPt.p
```

```
float BaseRotationAngle = UnityEditor.TransformUtils.GetInspectorRota  
float UArmRotationAngle = UnityEditor.TransformUtils.GetInspectorRota  
float LArmRotationAngle = UnityEditor.TransformUtils.GetInspectorRota  
float WRotationAngle = UnityEditor.TransformUtils.GetInspectorRotatio
```

```
sensor.AddObservation(BaseRotationAngle);  
sensor.AddObservation(UArmRotationAngle);  
sensor.AddObservation(LArmRotationAngle);  
sensor.AddObservation(WRotationAngle);
```



# Rewards



$$r = \begin{cases} -0.005 \\ -5 \\ +20 \end{cases} \quad d_1 \leq 0.5$$

$$r = \begin{cases} -0.005 & \text{per step} \\ -5 & \text{collision, out of range} \\ +20 & \text{goal, } d_2 \leq 0.5 \end{cases}$$

```
if (stage ==1 && PointTouch(EndTouchPlane, goalUpTouchPt, 0.1f))
{
    msg = trainingVE.name + " Goal 1! \n";
    Debug.Log(msg);
    stage = 2;
    AddReward(15.0f);
    goal.transform.parent = EndPivot.transform; //grab goal
}
```


```
else if (PointTouch(goalDownTouchPt, goal2UpTouchPt, 0.3f))
{
    msg = trainingVE.name + " Goal 2! \n";
    Debug.Log(msg);
    AddReward(100.0f);
    EndEpisode();
}
```

# I quit at 3M, looks promising


```
TouchCube. Step: 1400000. Time Elapsed: 1521.994 s. Mean Reward: 40.746. Std of Reward: 56.700. Training.
TouchCube. Step: 1450000. Time Elapsed: 1596.218 s. Mean Reward: 44.816. Std of Reward: 57.113. Training.
TouchCube. Step: 1500000. Time Elapsed: 1654.829 s. Mean Reward: 41.636. Std of Reward: 56.026. Training.
zation.py:93] Converting to results\1\TouchCube\TouchCube-1499998.onnx
zation.py:105] Exported results\1\TouchCube\TouchCube-1499998.onnx
TouchCube. Step: 1550000. Time Elapsed: 1710.662 s. Mean Reward: 53.289. Std of Reward: 58.197. Training.
TouchCube. Step: 1600000. Time Elapsed: 1772.814 s. Mean Reward: 55.664. Std of Reward: 57.972. Training.
TouchCube. Step: 1650000. Time Elapsed: 1839.616 s. Mean Reward: 66.710. Std of Reward: 57.145. Training.
TouchCube. Step: 1700000. Time Elapsed: 1905.361 s. Mean Reward: 68.004. Std of Reward: 56.702. Training.
TouchCube. Step: 1750000. Time Elapsed: 1974.348 s. Mean Reward: 71.199. Std of Reward: 55.970. Training.
TouchCube. Step: 1800000. Time Elapsed: 2034.998 s. Mean Reward: 75.024. Std of Reward: 54.928. Training.
TouchCube. Step: 1850000. Time Elapsed: 2093.756 s. Mean Reward: 77.129. Std of Reward: 54.333. Training.
TouchCube. Step: 1900000. Time Elapsed: 2153.092 s. Mean Reward: 80.611. Std of Reward: 52.868. Training.
TouchCube. Step: 1950000. Time Elapsed: 2214.371 s. Mean Reward: 82.048. Std of Reward: 52.347. Training.
TouchCube. Step: 2000000. Time Elapsed: 2273.202 s. Mean Reward: 83.202. Std of Reward: 51.604. Training.
zation.py:93] Converting to results\1\TouchCube\TouchCube-1999932.onnx
zation.py:105] Exported results\1\TouchCube\TouchCube-1999932.onnx
TouchCube. Step: 2050000. Time Elapsed: 2335.471 s. Mean Reward: 85.839. Std of Reward: 50.198. Training.
TouchCube. Step: 2100000. Time Elapsed: 2396.409 s. Mean Reward: 89.061. Std of Reward: 48.055. Training.
TouchCube. Step: 2150000. Time Elapsed: 2460.368 s. Mean Reward: 86.739. Std of Reward: 49.768. Training.
TouchCube. Step: 2200000. Time Elapsed: 2521.083 s. Mean Reward: 87.105. Std of Reward: 49.479. Training.
TouchCube. Step: 2250000. Time Elapsed: 2583.198 s. Mean Reward: 89.747. Std of Reward: 47.747. Training.
TouchCube. Step: 2300000. Time Elapsed: 2643.755 s. Mean Reward: 89.104. Std of Reward: 48.263. Training.
TouchCube. Step: 2350000. Time Elapsed: 2705.644 s. Mean Reward: 91.258. Std of Reward: 46.869. Training.
TouchCube. Step: 2400000. Time Elapsed: 2770.443 s. Mean Reward: 88.986. Std of Reward: 48.436. Training.
TouchCube. Step: 2450000. Time Elapsed: 2833.605 s. Mean Reward: 88.044. Std of Reward: 49.433. Training.
TouchCube. Step: 2500000. Time Elapsed: 2895.328 s. Mean Reward: 90.368. Std of Reward: 48.504. Training.
zation.py:93] Converting to results\1\TouchCube\TouchCube-2499995.onnx
zation.py:105] Exported results\1\TouchCube\TouchCube-2499995.onnx
TouchCube. Step: 2550000. Time Elapsed: 2956.370 s. Mean Reward: 89.639. Std of Reward: 48.371. Training.
TouchCube. Step: 2600000. Time Elapsed: 3028.253 s. Mean Reward: 90.657. Std of Reward: 47.175. Training.
TouchCube. Step: 2650000. Time Elapsed: 3095.190 s. Mean Reward: 90.670. Std of Reward: 47.247. Training.
TouchCube. Step: 2700000. Time Elapsed: 3157.692 s. Mean Reward: 90.483. Std of Reward: 47.232. Training.
TouchCube. Step: 2750000. Time Elapsed: 3226.099 s. Mean Reward: 91.113. Std of Reward: 46.861. Training.
TouchCube. Step: 2800000. Time Elapsed: 3285.665 s. Mean Reward: 91.785. Std of Reward: 46.426. Training.
TouchCube. Step: 2850000. Time Elapsed: 3349.785 s. Mean Reward: 92.851. Std of Reward: 46.171. Training.
TouchCube. Step: 2900000. Time Elapsed: 3413.532 s. Mean Reward: 91.497. Std of Reward: 46.714. Training.
TouchCube. Step: 2950000. Time Elapsed: 3480.496 s. Mean Reward: 89.342. Std of Reward: 48.568. Training.
TouchCube. Step: 3000000. Time Elapsed: 3545.361 s. Mean Reward: 90.762. Std of Reward: 47.138. Training.
```

# I quit at 3M, looks promising


Cumulative Reward  
tag: Environment/Cumulative Reward




Episode Length  
tag: Environment/Episode Length



Policy Loss  
tag: Losses/Policy Loss



Value Loss  
tag: Losses/Value Loss





# HW4(2)

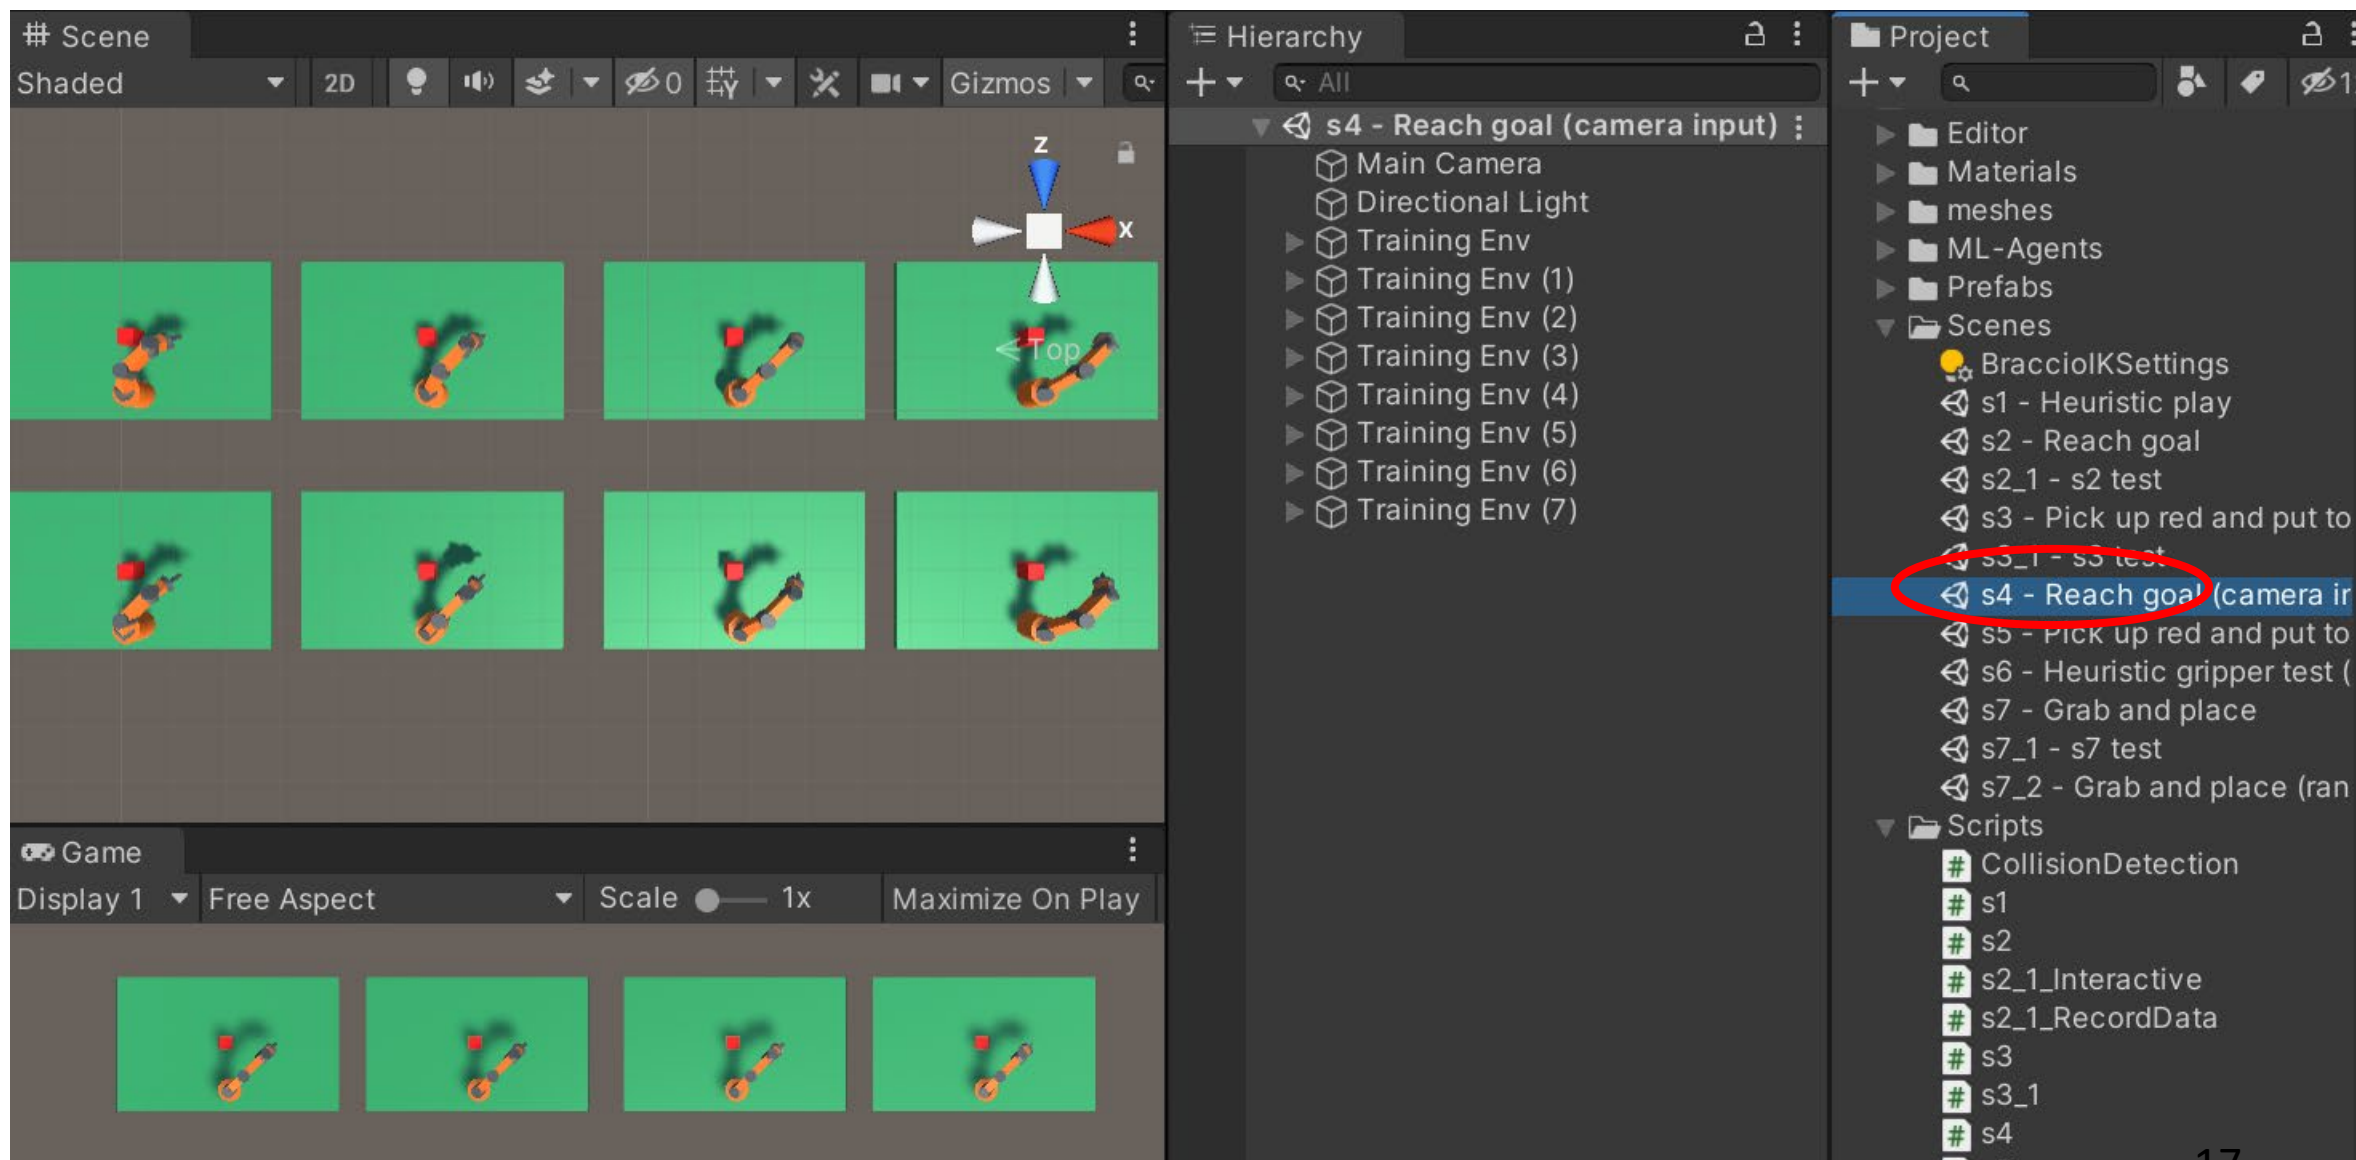
- Describe the training setting
- Show tensor board plots and discuss your training performance
- Describe your test performance



(3) Reach goal using camera sensor



# Open s4



# Training setting

$s$  = feature map vector from a CNN, size = ?

Input image to the CNN is captured by a camera from top, size = 84x84x3

$$a = (\Delta\theta_B, \Delta\theta_U, \Delta\theta_L, \Delta\theta_W)$$


$$r = \begin{cases} -0.005 & \text{per step} \\ -5 & \text{collision, out of range} \\ +20 & \text{goal, } d \leq 0.5 \end{cases}$$

No. of training environment = 9


Goal initialize = randomly positioned in polar system  $\theta = -80 \sim 80$ ,  $r = 0.8 \sim 1.5$

Arm initialize:  $(\theta_b = 0, \theta_u = 45, \theta_l = 45, \theta_w = 45)$

# Add camera sensor to the robot agent




# Add camera sensor to the robot agent



# Vector observation = 0

```
public override void CollectObservations(VectorSensor sensor)  
{  
    ...  
}
```




Looks good shortly (600K only)

```
TouchCube. Step: 50000. Time Elapsed: 222.493 s. Mean Reward: -1.205. Std of Reward: 12.602. Training.
TouchCube. Step: 100000. Time Elapsed: 417.921 s. Mean Reward: 5.743. Std of Reward: 13.921. Training.
TouchCube. Step: 150000. Time Elapsed: 618.958 s. Mean Reward: 12.166. Std of Reward: 11.773. Training.
TouchCube. Step: 200000. Time Elapsed: 819.051 s. Mean Reward: 15.312. Std of Reward: 9.549. Training.
TouchCube. Step: 250000. Time Elapsed: 1027.214 s. Mean Reward: 16.052. Std of Reward: 9.106. Training.
TouchCube. Step: 300000. Time Elapsed: 1230.887 s. Mean Reward: 18.197. Std of Reward: 5.736. Training.
TouchCube. Step: 350000. Time Elapsed: 1446.545 s. Mean Reward: 19.086. Std of Reward: 3.321. Training.
TouchCube. Step: 400000. Time Elapsed: 1662.651 s. Mean Reward: 19.512. Std of Reward: 1.555. Training.
TouchCube. Step: 450000. Time Elapsed: 1899.574 s. Mean Reward: 19.595. Std of Reward: 1.210. Training.
TouchCube. Step: 500000. Time Elapsed: 2154.839 s. Mean Reward: 19.635. Std of Reward: 0.976. Training.
zation.py:93] Converting to results\1\TouchCube\TouchCube-499992.onnx
ges\mlagents\trainers\torch\distributions.py:163: TracerWarning: Converting a tensor to a Python index might
n't record the data flow of Python values, so this value will be treated as a constant in the future. This
e to other inputs!
] * inputs.shape[0], axis=0)
ges\mlagents\trainers\torch\networks.py:352: TracerWarning: torch.Tensor results are registered as constant
is warning if you use this function to create tensors out of constant variables that would be the same eve
her case, this might cause the trace to be incorrect.
y_size]),
zation.py:105] Exported results\1\TouchCube\TouchCube-499992.onnx
TouchCube. Step: 550000. Time Elapsed: 2391.675 s. Mean Reward: 19.671. Std of Reward: 0.946. Training.
TouchCube. Step: 600000. Time Elapsed: 2647.336 s. Mean Reward: 19.707. Std of Reward: 0.785. Training.
```




# Good results only after 650K


Cumulative Reward  
tag: Environment/Cumulative Reward




Episode Length  
tag: Environment/Episode Length



Policy Loss  
tag: Losses/Policy Loss




Value Loss  
tag: Losses/Value Loss



# For simple environment, image input is easier than vector input


Cumulative Reward

tag: Environment/Cumulative Reward



Episode Length


tag: Environment/Episode Length



$s$  = feature map vector from a CNN, size = ?  
Input image to the CNN is captured by a camera from top, size = 84x84x3


Cumulative Reward

tag: Environment/Cumulative Reward



Episode Length

tag: Environment/Episode Length



$$s = (\Delta x, \Delta y, \Delta z, \theta_B, \theta_U, \theta_L, \theta_W)$$



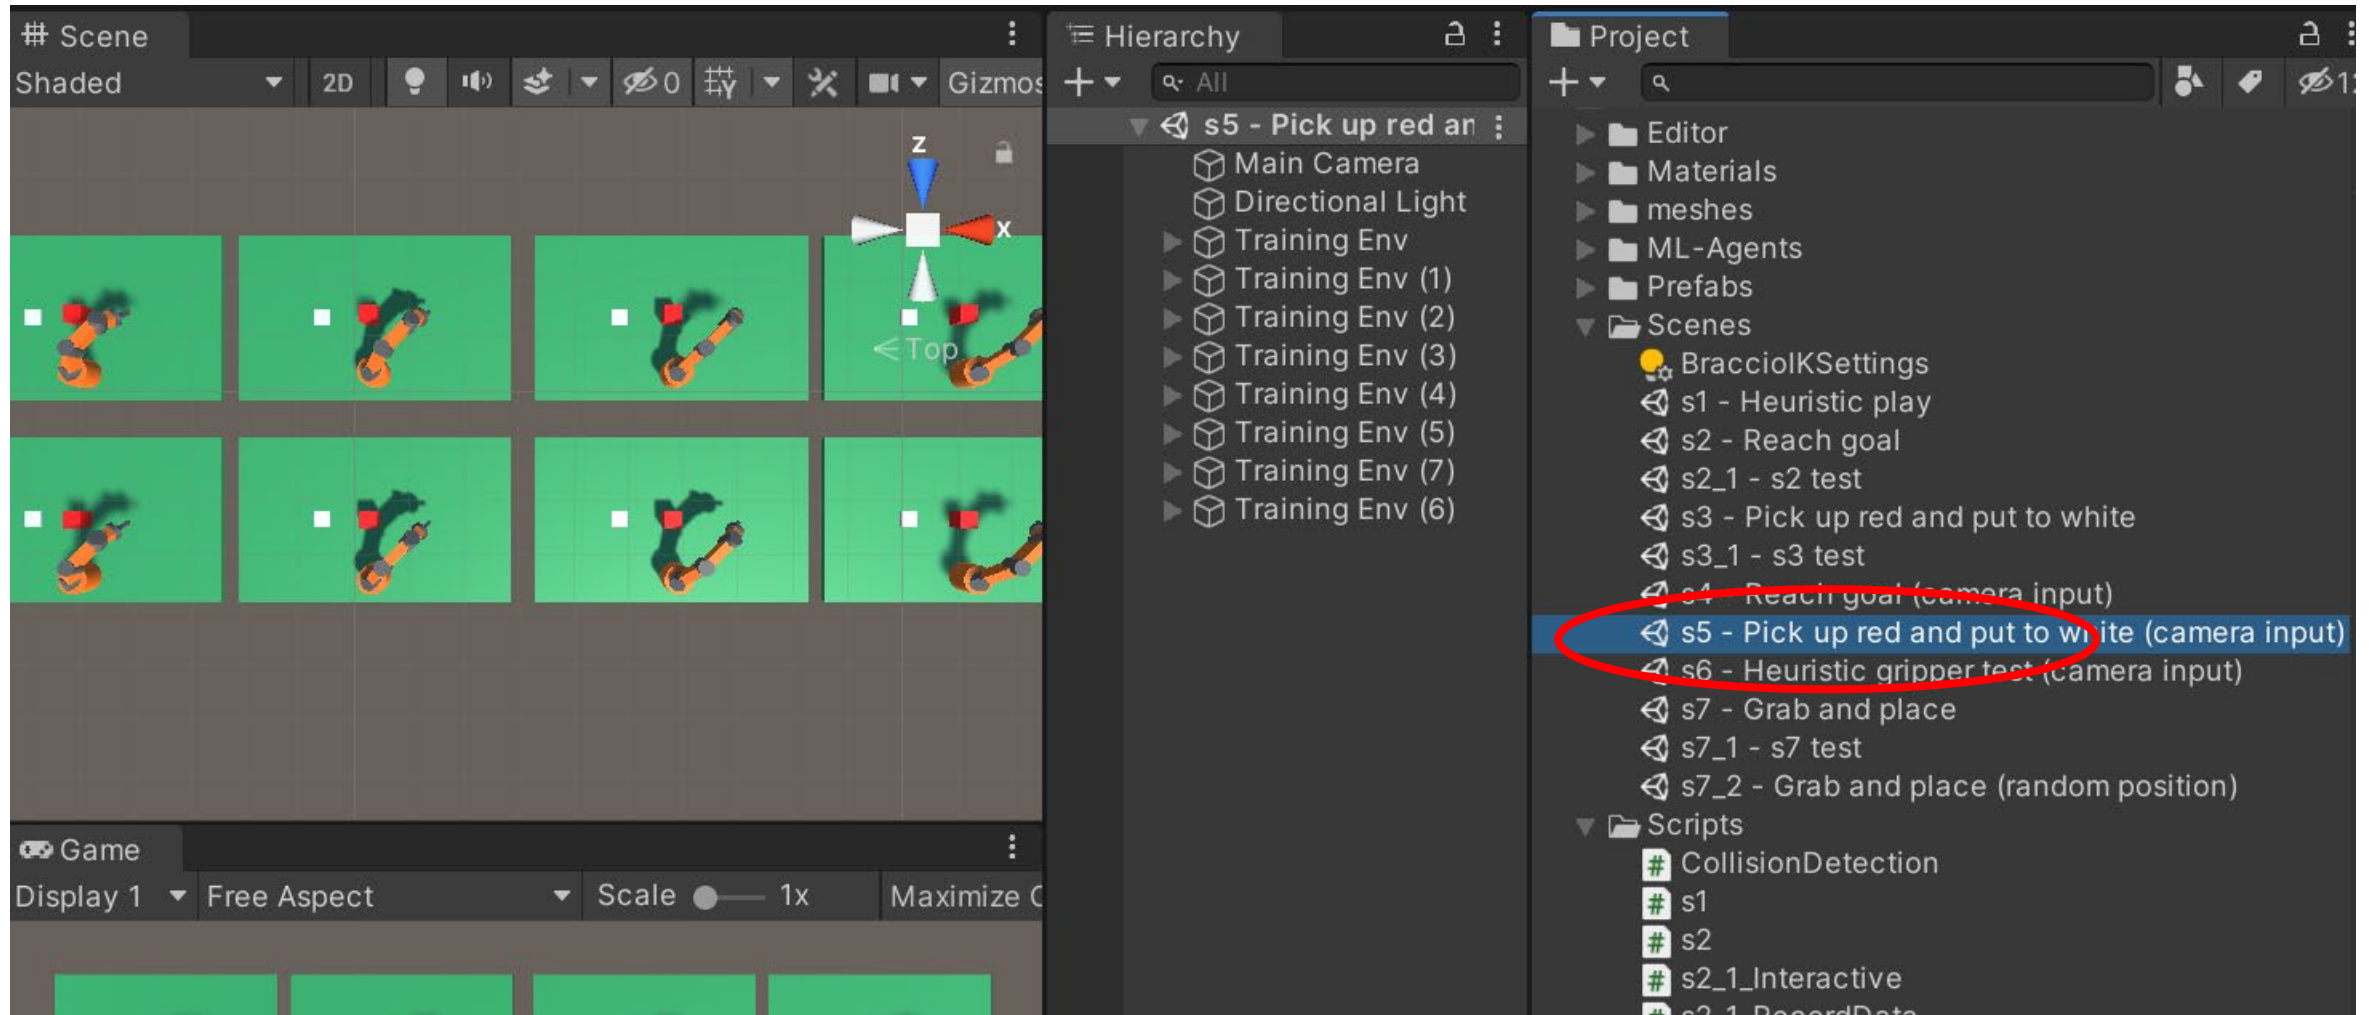
# HW4(3)

- Describe the training setting
- Show tensor board plots and discuss your training performance
- Describe your test performance




(4) Pick red cube and place it on top of the white cube using camera input

# Open s5



# Training setting



$s$  = feature map vector from a CNN, size = ?

Input image to the CNN is captured by a camera from top, size = 84x84x3

$$r = \begin{cases} -0.005 & \text{per step} \\ -5 & \text{collision, out of range} \\ +20 & \text{goal, } d_2 \leq 0.5 \end{cases}$$

$$a = (\Delta\theta_B, \Delta\theta_U, \Delta\theta_L, \Delta\theta_W)$$

No. of training environment = 9

Goal initialize = randomly positioned in polar system  $\theta = -80 \sim 80$ ,  $r = 0.8 \sim 1.5$

Goal2 initialize = same as goal 1

Arm initialize:  $(\theta_B = 0, \theta_U = 45, \theta_L = 45, \theta_W = 45)$

NN: ?-512-512-512-4

Time horizon = 2000

Buffer size = 20480

Batch size = 2048

# HW4(4)

- Describe the training setting
- Show tensor board plots and discuss your training performance
- Describe your test performance

